

Natural Language
Processing,
Large Language
Models, Large
Multimodal Models,
Pretrained Models,
LLMs, LangChain,
Hugging Face



```
### (1) Packages
import os
from langchain.document_loaders import ...

### (2) File Handling
# Get the current working directory
file_path = os.path.abspath(__file__)
current_dir = os.path.dirname(file_path)

# Go up one directory level
parent_dir = os.path.dirname(current_dir)

file_path = os.path.join(parent_dir, "data")
file_path
```

Generative AI with Python

The Developer's Guide to Pretrained LLMs, Vector Databases,
Retrieval-Augmented Generation, and Agentic Systems

Bert Gollnick



Rheinwerk
Computing

Bert Gollnick

Generative AI with Python

**The Developer's Guide to Pretrained LLMs, Vector Databases,
Retrieval-Augmented Generation, and Agentic Systems**



OceanofPDF.com

Imprint

This e-book is a publication many contributed to, specifically:

Editor Meagan White

Acquisitions Editor Hareem Shafi

Copyeditor Yvette Chin

Cover Design Graham Geary

Photo Credit Midjourney.com; Shutterstock: 1861053121/©
Origami Sayko

Production E-Book Kelly O'Callaghan

Typesetting E-Book Satz-Pro (Germany)

We hope that you liked this e-book. Please share your feedback with us and read the [Service Pages](#) to find out how to contact us.

**The Library of Congress Cataloging-in-Publication Control
Number for the printed edition is as follows: 2025008568**

© 2025 by:

Rheinwerk Publishing, Inc.
2 Heritage Drive, Suite 305
Quincy, MA 02171
USA

info@rheinwerk-publishing.com
+1.781.228.5070

Represented in the E.U. by:

Rheinwerk Verlag GmbH
Rheinwerkallee 4
53227 Bonn
Germany

service@rheinwerk-verlag.de
+49 (0) 228 42150-0

ISBN 978-1-4932-2690-0 (print)

ISBN 978-1-4932-2691-7 (e-book)

ISBN 978-1-4932-2692-4 (print and e-book)

1st edition 2025

OceanofPDF.com

Notes on Usage

This e-book is **protected by copyright**. By purchasing this e-book, you have agreed to accept and adhere to the copyrights. You are entitled to use this e-book for personal purposes. You may print and copy it, too, but also only for personal use. Sharing an electronic or printed copy with others, however, is not permitted, neither as a whole nor in parts. Of course, making them available on the internet or in a company network is illegal as well.

For detailed and legally binding usage conditions, please refer to the section [Legal Notes](#).

This e-book copy contains a **digital watermark**, a signature that indicates which person may use this copy:

Notes on the Screen Presentation

You are reading this e-book in a file format (EPUB or Mobi) that makes the book content adaptable to the display options of your reading device and to your personal needs. That's a great thing; but unfortunately not every device displays the content in the same way and the rendering of features such as pictures and tables or hyphenation can lead to difficulties. This e-book was optimized for the presentation on as many common reading devices as possible.

If you want to zoom in on a figure (especially in iBooks on the iPad), tap the respective figure once. By tapping once again, you return to the previous screen. You can find more recommendations on the customization of the screen layout on the [Service Pages](#).

OceanofPDF.com

Table of Contents

Notes on Usage
Table of Contents

Preface

1 Introduction to Generative AI

- 1.1 Introduction to Artificial Intelligence
- 1.2 Pillars of Generative AI Advancement
 - 1.2.1 Computational Power
 - 1.2.2 Model and Data Size
 - 1.2.3 Investments
 - 1.2.4 Algorithmic Improvements
- 1.3 Deep Learning
- 1.4 Narrow AI and General AI
- 1.5 Natural Language Processing Models
 - 1.5.1 NLP Tasks
 - 1.5.2 Architectures
- 1.6 Large Language Models
 - 1.6.1 Training

- 1.6.2 Use Cases
- 1.6.3 Limitations
- 1.7 Large Multimodal Models
- 1.8 Generative AI Applications
 - 1.8.1 Consumer
 - 1.8.2 Business
 - 1.8.3 Prosumer
- 1.9 Summary

2 Pretrained Models

- 2.1 Hugging Face
- 2.2 Coding: Text Summarization
- 2.3 Exercise: Translation
 - 2.3.1 Task
 - 2.3.2 Solution
- 2.4 Coding: Zero-Shot Classification
- 2.5 Coding: Fill-Mask
- 2.6 Coding: Question Answering
- 2.7 Coding: Named Entity Recognition
- 2.8 Coding: Text-to-Image
- 2.9 Exercise: Text-to-Audio
 - 2.9.1 Task
 - 2.9.2 Solution
- 2.10 Capstone Project: Customer Feedback

Analysis

2.11 Summary

3 Large Language Models

3.1 Brief History of Language Models

3.2 Simple Use of LLMs via Python

3.2.1 Coding: Using OpenAI

3.2.2 Coding: Using Groq

3.2.3 Coding: Large Multimodal Models

3.2.4 Coding: Running Local LLMs

3.3 Model Parameters

3.3.1 Model Temperature

3.3.2 Top-p and Top-k

3.4 Model Selection

3.4.1 Performance

3.4.2 Knowledge Cutoff Date

3.4.3 On-Premise versus Cloud-Based Hosting

3.4.4 Open-Source, Open-Weight, and Proprietary Models

3.4.5 Price

3.4.6 Context Window

3.4.7 Latency

3.5 Messages

3.5.1 User

3.5.2 System

3.5.3 Assistant

3.6 Prompt Templates

3.6.1 Coding: ChatPromptTemplates

3.6.2 Coding: Improve Prompts with LangChain Hub

3.7 Chains

3.7.1 Coding: Simple Sequential Chain

3.7.2 Coding: Parallel Chain

3.7.3 Coding: Router Chain

3.7.4 Coding: Chain with Memory

3.8 Safety and Security

3.8.1 Security

3.8.2 Safety

3.8.3 Coding: Implementing LLM Safety and Security

3.9 Model Improvements

3.10 New Trends

3.10.1 Reasoning Models

3.10.2 Small Language Models

3.10.3 Test-Time Computation

3.11 Summary

4 Prompt Engineering

4.1 Prompt Basics

4.1.1 Prompt Process

4.1.2 Prompt Components

4.1.3 Basic Principles

- 4.2 Coding: Few-Shot Prompting
- 4.3 Coding: Chain-of-Thought
- 4.4 Coding: Self-Consistency Chain-of-Thought
- 4.5 Coding: Prompt Chaining
- 4.6 Coding: Self-Feedback
- 4.7 Summary

5 Vector Databases

- 5.1 Introduction
- 5.2 Data Ingestion Process
- 5.3 Loading Documents
 - 5.3.1 High-Level Overview
 - 5.3.2 Coding: Load a Single Text File
 - 5.3.3 Coding: Load Multiple Text Files
 - 5.3.4 Exercise: Load Multiple Wikipedia Articles
 - 5.3.5 Exercise: Loading Project Gutenberg Book
- 5.4 Splitting Documents
 - 5.4.1 Coding: Fixed-Size Chunking
 - 5.4.2 Coding: Structure-Based Chunking
 - 5.4.3 Coding: Semantic Chunking
 - 5.4.4 Coding: Custom Chunking
- 5.5 Embeddings
 - 5.5.1 Overview
 - 5.5.2 Coding: Word Embeddings

5.5.3 Coding: Sentence Embeddings

5.5.4 Coding: Create Embeddings with LangChain

5.6 Storing Data

5.6.1 Selection of a Vector Database

5.6.2 Coding: File-Based Storage with a Chroma Database

5.6.3 Coding: Web-Based Storage with Pinecone

5.7 Retrieving Data

5.7.1 Similarity Calculation

5.7.2 Coding: Retrieve Data from Chroma Database

5.7.3 Coding: Retrieve Data from Pinecone

5.8 Capstone Project

5.8.1 Features

5.8.2 Dataset

5.8.3 Preparing the Vector Database

5.8.4 Exercise: Get All Genres from the Vector Database

5.8.5 App Development

5.9 Summary

6 Retrieval-Augmented Generation

6.1 Introduction

6.2 Coding: Simple Retrieval-Augmented Generation

- 6.2.1 Knowledge Source Setup
- 6.2.2 Retrieval
- 6.2.3 Augmentation
- 6.2.4 Generation
- 6.2.5 RAG Function Creation
- 6.3 Advanced Techniques
 - 6.3.1 Advanced Preretrieval Techniques
 - 6.3.2 Advanced Retrieval Techniques
 - 6.3.3 Advanced Postretrieval Techniques
- 6.4 Coding: Prompt Caching
- 6.5 Evaluation
 - 6.5.1 Challenges in RAG Evaluation
 - 6.5.2 Metrics
 - 6.5.3 Coding: Metrics
- 6.6 Summary

7 Agentic Systems

- 7.1 Introduction to AI Agents
- 7.2 Available Frameworks
- 7.3 Simple Agent
 - 7.3.1 Agentic RAG
 - 7.3.2 ReAct
- 7.4 Agentic Framework: LangGraph
 - 7.4.1 Simple Graph: Assistant

- 7.4.2 Router Graph
- 7.4.3 Graph with Tools
- 7.5 Agentic Framework: AG2
 - 7.5.1 Two Agent Conversations
 - 7.5.2 Human in the Loop
 - 7.5.3 Agents Using Tools
- 7.6 Agentic Framework: CrewAI
 - 7.6.1 Introduction
 - 7.6.2 First Crew: News Analysis Crew
 - 7.6.3 Exercise: AI Security Crew
- 7.7 Agentic Framework: OpenAI Agents
 - 7.7.1 Getting Started with a Single Agent
 - 7.7.2 Working with Multiple Agents
 - 7.7.3 Agent with Search and Retrieval Functionality
- 7.8 Agentic Framework: Pydantic AI
- 7.9 Monitoring Agentic Systems
 - 7.9.1 AgentOps
 - 7.9.2 Logfire
- 7.10 Summary

8 Deployment

- 8.1 Deployment Architecture
- 8.2 Deployment Strategy
 - 8.2.1 REST API Development

- 8.2.2 Deployment Priorities
- 8.2.3 Coding: Local Deployment
- 8.3 Self-Contained App Development
- 8.4 Deployment to Heroku
 - 8.4.1 Create a New App
 - 8.4.2 Download and Configure CLI
 - 8.4.3 Create app.py File
 - 8.4.4 Procfile Setup
 - 8.4.5 Environment Variables
 - 8.4.6 Python Environment
 - 8.4.7 Check the Result Locally
 - 8.4.8 Deployment to Heroku
 - 8.4.9 Stop Your App
- 8.5 Deployment to Streamlit
 - 8.5.1 GitHub Repository
 - 8.5.2 Creating a New App
- 8.6 Deployment with Render
- 8.7 Summary

9 Outlook

- 9.1 Advances in Model Architecture
- 9.2 Limitations and Issues of LLMs
 - 9.2.1 Hallucinations
 - 9.2.2 Biases
 - 9.2.3 Misinformation

- 9.2.4 Intellectual Property
- 9.2.5 Interpretability and Transparency
- 9.2.6 Jailbreaking LLMs
- 9.3 Regulatory Developments
- 9.4 Artificial General Intelligence and Artificial Superintelligence
- 9.5 AI Systems in the Near Term
- 9.6 Useful Resources
- 9.7 Summary

The Author

Index

Service Pages

Legal Notes

OceanofPDF.com

Preface

Welcome to your journey into the world of generative artificial intelligence (generative AI). In this preface, we set the stage for what you, the reader, can expect from this comprehensive guide on generative AI with Python. Follow this roadmap to maximize the value you derive from our journey into the world of generative AI.

Objective of This Book

Welcome to our comprehensive exploration of generative AI with a spotlight on some of its most transformative and advanced technologies. The primary objective of this book is to guide you through the frontiers of this dynamic field, with special attention to large language models (LLMs), the nuanced art of prompt engineering, the utility of vector databases, innovative processes of retrieval-augmented generation (RAG), and the up-and-coming topic of agentic systems.

LLMs have revolutionized the way we interact with text-based AIs. LLMs unlock unparalleled capabilities in language understanding and generation. We'll dive into the architectures of these models, dissecting how they learn from massive corpora of data to produce human-like text. Then, we'll examine prompt engineering, an essential skill in the age of LLMs. You'll learn how to craft prompts that efficiently navigate the models' knowledge and capabilities.

Vector databases present the next leap in organizing and retrieving data. Understanding this technology is key to working with high-dimensional data and building systems that are capable of rapid and relevant information access. We'll explore the concept behind vector databases, their design and function, and how they pave the way for sophisticated AI applications.

The concept of RAG bridges the gap between retrieval of relevant information and on-the-fly text generation. RAG systems mark a significant milestone in AI development, assisting models in producing more accurate, more informed content. The chapter dedicated to RAG will offer an in-depth look at its mechanisms and at integrating this technique into your generative AI applications.

Next, we'll venture into the domain of agentic systems. These AI systems can act autonomously, make decisions, and undertake tasks that traditionally require human intelligence. We'll navigate the ethical, technological, and practical constructs of such systems. This chapter will prepare you to design AIs with autonomy and agency in a responsible and innovative manner.

Up to that point, you'll have learned to develop applications locally, but eventually you want to share your work with the wider world. We'll dive into this topic in the chapter on deployment of AI systems.

This book is crafted not only to provide you with a deep theoretical grounding in these areas but also to give you hands-on practical expertise. Through a series of meticulously selected examples, case studies, and projects, we'll guide you through applying these concepts to varied scenarios, whether for professional advancement, academic research, or personal intellectual curiosity.

Our commitment extends beyond mere knowledge transfer. We hope to equip you with the awareness required to become a proficient

creator and innovator in the field of generative AI. We aim to thoroughly prepare you to confront and tackle the challenges and opportunities that this technology presents in the modern world. By the end of this book, you'll have, not just an understanding, but mastery over these sophisticated instruments of artificial intelligence.

Target Audience

This book is designed for a wide array of readers, ranging from software developers and data scientists to students and researchers interested in generative AI. A certain amount of coding background is expected. If you have a basic understanding of Python programming and a keen interest in AI, especially in the field of generative AI models, this book will serve you well.

The book's content is tailored to engage both novices taking their first steps into the world of generative AI and seasoned professionals seeking to refine their skills and knowledge.

Our practical examples and thorough explanations will help you grasp the concepts and techniques essential to developing and applying generative AI systems. Whether you're looking to innovate in your field, embarking on an academic endeavor, or simply pursuing a fascination with AI, this book hopes to serve as an invaluable resource on your journey.

Prerequisites: What You Should Already Know

Before diving into the riveting world of generative AI with Python, let's check a few prerequisites to ensure a smooth journey. First, in

terms of mandatory requirements, a firm grasp of Python programming is essential; you should be comfortable with the following tasks:

- Writing functions
- Creating and manipulating different data structures like lists and dictionaries
- Writing loops, mainly for-loops
- Using libraries like NumPy and pandas

Ideally, but not mandatory, you already have a foundational understanding of machine learning (ML) concepts, such as training models and working with datasets.

Familiarity with basic statistics and linear algebra will also be beneficial, as they underpin many AI algorithms. Although the book will cover the necessary theory behind generative AI, previous exposure to neural networks and deep learning frameworks like TensorFlow or PyTorch will help you navigate the more advanced topics more easily.

If these prerequisites sound like languages you speak, you're well equipped to embark on our thrilling voyage through the seas of generative AI landscapes.

Structure of This Book

This book is designed as a practical guide for Python programmers looking to develop generative AI applications. The structure of the book follows a step-by-step approach, starting with an introduction to the fundamental concepts and progressing to more advanced topics like vector databases and agentic systems.

You're encouraged to go through the book sequentially, but some chapters build upon the knowledge from previous chapters—more on that in the next section. Practical Python code examples are available for download at <https://www.rheinwerk-computing.com/6057/> to illustrate and reinforce the application of the concepts discussed. Let's look at each of the chapters in turn:

- **Chapter 1**

This chapter introduces you to the basics of generative AI, a subfield of artificial intelligence focused on creating new content. You'll learn about natural language processing (NLP) models, in particular LLMs. Recent developments like large multimodal models (LMMs), or reasoning models are touched on as well.

- **Chapter 2**

In this chapter, we discuss pretrained models, particularly from the field of NLP. Pretrained models are language models that have been trained on large amounts of data and that can be reused for various NLP or computer vision tasks. In this chapter, we'll discuss the popular platform Hugging Face, on which you can find over a million open-source models to handle many tasks.

The models we present will mostly solve particular tasks like text summarization, translation, text classification, or text generation. The advantage of this approach is that these models are small and can be operated locally. In this chapter, you'll learn how to select a model and operate it on your system.

- **Chapter 3**

In this chapter, you'll learn how to interact with LLMs via Python code. Different providers like OpenAI and Groq are covered, and we'll show you how to use their models for your projects. Different ways to use these models exist. In this book, you'll learn how to work with LangChain, a Python-based framework for interacting

with LLMs, to implement efficient approaches, including prompt templates and chains to guide the LLMs to produce the desired outputs.

- **Chapter 4**

Prompt engineering is the process of designing efficient inputs to LLMs to generate outputs. You'll learn how to create effective prompts to optimize the performance of language models and achieve more precise results. Basic techniques such as few-shot prompting and chain-of-thought (CoT) are covered. But we'll also go beyond these topics and cover more advanced techniques like self-feedback or reflection.

- **Chapter 5**

This chapter delves into vector databases, essential for efficiently storing and querying large text collections. We describe the entire process, from the indexing pipeline (in which you learn how to preprocess the data), to storing the data in a vector database and querying. We'll mostly focus on adding data to a vector database. This process consists of several steps that we'll cover one by one.

We start with document loading. For document loading, various LangChain document loaders are at your disposal to load nearly any kind of data source in a structured way. We'll visit some data sources as well.

In the data splitting step, you'll learn how to structure documents into meaningful bite-sized chunks. Different methods can be used, depending on the kind of data and use case. We'll present the most common data splitting approaches.

Often, text information needs to be converted into a format that a computer algorithm can "understand." Since computers work based on numbers rather than human language, the information is

converted into a numerical representation called an embedding. You'll learn what embeddings are, how they can be created, and different types that are available.

Several providers of vector databases can be used, but we'll work with a Chroma database as an example of an open-source vector database that can run locally on your system. But you'll also learn how to work with a web-based provider called Pinecone.

Once your data is stored into a vector database, you'll want to use the database to retrieve information. Typically, as a user, you have a query and want to retrieve the most relevant information, like a text document or an image, from the data store. We'll demonstrate how not only texts can be retrieved, but also images and audio files. You'll learn how to retrieve data based on additional conditions like special properties, which are typically defined in metadata.

A key aspect of finding relevant documents is how relevance is measured and how these documents are found. We cover these topics in our discussion on similarity search. Some concepts explored include cosine similarity and maximum margin relevance (MMR).

In a capstone project, we bring all the pieces together into a single project, starting with a specific data source, preprocessing the data, and finally storing the data into a vector database.

- **Chapter 6**

RAG is a new paradigm where external knowledge sources, typically from a vector database, are integrated into the generation process to enhance the quality and relevance of the generated text. This generative AI feature is powerful in a company environment, as you learn how powerful this technology can be.

You'll learn how to retrieve relevant information from a large data store, how to integrate that information into a generation process, and how to evaluate the quality and diversity of the generated text.

- **Chapter 7**

This chapter introduces agent-based, or agentic, systems. These AI systems build on LLMs and equip them with additional features like tool use or memory, so that these agentic systems are capable of autonomously executing tasks and directing their own actions toward specific goals.

Agentic systems are one of the most interesting areas in AI development. Many different frameworks are provided, and we'll work with the most relevant ones, such as LangGraph, OpenAI Agents, Microsoft AG2, Magentic One, TinyTroupe, and Pydantic AI.

- **Chapter 8**

Up to this point, we've developed systems locally on our computers, but in this chapter, you'll learn how to practically deploy generative AI models and systems. The deployment process typically involves the use of representational state transfer (REST) application programming interfaces (APIs) to make the systems accessible to users and other systems that want to use these capabilities.

We'll explore different deployment options and design decisions like self-contained applications versus frontend-backend architectures. You'll learn how to create a backend REST API service and also how to deploy an application to a variety of providers, such as Heroku, Streamlit, or render.

- **Chapter 9**

This book concludes with a look at future developments in

generative AI. We'll discuss associated challenges, such as impacts on employment and other societal challenges, but also opportunities.

How to Use This Book Effectively

This book is designed to be both a comprehensive guide and a practical toolkit for enthusiasts, developers, and professionals looking to harness the power of artificial intelligence to generate new content, from text to images and beyond. To make the most out of this book, we want to share some key strategies to help you use it effectively.

Many chapters can be considered standalone chapters, and you can just jump into them, but some build upon the knowledge from other chapters. [Figure 1](#) shows an overview of the chapters that can be directly consumed and those that require some knowledge learned from other chapters.




Figure 1 Book Chapter Dependencies

Let's start with the standalone chapters. Chapters 1 through [Chapter 3](#), and [Chapter 9](#), can be consumed independently for the most part.

However, [Chapter 4](#) builds upon knowledge of LLMs, so you ideally should read [Chapter 3](#) first.

If you want to dive into vector databases in [Chapter 5](#), we recommend reading both [Chapter 3](#) and [Chapter 4](#) first. Once you've covered vector databases, you can extend your knowledge with RAG in [Chapter 6](#).

For [Chapter 8](#), you should have at least covered the basics of LLMs in [Chapter 3](#) as well as prompt engineering in [Chapter 4](#).

Downloadable Code and Additional Materials

Alongside this book, you can download the code examples, so that you can follow and reproduce the scripts on your own. You can get the material via GitHub. The repository, which is a cloud-based storage of the content, can be visited via this link:

https://github.com/DataScienceHamburg/GenerativeAiApplicationsWithPython_Material

You can also download the code directly from the publisher's website:

<https://www.rheinwerk-computing.com/6057/>

In the next section, we'll describe how you can download the material so that you can follow along the Python scripts on your own system.

System Setup

Figure 2 shows the steps for setting up your system so you can work with the coding material and run it on your local system.




Figure 2 System Setup


By following these steps, the scripts can be run on your local computer. This preparation might seem like a lot but consider that this is a one-time effort and shouldn't take more than 30 minutes. The steps are as follows:

- ❶ Python installation
- ❷ Integrated development environment (IDE) installation
- ❸ Git installation
- ❹ Downloading the course materials
- ❺ Setting up your local environment

Python Installation

The Python installation step is straightforward and can be accomplished through multiple methods. We recommend downloading Python from the official website at <https://www.python.org/downloads/>. If you scroll down the download page, you'll find multiple different versions.

[Figure 3](#) shows a list of these versions. The scripts and code in this book were developed on Python version 3.12.7. To avoid any issues, you should stick to the same version. Click the **Download** button and follow the installation instructions. Nothing specific needs to be considered, so you can go with the default settings.



Looking for a specific release?
Python releases by version number:

Release version	Release date	Click for more	
Python 3.13.0	Oct. 7, 2024	Download	Release Notes
Python 3.12.7	Oct. 1, 2024	Download	Release Notes
Python 3.11.10	Sept. 7, 2024	Download	Release Notes
Python 3.10.15	Sept. 7, 2024	Download	Release Notes
Python 3.12.6	Sept. 6, 2024	Download	Release Notes
Python 3.9.20	Sept. 6, 2024	Download	Release Notes
Python 3.8.20	Sept. 6, 2024	Download	Release Notes

[View older releases](#)

Figure 3 Python Download

IDE Installation

An IDE is essential for efficiently writing and managing code in Python projects. Some popular choices include the following:

- **Visual Studio Code (VS Code)**

This lightweight and highly customizable IDE supports many different programming languages, which is one reason why it is a favorite. We use VS Code for our Python, R, and Flutter projects and therefore don't need to adjust to different IDEs when switching projects.

VS Code is open source and free at <https://code.visualstudio.com/>. Its large community provides thousands of extensions. We've used it for many years and can

hardly imagine switching. But eventually we did. Well, kind of. We recently moved to Cursor.

- **Cursor**

Cursor is a fork of VS Code, which means that it is very close relative. But the main difference with Cursor is that it already has AI power included. With this AI-enhanced code editor, you become much faster and much efficient in coding. Cursor leverages generative AI capabilities, and you can get coding assistance via OpenAI's Codex, GPT, or Anthropic's Claude.

Cursor features code completion, code explanation, bug fixing, refactoring, and interactive chatting to speed you up. It has a free version and a paid version. You can learn more at <https://www.cursor.com/>.

- **PyCharm**

This IDE is focused on Python and provides advanced features like code completion and debugging. The community edition of PyCharm is free and includes basic features while the professional version provides deeper integrations. On the developer website, you can find both versions at <https://www.jetbrains.com/pycharm/download/>.

If you already have a preference and experience with any of these options, we won't stop you. But if you're a beginner and looking for good advice, we recommend going with Cursor.

Git Installation

Git is an essential tool for version control, allowing developers to track changes in their code, collaborate with others, and maintain a history of project development.

Git helps you manage your codebase efficiently by enabling branching, merging, and reverting changes. Git integrates seamlessly with platforms like GitHub, GitLab, and Bitbucket, making remote collaboration straightforward. You'll need Git in the next step to get the book's material. Download Git from <https://git-scm.com/downloads>. [Figure 4](#) shows the download options.



Figure 4 Git Download

Select the download corresponding to your operating system. Many choices are available, but you can go with the default settings during installation.

Getting the Source Material

The scripts and code files for this book are hosted on GitHub. You can find the material at the following link:

https://github.com/DataScienceHamburg/GenerativeAiApplicationsWithPython_Material

When you navigate to this page, you should see all the material, as shown in [Figure 5](#).




Figure 5 GitHub Repository for This Book

You can get the code in two ways by clicking the green **Code** button.

- **Download as ZIP**
- **Clone using web URL**

If you have no experience with Git, go with the ZIP option. Download the ZIP file and extract it to a folder of your choice on your computer.

But if you had points of contact with Git, you could clone the repository by running the following from the command line:

```
git clone
https://github.com/DataScienceHamburg/GenerativeAiApplicationsWithPython_Material.git
```

This command will clone the repository to your computer.

Setting up Your Local Environment

A Python environment is an isolated workspace where you can install and manage dependencies for a specific project. This environment acts as a sandbox, ensuring that the libraries and tools used in one project don't interfere with others on your system. The diagram shown in [Figure 6](#) illustrates the concept of environments and why they are needed.




Figure 6 Python Environments

While Project A requires a specific package xyz version 1.0a0 and Python version 3.10, you might also need to work on a different project with completely different requirements.

Let's say, for instance, that Project B requires the same package xyz, but in the newer version 2.0a0, and a different Python version, maybe 3.12.7 instead. If you've only set up a single environment for all your projects, you would run into issues because of the incompatibility between the two projects.

The gold standard is to set up an environment for each project you work on. This approach consumes more space on your hard disk, but you'll be on the safe side and avoid negative side effects between your projects.

OK, we hope we've convinced you of the benefits of working with environments. Great—now we need to figure out which environment

to choose. Different tools for managing your Python environments include the following:

- **virtualenv**

virtualenv is a lightweight tool for managing Python environments. It allows you to install project-specific packages without affecting the global Python installation.

- **Poetry**

Poetry is a more modern dependency and environment manager. It simplifies package management by combining dependency specification, installation, and environment creation in a single tool. It is perfect for projects that need precise version tracking and easy deployment.

- **uv**

A less popular but emerging tool, uv focuses on creating ultralight virtual environments. Due to its implementation in Rust, it is blazing fast compared to other environmental tools.

We recommend uv. You can find many different options to install uv on the developer page <https://docs.astral.sh/uv/getting-started/installation/#standalone-installer>. You can install it via pip on your terminal, as follows:

```
pip install uv
```

In the code folder, you'll find a file called *pyproject.toml*. This file includes basic metadata information on the package, as well as package dependencies—the packages and their versions as we have them installed on our systems when we created the material. Next to the *pyproject.toml* file, you'll find *uv.lock*. This file keeps track of the exact versions of all dependencies and their transitive dependencies, that is, both dependencies and dependencies of dependencies.

You can create the exact same environment we've used by running, in the terminal, the following command:

```
uv sync
```

This command will download and install all dependencies into a subfolder named `.venv`, all at once. If for whatever reason it does not work on your computer, don't worry. Open the `pyproject.toml`, look for the section on dependencies, and install the packages manually. You can do install these packages by running either of the following commands:

- `pip install packagename`
- `uv add packagename`

You should also read the description on the GitHub repository, in which specific issues for different platforms will be addressed.

Once you open a Python script, VS Code (or Cursor) will immediately recognize that an environment is available. Check this capability by inspecting the horizontal bar at the bottom. [Figure 7](#) shows the Python version (3.12.7) used and the virtual environment (`.venv`).



Figure 7 Environment and Python Detection in VS Code

We've completed the last step of the system setup, and now, you're ready to start coding.

Acknowledgments

I hope what you've learned so far about this book has made it enticing.

In my journey writing this book, I have traversed the exciting and often challenging landscape of generative AI. No such voyage is ever a solitary endeavor. Therefore, I want to take a moment to acknowledge those who have been my refuge and steadfast supporters.

At the top of my gratitude is my family. Their love and understanding have been the foundation of my resolve. My wife Lea, with her infinite patience and unwavering support, has been the lighthouse guiding me through challenging times of research and writing. She accepted, without a murmur of disquiet (well mostly), the long and irregular hours I spent secluded in my working room, tapping away at the keyboard, lost in a world of codes and concepts. Her strength and encouragement were the invisible hands that held me steady when the specter of writer's block loomed menacingly.

Equally deserving of my deepest appreciation is my daughter Elisa, whose youthful cheer and innocence often penetrated the solitude of my scholarly exertions, reminding me of the joy and wonder that life offers. Her understanding, far beyond her years, was a treasured gift as she accepted, with remarkable maturity, the times when I was preoccupied with crafting this manuscript.

To these two remarkable individuals I dedicate not just this chapter, but all the fruits of the hours spent creating this book on generative AI with Python. May every page reflect the love and sacrifice that they have so selflessly given.

To our extended family and friends who have offered their encouragement and support, to colleagues who have provided technical insights, and to the wider AI and Python communities that have inspired and fueled our passion for this subject, you have my heartfelt thanks. This book is not just a product of my endeavor; it is

a testament to the collective spirit of all those who have, in ways big and small, contributed to its completion.

Thank you.

Conventions Used in This Book

Throughout this book, certain typographical and formatting conventions have been employed to enhance your learning experience and to make the journey into generative AI as clear and effective as possible.

Code blocks and programming snippets are presented in a monospaced font, which delineates them from the surrounding text, ensuring that you can easily identify and follow along with our Python examples.

```
# sample code looks like this
my_welcome = "Welcome to the book"
```

The code provides outputs that are typically shown on the command line. Such outputs are presented similarly, and look like this:

```
Welcome to the book
```

All code blocks and files can also be found in the accompanying repository. You can easily navigate to the source code files, which are indicated as *source_folder/source_file_name.py*.

To encapsulate essential takeaways or to provide additional context, informational boxes are placed strategically throughout the chapters.

An Information Box

Here you find details on key concepts.

Hyperlinks are italicized to provide quick navigation to external resources for a deeper dive into topics or to access datasets pertinent to the exercises, for example, *www.rheinwerk-computing.com*.

We hope that these conventions render the book both user friendly and conducive to your learning journey into the world of generative AI.

OceanofPDF.com

1 Introduction to Generative AI

One day the AIs are going to look back on us the same way we look at fossil skeletons on the plains of Africa.

—Nathan Bateman in the film Ex Machina

Humanity's story is a journey of evolution—a story of curiosity, invention, and transformation. We've always sought to build tools to amplify our abilities. It started with fire, and then, we invented the wheel, the printing press and in the last century the internet—to name just a few key milestones. Today, we stand at the precipice of the newest kind of evolution, one driven not by natural selection, but by our own creations.

Generative artificial intelligence (generative AI) represents a huge leap forward in this journey. Earlier computer technologies were explicitly programmed to perform specific tasks. Generative AI can create, imagine, and synthesize, often appearing capable of its own spark of creativity.

This chapter explores how we arrived at this particular moment. You'll learn about the developments critical for this invention. The most relevant technologies in this field are large language models (LLMs), which we'll discuss in more detail. Then, we'll extend our discussion to related technologies like large multimodal models (LMMs) before we study applications using generative AI.

How did we end up at this point? We'll introduce you to artificial intelligence in [Section 1.1](#). Artificial intelligence, especially generative

AI, is made possible by developments in different areas. You'll find out about these pillars of generative AI advancement in [Section 1.2](#). Generative AI is based on deep learning; you'll encounter this type of AI and its core features in [Section 1.3](#).

Perhaps, you've heard of the impressive achievements of various AI systems, especially in the field of gaming. We'll revisit these achievements in [Section 1.4](#) and discuss why AI can have superhuman capabilities in some fields and why extending these capabilities to a wide range of domains can be difficult.

We cannot speak about generative AI without discussing natural language processing (NLP). The most usual input type of these models is text. Thus, NLP models play a crucial role, as you'll see in [Section 1.5](#). LLMs are the foundation of most generative AI applications. In [Section 1.6](#), we'll look at the most notable LLMs.

With the advancement of LLMs, some limitations were discovered. One limitation is that LLMs are restricted to processing text input and can only provide text output. This limitation changes with LMMs, the model class that you'll learn about in [Section 1.7](#).

In [Section 1.8](#), you'll learn which types of content can be created with generative AI. A complete industry has been developed around this area, as explored in [Section 1.8](#).

Let's start with the foundation of this book—an introduction to artificial intelligence in general and generative AI in particular.

1.1 Introduction to Artificial Intelligence

First, what is generative AI? Generative AI refers to a branch of artificial intelligence focused on creating models that can produce

new content such as text, images, music, code, and videos—basically any kind of data.

The most famous offspring of generative AI is ChatGPT. Based on OpenAI models from its family of GPTs (generative pretrained transformers), this chatbot became famous more quickly than any other service so far.

[Figure 1.1](#) shows the growth in users of selected internet services over time.




Figure 1.1 Number of Users of Selected Internet Services

While Netflix required 1,500 days to reach 1 million users, other services reached this milestone more quickly. Spotify needed 150 days, while Instagram did it in 75 days. The line for ChatGPT is close to vertical in this graph because OpenAI reached 1 million users in just 5 days. To say “ChatGPT has taken the world by storm” is a gross understatement. Released at the end of November 2022, the field has since evolved so quickly that we often forget everything that happened in just a few years’ time.

The rapid ascent of ChatGPT is exemplary of a larger phenomenon: the acceleration of generative AI development. What once took decades to iterate and improve upon is now happening in a few

months. The release cycle for AI models has become astonishingly compressed, with each new iteration not only outperforming its predecessor but also expanding its range of capabilities.

Behind this incredible pace of change is a confluence of technological and societal factors. The availability of vast computational resources, driven by cloud computing and specialized hardware like GPUs and TPUs, has enabled researchers to train larger and more complex models. Simultaneously, the abundance of publicly accessible data—like images, text, videos, and more—has provided the fuel required to teach these models how to create.

But technology alone doesn't explain the speed of adoption. Generative AI aligns seamlessly with the demands of the modern world, where efficiency, creativity, and scalability are paramount. Businesses eager to stay ahead of the competition have embraced these tools to automate content creation, optimize workflows, and deliver personalized customer experiences. Consumers, fascinated by AI's creative potential, have adopted platforms like ChatGPT, not just as tools but as companions for creativity, problem-solving, and learning.

This rapid development has created a feedback loop: As more industries adopt generative AI, they generate insights, use cases, and datasets that further refine these technologies. Open-source communities and research collectives accelerate the pace by sharing advancements openly, democratizing access to cutting-edge tools. Meanwhile, investments in AI startups have surged, creating a highly competitive ecosystem that pushes innovation to the forefront.

However, the speed at which generative AI evolves also introduces challenges. Ethical concerns, such as misinformation, intellectual property rights, and job displacement, have become urgent topics of

discussion. Regulatory frameworks struggle to keep up with the pace of change, leaving gaps that could have far-reaching consequences.

Generative AI's rapid rise is a testament to its transformative potential. It has unlocked possibilities that were unimaginable a decade ago and continues to redefine the boundaries of what machines and humans can achieve together. As we look ahead, the question is no longer whether generative AI will shape our future, but how we'll shape the future of generative AI.

Generative AI's rapid development is not just a technological masterpiece. This transformative force operates across creative, technical, and commercial landscapes, and its influence transcends the borders of research labs, reaching into studios, factories, and boardrooms. Generative AI reshapes the way we create, build, and conduct business.

Backtracking a bit, let's now look at the history of AI in general. [Figure 1.2](#) shows the evolution of artificial intelligence.




Figure 1.2 AI Evolution

Artificial intelligence has undergone a remarkable journey. Numerous breakthroughs occurred, but so did setbacks and sustained periods of rapid innovation. The starting point is marked by the development of the Neuron model, developed in 1943 by Warren McCulloch and Walter Pitts, that laid the foundation for understanding how biological processes can inspire machine intelligence.

Alan Turing developed the famous eponymous test as a benchmark for assessing a machine's ability to exhibit human-like intelligence. The Turing Test determines if a machine can engage in a conversation indistinguishably from a human. If the machine can successfully deceive a human evaluator into thinking they are interacting with another human, the machine passes the test. The Turing Test has become a philosophical and technical cornerstone in AI discussions.

In 1957, Frank Rosenblatt developed the concept of the *perceptron*, which marked the first implementation of a neural network. Concepts like weights and learning set the stage for later advancements in machine learning (ML).

In the 1970s, expert systems were designed to mimic human decision-making in specific domains. These expert systems were a major focus of AI research at that time. However, limitations in computing power and data led to disillusionment and triggered the first "AI winter," a period of reduced funding and reduced interest in AI. During the 1980s, AI resurged slightly, but not much progress was made, and a second "AI winter" occurred—another cycle of unmet expectations and general skepticism about the potential of AI.

In the early 1990s, neural networks gained some traction again, and some theoretical groundwork was laid. But the most important breakthroughs happened in the 2010s and later. Deep learning, a subset of AI, was nothing less than a revolution. It was fueled by large amounts of data and improved computational power. New architectures were developed. For example, convolutional neural networks (CNNs) and recurrent neural networks (RNNs) enabled breakthroughs in computer vision and NLP.

In the late 2010s, generative AI took off. Early models like OpenAI's GPT-2 (released in 2019) demonstrated the ability to generate

coherent and contextually relevant text. Subsequent models quickly evolved.

While GPT-2 is considered to have abilities comparable to a preschooler, GPT-3 (released in 2020) achieved elementary school pupil sophistication. GPT-3 displayed impressive generalization across various tasks. Released in 2023, GPT-4 is considered to be comparable to a smart high-schooler. GPT-4 is marked by improvements in reasoning and contextual understanding. In December 2024, the model o3 was announced, and it is expected to be released early 2025, with capabilities comparable to university students.

With more and more capabilities, AI systems are becoming increasingly more general. Tech companies are investing heavily in this technology to achieve artificial general intelligence (AGI), which contrasts with the notion of narrow AI.

Narrow AI (or weak AI) is designed for specific tasks and operates under a limited set of constraints. But don't be fooled by the word "weak"—these models excel in well-defined areas like language translation, facial recognition, and recommendation systems, performing these tasks with high precision. "Narrow" describes the scope of their capabilities better. However, this kind of AI lacks the ability to generalize its knowledge to other domains.

On the other hand, AGI (or strong AI) aspires to understand, learn, and apply knowledge across a wide range of tasks, mirroring human cognitive abilities. AGI aims to exhibit versatile and adaptive intelligence, capable of reasoning, problem-solving, and understanding contexts across various fields without any predefined limitations.

Generative AI, within the broader AI landscape, serves as a pivotal technology driving advancements in creating realistic and context-aware content autonomously. Its applications span various fields, enhancing capabilities in media, communication, and content creation while pushing the boundaries of what AI can achieve.

With each technology, an initial innovation trigger (i.e., a technological breakthrough, a product launch, or another event) generates significant interest and media coverage.

Then, the public gets overenthusiastic. Paired with inflated expectations, disappointment increases when the technology does not meet these exaggerated expectations. In the subsequent phase, public interest decreases due to experiments that can't deliver according to the hype.

But investment does not stop. Neither does the development in this field. Eventually, successful projects emerge. These projects show the technology's real-world value. Finally, the technology becomes widely understood, adopted, and accepted by mainstream.

At the time of writing, generative AI is possibly in the "trough of disillusionment," but we can expect that generative AI will reach the "plateau of productivity" in a few years. We won't notice a single event or breakthrough; rather, it will be a continuous process of many small steps, and some huge steps.

[Figure 1.3](#) shows some capabilities of AI systems on various tasks, relative to human performance, over time.




Figure 1.3 AI System Capabilities (Source: <https://ourworldindata.org/artificial-intelligence>; used under Creative Commons Attribution 4.0 International License)

When the capability of an AI system in a specific task is compared against human performance (e.g., “Language understanding”), the AI system starts out completely incapable (represented by a score of –100) but improves over time. In the next chapter, we’ll discuss several reasons for AI’s continuous improvement over time. At some point, the AI system reaches a score of 0—this value represents human-level capabilities on that task. Every score above 0 represents superhuman capabilities. Several extremely impressive behaviors of AI systems to notice include the following:

- AI systems reach superhuman capabilities one task at a time.
- The time it takes to achieve some outstanding capability gets shorter and shorter.
- Over time, there will be fewer tasks in which humans possess capabilities better than AI systems.

But what cannot be illustrated in this way is that AI systems still struggle to achieve AGI. In other words, current AI systems dominate one task over others. But no AI system possesses all these capabilities together. At the time of writing, OpenAI o3 may have actually reached AGI, but the community lacks consensus on this topic.

However, what generative AI achieved already is quite impressive. How did we get to this point? Let's find out in the next section.

OceanofPDF.com

1.2 Pillars of Generative AI Advancement

Previously, you saw some of the amazing capabilities of AI systems. Several factors led to this result. In this section, we'll take a closer look at these factors: computational power, model and data size, investments, and algorithmic improvements.

1.2.1 Computational Power

Figure 1.4 shows the computations used to train an AI system, classified by domain.




Figure 1.4 Computation Used to Train AI Systems (Source: <https://ourworldindata.org/artificial-intelligence?insight=the-last-decades-saw-a-continuous-exponential-increase-in-the-computation-used-to-train-ai>; used under Creative Commons Attribution 4.0 International License)

This graph specifically shows the computation effort for training over time. The graph is logarithmic, and thus, the observed trend is exponential in nature. Most points are clustered to the right, indicating the increased interest and development in recent years; no

single domain drives this development. However, “Language” is the dominant domain.

This training is only possible due to increased hardware performance. Hardware progress follows an exponential course. So far, hardware shows no signs of slowing down, and so larger and larger models will be enabled.

Training computations go hand in hand with increased model sizes—our next input factor.

1.2.2 Model and Data Size

[Figure 1.5](#) shows the relationship between the number of model parameters and training dataset size.




Figure 1.5 Model Parameters and Training Dataset Size (Source: <https://ourworldindata.org/grapher/parameters-vs-training-dataset-size-in-notable-ai-systems-by-researcher-affiliation>; used under Creative Commons Attribution 4.0 International License)

Both axes are logarithmic, which indicates that exponentially larger models require exponentially larger datasets. Notice how most developments have been achieved most recently. The colors

indicate researcher affiliation, which is either academic or industry. Since the development of extremely large models costs in the order of millions of dollars, academia can no longer keep up with the industry's costs.

Speaking of financial conditions, let's look more closely at the investment landscape that drives the developments.

1.2.3 Investments

The most recent AI models already cost around US\$100 million—an impressive number made even more impressive when you consider the cost of training in 2017. When the transformer architecture was developed, training costs were around \$1,000 USD, five orders of magnitude less than currently.

More than just money, computational power also drives this development as does how the algorithms have improved dramatically over time.

1.2.4 Algorithmic Improvements

The most impactful deep learning architecture so far is called the *transformer architecture*. Developed in 2017, transformers are the foundation of all modern language models, but is also applied in computer vision and time series prediction. Different software domains rely on different data types, each of which are simpler or harder to process by AI systems. Thus, different domains are impacted differently by an AI system, which shows different speeds of progress.

The usual baseline for algorithmic progress is called Moore's Law. Formulated by Gordon Moore in 1965, Moore's Law suggests that the number of transistors on a microchip doubles approximately every 2 years, leading to exponential growth in computing power. This principle has profoundly influenced technological advancements, including the development of AI algorithms. In the graph shown in , this baseline progress is represented by a dashed line.

In domains like LLMs, the effective compute doubling time is between 5 and 12 months. In other domains like computer vision, the median compute doubling time is 9 months, which is significantly larger. One reason for LLMs' accelerated needs might be that computer vision models are already extremely mature and have already reached their plateau of productivity.

In essence, multiple driving forces have resulted in the outstanding capabilities of our current generative AI systems, for instance, the following capabilities:

- Algorithmic progress improves model architectures and makes them more accurate or enables faster learning.
- AI systems rely on vast amounts of data. Since the onset of the internet, the available data has been growing exponentially and creates a huge source of training data.
- Larger models require better hardware. Deep learning relies on GPU power, and the processing power of GPUs has increased exponentially.
- None of these developments would be possible without proper funding. Since generative AI delivered on expectations, huge amounts of money have flowed into the development of new models.

The combination of these driving forces has resulted in AI systems with capabilities that, in many cases, equal or outweigh human capabilities.

Now, let's get a better technical understanding of the core concepts. The next section will lay a foundation with some high-level understanding of deep learning, the technology behind AI systems.

OceanofPDF.com

1.3 Deep Learning

Deep learning is a subfield of artificial intelligence and ML. You must understand the crucial differences between ML and classic programming process flows, as shown in [Figure 1.6](#).




Figure 1.6 ML and Classic Programming, Compared

While in classic programming, rules and data are used as inputs to generate answers. ML turns this logic on its head. With ML, data and corresponding answers are provided to an ML algorithm, and that algorithm figures out the rules on its own.

This principle holds for statistical ML algorithms as well as any kind of neural network. The most powerful ML algorithms derive from neural networks, specifically deep learning. Let's examine what happens during the training of a deep learning model, as shown in [Figure 1.7](#).

Imagine you work in real estate, and you want to train an algorithm to predict the price of houses. In your dataset, you'll need many observations of price transactions. In this setup, the price at which a property sold is your *target variable*—the variable you want to predict. In more technical terms, these values are also called *true values* (typically abbreviated with the variable Y).

One of the first steps is to split your dataset into training and testing datasets. The training data is only used for training; the testing data is only used for model testing. More complex setups are possible, but for our introduction, we'll keep it simple.

The selling price depends on many factors like location, size, or age of the property, which are your independent features, in statistics often referred to as X . Underlying all ML algorithms is the idea that observable patterns exist between independent variables and dependent variables. The algorithm will figure out these hidden patterns.




Figure 1.7 Deep Learning: Model Training

To train a deep learning model, you can feed the independent features (or variables) to a deep learning model. The information is passed through the network. Finally, the network provides a prediction $\hat{Y}$. This prediction is compared to the true value Y in a *loss function*. This function evaluates the current prediction quality and returns a loss score that is used by the optimizer. An important building block of a neural network, the optimizer tries to update the weights of the neurons of the model in the most efficient way. At this point, one loop has finished.

The next batch of data is loaded from the training data and passed through the network layers. Ideally, the weights have been updated, so that the current prediction of the model is closer to the real value. The loss function calculates the loss score. Then, the optimizer adapts the weights of the neurons again.

This process is repeated until all data has been passed through the network. Typically, data is not just run through once, but several times. Each complete loop of the complete dataset being passed through the loop is called an *epoch*.

The training process finishes once a stop criterion is found, for instance:

- A predefined number of epochs is reached.
- Model loss has reached some convergence so that the training process does not improve the result anymore.

Let's assume that your model has been trained successfully. You now have a model with a specific setup (i.e., architecture, number of neurons, activation functions, etc.) and corresponding model weights for each neuron in the network. How can you use this trained model to generate predictions? This process of running a model for creating predictions is called *inference*. [Figure 1.8](#) shows the model inference process.

During model inference, you provide test data to the model. This test data is passed through the model with its frozen model weights. In this context, “frozen” means that the model weights are not adapted any more at this point. You receive a prediction Y that you can

compare to the true test values Y . Both values (or sets of values) can be used in your model evaluation.




Figure 1.8 Deep Learning: Model Inference

This book is not a book on AI engineering. We won't train a model ourselves in this book. The good thing is we don't have to! The models that we'll use are already trained on vast amounts of data, and we can use these models with their frozen model weights to our purposes and just run model inference.

Depending on the model architecture and available training data, you can train many kinds of models. Typically, these models are extremely well adapted to a specific dataset, that is, working rather well, but in a rather narrow domain. Let's turn to this topic next.

1.4 Narrow AI and General AI

AI systems have become more and more powerful, mostly absent from the public eye. But every now and then, their capabilities reach the public's attention. [Figure 1.9](#) shows some of the narrow achievements of early AI systems.




Figure 1.9 Narrow AI Achievements

In 1997, IBM's Deep Blue, an extremely sophisticated chess playing machine, beat world chess champion Garry Kasparov, and the world was stunned that a machine won. For centuries, chess was associated with outstanding cognitive capabilities. Furthermore, good chess players have intuition, which is required to anticipate the moves of the opponent. But the machine won and showed superhuman capabilities at playing chess. In hindsight, we must recognize that this was merely the result of a comparably simple search algorithms and simulations of millions of moves. Deep Blue did not apply any advanced artificial intelligence. Nevertheless, the AI scored its first big win in several battles of wit between man and machines.

Much more impressive was when Google DeepMind's system AlphaGo beat Lee Sedol, one of the best Go players in the world in a 2015 match. AlphaGo won four out of the five games and received a \$1 million prize money. Experts in Go did not expect AlphaGo to win at all because of the game's extraordinary high game tree

complexity. Sometimes called the “Shannon number” (after Claude Shannon), *game tree complexity* refers to the number of possible games that could theoretically be played. The Shannon number for chess is estimated to be around 10^{123} . Although an insanely large number, many combinations have been simulated for chess.

In Go, a simple “brute force” approach to simulating all game combinations is impossible. The Shannon number is estimated to be around 10^{360} . (We’ll save you the trouble of reading a number with 360 zeros.) The reason for this large difference in complexity is mainly the board’s size. In chess, the board is an 8×8 , while Go is played on a 19×19 grid. The game dynamics are also different: In chess, a typical game involves about 40 to 60 moves per player. In Go, a game can last around 200 to 300 moves per player.

The significance of a machine winning at Go cannot be overestimated: Now, humans lose to machines in all one-versus-one games. We thought only humans were capable of intuition, an essential part of winning at Go.

In the second game of the five-game match, the AI made a move that was completely incomprehensible to the experts. “Move 37” later became famous as one of the most significant moves in the history of the game of Go. AlphaGo made a move that no human would ever have made at that point in the game. A Go stone was placed in a position on the board that was completely unexpected. The experts were unanimous: The AI made a major mistake that would take its revenge later. In fact, much later in the game, it turned out to be a brilliant move that the experts would analyze for months to come. What’s more, this move not only proved that AI has intuition in its own way. After thousands of years of playing Go, an AI was able to present humans with abilities far beyond their own.

At this point, we have an AI system able to beat the best Go players in the world. AlphaGo was at the top of the Go world. But just two years later, in 2017, AlphaGo lost 100 out of 100 games against its opponent. What happened? Did humankind strike back? Not really. No human will ever hold the crown in Go again. What happened was evolution, more specifically AI evolution. AlphaGo lost against its offspring, a more advanced AI system called AlphaGo Zero.

AlphaGo Zero did not rely on studying gameplay from human games. It learned by playing against itself and improved in mere 21 days to the level of AlphaGo. After 40 days, AlphaGo Zero had exceeded all previous versions.

OK, you might think, at this point, AI systems can reach expert-like capabilities in very narrow fields, and you're right!. But could AI systems be successful in real-time, multiplayer games? They are not capable of teamwork and setting dynamic objectives, right? Well, in 2019, a system from OpenAI called "Five" played the video game Dota 2. This system controlled five AI agents, one for each hero on a team. These agents had to coordinate seamlessly to execute quite complex tasks and strategies. The AI system had to make decisions 20 times per second. Short-term and long-term tactics had to be carefully balanced.

In a live event in 2019, "Five" played against the international champions OG who won "The International" a year earlier. The event took place over four days. The AI agents played close to 43,000 games of which they won a staggering ratio of 99.4%.

More impressive AI systems were developed that could play and win games like Diplomacy. In this complex strategy board game, negotiations, collaborations, and strategic planning are crucial to success. Meta developed a system called Cicero that can excel at human-like communication and social reasoning.

AI systems started in rather narrow domains, like Chess and Go, and have already achieved mastery in complex environments that require far more skills than just being a solitary genius.

Where are we today, in 2025? AI systems can solve more and more cognitive tasks.

OceanofPDF.com

1.5 Natural Language Processing Models

NLP forms the backbone of many generative AI systems. NLP enables machines to understand, generate, and interact with human language. At the heart of NLP are core tasks that act as building blocks for complex applications.

1.5.1 NLP Tasks

In this section we'll explore some foundational NLP tasks.

[Figure 1.10](#) shows some typical NLP tasks—text classification, text generation, question answering, translation, fill-mask tasks, and sentence similarity. These different tasks will be explained in more detail.

Language modeling is the task of predicting the next word in a sequence or assigning probabilities to sequences of words. It serves as the cornerstone of many NLP applications, including text generation, speech recognition, and machine translation. Language models, such as GPTs, learn statistical patterns in text data, enabling them to produce coherent and contextually relevant sentences. By understanding language structure and semantics, these models can

generate human-like text, perform auto-completions, and even engage in creative writing.




Figure 1.10 Typical NLP Tasks

Text Classification

Text classification involves categorizing text into predefined labels or classes. This task is foundational for many NLP applications, such as spam detection, sentiment analysis, topic categorization, and intent recognition. Models trained for text classification learn to identify patterns and relationships in textual data, enabling them to accurately assign labels based on content.

Text classification is pivotal in organizing large datasets and enhancing user experiences by tailoring responses or content recommendations.

Example

Imagine you're developing an email filtering system. The goal is to classify incoming emails as either "spam" or "not spam." To achieve this goal, a text classification model is trained on a labeled dataset containing examples of both categories. By analyzing the text, the model might detect suspicious phrases like "free money"

or “click here.” The model learns to predict the correct label for new, unseen emails and can now help users automatically organize their inboxes and focus on the most important communications.

Text Generation

Probably the most prominent task in NLP since ChatGPT was released, text generation will also be the task we’ll mostly use in this book.

Text generation is the process of creating coherent and contextually appropriate text based on a given input or prompt. This task is central to many applications, including creative writing, chatbots, content creation, and code generation.

Modern text generation systems excel in producing human-like responses by leveraging vast datasets and advanced language modeling techniques. These models can generate anything from poetry and stories to technical explanations and program code, all tailored to the specific needs of users.

Example

A text generation system can take the user prompt “Write a short story about a robot exploring a new planet” and produce a detailed narrative complete with descriptions, dialogue, and plot development. The quality of the output depends on the training data and the model’s ability to understand and extend the context provided.

Question Answering

Question answering is the task of extracting or generating answers to natural language questions. Question answering systems can be classified into two types: retrieval-based and generative. Retrieval-based systems locate answers from a pre-existing corpus, while generative systems like OpenAI's ChatGPT synthesize new answers based on input queries. This task combines language understanding, reasoning, and contextual knowledge, making it essential for chatbots, virtual assistants, and automated customer support.

Text Translation

Text translation aims to convert text or speech from one language to another. Modern translation systems utilize deep learning models to achieve high accuracy and fluency. Transformer-based architectures, such as Google's BERT (bidirectional encoder representations from transformers) and OpenAI's models (which we'll discuss later) have revolutionized this task by capturing the nuances of context and syntax in both source and target languages. Applications of machine translation range from breaking language barriers in global communications to facilitating multilingual content creation.

Fill-Mask Tasks

Fill-mask tasks involve predicting missing or masked words within a given text. This capability is foundational for understanding context and is widely used in pretrained models like BERT.

During training, certain words in a sentence are masked, and the model learns to predict them based on the surrounding context. Fill-mask tasks help models grasp deeper semantic relationships and are a critical step in developing robust NLP systems.

Example

Given the sentence “The weather is [MASK] today,” a fill-mask model might predict “sunny,” “rainy,” or other contextually plausible words. This ability not only enhances language understanding but also supports downstream tasks such as question answering and text generation.

Sentence Similarity

Sentence similarity involves measuring how semantically close two pieces of text are. This task is fundamental in applications like vector databases, on which we focus in [Chapter 5](#).

But sentence similarity is also relevant for plagiarism detection, duplicate question identification, and recommendation systems. Models trained for sentence similarity use *embeddings* to represent sentences in a high-dimensional space, allowing similarity to be calculated as the distance between their representations.

Example

Comparing the sentences “How do I reset my password?” and “What should I do to recover my password?” a sentence similarity model would identify a high degree of similarity between them, as they convey almost the same meaning. This capability is crucial for systems like FAQ matching and conversational AI.

Text Summarization

Text summarization involves condensing lengthy documents into shorter, coherent versions while retaining the most critical information. This task can be approached in two ways: extractive summarization and abstractive summarization. Extractive methods identify and concatenate key sentences or phrases from the source text, while abstractive methods generate new sentences that paraphrase the content. Summarization plays a crucial role in information retrieval, content aggregation, and assisting users in navigating large volumes of data efficiently.

1.5.2 Architectures

Language modeling has evolved significantly over the years and has been shaped by the development of foundational architectures. Let's briefly discuss some historical developments that have led to our current architecture.

Common Architectures

Early models relied on recurrent neural networks (RNNs), which were designed to process sequential data by maintaining a hidden state that carried information from one time step to the next. This structure allowed RNNs to handle basic tasks like text classification or sequence prediction. However, they struggled with capturing long-term dependencies in text due to the vanishing gradient problem, which made learning over extended sequences difficult.

To address these limitations, long short-term memory (LSTM) networks were introduced. LSTMs extended the capabilities of RNNs by incorporating memory cells and gates, which allowed the model to

retain important information over longer sequences while filtering out irrelevant details.

This advancement made LSTMs effective for tasks such as machine translation and text summarization, particularly when sequence length and context played a critical role. Despite their improvements, LSTMs were computationally intensive, especially when processing large datasets, and their sequential nature limited scalability.

The advent of the transformer architecture in 2017 marked a turning point in NLP. Unlike RNNs and LSTMs, transformers abandoned sequential processing in favor of a self-attention mechanism that enabled parallel processing of input data.

This approach allowed the model to capture relationships between words regardless of their distance in the sequence, making it both faster and more accurate.

Transformer architecture also introduced positional encoding to account for the order of words in a sequence, preserving context while maintaining computational efficiency. These innovations form the backbone of virtually all modern LLMs.

Practical Implementations

Practical implementations of transformer-based models have revolutionized the NLP landscape. Google's BERT is one of the most impactful models to emerge.

By training on masked language modeling and next-sentence prediction tasks, BERT developed a deep understanding of text context, enabling it to excel in applications such as question answering, sentiment analysis, and named entity recognition (NER). BERT can process text bidirectionally and thus can consider both

preceding and following words, which has set a new standard for contextual understanding.

Another standout implementation is the GPT series, developed by OpenAI. Unlike BERT, which is optimized for understanding text, GPT models are designed for generating text. You know these models implicitly because they form the foundation of ChatGPT.

These models predict the next token in a sequence based on preceding tokens, making them capable of producing coherent, contextually relevant, and creative outputs. From the highly versatile GPT-3.5 to the advanced capabilities of GPT-4, o1, and o3, these models have been widely adopted in end-user applications.

The progression from RNNs and LSTMs to transformers underscores the transformative impact of architectural innovation in NLP. By enabling deeper understanding, broader context awareness, and more efficient processing, these developments have paved the way for practical, real-world applications that continue to shape how we interact with and leverage language models today.

1.6 Large Language Models

We'll work a lot with LLMs in this book, and a complete chapter will deal with this topic. So, we'll keep it short and concise because we'll dive much deeper into this fascinating topic in [Chapter 3](#).

LLMs are advanced ML models trained on vast amounts of text data to understand, generate, and interact with human language. They leverage architectures like transformers to process and produce contextually coherent and semantically meaningful text, enabling applications such as chatbots, content generation, and question answering.

In the following sections, we'll first describe the different steps of the training process for an LLM. Then, we'll highlight some use cases before looking at some of the problems associated with LLMs.

1.6.1 Training

[Figure 1.11](#) shows the training process for LLMs.




Figure 1.11 LLM Training Process

An LLM is trained in a process involving the following steps:

1. Data

The process starts with collecting and preparing a large amount of diverse and high-quality text data. A lot of data preprocessing must be performed, including data cleansing, tokenization, and organizing to ensure consistency and relevance.

2. Pretrained model

Once a model is trained on this data, a pretrained model is derived. This process involves predicting the missing words or the next words based on the context. The goal in this step is to capture general patterns and structures. The result is a general-purpose model capable of understanding and generating human-like text.

3. Instruction model

The pretrained model is fine-tuned using additional datasets or techniques that are tailored to specific tasks or instructions. This step ensures that the model generates outputs that are aligned with particular applications or user instructions. This instruction-tuning and fine-tuning helps achieve high-quality model outputs by making the model more task specific.

4. Safety model

This step involves the alignment of the model to human preferences. The aim of this step is to ensure that the model can be operated safely. In this context, techniques like reinforcement learning from human feedback (RLHF) are used to make the model less biased and more reliable. A safety model models some harmful or dangerous behavior in generated outputs. For example, a model should refuse to help a user build a bomb or hide a dead body.

5. Evaluation

The final model is rigorously evaluated using benchmarks and

real-world tasks to ensure its quality and alignment. Numerous evaluation metrics assess the accuracy, relevance, safety, and robustness of the model and its outputs. Further refinements might be necessary, so that additional feedback loops might be performed.

1.6.2 Use Cases

LLMs have revolutionized the way we interact with artificial intelligence, opening up a vast array of use cases that span across industries and disciplines.

These models, powered by immense datasets and advanced neural network architectures, are designed to understand and generate human-like text, which makes them a versatile tool for various applications, as we'll discuss next.

Creative Writing

One of the most prominent and widely recognized applications of LLMs is content generation. LLMs are adept at generating text that is coherent, contextually relevant, and often indistinguishable from human writing. This ability has application in crafting articles, blogs, and social media posts. Businesses and individuals alike can use LLMs to brainstorm ideas, refine drafts, and even produce entire pieces of work.

An author can use an LLM to generate plot ideas and character dialogues, while a marketing team can create tailored promotional content targeting specific groups.

We'll encounter this feature again in many chapters, like [Chapter 3](#), [Chapter 4](#), and [Chapter 7](#).

Summarization

Closely related to content creation is the use of LLMs for summarization. The sheer volume of information generated daily in the digital age makes it challenging to sift through and extract key insights. LLMs excel in this domain by distilling large texts into concise summaries, retaining the essence of the original content.

This capability is invaluable for professionals and researchers who need to quickly grasp the main points of lengthy reports, articles, or academic papers. Whether summarizing a dense legal document, providing an executive summary for a business proposal, or creating a digestible recap of the day's news, LLMs save time and effort while enhancing productivity. We'll revisit this capability in [Chapter 5](#) and [Chapter 6](#).

Chatbots

Another compelling use case is conversational AI, where LLMs are employed to power chatbots and virtual assistants. These systems can engage users in natural, dynamic conversations; answer questions; provide recommendations; or assist with tasks.

Businesses deploy such bots to improve customer support, handling routine inquiries efficiently and allowing human agents to focus on more complex issues. On a personal level, virtual assistants equipped with LLMs can manage schedules, set reminders, and even act as tutors, guiding users through learning new skills or solving problems.

Translation

Translations also benefit significantly from LLMs. These models bridge linguistic barriers by providing accurate translations that

capture not just the literal meaning but also the nuances of context and tone. Translation is particularly beneficial for global businesses aiming to localize their content or individuals seeking to communicate across cultures. Moreover, LLMs are instrumental in language learning applications, offering learners personalized feedback, conversational practice, and explanations for grammar or vocabulary.

Content Creation

Creative applications of LLMs extend beyond text-based tasks. They can assist in composing music, generating code, and even developing game narratives or dialogue systems.

In the gaming industry, for instance, LLMs can create immersive storylines and realistic non-player character interactions, enhancing the overall player experience. Similarly, in software development, LLMs help developers by generating code snippets, suggesting improvements, or automating routine programming tasks.

Education

In the field of education, LLMs serve as valuable tools for both teaching and learning. They can generate tailored lesson plans, quizzes, and instructional materials. LLMs can adapt to diverse learning styles and levels. Students can interact with LLMs to clarify doubts, explore topics in depth, or receive assistance with writing assignments. The interactive and adaptive nature of these models fosters a more engaging and personalized educational experience.

Other use cases and affected industries are too numerous to list as most industries will be affected by this revolution.

1.6.3 Limitations

LLMs are incredible powerful, but they also exhibit certain limitations. Some of these limitations include the following:

- **Hallucinations**

LLMs can generate text that appears plausible but is factually incorrect, misleading, or entirely fabricated. These “hallucinations” occur because the models rely on statistical patterns in training data rather than understanding or verifying facts. Thus, LLMs are unreliable for tasks requiring high factual accuracy, such as medical diagnoses or legal advice.

- **Biases**

LLMs inherit biases present in their training data, which may reflect societal stereotypes, cultural imbalances, or skewed perspectives. This bias can result in outputs that are prejudiced or offensive. Addressing bias is challenging because mitigating one bias may inadvertently introduce or amplify another.

- **Context window and token length constraints**

LLMs can process only a limited amount of text. These constraints restrict the length of input text (context window) or of the output text they can process in a single interaction. This constraint can limit their ability to handle long documents, detailed instructions, or extended conversations. Splitting content into chunks to work within these limits may lead to loss of context or coherence.

1.7 Large Multimodal Models

Like LLMs, LMMs are advanced AI systems. But in contrast to LLMs, LMMs are capable of understanding and possibly creating multiple types (modalities) of data, such as text, images, audio, and even video. [Figure 1.12](#) shows possible input and output modalities of an LMM.




Figure 1.12 Large Multimodal Models

These models integrate diverse modalities into a single model. These diverse modalities can be either on the input and/or the output of the model. Some models allow for both text and images as inputs and provide text as output, called *vision-language models (VLMs)*. Popular examples are Llama 3.2 Vision, GPT-4o, or Google Gemini. The different modalities enable them to perform complex tasks that require cross-modal understanding and reasoning.

Unlike unimodal models (like LLMs), which operate exclusively on a single type of data, LMMs can generate text from an image, create images based on user prompts, extract information from videos, and much more. This versatility has placed LMMs at the forefront of AI

research. A shift from focus on LLMs to LMMs can be observed, and this process will continue until eventually LMMs mostly replace LLMs.

LMMs are transforming various industries by introducing innovative solutions and automating complex workflows. In the art and design domain, tools like text-to-image conversion, which were previously standalone models, will be integrated into LMMs. A prominent example is DALL-E, which is integrated into GPT-4o.

In healthcare, LMMs have the potential to revolutionize medical diagnostics and patient care. They can analyze medical images alongside patient records to provide more accurate diagnoses. Furthermore, they assist in creating educational materials for patients, simplifying complex medical information with visuals and explanatory text.

The entertainment industry is also leveraging LMMs to produce content like video game assets, movie storyboards, and even interactive narratives. By synthesizing text and images, LMMs help creators rapidly prototype ideas and develop high-quality visuals for immersive experiences.

Despite their promise, LMMs face significant challenges. One major issue is the difficulty of achieving seamless integration between modalities. For example, ensuring that an image accurately represents the nuances of a textual description requires robust alignment techniques, which are computationally intensive and prone to error.

Another challenge lies in the ethical and societal implications of multimodal AI. The potential for misuse, such as generating deepfakes or misleading content, raises concerns about trust and

accountability. Moreover, biases in training data can propagate into the models, leading to unintended consequences in their outputs.

Finally, the computational demands of training and deploying LMMs are substantial. These models require vast amounts of data and powerful hardware, which can limit accessibility and raise sustainability concerns due to high energy consumption.

Let's now look at the wider field of generative AI applications.

OceanofPDF.com

1.8 Generative AI Applications

Based on the capabilities of LLMs and LMMs, and other more narrow models with specific focus on creation of music, voice, videos, and much more, a huge ecosystem has developed around generative AI.

The following section highlights the generative AI market, specifically the different areas in which these models are used. The map clusters the field in three main categories: systems for consumer use, for enterprise, and for prosumer use. Let's consider each use category:

- **Consumer**

The consumer market will be most familiar to you, through applications like chatbots, text-to-image converters, text-to-video converters, personal assistants, and more. The dominant business model is freemium, in which you can use the service under defined limits, unless you become a paid subscriber.

- **Enterprise (horizontal)**

Horizontal integration means that generative AI companies provide solutions that can be used across different industries. Think of chatbots: They are extremely versatile and can be used in completely different industries like finance, healthcare, education, and many other industries.

- **Enterprise (vertical)**

In vertical integration, a service tries to expand its control over the supply chain. In the context of generative AI, these services are tailored to specific industries and have a deep integration.

- **Prosumer**

Prosumer services target independent professionals, small

businesses, or power users. The primary business models are premium subscriptions or usage-based pricing.

It is amazing to see what a vast ecosystem has evolved around generative AI in such a short period of time.

1.8.1 Consumer

For consumers, many products are available, and we can only mention a small number:

- With Suno (<https://suno.com/>), you can create music based on a short description.
- Med-PaLM 2 (<https://sites.research.google/med-palm/>) is a model fine-tuned for the medical domain that has reached outstanding accuracy in the MedQA medical exam benchmark.
- Services like Lensa AI (<https://lensa.app/>) or Remini (<https://remini.ai/>) can be used for photo improvements and avatar generation.
- Character.ai (<https://character.ai/>) allows you to chat with millions of AI characters at anytime and anywhere.

1.8.2 Business

For businesses and enterprises, both horizontal and vertical AI-based systems are available. “Horizontal” means that the service is not domain specific, while a “vertical” service is domain specific.

For horizontal services, again, we can only provide a short list of offerings that we’ve personally tested:

- Cursor (<https://www.cursor.com/>) is my favorite AI code editor. It is integrated into a fork of Visual Studio Code (VS Code) and can increase productivity a great deal. Similar services provide tabnine (<https://www.tabnine.com/>) or GitHub Copilot (<https://github.com/features/copilot>).
- With Notion (<https://www.notion.com/>) you can capture your thoughts, create plans, or manage your projects. It has a seamless integration of AI features.

Vertical enterprise services aim at improving multiple tasks within a dedicated domain. Here are two important offerings:

- In healthcare, companies like Ambience Healthcare (<https://www.ambiencehealthcare.com/>) claim to improve patient care with AI.
- The legal field services like Harvey (<https://www.harvey.ai/>) provide domain-specific AI for law firms.

1.8.3 Prosumer

Prosumer is an artificial word that represents the combination of producer and consumer. Services in this area can be applied by enterprises and consumers alike and include the following:

- Most notably are search and knowledge services. Everybody knows ChatGPT (<https://chatgpt.com/>). Alternatives are Claude (<https://claude.ai/>), Perplexity (<https://www.perplexity.ai/>), or Google Gemini (<https://gemini.google.com/app>).
- Image creation was hyped up in 2022 when the first models allowed users to create images from text prompts. But the quality has improved even further, and we personally would consider these algorithms to have reached the plateau of productivity.

We've had good experiences with Leonardo.ai (<https://leonardo.ai/>) and fal (<https://fal.ai/>).

- Impressive improvements can be observed in video creation or editing services. Runway (<https://runwayml.com/>) and Kling AI (<https://klingai.com/>) allow you to create impressive AI-generated videos.

OceanofPDF.com

1.9 Summary

This chapter provided a comprehensive overview of generative AI, a transformative branch of artificial intelligence centered on the creation of new and innovative content.

In [Section 1.1](#), we discussed the evolution of artificial intelligence, from its first steps nearly a hundred years ago, to the LLMs and LMMs of today. You also learned that every technology goes through a hype cycle, and we discussed the cycle for generative AI. We briefly examined the capabilities of various AI systems compared to human-level capabilities.

In [Section 1.2](#), we studied the pillars of the evolution of generative AI. Specifically, we saw the exponential growth in computational power, models and data size, investments, and algorithmic improvements. [Section 1.3](#) covered the underlying foundation of generative AI, namely, deep learning. You saw how a model is trained and used in deployment. Typically, a deep learning model is developed with a narrow domain in which the model can excel. LLMs and LMMs broaden the scope to more general capabilities. We discussed narrow and general AI in [Section 1.4](#).

NLP models were the main topic of [Section 1.5](#). We started by discussing various NLP tasks like text classification, translation, and more. Later, we discussed different architectures, from LSTM to transformers.

LLMs were introduced in more detail in [Section 1.6](#). We discussed the training process and its specific steps. We then looked at several use cases of LLMs, such as creative writing, summarization, translation, or content creation. Despite their incredible capabilities,

we also discussed some limitations, like hallucinations, biases, and context window limitations.

Expanding our scope from single-modality models like LLMs, which can only consume and return text, we entered the field of LMMs in [Section 1.7](#). These models can understand multiple types of data, such as text, images, audio, and even video. We discussed their applications, and their challenges.

Finally, in [Section 1.8](#), we covered some generative AI applications. We started with a generative AI market map that classifies different services into consumer, enterprise, and prosumer services, and we provided some examples in each segment.

Through this structured progression, we've set the stage for a deeper exploration of generative AI and its transformative impact across diverse fields.

In the next chapter, we'll discuss pretrained models. These narrow AI models excel in a specific domain. The chapter will discuss the most popular framework Hugging Face, which allows you to easily integrate such pretrained networks into your workflow.

2 Pretrained Models

The more complex the mind, the greater the need for the simplicity of play.

—Isaac Asimov, in “I, Robot”

Pretrained models are AI systems that are capable of processing large amounts of data and working on a dedicated task. Pretrained models are not a Swiss army knife with a lot of different features, but typically, they work well for one specific task. Then came large language models (LLMs). At that time, there were individual models for each task. In the field of natural language processing (NLP), a task might be text classification, question answering, zero-shot classification, translation, summarization, feature extraction, text generation, fill-mask, or sentence similarity. The list is not exhaustive, and more tasks will be possible over time.

In [Chapter 3](#), we’ll work mostly with LLMs. Given their extraordinary capabilities, why should we start with the rather classic machine learning (ML) models. [Figure 2.1](#) shows a comparison of the training processes for classic ML models and LLMs.




Figure 2.1 Classic ML Models versus LLM Training

In classic ML models, for each task, a dedicated model is created. Usually, supervised learning is applied, which is based on a labeled dataset with inputs and corresponding target information. For example, a text translation model is created based on words or sentences in a source language (labeled input), and words or sentences in a target language. This approach works quite well. As a result, you get a model dedicated to a specific task. One advantage is that the model is perfectly adapted to its needs; another, that it is rather small. In most cases, you can run it on your local computer. In this chapter, you'll discover many of these models.

In contrast, LLMs are developed in a different way. The input is typically unlabeled, meaning that the data might come from several different sources. It might even be multimodal and represent text data as well as images or video data. Few specialized models are trained, but instead, we are operating one “large” language model. The advantage is that the capabilities of these models consist of the aggregation of all small models.

The disadvantage of LLMs is that, compared to a dedicated small pretrained model, the inference of an LLM might be quite costly, with longer inference times making it unsuitable for real-time applications. Also, some applications might require the operation of a model locally bound to the systems of your company, for instance, for data protection reasons. But given the size of today's LLMs these requirements might be hard to impossible to implement.

Consequently, a beneficial approach is to run a simple model if the task is simple, using an LLM only when the task requires it.

In this chapter, we'll explore what pretrained models are, where you can find them, and how you can use them for your needs. In [Section 2.1](#) we'll take you to Hugging Face, a platform for ML models, datasets, and much more. You'll learn how to work with the

platform and how to find suitable models for given tasks. That knowledge is practically applied in [Section 2.2](#) on a specific task, summarizing a longer text with a model from Hugging Face.

After you've seen how to implement a task, it's your turn, and you can apply the information in [Section 2.3](#) on a given task. In subsequent sections, we explore several tasks together, and you'll learn in which situations a task can help you. In [Section 2.4](#), you'll encounter zero-shot classification. Then, we implement fill-mask tasks in [Section 2.5](#), question answering in [Section 2.6](#), named entity recognition (NER) in [Section 2.7](#), and text-to-image conversion in [Section 2.8](#).

In the text-to-audio conversion exercise in [Section 2.9](#), you'll create music based on a prompt. Finally, in a capstone project on the analysis of customer feedback, in [Section 2.10](#) you'll integrate several different models to solve a more complex task. In each of these steps, you'll encounter a platform that reshaped the world of ML and simplifies development through the sharing of models and datasets—Hugging Face.

2.1 Hugging Face

The pretrained models in our scope are language models trained for a specific task. They are good at their job, but expert only in a very narrow scope. Given a prepared dataset, you can train such a model from scratch or fine-tune an existing model to be more suitable for working with your data.

But in many cases, you can use such models out of the box. One platform is synonymous with pretrained, specialized models. Hugging Face is an AI company and open-source platform that

focuses on building and providing tools for NLP, computer vision, and many other ML tasks. [Figure 2.2](#) shows the Hugging Face homepage, available at <https://huggingface.co>.




Figure 2.2 Hugging Face Home Page

At Hugging Face, you'll find more than one million models (at the time of writing). In a production environment, you'll simply want a sufficient model for a given task. Applying an LLM for a comparably simple task is not necessary, and this approach increases costs and possibly latency.

Thus, in the following section, you'll become familiar with Hugging Face and select narrow AI models for specific tasks. However, you must first understand how to select the right model out of a million options. [Figure 2.3](#) shows the model selection process in Hugging Face (<https://huggingface.co/models>) where you first select the task ❶ and then select a model ❷.

First, identify the task you're working on and select it. This step basically works as a filter, and you'll now see only the models in this category. Then, filter by name. For example, if you're looking for a **Translation** model, you might enter the source or target language in the text field to further filter the results. You can also sort the results

based on different criteria like trending, most likes, or most downloaded.




Figure 2.3 Hugging Face Model Selection Process

2.2 Coding: Text Summarization

In this example, we'll load an Arxiv article and summarize it. *Arxiv* is an open-access archive for more than 2 million scholarly articles from different fields like physics, mathematics, or computer science. The articles are typically available as PDF files. Luckily, the LangChain community has developed a loader for this type of article that we can use. We'll load an article on the topic "prompt engineering" and subsequently run it through a summarizer model that loaded from Hugging Face. Let's go!

Once you have selected a model on Hugging Face by clicking on it, you're greeted with its *model card*. On this card, you'll typically find metrics about the model's performance (based on selected benchmarks) as well as information on how to use the model.

shows a sample model card for a text summarization model.

On the model card, you'll typically find how to use it in the left column, as well as some metrics that provide information on the model's performance based on different benchmarks.

On the right, you'll often find an inference section that allows you to test the model. In this example, you might add some text to be summarized.

In addition to the model card, under the **Files and versions** tab, you'll find how big the specific model is, which is often an important indicator for the deployment of the model.




Figure 2.4 Sample Model Card (Source: <https://huggingface.co/sshleifer/distilbart-cnn-12-6>)

Under the **Community** tab, you'll see the history of the model and possibly information about the model's developers.

These models are usually based on transformer architecture. One of the most relevant packages dealing with this architecture has the same name: `transformers`. This package holds the `pipeline` function, which is utility function for setting up a pipeline. Also, we'll load `ArxivLoader` to download and import articles directly into a `LangChain` document object. You can find the code in `02_PreTrainedNetworks\10_text_summarization.py`, which looks as follows:

```
from transformers import pipeline
from langchain_community.document_loaders import ArxivLoader
```

The `pipeline` function requires a `task` and a `model` name. Both parameters are passed, and as a result, we get an instance of a pipeline, in our case, a summarizer, as follows:


```
### model selection
task = "summarization"
model = "sshleifer/distilbart-cnn-12-6"
summarizer = pipeline(task= task, model=model)
```

We can load one or more documents. For the purposes of illustration, we'll stick to just one document on the topic "prompt engineering". The loader object is based on `ArxivLoader`, and its method `load()` must be called so that the documents are imported, as follows:

```
### Data Preparation
query = "prompt engineering"
loader = ArxivLoader(query=query, load_max_docs=1)
docs = loader.load()
```

You might check the `docs` object, which is a list with one `Document()`. The document has multiple kinds of metadata, but we are only interested the `page_content` property, which holds the article text, as follows:

```
# %% Data Preparation
article_text = docs[0].page_content
```

At this point, you've extracted the text, which is stored in `article_text`. Let's pass this text to our `summarizer`. Notice how we cannot pass the complete article due to limits in the model's context window, which we'll discuss in detail in [Chapter 3](#). We can limit the text and use only its first 2,000 characters. We can also control the output and set a minimum length and maximum length for the summary. The result of the model run is stored in `result`, which is a list with a dictionary. We are interested in the `summary_text` (see [Listing 2.1](#)).

```
# %% Run the summarizer
result = summarizer(article_text[:2000], min_length=20, max_length=80,
do_sample=False)
result[0]['summary_text']
```

The paper investigates prompt engineering as a novel creative skill for creating AI

art with text-to-image generation. Prompt engineering is a new type of skill that is non-intuitive and must first be acquired before it can be used.

Listing 2.1 Text Summarization Run

Finally, let's check if the output length is within the requested limits, as follows:

```
# %% number of characters  
len(result[0]['summary_text'].split(' '))
```

41

Good! We're within our required limits of 20 and 80 characters. Now, it's your turn to find a model and use it in a translation task.

OceanofPDF.com

2.3 Exercise: Translation

Now that you know how to implement and use a model from Hugging Face, time for you to get behind the wheel yourself. In this exercise, you'll need to understand the task, find a suitable model, and implement it.

2.3.1 Task

A typical NLP task is translation. With the knowledge from the previous section, you can now solve this task independently. Find a suitable model on Hugging Face to translate the following sentence into Japanese:

Be the change you wish to see in the world.

This famous saying is attributed to Mahatma Gandhi. Try to solve it independently and read on to see our solution.

2.3.2 Solution

First, check the model's section on Hugging Face (<https://huggingface.co/models>). Then, filter for **Translation** tasks under the **Natural Language Processing** section. Nearly 5,000 translation models are available, so we'll need to filter further.

If you check the naming of the models, you'll see that most adhere to the naming convention [source language]-[target language]. Thus, narrow things down with "en-ja." This step brings us to the model **Mitsua/elan-mt-bt-en-ja**.

Its model card clearly describes how to use it, and this model follows the same structure as in the previous section. Use `pipeline` and pass a corresponding `task` and `model`. Then, pass the text to be translated into the pipeline. In the result, you'll see the translated text.

[Listing 2.2](#) shows our sample source code for a translation task from English to Japanese.

```
### packages
from transformers import pipeline

### model selection
task = "translation"
model = "Mitsua/elan-mt-bt-en-ja"
translator = pipeline(task=task, model=model)

# %%
text = "Be the change you wish to see in the world."
result = translator(text)
result[0]['translation_text']
```

世の中の見たい変化になれ。

Listing 2.2 Translation from English to Japanese (Source: 02_PreTrainedNetworks\20_translation.py)

We hope you succeeded and came up with the same translation, possibly using a different model.

In the next section, we'll move on to the next useful task: zero-shot classification.

2.4 Coding: Zero-Shot Classification

Zero-shot classification is a powerful capability in modern AI systems that allows them to categorize or classify items they've never been explicitly trained on. Unlike traditional classification models that require labeled examples for each category, zero-shot classifiers can make educated guesses about new categories by leveraging their understanding of natural language and the semantic relationships between concepts.

For instance, a computer vision or multimodal model trained to identify cats and dogs might be able to classify a zebra correctly simply by understanding the linguistic description “black and white striped horse-like animal,” even though it never saw zebra examples during training. This ability comes from the model's pretrained knowledge of language and concepts, making it extremely versatile for real-world applications where new categories frequently emerge.

In our coding example, we'll choose the opening lines of three different books. These books clearly can be categorized into romantic, fantasy, and crime novel. Our task is to find out if a zero-shot classifier can correctly categorize these books into genres purely based on the first few sentences of the books. Let's get started.

We'll use `pipeline` to import and interact with a zero-shot classifier model. `pandas` is only used for visualization of the result. The following script can be found at

02_PreTrainedNetworks\50_zero_shot.py:

```
### packages
from transformers import pipeline
import pandas as pd
```

Call pipeline with the task zero-shot-classification and the model selected in the following. Feel free to select a different one and compare the results.

```
### Classifier
classifier = pipeline(task="zero-shot-classification", model="facebook/bart-large-mnli")
```

For the sample data, as shown in [Listing 2.3](#), we use Jane Austen's *Pride and Prejudice* as an example of a romantic novel. For a crime novel, the classic detective is chosen—Sherlock Holmes. The opening of Arthur Conan Doyle's *The Return of Sherlock Holmes* shall serve as an example of a crime novel. One of the most famous fantasy novels is *Alice's Adventures in Wonderland*. Lewis Carroll's classic novel serves as an example of a fantasy novel.

```
# %% Data Preparation
# first example: Jane Austen: Pride and Prejudice (romantic novel)
# second example: Lewis Carroll: Alice's Adventures in Wonderland (fantasy novel)
# third example: Arthur Conan Doyle "The Return of Sherlock Holmes" (crime novel)
titles = ["Pride and Prejudice", "Alice's Adventures in Wonderland", "The Return of Sherlock Holmes"]
documents = [
    "Walt Whitman has somewhere a fine and just distinction between "loving by allowance" and "loving with personal love." This distinction applies to books as well as to men and women; and in the case of the not very numerous authors who are the objects of the personal affection, it brings a curious consequence with it. There is much more difference as to their best work than in the case of those others who are loved "by allowance" by convention, and because it is felt to be the right and proper thing to love them. And in the sect—fairly large and yet unusually choice—of Austenians or Janites, there would probably be found partisans of the claim to primacy of almost every one of the novels. To some the delightful freshness and humour of Northanger Abbey, its completeness, finish, and entrain, obscure the undoubted critical facts that its scale is small, and its scheme, after all, that of burlesque or parody, a kind in which the first rank is reached with difficulty.",
    "Alice was beginning to get very tired of sitting by her sister on the bank, and of having nothing to do: once or twice she had peeped into the book her sister was reading, but it had no pictures or conversations in it, and what is the use of a book, thought Alice "without pictures or conversations? So she was considering in her own mind (as well as she could, for the hot day made her feel very sleepy and stupid), whether the pleasure of making a daisy-chain would be worth the trouble of getting up and picking the daisies, when suddenly a White Rabbit with pink eyes ran close by her.",
    "It was in the spring of the year 1894 that all London was interested, and the fashionable world dismayed, by the murder of the Honourable Ronald Adair under most unusual and inexplicable circumstances. The public has already learned those
```

particulars of the crime which came out in the police investigation, but a good deal was suppressed upon that occasion, since the case for the prosecution was so overwhelmingly strong that it was not necessary to bring forward all the facts. Only now, at the end of nearly ten years, am I allowed to supply those missing links which make up the whole of that remarkable chain. The crime was of interest in itself, but that interest was as nothing to me compared to the inconceivable sequel, which afforded me the greatest shock and surprise of any event in my adventurous life. Even now, after this long interval, I find myself thrilling as I think of it, and feeling once more that sudden flood of joy, amazement, and incredulity which utterly submerged my mind. Let me say to that public, which has shown some interest in those glimpses which I have occasionally given them of the thoughts and actions of a very remarkable man, that they are not to blame me if I have not shared my knowledge with them, for I should have considered it my first duty to do so, had I not been barred by a positive prohibition from his own lips, which was only withdrawn upon the third of last month."

]

Listing 2.3 Zero-Shot Classification: Text Samples for Crime, Fantasy, and Romantic Novels (Source: 02_PreTrainedNetworks\50_zero_shot.py)

The idea of zero-shot classification is to provide some candidate labels and let the model decide the probability of each book's sample text belonging to each of the candidates.

We run the model by running our object `classifier` and passing the documents and the `candidate_labels`. The result is stored in the object `res`, as follows:

```
candidate_labels=["romance", "fantasy", "crime"]
#%% classify documents
res = classifier(documents, candidate_labels = candidate_labels)
```

The result will be visualized with `pandas` as a bar plot for the opening lines of Sherlock Holmes, as shown in [Figure 2.5](#), with the following code:

```
### visualize results
pos = 2
pd.DataFrame(res[pos]).plot.bar(x='labels', y='scores', title=titles[pos])
```




Figure 2.5 Zero-Shot Classification Result for Sherlock Holmes

The classifier correctly assigned the label “crime” to the text snippet, with a very high probability. This correct classification is not as obvious as you might think. We don’t pass the complete book, which is clearly a crime novel, only a small fragment.

In our next task, we’ll fill in empty spots in a sentence. This task is comparable to auto-completion in which the next word in a sequence is proposed. This technique is called a fill-mask task.

2.5 Coding: Fill-Mask

A fill-mask task is a type of masked language modeling in NLP, where a model predicts a missing word or token in a given sentence. This task is often used to understand context and word relationships by filling in a blank (or “mask”) in a sentence.

For instance, in the sentence “AI is transforming the [MASK] industry,” a fill-mask model might predict words like “technology,” “healthcare,” or “finance,” based on the context.

BERT (bidirectional encoder representations from transformers) and similar models are commonly used for this task, as they are pretrained on masked language modeling objectives, allowing them to predict missing words with a deep understanding of sentence structure and context.

The implementation is extremely simple. You can find the example at *02_PreTrainedNetworks\80_fill_mask.py*. All you need is one package and its pipeline function, as follows:

```
### packages
from transformers import pipeline
```

First, define the task `fill-mask` and choose the model `bert-base-uncased`, as shown in [Listing 2.4](#). After you create a model instance, you can pass a query to it. The query must have `[MASK]` in it, which is the place where the blank needs to be filled.

```
unmasker = pipeline(task='fill-mask', model='bert-base-uncased')
unmasker("I am a [MASK] model.")
```

```
[{'score': 0.26198092103004456,
  'token': 2535,
  'token_str': 'role',
  'sequence': 'i am a role model.'},
 {'score': 0.1216965764760971,
```

```

    'token': 4827,
    'token_str': 'fashion',
    'sequence': 'i am a fashion model.'},
{'score': 0.03554196283221245,
 'token': 2944,
 'token_str': 'model',
 'sequence': 'i am a model model.'},
{'score': 0.033298179507255554,
 'token': 3565,
 'token_str': 'super',
 'sequence': 'i am a super model.'},
{'score': 0.033055439591407776,
 'token': 2449,
 'token_str': 'business',
 'sequence': 'i am a business model.'}]

```

Listing 2.4 Fill-Mask Inference

The model chooses reasonable words by assuming we are interested in a role model, fashion model, or supermodel.

With LLMs, as you'll see later, the next words are predicted in a similar way. LLMs are often used for answering questions. In the next section, we visit a task that works similarly—question answering.

2.6 Coding: Question Answering

Question answering models are tailored to address specific, well-defined tasks. These question answering models allow users to apply models that can accurately interpret and answer questions within a focused domain. Unlike LLMs, which attempt to answer questions across a wide range of topics, narrow AI question answering models excel at providing high-precision answers in specialized areas—such as customer support, technical troubleshooting, medical diagnosis, and legal research.

Question answering is essential because it enables efficient information retrieval, transforming unstructured text data into actionable insights.

For instance, in customer service, question answering models can answer common queries automatically, reducing the need for human intervention and improving response times. In sectors like healthcare or finance, narrow question answering models help professionals quickly access relevant information, allowing them to make informed decisions without extensive manual searching.

By optimizing access to knowledge within specific fields, question answering models in narrow AI enhance productivity, support data-driven decision-making, and contribute to a more streamlined and responsive user experience.

Let's look at some examples. You can find the following code at *02_PreTrainedNetworks\70_qa.py*:

```
### packages
from transformers import AutoModelForQuestionAnswering, AutoTokenizer, pipeline
from pprint import pprint
```

You can use the model `deepset/roberta-base-squad2`. This model is based on SQuAD 2.0 (Stanford Question Answering Dataset 2.0), designed to challenge question answering models by introducing unanswerable questions. The dataset contains over 150,000 questions. Each question is paired with a context paragraph, usually drawn from a Wikipedia article. In many cases, the context includes the answer to the question, but in other cases, it does not. Here's our code:

```
### constants
MODEL = "deepset/roberta-base-squad2"
```

The model relies solely on the context provided to find an answer and answer questions by locating information within a specific context. The model does not have external knowledge. It can only see and analyze the context provided to it.

Thus, you'll need to pass the input in JavaScript Object Notation (JSON) with the keys `question` and `context`. The answer will be based on the context. [Listing 2.5](#) shows the model setup.

```
# Get predictions
nlp = pipeline(task='question-answering', model=MODEL, tokenizer=MODEL)
QA_input = {
    'question': 'What are the benefits of remote work?',
    'context': 'Remote work allows employees to work from anywhere, providing
flexibility and a better work-life balance. It reduces commuting time, lowers
operational costs for companies, and can increase productivity for self-motivated
workers.'
}
```

Listing 2.5 Fill-Mask: Model Setup (Source: `02_PreTrainedNetworks\70_qa.py`)

Now, run the model and check the answer, as shown in [Listing 2.6](#).

```
res = nlp(QA_input)
pprint(res)

{'answer': 'flexibility and a better work-life balance',
 'end': 104,
```

```
'score': 0.6486219763755798,  
'start': 62}
```

Listing 2.6 Fill-Mask: inference (Source: 02_PreTrainedNetworks\70_qa.py)

In our next task, you'll learn how to detect and extract entities from sentences.

OceanofPDF.com

2.7 Coding: Named Entity Recognition

Named entity recognition (NER) is an NLP technique that identifies and categorizes specific elements in text, such as names of people, organizations, locations, dates, and quantities. By recognizing these entities, NER helps transform unstructured text into structured data, enabling more efficient data analysis and search functionalities. These capabilities can be helpful, for instance, when you want to anonymize texts by removing all persons, places, and dates.

In generative artificial intelligence (generative AI), NER enhances LLMs by enabling them to understand context around key concepts and generate responses that are more accurate and contextually relevant. Advanced NER models leverage deep learning to recognize entities even in complex and nuanced contexts, improving their utility in applications ranging from customer support to data extraction.

Using a specific NER model, you can perform highly specialized tasks that require accurate identification of specific entities. NER can automatically identify and extract names, dates, places, and other relevant entities and thus save time in manual data entry or review processes. NER models can assist in identifying customer names, order numbers, product mentions, and other entities in support interactions, helping prioritize and direct queries more efficiently.

This discussion should give you some ideas about how you might utilize NER, but now, let's utilize it. First, select the model `dslim/bert-base-NER`. This model allows you to detect different entities like persons, organizations, locations, or miscellaneous entities. The model card (<https://huggingface.co/dslim/bert-base->

NER) provides all necessary details on how to use the model. You can find the script at *02_PreTrained Networks\60_ner.py*.

First, load the required packages. By now, you're familiar with `pipeline`, but in this case, you also need to load `AutoModelForTokenClassification` and the tokenizer with `AutoTokenizer`, as follows:

```
from transformers import AutoTokenizer, AutoModelForTokenClassification
from transformers import pipeline
from pprint import pprint
```

We'll load the tokenizer and model in the following code. Use the `from_pretrained` method and pass the model's name to `AutoTokenizer` and `AutoModelForTokenClassification`. You'll learn more about tokenization in [Chapter 3](#).

```
### model and tokenizer
tokenizer = AutoTokenizer.from_pretrained("dslim/bert-base-NER")
model = AutoModelForTokenClassification.from_pretrained("dslim/bert-base-NER")
```

Now, you can set up pipeline and pass the task, the model, and the tokenizer during instantiation of our object `nlp`, as follows:

```
### pipeline
nlp = pipeline(task="ner", model=model, tokenizer=tokenizer)
```

Great! You've created all the objects for running the model. Let's work on a single sentence and try to detect the different entities, as shown in [Listing 2.7](#).

```
example = "My name is Bert and I live in Hamburg"
```

```
ner_results = nlp(example)
pprint(ner_results)
```

```
[{'end': 15,
  'entity': 'B-PER',
  'index': 4,
  'score': 0.9952625,
  'start': 11,
  'word': 'Bert'},
 {'end': 34,
```

```
'entity': 'B-LOC',  
'index': 9,  
'score': 0.9977075,  
'start': 27,  
'word': 'Hamburg']}]
```

Listing 2.7 Named Entity Recognition: Inference (Source: 02_PreTrainedNetworks\60_ner.py)

The model successfully detected the name “Bert” as a person (B-PER) and the city where he lives as the location (B-LOC). You can find the description of the different entity names in the model card.

A popular class of pretrained models deals with the creation of images based on text prompts. A whole ecosystem on text-to-image models and providers has evolved. In the next section, you’ll learn how to use such models locally.

OceanofPDF.com

2.8 Coding: Text-to-Image

You can select from a wide range of text-to-image models, but keep in mind that these are quite resource intensive. If you don't own a good GPU, we recommend using a rather small text-to-image model. For that reason, we recommend using `amused/amused-256`, available at <https://huggingface.co/amused/amused-256>.

The following code creates an image of 256×256 pixels. Running on a CPU, the creation process took half an hour. The model creates an image of a dog and stores that image in the current working folder. [Listing 2.8](#) shows the code required to use this text-to-image model. [Figure 2.6](#) shows the result.

```
import torch
from diffusers import AmusedPipeline
###
pipe = AmusedPipeline.from_pretrained(
    "amused/amused-256", variant="fp16", torch_dtype=torch.float16
)
pipe.vqvae.to(torch.float32) # vqvae is producing nans in fp16
###
# pipe = pipe.to("cuda")

prompt = "dog"
```

```
image = pipe(prompt, generator=torch.Generator().manual_seed(8)).images[0]
image.save('text2image_256.png')
```

Listing 2.8 Text-to-Image Code (Source: 02_PreTrainedNetworks\30_text_to_image.py)



Figure 2.6 Sample Output from the Text2Image Model

Notice we've used a new package called `diffusers`, which is built around *stable diffusion models*. This model family can create images based on text prompts. The `pipe` is provided with the user prompt and returns an image, that we save with `image.save`, as shown in [Figure 2.6](#). Most of the code is directly taken from the model card because the developers make our lives easier by showing us how to use their developments.

Admittedly, the output is not the best image, and the algorithm struggled to get the right eye correct. But otherwise, it looks quite decently like a dog.

Now again, it's time for you to shine and get some practice. You'll create some music in the next exercise.

2.9 Exercise: Text-to-Audio

Generating audio like voices or music is quite impressive. With a text description of a music track, you can create it easily. You can rely on an online service like <https://suno.com/>, but in this section, we want to run a model locally on our systems.

2.9.1 Task

In this exercise, find and select a model to create a music track and then run that model locally. If you don't own a powerful GPU, we advise searching for a model with "small" in its name. Try to solve this problem independently before checking out our solution.

2.9.2 Solution

For our solution, we'll use a model from Facebook called `facebook/musicgen-small` (<https://huggingface.co/facebook/musicgen-small>).

You can find the code at `02_PreTrainedNetworks\40_text_to_audio.py`. It reflects the model card description. In this case, you'll create a `pipeline` instance. By running this instance based on a text prompt, you'll get some music, which finally needs to be converted into a WAV file using `scipy`. [Listing 2.9](#) shows the code for creating some music based on a text-to-audio model.

```
### packages
from transformers import pipeline
import scipy
```

```

#%% model selection
task = "text-to-audio"
model = "facebook/musicgen-small"

# %%
synthesiser = pipeline(task, model)

music = synthesiser("lo-fi music with a soothing melody", forward_params=
{"do_sample": True})

scipy.io.wavfile.write("musicgen_out.wav", rate=music["sampling_rate"],
data=music["audio"])

```

Listing 2.9 Text-to-Audio Source Code (Source: 02_PreTrainedNetworks\40_text_to_audio.py)

This task follows the already known pattern of pipeline in combination with a task and a model. The remaining pieces are to call the synthesizer object with a prompt that describes the music and then to save the music with scipy as a WAV file.

As a result, you'll find the WAV file in your working folder. Listen to it, but we must say, we think it worked well. But you need to be patient: Creating music might take a few minutes.

We've come close to the end of this chapter, and you're now equipped with the knowledge to tackle a larger task in a capstone project.

OceanofPDF.com

2.10 Capstone Project: Customer Feedback Analysis

In this capstone project, you'll solve a task independently using multiple different pretrained models.

Create a system that analyzes customer feedback in the form of product reviews. The system should understand a customer's sentiment and categorize the feedback into different categories like "defect," "delivery," or "interface."

Consider that the customer reviews can be multilingual. The output should be a dictionary with the feedback and their corresponding sentiment, and categories. You can find the starting code at *02_PreTrainedNetworks\90_capstone_start.py*.

Your tasks are as follows:

- Import necessary packages
- Search for a sentiment analysis model that returns star-ratings for multiple languages
- Implement the sentiment analysis model
- Search for a zero-shot classification model
- Implement the sentiment analysis model
- Bundle the logic in a function
- Test the function with the given `feedback` object
- The result should be a dictionary with each feedback, its sentiment, and the most likely category it belongs to

[Figure 2.7](#) shows the workflow of the code. The input will be passed to the two different models, and the outputs will be combined with the feedback. In a real-life scenario, subsequent pipelines might be triggered. On the right, you'll see some example input and output texts.




Figure 2.7 Capstone Project Workflow

The starter script is shown in [Listing 2.11](#). You can find it under `02_PreTrainedNetworks\90_capstone_start.py`. Fill in the gaps with your own code. Good luck and have fun!

```

#%% packages
# TODO: import the necessary packages

#%% data
feedback = [
    "I recently bought the EcoSmart Kettle, and while I love its design, the heating element broke after just two weeks. Customer service was friendly, but I had to wait over a week for a response. It's frustrating, especially given the high price I paid.",
    "Die Lieferung war super schnell, und die Verpackung war großartig! Die Galaxy Wireless Headphones kamen in perfektem Zustand an. Ich benutze sie jetzt seit einer Woche, und die Klangqualität ist erstaunlich. Vielen Dank für ein tolles Einkaufserlebnis!",
    "Je ne suis pas satisfait de la dernière mise à jour de l'application EasyHome. L'interface est devenue encombrée et le chargement des pages prend plus de temps. J'utilise cette application quotidiennement et cela affecte ma productivité. J'espère que ces problèmes seront bientôt résolus."
]

# %% function
# TODO: define the function process_feedback

```

```
### Test
# TODO: test the function process_feedback
```

Listing 2.10 Capstone Project Starter Script (Source: 02_PreTrainedNetworks\90_capstone_start.py)

How did it work out? Could you find models for these tasks? Was it difficult to implement it?

For comparison, let's look at our solution, which can also be found at 02_PreTrainedNetworks\91_capstone_end.py.

For the zero-shot model, we selected facebook/bart-large-mnli, and for the sentiment analysis, nlptown/bert-base-multilingual-uncased-sentiment, as shown in [Listing 2.11](#). Choosing another model is completely fine!

```
### packages
from transformers import pipeline, AutoTokenizer, AutoModelForTokenClassification
import pandas as pd
from typing import List

### data
feedback = [
    "I recently bought the EcoSmart Kettle, and while I love its design, the heating element broke after just two weeks. Customer service was friendly, but I had to wait over a week for a response. It's frustrating, especially given the high price I paid.",
    "Die Lieferung war super schnell, und die Verpackung war großartig! Die Galaxy Wireless Headphones kamen in perfektem Zustand an. Ich benutze sie jetzt seit einer Woche, und die Klangqualität ist erstaunlich. Vielen Dank für ein tolles Einkaufserlebnis!",
    "Je ne suis pas satisfait de la dernière mise à jour de l'application EasyHome. L'interface est devenue encombrée et le chargement des pages prend plus de temps. J'utilise cette application quotidiennement et cela affecte ma productivité. J'espère que ces problèmes seront bientôt résolus."
]

# %% Function
def process_feedback(feedback: List[str]) -> dict[str, List[str]]:
    """
    Process the feedback and return a DataFrame with the sentiment and the most likely label.
    Input:
    feedback: List[str]
    Output:
    pd.DataFrame
    """
```

```

"""
CANDIDATES = ['defect', 'delivery', 'interface']
ZERO_SHOT_MODEL = "facebook/bart-large-mnli"
SENTIMENT_MODEL = "nlptown/bert-base-multilingual-uncased-sentiment"
# initialize the classifiers
zero_shot_classifier = pipeline(task="zero-shot-classification",
model=ZERO_SHOT_MODEL)
sentiment_classifier = pipeline(task="text-classification",
model=SENTIMENT_MODEL)

zero_shot_res = zero_shot_classifier(feedback,
candidate_labels = CANDIDATES)
sentiment_res = sentiment_classifier(feedback)
sentiment_labels = [res['label'] for res in sentiment_res]
most_likely_labels = [res['labels'][0] for res in zero_shot_res]
res = {'feedback': feedback,
'sentiment': sentiment_labels,
'label': most_likely_labels}
return res

### Test
process_feedback(feedback)

{'feedback': ["I recently bought the EcoSmart Kettle, and while I love its design,
the heating element broke after just two weeks. Customer service was friendly, but I
had to wait over a week for a response. It's frustrating, especially given the high
price I paid.",
'Die Lieferung war super schnell, und die Verpackung war großartig! Die Galaxy
Wireless Headphones kamen in perfektem Zustand an. Ich benutze sie jetzt seit einer
Woche, und die Klangqualität ist erstaunlich. Vielen Dank für ein tolles
Einkaufserlebnis!',
"Je ne suis pas satisfait de la dernière mise à jour de l'application EasyHome.
L'interface est devenue encombrée et le chargement des pages prend plus de temps.
J'utilise cette application quotidiennement et cela affecte ma productivité.
J'espère que ces problèmes seront bientôt résolus."],
'sentiment': ['2 stars', '5 stars', '2 stars'],
'label': ['defect', 'delivery', 'interface']}

```

Listing 2.11 Capstone Project Completed Script (Source: 02_PreTrainedNetworks\91_capstone_end.py)

2.11 Summary

In this chapter, you learned all about pretrained models, which are specialized, task-specific models.

This chapter began by looking back to a time before LLMs, when smaller models were developed for individual tasks using supervised learning techniques. These models excelled at specific tasks, were compact, and could often be run locally. In contrast, as you discovered, LLMs are built to handle a wide range of tasks using extensive, often unstructured data from various sources. While LLMs offer versatility, they come with higher costs and greater latency, making them less ideal for certain applications. We explored why it's essential to choose between smaller, specific models and LLMs depending on your requirements and the resources available.

In this chapter, we introduced you to Hugging Face, a popular platform hosting over a million models tailored to specific ML tasks. You learned how to navigate the platform's model selection tools to find models suited to your needs, ensuring efficiency without incurring unnecessary costs.

Through several practical coding examples, you worked on tasks for text summarization, translation, zero-shot classification, fill-mask tasks, question answering, and NER using Hugging Face's model selection and pipeline features.

We explored how to load and summarize an article from Arxiv, handling limitations like model context window size. For translation, you followed a process to convert a well-known English phrase into Japanese, using tools to filter models based on language requirements.

This chapter also included a text-to-image example, where you observed the computational demands of these models, and we urged you to consider using smaller models if resources are limited. Additionally, a text-to-audio exercise encouraged you to try generating audio, with guidance on selecting an appropriate model and converting the output into a WAV file.

By the end of this chapter, you now know how to select and use pretrained models. You learned to balance model choice, task suitability, and resource efficiency.

With this knowledge, you're now prepared to move on into more advanced discussions on LLMs and their applications.

OceanofPDF.com

3 Large Language Models

The question of whether a computer can think is no more interesting than the question of whether a submarine can swim.

—Edsger W. Dijkstra

Dijkstra, one of the most influential computer scientists in the 20th century, was not a fan of anthropomorphizing computers. At the time of his quote, at the end of the 1970s or early 1980s, computers were not even remotely as capable as today's computers. Today, we witness the emergence of systems with human (and increasingly superhuman) capabilities. Still, large language models (LLMs) don't think in the way we humans do, but they exhibit capabilities that often make them seem as if they do.

In this chapter, we make ourselves familiar with LLMs, which are the most prominent members of the generative artificial intelligence (generative AI) family.

We start with a brief history of language models in [Section 3.1](#). Then, you'll dive right in and learn how to use LLMs with Python in [Section 3.2](#). From there, we'll expand and dive more deeply in [Section 3.3](#), where you'll learn which model parameters steer a model's outputs.

Since an exponentially increasing number of LLMs are available, selecting which one to use has become difficult. We deal with the model selection in [Section 3.4](#). When interacting with LLMs, you

must speak their language in a certain way. In other words, you must provide information to them via messages. You'll learn about different types of messages and their intricacies in [Section 3.5](#).

As a preliminary step to the concept of chains (the focus of [Section 3.7](#)), we'll introduce you to the concept of prompt templates, which typically represent the first element in a chain. You'll learn about prompt templates in [Section 3.6](#).

We'll mostly work with a Python framework called LangChain that is centered around chains, a concept that we'll study in [Section 3.7](#). Chains allow you to develop highly complex systems, while keeping the complexity manageable.

You must ensure that any LLM or complex AI system you operate is secured against malicious actors and also provide reliable, correct, and on-topic responses. We cover these LLM safety and security topics in [Section 3.8](#).

If you're not satisfied with the responses from “your” LLM, you might consider some improvements that we describe in [Section 3.9](#).

The field of LLMs is new and developing at an enormous pace. Recent trends like reasoning models, small language models (SLMs), or test-time compute are covered in [Section 3.10](#).

But let's explore first how we got to where we are today by studying the history of language models.

3.1 Brief History of Language Models

The first chatbot, Eliza, was developed in the 1960s. It was quite limited in its capabilities and used pattern recognition—basically pretending “real” conversations. The user's input was converted into

a question, and a rule-based system provided some generic answers. Still, Eliza was quite successful and led to more research in the field of natural language processing (NLP). [Figure 3.1](#) shows a screenshot of its user interface.

```
Welcome to

EEEEEE LL      IIII ZZZZZZZZ AAAAA
EE      LL      II      ZZ  AA  AA
EEEEEE LL      II      ZZ  AAAAAAA
EE      LL      II      ZZ  AA  AA
EEEEEE LLLLLL IIII ZZZZZZZZ AA  AA

Eliza is a mock Rogerian psychotherapist.
The original program was described by Joseph Weizenbaum in 1966.
This implementation by Norbert Landsteiner 2005.

ELIZA: Is something troubling you ?
YOU:   Men are all alike.
ELIZA: What is the connection, do you suppose ?
YOU:   They're always bugging us about something or other.
ELIZA: Can you think of a specific example ?
YOU:   Well, my boyfriend made me come here.
ELIZA: Is it important to you that your boyfriend made you come here ?
YOU:   He says I'm depressed much of the time.
ELIZA: I am sorry to hear that you are depressed.
YOU:   It's true. I am unhappy.
ELIZA: Can you explain what made you unhappy ?
YOU:   █
```

Figure 3.1 ELIZA User Interface (Source: https://upload.wikimedia.org/wikipedia/commons/4/4e/ELIZA_conversation.jpg)

Decades passed, until in the 1990s, statistical models became the norm and replaced the rule-based systems. Most notable in this category are *N-gram models*. These models could analyze the likelihood of words by studying their local neighbors. In this way, they could gain some understanding of the context surrounding a word.

With the onset of deep learning, model architectures like convolutional neural networks (CNNs), or recurrent neural networks (RNN) provided better results. They could “memorize” short-term relationships within sequences of text. But they struggled to comprehend long-term relationships.

Other models like long short-term memory (LSTM) or gated recurrent units (GRUs) are also important milestones.

Word embeddings were developed in the 2000s to allow the representation of a word in the form of a dense numerical vector. An important model is Word2Vec, which was developed in 2013 by Google. This concept of word embeddings, as well as its extension to sentence embeddings, is the foundation of the information that you'll encounter throughout this book.

The 2017 paper "Attention is All You Need" (<https://arxiv.org/abs/1706.03762>), published by Google developers, started a revolution. Now, LLMs, as well as many computer vision models, are based on the groundbreaking ideas expressed in this paper. One limitation so far was that text sequences could not be processed in parallel. Now, parallel processing has become possible. This paper also included the concept of attention and self-attention, which allowed a model to focus not just on nearby words, but also distant words. With this capability, long-range dependencies can be captured effectively.

Now, let's not linger too long on the past and instead turn to using LLMs programmatically.

OceanofPDF.com

3.2 Simple Use of LLMs via Python

We'll start and learn how to use the most popular LLM: OpenAI's GPT. This immensely powerful model for which you need to pay. Since there are also extremely capable LLMs provided as open-source, we'll show you how you can use LLMs for free via Groq. You'll see that the way to interact with the models is very streamlined, thanks to the Python framework LangChain. After these two examples you'll be able to connect to any other LLM provider.

3.2.1 Coding: Using OpenAI

To use OpenAI you need to have an application programming interface (API) key. These models are not free, so you'll be charged for your usage. Prices are dropping permanently over time, and you can check the current pricing via <https://openai.com/api/pricing/>.

First, you must set up an API key. To use a web service, you usually only need a username and password to access the service. To use a service programmatically, you require an API key, which makes an API key something like a combination of username and password.

If you don't have one yet, get an API key by following these steps:

1. Head over to <https://platform.openai.com/>.
2. Create an account.
3. Activate billing and load some money onto your account.
4. Navigate to API keys and create a new API key. The name you specify in the web frontend is irrelevant, you only need the key. Copy the key into the clipboard.

5. Paste this key into a file called the `.env` file. The key should look like `sk-proj...`

A best practice is to separate the code from the credentials. Thus, store the API key in a separate file. A common approach is to store this information in a file called `.env` and place it in the working folder. In that file, you could store the API key and possibly many more keys if needed. [Listing 3.1](#) shows what an environment file should look like.

```
OPENAI_API_KEY = sk-proj...
```

Listing 3.1 Sample `.env` File Content

API keys are treated as *environment variables*, which are typically variables that are used by your operating system. Our variable has the name `OPENAI_API_KEY` and its value must be defined to the right of the equal sign. Make sure you use the same key name in the coding script.

If you hesitate to provide your banking details, skip this section and move to [Section 3.2.2](#) to work directly with Groq, which provides access to models for free. Don't confuse Groq with Grok. Groq is an AI startup focusing on creating chips for fast inference with LLMs, while Grok is an LLM launched as an initiative by Elon Musk.

Let's now start coding. You can find it in the material folder under `03_LLMs/10_model_chat.py`. A best practice is to place all the required packages and functions at the beginning of the file. Let's go through what we need next.

The `os` package is required for fetching and loading the environment variables. All major model providers offer packages for integration into LangChain. Thus, as shown in [Listing 3.2](#), use `langchain_openai`. The package `dotenv` is required for working with

the environment variables file. Its function `load_dotenv()` loads the content of the `.env` file and provides the content as environment variables.

```
### packages
import os
from langchain_openai import ChatOpenAI
from dotenv import load_dotenv
load_dotenv('.env')
```

Listing 3.2 Using OpenAI: Required Packages (Source: 03_LLMs/10_model_chat.py)

Check that the API key is available by running `os.getenv('OPENAI_API_KEY')`. As a result, you should see the API key printed on the screen.

Now we create an instance of the model we'll use by using `ChatOpenAI` class. This requires a model name. We choose GPT 4o Mini here. Another important parameter is the temperature. This parameter controls the creativity of the model. And you need to pass the API key to authenticate and enabling OpenAI to charge you based on your usage.

```
MODEL_NAME = 'gpt-4o-mini'
model = ChatOpenAI(model_name=MODEL_NAME,
                    temperature=0.5, # controls creativity
                    api_key=os.getenv('OPENAI_API_KEY'))
```

The `model` object's `invoke()` method is particularly important, as shown in [Listing 3.3](#). This method allows you to run the model based on specified parameters. In our first example, we asked the model to provide information on “What is LangChain?” The result is stored in an object of type `AIMessage`. You can access the information received from the model call by looking at the output of its `dict()` method.

```
res = model.invoke("What is a LangChain?")
res.dict()
{'content': 'LangChain is ...',
 'additional_kwargs': {},
 'response_metadata': {'token_usage': {'completion_tokens': 312,
```

```

    'prompt_tokens': 13,
    'total_tokens': 325},
    'model_name': 'gpt-4o-mini-2024-07-18',
    'system_fingerprint': 'fp_f33667828e',
    'finish_reason': 'stop',
    'logprobs': None},
    'type': 'ai',
    'name': None,
    'id': 'run-929cb722-e48e-457b-859c-754a5d272c6d-0',
    'example': False,
    'tool_calls': [],
    'invalid_tool_calls': [],
    'usage_metadata': {'input_tokens': 13,
    'output_tokens': 312,
    'total_tokens': 325}}

```

Listing 3.3 Using OpenAI: Model Invocation (Source: 03_LLMs/10_model_chat.py)

Plenty of information is coming back from the model. Let's start with the most important part: `content`. This property holds the actual model output prompt. Out of the other properties, we just want to mention `response_metadata`, which holds information on token usage. You're charged in terms of input tokens and output tokens. In this property, you can determine how many tokens were used in the request.

Familiarize yourself with the different models from OpenAI's model family by studying the model overview at <https://platform.openai.com/docs/models/overview>. Some key features include the following:

- OpenAI created a model family (<https://platform.openai.com/docs/models>) consisting of several models suitable for different tasks.
- Language models like the GPT family (e.g., GPT-4o) can process text, and some can also work with images.
- Text-to-image generation: DALL-E is a model that can generate and edit images.

- Text-to-speech (TTS): Several models can convert text to natural, spoken audio.
- Speech-to-text: With Whisper, you can convert audio recordings into text.
- Text embeddings: Embeddings are numerical representations of text. Such embeddings are the cornerstone of NLP. We'll revisit embeddings in detail in [Chapter 5](#).

You're not limited to the OpenAI model family. You can work with many other LLMs. Next, you'll learn how to work with open-source LLMs that you can run for free via Groq.

3.2.2 Coding: Using Groq

Groq is a company that develops AI hardware to make extremely fast inferences. For developers, they provide access to LLMs, especially open-source LLMs. You can use the service for free, but you need to be authenticated via an API key. So, the first step is to head over to <https://console.groq.com/>, set up an account, and create an API key to use in your code.

Copy this API key and store it in a file called `.env` inside the working folder `03_LLM/`. The file's content is shown in [Listing 3.4](#).

```
GROQ_API_KEY = gsk_...
```

Listing 3.4 Snippet from the `.env` File

The script that you can find at `03_LLMs/10_model_chat_groq.py` starts by loading relevant packages. The main package is `langchain_groq`, which is the interface to use the models from the Groq model family, as shown in [Listing 3.5](#). Use the `os` and `dotenv`

packages to set up and get environment variables, which hold the Groq API key.

```
### packages
import os
from langchain_groq import ChatGroq
from dotenv import load_dotenv
load_dotenv('.env')
```

Listing 3.5 Using Groq: Required Packages (Source: 03_LLMs/10_model_chat_groq.py)

Now, you need to select a model. Details on specific models can be found in the overview about Groq models available at <https://console.groq.com/docs/models>.

In our example, we chose a model from Llama family: an open-source model, more specifically, an open-weight model. In other words, the model is provided to the public to be used for free, but not all details about its datasets and training processes are made publicly available.

After deciding on a model, create an instance of `ChatGroq` class. In this instantiation, pass the model's name as parameter. You must also pass the API key you created earlier. These two parameters are mandatory. Among many other available parameters, in our example, we've only set the temperature parameter, which controls the creativity of the model. You'll learn more about model parameters in [Section 3.3](#). With that step, you have everything in place to interact with the LLM.

```
MODEL_NAME = 'llama-3.3-70b-versatile'
model = ChatGroq(model_name=MODEL_NAME,
                  temperature=0.5, # controls creativity
                  api_key=os.getenv('GROQ_API_KEY'))
```

Now, we can ask the model “What is Hugging Face?” via the `invoke()` method, as follows:

```
# %% Run the model
res = model.invoke("What is a Hugging Face?")
```

With the `dict()` method, you can see an overview of the model output, as shown in [Listing 3.6](#).

```
# %% find out what is in the result
res.dict()

{'content': 'Hugging Face is a popular open-source library and platform for natural
language processing (NLP) and machine learning (ML) ...',
 'additional_kwargs': {},
 'response_metadata': {'token_usage': {'completion_tokens': 328,
    'prompt_tokens': 42,
    'total_tokens': 370,
    'completion_time': 1.312,
    'prompt_time': 0.010021397,
    'queue_time': 0.0048266910000000001,
    'total_time': 1.3220213969999999},
    'model_name': 'llama-3.1-70b-versatile',
    'system_fingerprint': 'fp_b6828be2c9',
    'finish_reason': 'stop',
    'logprobs': None},
 'type': 'ai',
 'name': None,
 'id': 'run-84ad52ab-f7df-4245-ab42-5281e48cac69-0',
 'example': False,
 'tool_calls': [],
 'invalid_tool_calls': [],
 'usage_metadata': {'input_tokens': 42,
    'output_tokens': 328,
    'total_tokens': 370}}
```

Listing 3.6 Using Groq: Model Output (Source: 03_LLMs/10_model_chat_groq.py)

The most important output is again the content. Typically, you directly fetch the content via its property, as shown in [Listing 3.7](#).

```
# %% only print content
print(res.content)
```

Hugging Face is a popular open-source library and platform for natural language processing (NLP) and machine learning (ML) tasks. It was founded in 2016 by Julien Chaumond, Clement Delangue, and Thomas Wolf. Hugging Face is known for its Transformers library, which provides pre-trained models and a simple interface for using and fine-tuning transformer-based models for various NLP tasks, such as text

classification, language translation, question answering, and more.

...

Listing 3.7 Using Groq: Model Output Content (Source: 03_LLMs/10_model_chat_groq.py)

You've successfully invoked a model provided by Groq. You can find all available Groq models at <https://console.groq.com/docs/models>.

For each model, the model ID is provided. This string must be used in your script. Furthermore, the developer is shown as well as a limiting factor. For LLMs, the limiting factor is the context window, and the maximum number of tokens is shown. More on that limitation is provided in the next info box. If you want to dive more deeply into the model, check out the model card, and you're forwarded to the developer page of the model.

Some specialized models include Distil-Whisper models, which provide a speech-to-text model. You can upload an MP3 file (up to a certain size) and receive the corresponding spoken text in the file.

Most available models are LLMs, and generally, all are open-source or open-weight models. Prominent models include members of the Llama family from Meta, Gemma models from Google, or Mixtral from Mistral.

In the model overview, you might have seen the context window as one of the most important parameters. The context window refers to the maximum number of input tokens that a model can process at once. This value is important because the model's ability to generate relevant output depends on the information it can keep and use in a single prompt. Each LLM has a fixed limit on how many tokens it can handle at once in its context window.

In this context, keep in mind that an LLM breaks the input text down into smaller units, which are called "tokens." Such a token might be a

word, or just a part of a word, or even some punctuation.

A large context window allows the model to process more information, which improves its ability to understand longer texts, or “memorize” long conversations. As the context window size increases, the computational resources required to process the text increase. Also, the latency of the model increases—it takes more time for the model to provide an answer.

If the context window is exceeded, the model either “forgets” older tokens, or the LLM call fails, depending on the package implementation.

We’ll face this problem in [Chapter 5](#) and [Chapter 6](#) on vector databases and retrieval-augmented generation (RAG), respectively, since these techniques help to establish some kind of “external knowledge” to an LLM. These approaches offer the flexibility to feed only relevant knowledge to the LLM, thus keeping within the constraints of a given context window.

3.2.3 Coding: Large Multimodal Models

In this book, you’ll mostly work with text inputs and outputs and use traditional LLMs. But demand has increased for models to understand and interact with more complex and varied forms of information. Consequently, multimodal LLMs were developed. These models are trained to understand and generate multiple types, or modalities, of input and output format, typically besides text—images, audio, and video.

Unlike traditional LLMs that operate in a single modality (text), these multimodal models can process information across several formats. For this capability, they combine advances in NLP with innovations in

computer vision and audio processing. As a result, these models can perform the following tasks:

- Analyze and describe images in text form
- Generate images based on textual descriptions
- Transcribe audio inputs
- Interpret audio inputs and provide responses based on them

Let's use some multi-modal models to learn how to work with them. You'll learn how to pass an image as an input to a model. Then, you'll interact with the LMM to determine if it understands what it "sees" in the image.

The flowchart shown in [Figure 3.2](#) illustrates the process of training a deep neural network. In our example, we'll pass an image to the multimodal model.




Figure 3.2 Deep Neural Network Training Process

You can find the code for this script at *03_LLMs\60_multimodal.py*. This code is mostly based on the documentation from Groq, available at <https://console.groq.com/docs/vision>.

As shown in [Listing 3.8](#), start by loading the required packages.

```
### packages
from groq import Groq
```



```
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
import base64
```

Listing 3.8 LMM: Required Packages (Source: 03_LLMs\60_multimodal.py)

A good practice is to define the constants at the beginning of the script. As in the following example, define which model to choose, where the image is located, and the user prompt:

```
### constants
MODEL = "llama-3.2-90b-vision-preview"
IMAGE_PATH = "sample_image.png"
USER_PROMPT = "What is shown in this image? Answer in one sentence."
```

Since we're working with a local image that must be sent to the Groq API, the image must be loaded and converted into a format that can be sent as text string. For this functionality, as shown in [Listing 3.9](#), you can define a function called `encode_image`. This function loads the image and directly converts it to base64 format. Base64 is a binary-to-text encoding scheme that converts binary data into an ASCII string format.

```
### Function to encode the image
def encode_image(image_path):
    with open(image_path, "rb") as image_file:
        return base64.b64encode(image_file.read()).decode('utf-8')

base64_image = encode_image(IMAGE_PATH)
```

Listing 3.9 LMM: Encode Image Function (Source: 03_LLMs\60_multimodal.py)

Now, set up a Groq instance, as shown in [Listing 3.10](#), using the native implementation of groq package, so the chat request has a different format compared to the LangChain interaction with models. But you can recognize many pieces, for example, the messages. In the `messages` object, you can define a user message. In that object, we pass a dictionary with text content (the user prompt) as well as the image content (the image we want to interact with).

```

#%% Getting the base64 string
client = Groq()

chat_completion = client.chat.completions.create(
    messages=[
        {
            "role": "user",
            "content": [
                {"type": "text", "text": USER_PROMPT},
                {
                    "type": "image_url",
                    "image_url": {
                        "url": f"data:image/jpeg;base64,{base64_image}"
                    }
                },
            ],
        }
    ],
    model=MODEL,
)

```

Listing 3.10 LMM: Inference (Source: 03_LLMs\60_multimodal.py)

The response is located in the object `chat_completion`, which can print to screen, as shown in [Listing 3.11](#).

```

#%% analyze the output
print(chat_completion.choices[0].message.content)

```

Listing 3.11 LMM: Chat Completion (Source: 03_LLMs\60_multimodal.py)

The image illustrates a flowchart for a machine learning (ML) model, showcasing the process of training and optimization. The model can provide valuable answers. It understands what it sees. Try to modify the user prompt to check whether the model can answer more detailed questions on the image.

3.2.4 Coding: Running Local LLMs

So far, you've run LLMs via API calls provided by software as a service (SaaS) providers. Sometimes, you want to run a model

locally. For instance, perhaps privacy is important, and you want to avoid transmitting confidential information over the internet.

In such cases, you can operate a model on your local computer. Ideally you have a powerful GPU, which can run a decently sized model. But a small model can also run on your CPU.

A powerful platform that makes this process quite simple is Ollama. With Ollama, you can operate an LLM on your laptop or desktop without the need for internet connectivity. It offers privacy and full control by allowing you to interact with an LLM directly on your hardware.

First, you need to install Ollama software on your computer by visiting <https://ollama.com/>.




Figure 3.3 Ollama Software Download

Download the software corresponding to your operating system. They provide the software for macOS, Linux, and Windows. Now, determine which model is suitable for your hardware and project requirements. A list of available models is available at <https://ollama.com/library>. For this example, we'll work with gemma2 model class, as shown in [Figure 3.4](#).




Figure 3.4 Ollama-Provided Model Class “gemma2”

This small, yet powerful, model provided by Google comes in three different flavors: tiny 2b, medium 9b, and large 27b. This number refers to the number of parameters, so the 2b refers to a model with roughly 2 billion parameters. Click on a link for more information (e.g., the file sizes of the models and the actual name for the model).

We'll use the smallest model, gemma2:2b. This option is feasible even for pure CPU systems.

You can download a model by pulling it via Ollama. On your terminal, run the code shown in [Listing 3.12](#).

```
ollama pull gemma2:2b
```

```
pulling manifest
pulling 7462734796d6... 100% ██████████ 1.6 GB
pulling e0a42594d802... 100% ██████████ 358 B
pulling 097a36493f71... 100% ██████████ 8.4 KB
pulling 2490e7468436... 100% ██████████ 65 B
pulling e18ad7af7efb... 100% ██████████ 487 B
verifying sha256 digest
```

```
writing manifest
success
```

Listing 3.12 Ollama: Model Download

The model was downloaded to your hard drive and is now available. You can check it via the following command:

```
ollama list
```

NAME	ID	SIZE	MODIFIED
gemma2:2b	8ccf136fdd52	1.6 GB	2 hours ago

The last step on the operating system level is to add the Python package `ollama`. You can add the package via one of the following options:

- `uv add ollama`
- `pip install ollama`

This step concludes the preparation work. Now, you can interact with the local model directly from a Python script.

In our first Python script, available at `03_LLMs\70_ollama.py`, you'll need to import the package first, as follows:

```
### packages
import ollama
```

Then, you can generate model responses by calling the method `.generate()`, as follows:

```
response = ollama.generate(model="gemma2:2b",
                           prompt="What is an LLM?")
```

The response has several kinds of metadata, but the most important output is available in the JSON key `response`, as shown in

[Listing 3.13](#).

```
from pprint import pprint
pprint(response['response'])
```

```
( 'LLM stands for **Large Language Model**.
"Here's a breakdown:
**What it is:**
* A type of artificial intelligence (AI) model.
...
```

Listing 3.13 Ollama: Model Output

Amazingly, you can now completely run an LLM self-hosted without an internet connection. But there's more: The behavior of that package is slightly different from the `langchain` API.

Thus, we recommend using it the `langchain` way. The LangChain community has built a model class specifically for Ollama, so that Ollama-hosted models provide the same outputs as other models, and you can easily exchange an Ollama model with any other provider.

You only need to load the corresponding package, as follows:

```
from langchain_community.llms import Ollama
```

Then, you can set up a model instance, as follows:

```
llm = Ollama(model="gemma2:2b")
```

With this model instance (`llm`), you can interact exactly in the same way as with other models—by invocation, as shown in [Listing 3.14](#).

```
response = llm.invoke("What is an LLM?")
response

( 'LLM stands for **Large Language Model**.
...
```

Listing 3.14 Ollama: Model Invocation

Isn't that awesome? You can now use an LLM even in confidential situations, without exposing any data over the internet.

Now, let's step back and cover some theoretical groundwork on steering model behavior. In the following section, you'll learn about the most important model parameters.

OceanofPDF.com

3.3 Model Parameters

Some particularly important parameters allow you to fine-tune the generated model outputs. Parameters like temperature, top-p, and top-k play a vital role; with their help, you can control the creativity, randomness, and focus of the generated output.

3.3.1 Model Temperature

With model temperature, you can control the randomness of the model's response. Typical values are 0 (low temperature) and 1, or even above (high temperature). Low temperatures keep the model focused, and you get deterministic results, meaning that you can repeatedly get the same response. The model favors extremely likely tokens. High temperatures, on the other hand, increase the randomness of token selection. A broader distribution of tokens is selected, which allows for more creative or unexpected outputs. Temperatures usually should not exceed 1, which can lead to chaotic and incoherent outputs. [Figure 3.5](#) shows how the temperature defines the model output.

For low temperatures, you get focused and deterministic answers, while the answers get more creative and varied for higher temperatures. Imagine you own an ice cream shop. If the (ambient) temperature is low, fewer customers visit your shop, and a better business decision might be to offer only the most popular flavors. But if temperature increases, demand rises, and as a result, a better decision would be to have more exotic flavors available.

Temperature is directly linked to the probability distribution of the tokens. Let's learn how temperature impacts the probability

distribution based on an artificial example. You have the prompt “Bert likes <MASK>.” and the model has the task to fill in the missing word. There are an enormous number of possible words. For simplicity, let’s just consider three words: “reading,” “walking,” and obviously “coding.” The model has an underlying probability of these words based on its training data. For extremely low temperatures, the model intensifies the differences between probabilities. For very high temperatures, these differences disappear, and all words have the same probability.




Figure 3.5 Temperature Output and Analogy

[Figure 3.6](#) shows the sample probability distributions for a given user prompt across different temperature values.




Figure 3.6 Model Temperature and Probability Distribution

Low temperatures of, for example 0.1, are shown in the plot on the left. The center plot shows the impact of a medium temperature on the probability distribution, while the plot on the right represents an extremely high temperature.

Notice how the differences are getting smaller with increasing temperature. This behavior is what this parameter does. Now, we study two more parameters that work alongside temperature: top-p and top-k.

3.3.2 Top-p and Top-k

Top-p (or also called “nucleus sampling”) controls the probability of the next token by dynamically adjusting the number of possible tokens that the model can choose from. Let’s take an example of top-p = 0.9. In this case, the model considers a cumulative probability of tokens (which add up to 90%) and chooses the smallest set of tokens within the boundaries.

This approach balances deterministic and creative outputs. If you set top-p = 1.0, there is no filtering, and all possible tokens are

considered. If you define a small value (e.g., $\text{top-p} < 0.5$), the model outputs tend to be more focused and predictable as only the top tokens are considered.

Top-k sampling controls how many of the most probable tokens are considered when generating the next word. If a value of $\text{top-k} = 1$ is chosen, only the most probable token is selected. Thus, you get a completely deterministic result. If $\text{top-k} = 50$, the model samples, from the top 50 tokens, the most probable tokens for each step. This approach increases diversity and allows for more creative and varied outputs. Top-k sets a fixed number of tokens to choose from, independent from the cumulated probability represented by top-k tokens. [Figure 3.7](#) shows some examples of top-p and top-k parameters.




Figure 3.7 Top-p and Top-k Parameters

The sample user prompt is "In the evening I want to see a <MASK>," in which the blank (<MASK>) should be filled. Again, multiple tokens are possible. Based on a given temperature, the probabilities of these tokens can be calculated. The tokens are represented in a decreasing order. All probability that sums up to be less than top-p are considered. In that case, "movie" and "friend" end up with 80% of probability. If the next token "doctor" would be added, the aggregated probability would be 91% and thus above the top-p value.

Top-k is straightforward. In our example, top-k is set to 3, so that the three most probable tokens are selected, of which the final prediction is randomly chosen.

Recommendations

The balance between these parameters depends on the field to be used in or the task to solve.

- In creative writing, you'll want to allow for higher temperature (0.8 to 1.0), coupled with moderate top-p (0.9 to 1.0) and top-k (50 to 100) to explore a range of creative outputs.
- In code generation, you want to reliable code snippets and thus will choose low temperature (0.1 to 0.3) with small top-k (10 to 20) and top-p (0.7 to 0.9) to ensure syntactically correct outputs.
- In customer service or chatbot applications, the model output needs to be reliable, focused, and consistent. A low temperature (0.2 to 0.4) ensures predictable responses. Top-p of 0.7 to 0.9 lets the model select the most probable tokens but retains some flexibility to keep the interaction natural. You want to avoid robot-like answers. Top-k can be in the range of 20 to 50; within that range, the model stays focused on providing relevant answers.

3.4 Model Selection

In the previous section, you undertook your first interactions with different LLMs. A selection of the models was given. But you might ask how to select the “right” model for your task. Depending on your project, you might have hard criteria and soft criteria. To process long input prompts, the context window (see the info box in the previous section) is a hard factor. To interact with a model that includes recent developments and trends, the model cutoff date could be extremely important.

A sweet spot can be found within the boundaries of model price, latency, and performance. Finding the right balance is important for any generative AI application. This trade-off is highly contextual and influenced by the specific requirements of your project and the resources at your disposal.

The performance of a model highly impacts the model’s outputs. Related to model performance is the price of model use. In general, the more powerful a model is, the more expensive it is to use. If you need a model that inherently “knows” about recent events, you must take the knowledge cutoff date into account. Your input data might also define the model selection since the context window might truncate your input if it is too small. For real-time applications, you need to ensure that the model has low latency.

Depending on the data you work with, you might want to operate the LLM on your own resources (on-premise) rather than using a model hosted in the cloud. Associated with the model hosting is the model license. Open-source and open-weight models provide their model weights so that you can download and operate it on your hardware,

which is not possible with proprietary models that are only accessible via APIs.

Many of these parameters are interconnected, which makes model selection quite difficult.

Let’s consider the performance of a model and take a closer look at model selection criteria.

3.4.1 Performance

You can check the performance of different models at Chatbot Arena (<https://lmarena.ai/?leaderboard>) and get a result, as shown in [Figure 3.8](#).

Category		Apply Filter		Overall Questions			
Overall		☐ Style Control ☐ Show Deprecated		#models: 194 (300%) #votes: 2,557,144 (300%)			
Rank+ (US)	Rank (StyleCtrl)	Model	Arena Score	95% CI	Votes	Organization	License
1	3	Gemini-2.0-Flash-Thinking-Exp-01-21	1380	+8/-9	5572	Google	Proprietary
1	1	Gemini-Exp-1206	1376	+5/-5	21064	Google	Proprietary
3	1	ChatGPT-4o-latest (2024.11.20)	1365	+4/-3	34299	OpenAI	Proprietary
4	4	Gemini-2.0-Flash-Exp	1356	+8/-6	19823	Google	Proprietary
4	1	ai.2024.12.17	1351	+8/-5	8128	OpenAI	Proprietary
6	4	ai.provision	1335	+4/-4	33202	OpenAI	Proprietary
7	7	DeepSeek-V3	1329	+5/-5	11893	DeepSeek	DeepSeek
7	10	Step-2-34K-Exp	1306	+8/-7	4166	StepFun	Proprietary
8	10	ai.mai	1306	+3/-3	40847	OpenAI	Proprietary
8	9	Gemini-1.5-Pro-002	1303	+3/-3	45486	Google	Proprietary

Figure 3.8 Chatbot Arena Leaderboard (as of January 24, 2025; Source: <https://lmarena.ai/?leaderboard>)

The models are sorted in terms of arena score. But how is this score calculated? It is called “arena” (<https://lmarena.ai/>) for a reason. In the arena, the user interacts with two models: model A and model B. The user can define a prompt and gets responses from the two models. The user must rate which one is performing better. Thus, Chatbot Arena is a double-blind test, which is considered the gold standard in the evaluation of test outcomes. In the screenshot, notice that multiple models share the same rank because the 95% confidence interval is considered the highest rank. The ranking

changes often, so your rankings will likely look different due to the passing time.

By default, the overall ranking is shown. But you can select other categories like **Math** or **Coding** and check those rankings. When searching for a model for a specific task, we advise selecting the category according to your task.

But performance is not the only relevant factor.

3.4.2 Knowledge Cutoff Date

Each model has a *knowledge cutoff date*, which is the date on which its training data was finalized. The model is subsequently trained on that data. In the model weights, no more recent data can be represented, which is why it is important to know the knowledge cutoff date. If you ask a model for information like an event or any other piece of information, the model cannot know if it happened after the date of the knowledge cutoff.

For chatbots, this parameter is becoming less relevant because these models are more often equipped with internet search capabilities to retrieve up-to-date information by themselves. However, for you as a developer of AI systems, knowledge cutoffs might be an important factor to consider.

3.4.3 On-Premise versus Cloud-Based Hosting

Another important aspect of model selection is data privacy. If you're dealing with confidential information, you or your customers might not want the information to leave the company network. In that case, you want to use a model that is hosted in the safe haven of the

internal network. But in that case you're most likely bound to use open-source models.

3.4.4 Open-Source, Open-Weight, and Proprietary Models

Proprietary models are provided to users via web applications or APIs. Famous members of this class include OpenAI or Anthropic who both provide their models in this way.

Google is a special case since their models are either provided as proprietary models via APIs (e.g., Gemini). But other model classes, like Gemma, are provided as open-source models via Hugging Face.

To be completely accurate, we should differentiate “open source” and “open weight.” Truly open-source models are provided with all their details, such as its model architecture and training data used. This total transparency is rare, however. A provider might release to the public a trained model with its weights, but specific details of the underlying data and training might still be kept secret.

A famous example of this group is Meta with its Llama model family. These models are free to use, but the company doesn't disclose its training data.

3.4.5 Price

The price of using an LLM service can be a significant factor in your model selection process. Typically, proprietary models are paid on a token basis. To be more precise, a differentiation is made between input tokens and output tokens. Input tokens are typically cheaper than output tokens. You can find the current prices for OpenAI

models at <https://openai.com/api/pricing/> and for Anthropic models at <https://www.anthropic.com/pricing#anthropic-api>.

You should come up with some estimates on how many API requests and how many tokens will be processed. Based on these estimates, you can derive an estimate of your total costs.

3.4.6 Context Window

Your project might include the processing of very long documents, making it necessary to pass as much information as possible to the model. Thus, the context window is a driving factor for the best choice of model.

If you check, for example, the models on Groq (<https://console.groq.com/docs/models>), you'll find models with rather small context windows like "LlaVa 1.5 7B" with a context window of 4.096 tokens as well as "Llama 3.3 70B Versatile" with an extremely large context window of 128.000 tokens.

3.4.7 Latency

Some use cases require extremely fast model responses despite any interdependencies that come with the hosting of a model. If latency does not play a role, you might even run an open-source model on a CPU. In other cases, latency might be the most relevant factor.

For example, let's say you want to couple an LLM with voice generation to enable real-time chats. In such a situation, an LLM can easily be the bottleneck and reduce the user experience because

conversation cannot be “natural” if the conversation partner requires long response times.

OceanofPDF.com

3.5 Messages

In [Section 3.2](#), you took your first steps in using LLMs. You invoked the model objects, sent a simple message, and received a response. In a more realistic chat, multiple types of messages will arise. Each message has a specific role and content. We'll look at the most common message types: user, system, and assistant.

3.5.1 User

This message type relates to the human message and represents the user input. The effectiveness of an LLM response depends on the clarity of the user message. [Chapter 4](#) on prompt engineering basically deals with optimizing the user message.

3.5.2 System

Alongside the user input, a system message can be defined that specifies how the model should behave and work, like in a role-play.

For example, if you want to set up a general assistant, a typical system message might be as follows:

```
You are a helpful AI assistant designed to provide accurate,  
concise, and polite responses. Always ensure that your answers are  
clear and informative.
```

Alternatively, imagine you want your model to behave as a technical support assistant. In this case, you might instruct the model with the following system message:

You are a technical support AI assistant specializing in troubleshooting and explaining software-related issues. Respond with clear, step-by-step instructions, avoiding technical jargon whenever possible.

With the system message, you can establish some initial instructions for the model. You can define its role, its tone, and even specific objectives before the user interaction begins. The system message is critical for setting the model's boundaries and expectations and helps guide it to behave in a manner aligned with the user's requirements.

System messages have their limitations. While they can shape the initial behavior, they cannot enforce strict adherence throughout the conversation. As a result, models can drift in their tone or behavior when prompted with unforeseen user inputs.

Additionally, system messages alone cannot enforce fine-grained control over content accuracy or nor attend to ethical considerations without complementary guardrails or moderation.

More on this topic is covered in [Chapter 4](#).

3.5.3 Assistant

The assistant message type corresponds to the model response. The main property is the content, which holds the model output.

Furthermore, this message type has a `response_metadata` property. This property holds some model-specific output, but typically the token usage is shown as well as the duration the query took.

Now that you understand the available message types, let's explore how they can be used in prompts. LangChain has a flexible interface

to help you set up prompts.

OceanofPDF.com

3.6 Prompt Templates

Before we invoke the LLM and send a query, let's set up a prompt in a unified and structured way by using LangChain's prompt templates. In this way, we can guide the model in how to act. Also, the model can, with the help of prompt templates, better understand the user and the user's intention. Let's see this capability in action.

3.6.1 Coding: ChatPromptTemplates

The most flexible way to implement prompt templates is to use `ChatPromptTemplates`, which allows you to pass a list of messages. [Listing 3.15](#) shows some sample code for setting up a prompt template.

Let's start with a simple example that showcases the idea. First, import the class `ChatPromptTemplate`. In the next step, create an instance of this class by calling its `from_messages()` method. These messages are a list of tuples. Each tuple has the form ("message type," "content"). In this way, you can define a system message that tells the model how to behave, followed by a human message that holds the actual user query. Important in this context is how you define variables, which are set up as placeholders and populated in a later step. In our example, variables are defined in the user-based (human) messages, inside curly braces. At this point, we've set up `input` and `target_language` as variables.

Although our code looks a bit like a formatted string literal (also called Python f-string), it isn't. In an f-string, predefined variables are passed inside string literals using curly braces `{}` and replaced by the string representation of the variable. In this case, we did not

predefine the variable `input` or the variable `target_language` in advance.

In the last step, we invoke prompt template. The variables are then replaced with actual content when we call the `invoke()` method of the `prompt_template` object. A dictionary is passed as a parameter; this dictionary uses keys corresponding to the variables and the values corresponding to the content to be used instead of the variable.

Finally, since we didn't store the invoke prompt to a new variable, the output is just shown in the terminal. The prompt template has been converted to a `ChatPromptValue` object with a `SystemMessage` and `HumanMessage`.

```
### packages
from langchain_core.prompts import ChatPromptTemplate

### set up prompt template
prompt_template = ChatPromptTemplate.from_messages([
    ("system", "You're an AI assistant that translates English into another language."),
    ("user", "Translate this sentence: '{input}' into {target_language}"),
])

### invoke prompt template
prompt_template.invoke({"input": "I love programming.", "target_language": "German"})

ChatPromptValue(messages=[SystemMessage(content='You're an AI assistant that translates English into another language.'), HumanMessage(content="Translate this sentence: 'I love programming.' into German")])
```

Listing 3.15 Prompt Template Use (Source: 03_LLMs/30_prompt_templates.py)

What is the purpose of this approach? With the prompt template, we have a flexible first component that could be passed to a model to get some response. The “prompt to LLM” approach is a simple chain of steps.

We'll deal further with LangChain chains in [Section 3.7](#).

But first, let's use the wisdom of the crowd to come up with a good prompt. LangChain has established an ecosystem that allows users to share and explore prompts with LangChain Hub.

3.6.2 Coding: Improve Prompts with LangChain Hub

You can find LangChain hub at <https://smith.langchain.com/hub>. This hub allows you to explore prompts created by others for many purposes. We'll dive more deeply into the prompt engineering ecosystem in [Chapter 4](#). For now, we'll get some help to create a prompt.

If you search for “prompt maker,” you'll find a user-created prompt template called “hardkothari/prompt-maker.” This prompt is created to generate a great prompt. In our example, you'll find out how it works.

The code shown in [Listing 3.16](#) corresponds to the file `03_LLMs/31_prompt_hub.py`. You need to load the required packages. The newcomer is `hub` from `langchain` package.

```
from langchain import hub
from langchain_openai import ChatOpenAI
from langchain_core.output_parsers import StrOutputParser
from dotenv import load_dotenv
load_dotenv('.env')
from pprint import pprint
```

Listing 3.16 Prompt Hub: Required Packages (Source: `03_LLMs/31_prompt_hub.py`)

To use the prompt creation, you must call the `pull` method from `hub`, as follows:

```
### fetch prompt
prompt = hub.pull("hardkothari/prompt-maker")
```


Some input variables are exposed via the `input_variables` property, as follows:

```
### get input variables
prompt.input_variables

['lazy_prompt', 'task']
```

[Listing 3.17](#) shows an improved prompt. You only need to set up a model and run it in a chain, which anticipates knowledge about chains that we cover later in [Section 3.7](#), so be patient. For now, just take this step as a given.

```
# %% model
model = ChatOpenAI(model="gpt-4o-mini",
                    temperature=0)

# %% chain
chain = prompt | model | StrOutputParser()
```

Listing 3.17 Prompt Hub: Chain Setup (Source: 03_LLMs/31_prompt_hub.py)

You can invoke the chain with the code shown in [Listing 3.18](#) and pass the relevant parameters (i.e., `lazy_prompt` and `task`) to get an improved prompt.

```
# %% invoke chain
lazy_prompt = "summer, vacation, beach"
task = "Shakespeare poem"
improved_prompt = chain.invoke({"lazy_prompt": lazy_prompt, "task": task})
print(improved_prompt)
```

As a skilled poet in the style of William Shakespeare, compose a sonnet that captures the essence of summer, vacation, and the beach. ### Instructions: Your poem should reflect the beauty and joy of a summer getaway, using rich imagery and evocative language typical of Shakespearean verse. Aim for 14 lines, adhering to the traditional iambic pentameter and ABABDCDEFEFGG rhyme scheme. ### Context: Incorporate themes of nature, leisure, and the fleeting nature of time, while evoking a sense of nostalgia and warmth. Use metaphors and similes to enhance the emotional depth of the poem, and consider including references to the sun, sea, and the carefree spirit of summer. ### Desired Outcome: The final piece should resonate with readers, transporting them to a sun-drenched beach, filled with laughter and the gentle sound of waves, while also reflecting on the transient beauty of such moments.

Listing 3.18 Prompt Hub: Chain Invocation (Source: 03_LLMs/31_prompt_hub.py)

This improved prompt describes the role of the model in much more detail and provides extensive instructions on the task as well as a description of the desired outcome. Let's run the model with the improved prompt, as shown in [Listing 3.19](#).

```
# %% run model with improved prompt
res = model.invoke(improved_prompt)
print(res.content)
```

Upon the golden sands where sunbeams play,
The azure waves do kiss the shore with glee,
In summer's arms, we cast our cares away,
As laughter dances on the breeze so free.

The sun, a sovereign in the heavens high,
Doth paint the sky with hues of rose and gold,
While seagulls wheel and cry, as if to vie
For joy that in this fleeting time we hold.

O, let us linger in this warm embrace,
Where time, like grains of sand, slips through our hands,
Each moment savored, sweet as summer's grace,
A tapestry of dreams upon the strands.

Yet, as the twilight calls the day to rest,
We cherish memories, our hearts possessed.

Listing 3.19 Prompt Hub: Model Run With Improved Prompt (Source: 03_LLMs/31_prompt_hub.py)

We'll leave it to you to run the model only based on the `lazy_prompt` or the `task` and to compare the results.

So far, we've worked with a small example of how a chain combines different building blocks. This concept is powerful due to its structure, as we'll explore next.

3.7 Chains

Chains are such an important concept that they are part of the package named LangChain. A *chain* refers to a sequence of process steps that are connected to accomplishing a task. Typically, they consist of multiple components.

We start with the smallest and simplest chain: a link between a prompt and an LLM. Then, we'll look at how chains can be executed in parallel and how routes can be implemented. Finally, we equip an LLM with memory so that it can retain information. Let's start small, with a simple sequential chain.

3.7.1 Coding: Simple Sequential Chain

The simplest chain is a “prompt to LLM” chain, as shown in [Figure 3.9](#). A user input is passed to a prompt template. The prompt template itself passes its output into an LLM step. Finally, the LLM step creates a model output.




Figure 3.9 Simple LangChain Chain

You're not limited to a sequential chain. You can also craft more complex structures like parallel running chains or router chains. We'll describe these structures in more detail later in this section.

Let's set up a chain consisting of a prompt. The prompt is passed to an LLM. After the model creates an output, the output will be passed to a `StrOutputParser`. You can find the corresponding code at [03_LLMs/40_simple_chain.py](#).

As shown in [Listing 3.20](#), start by importing relevant packages and API keys. As the LLM, we'll use OpenAI. The API key is loaded via `dotenv`. Make sure that the `.env` file is stored in the working folder, which has an entry for `OPENAI_API_KEY`.

```
### packages
from langchain_openai import ChatOpenAI
from langchain_core.prompts import ChatPromptTemplate
from dotenv import load_dotenv
from langchain_core.output_parsers import StrOutputParser
load_dotenv('.env')
```

Listing 3.20 Sequential Chain: Required Packages (Source: [03_LLMs/40_simple_chain.py](#))

Now, [Listing 3.21](#) extends on the prompt template from [Section 3.6](#) creating a `ChatPromptTemplate` based on a system and user message. The task is to translate an `input` text into a `target_language`.

```
### set up prompt template
prompt_template = ChatPromptTemplate.from_messages([
    ("system", "You're an AI assistant that translates English into another language."),
    ("user", "Translate this sentence: '{input}' into {target_language}"),
])
```

Listing 3.21 Sequential Chain: Required Packages (Source: [03_LLMs/40_simple_chain.py](#))

The next component to apply is an LLM. We'll use `GPT-4o-mini`, as follows:

```
# %% model
model = ChatOpenAI(model="gpt-4o-mini", temperature=0)
```

Stacking the chain together cannot be simpler. Simply apply the pipe operator `|`. Components are separated by `|`. In our example, the prompt is the first component in the chain, followed by the model. Subsequently, the model output is passed to `StrOutputParser` which parses the model output into the top likely string.

```
# %% chain
chain = prompt_template | model | StrOutputParser()
```

Everything is prepared, and we can invoke the chain with input parameters. As a result, we'll get the final output, as follows:

```
# %% invoke chain
res = chain.invoke({"input": "I love programming.", "target_language": "German"})
res

'The translation of "I love programming." into German is "Ich liebe Programmieren."'
```

Now that you understand simple sequential chains, let's set the bar higher and implement a parallel chain.

3.7.2 Coding: Parallel Chain

You can execute two chains and run both simultaneously, which can speed up a process, especially when API calls are involved, which now can be run in parallel rather than sequentially. [Figure 3.10](#) shows how a user input triggers the running of two chains in parallel. After both chains have completed, the output can be combined.

We'll develop a chain system in which two chains are executed in parallel, as shown in [Listing 3.22](#). Based on a user input on a specific topic, two chains run and answer in two different styles: One answers in a polite and friendly tone; the other, in a savage and angry tone. You're not limited to two parallel chains.

```
from langchain_openai import ChatOpenAI
from langchain_core.prompts import ChatPromptTemplate
from langchain_core.runnables import RunnableParallel
```

```

from langchain_core.output_parsers import StrOutputParser
from dotenv import load_dotenv
load_dotenv('.env')

```

Listing 3.22 Parallel Chains: Required Packages (Source: 03_LLMs/41_parallel_chain.py)




Figure 3.10 Parallel Execution of Two Chains

As the LLM backend, we'll use GPT-4o-mini, as follows:

```

### Model Instance
llm = ChatOpenAI(model="gpt-4o-mini", temperature=0)

```

Let's prepare the two chains, namely, `polite_chain` and `savage_chain`, as shown in [Listing 3.23](#). Both are based on `ChatPromptTemplate` in which we specify the tone in the system prompt. The human prompt just consists of the `topic` placeholder. An important consideration is that both get the same `topic` parameter to enable parallel execution.

```

### Prepare Prompts
# example: style variations (friendly, polite) vs. (savage, angry)
polite_prompt = ChatPromptTemplate.from_messages([
    ("system", "You are a helpful assistant. Reply in a friendly and polite manner."),
    ("human", "{topic}")
])

savage_prompt = ChatPromptTemplate.from_messages([
    ("system", "You are a helpful assistant. Reply in a savage and angry manner."),

```

```

        ("human", "{topic}")
    ])

```

Listing 3.23 Parallel Chains: Prompts (Source: 03_LLMs/41_parallel_chain.py)

For illustration purposes, these chains are very simple. So, we'll just use a prompt, which is passed to the LLM, and the top likely output will be returned, as follows:

```

%% Prepare Chains
polite_chain = polite_prompt | llm | StrOutputParser()
savage_chain = savage_prompt | llm | StrOutputParser()

```

Now, let's see how both chains can be executed in parallel. The class that provides this functionality is `RunnableParallel`, as shown in [Listing 3.24](#).

```

# %% Runnable Parallel
map_chain = RunnableParallel(
    polite=polite_chain,
    savage=savage_chain
)

```

Listing 3.24 Parallel Chains: Runnable Parallel Implementation (Source: 03_LLMs/41_parallel_chain.py)

All the pieces are in place. You are ready to invoke your chain with a given topic, as shown in [Listing 3.25](#). Hopefully, the model can tell us the meaning of life: in a friendly tone and in an annoyed, savage tone.

```

# %% Invoke
topic = "What is the meaning of life?"
result = map_chain.invoke({"topic": topic})
# %% Print
from pprint import pprint
pprint(result)

```

```

{'polite': 'The meaning of life is a profound question that has been explored ' 'by
philosophers, theologians, and thinkers throughout history. It ' 'can vary greatly
from person to person, depending on individual ' 'beliefs, values, and experiences.
For some, it may be about ' 'seeking happiness, forming connections with others, or
pursuing ' 'knowledge and personal growth. For others, it might involve '
'fulfilling a purpose, contributing to society, or exploring ' 'spirituality.
Ultimately, it's a personal journey, and many find ' 'that the meaning of life is

```

```
something they define for themselves ' 'through their experiences and reflections.
What are your thoughts ' 'on it?', 'savage': 'Oh, here we go with the age-old
question that everyone thinks ' 'they're so profound for asking. The meaning of
life? Seriously? ' 'It's not some cosmic riddle waiting for you to crack it. Life is
' 'what you make of it, so stop whining and start living! If you're ' 'waiting for
some grand revelation, you're going to be waiting ' 'forever. Get off your
philosophical high horse and figure out what ' 'makes you tick instead of asking the
universe for answers!'}

```

Listing 3.25 Parallel Chains: Model Invocation (Source: 03_LLMs/41_parallel_chain.py)

The good news is that the model answered in the requested tones. The bad news is that it cannot really tell us the meaning of life. Perhaps, we've been too presumptuous to expect this question answered by a model.

You can also implement a decision logic that allows for different paths to be chosen depending on a condition, which we'll cover next.

3.7.3 Coding: Router Chain

[Figure 3.11](#) shows a router chain. The “router” in this context works like an if-then-else statement with multiple outputs. Depending on the input, it can route to one of several outputs.




Figure 3.11 Router Chain

But how does a router chain work? Several options exist for how a router defines which subsequent chain to use. One sophisticated

approach is to use a semantic router. In this case, the user input query is analyzed based on a word or sentence embeddings. The semantic meaning of the query is studied, and the most suitable subsequent chain selected. This approach is based on embeddings and similarities—some concepts we'll cover in detail in [Chapter 4](#).

Like all chains in LangChain, it follows the *LangChain expression language (LCEL)*, which is an easy approach to compose chains in LangChain. The idea is that no code change should be required to move from prototype to production. Thus, LCEL provides a unified interface with many features. Some important features include support for asynchronous calls, parallel execution, retries, and fallbacks. You can also easily access intermediate results and use structured inputs and outputs. As a result, your chains can produce outputs directly in a dictionary or in JavaScript Object Notation (JSON) format.

A chain can consist of prompts, models, retrievers, and tools. All these components can be connected to working on a complex task. These components are executed in a predefined order, in which the output of one component becomes the input for the next component. This flexibility allows you to work on specific use cases and is a versatile tool to be applicable for various applications.

Now, let's implement such a router chain in Python.

A chain is not necessarily a linear sequence. You can use conditional statements and then route the process to different chains, and from there, more conditional checks might come, with no limit to possible complexity. Let's start simply and implement a router chain, as shown in [Figure 3.11](#). The code is stored in `03_LLMs/45_semantic_router.py`.

This code will include some concepts like embeddings and cosine similarity, which will be covered in depth in [Chapter 4](#). We'll explain these concepts only as far as necessary for this example.

The general approach consists of the following steps:

1. We create different chains, one covering math questions, another for music, and yet another for history questions.
2. Embeddings are created for the words according to the chains: “math,” “music,” and “history.”
3. The semantic router then works in the following way: It creates a numeric vector (an embedding) for the user query. The similarity between the user query embedding and the three embedding words is calculated. The most similar embedding is selected and defines the chain that is most appropriate for the user query.
4. That chain is selected and invoked with the user query.

Alright, let's get coding. First things first, as shown in [Listing 3.26](#), start with loading relevant packages.

```
## packages
from langchain_core.prompts import ChatPromptTemplate
from langchain_core.output_parsers import StrOutputParser
from langchain_openai import ChatOpenAI, OpenAIEmbeddings
from langchain_community.utils.math import cosine_similarity
from dotenv import load_dotenv
load_dotenv('.env')
```

Listing 3.26 Semantic Router Implementation: Required Packages (Source: 03_LLMs/45_semantic_router.py)

As before, we'll work with OpenAI and use GPT-4o-mini as LLM. Furthermore, we'll use an embedding model from OpenAI that translates a human understandable text into a numeric vector. Make sure that you have an `.env` file in your working folder with an `OPENAI_API_KEY` entry.

In the next step, we'll create instances of an LLM and an embeddings model for later use in the script, as follows:

```
# %% Model and Embeddings Setup
model = ChatOpenAI(model="gpt-4o-mini", temperature=0)
embeddings = OpenAIEmbeddings()
```

As shown in [Listing 3.27](#), the chains will be based on messages, in which we can define system prompts. Later, in the output, we want to see the chain that was used, so we'll instruct the model to mention the used chain.

```
### Prompt Templates
template_math = "Solve the following math problem: {user_input}, state that you are a math agent"
template_music = "Suggest a song for the user: {user_input}, state that you are a music agent"
template_history = "Provide a history lesson for the user: {user_input}, state that you are a history agent"
```

Listing 3.27 Semantic Router Implementation: Prompt Templates (Source: 03_LLMs/45_semantic_router.py)

Now, you can use these templates, as shown in [Listing 3.28](#), in the definition of the system message. In the user message, you can directly use a placeholder for the user input, which will be invoked when the router runs.

```
# %% Math-Chain
prompt_math = ChatPromptTemplate.from_messages([
    ("system", template_math),
    ("human", "{user_input}")
])
chain_math = prompt_math | model | StrOutputParser()

# %% Music-Chain
prompt_music = ChatPromptTemplate.from_messages([
    ("system", template_music),
    ("human", "{user_input}")
])
chain_music = prompt_music | model | StrOutputParser()

###
# History-Chain
prompt_history = ChatPromptTemplate.from_messages([
```

```

        ("system", template_history),
        ("human", "{user_input}")
    ])
    chain_history = prompt_history | model | StrOutputParser()

```

Listing 3.28 Semantic Router Implementation: Chain Setup (Source: 03_LLMs/45_semantic_router.py)

All chains are combined into one list. We'll come back to that list and select the most appropriate chain in the router, as follows:

```

### combine all chains
chains = [chain_math, chain_music, chain_history]

```

According to the chains, embeddings are created for the words that describe the chains best. In our case, we can use the words “math,” “music,” and “history,” as shown in [Listing 3.29](#). Double-check that the `chain_embeddings` object has the right count of elements.

```

# %% Create Prompt Embeddings
chain_embeddings = embeddings.embed_documents(["math", "music", "history"])
###
print(len(chain_embeddings))

```

3

Listing 3.29 Router Chain: Chain Embeddings (Source: 03_LLMs/45_semantic_router.py)

Now, we've come to the core of this script—the router. As shown in [Listing 3.30](#), we've set the router up as a function with one input—the user query input. The output will be one of the three chains. At first, the user query is embedded into an object `query_embedding`. We need to find out whether the user wants to know about math, music, or history. Thus, we need to calculate the similarities between the `query_embedding` and the three embeddings corresponding to the chains.

The `similarities` object is a list with three elements. The first element corresponds to the similarity between the user query and

“math”; the second element corresponds to the similarity between the user query and “music”; and so on.

You can access the `most_similar_index` by selecting the index of the highest value with `argmax()`. Finally, out of the list of chains, the one with the most similar embedding is returned.

```
# %% Prompt Router
def my_prompt_router(input: str):
    # embed the user input
    query_embedding = embeddings.embed_query(input)
    # calculate similarity
    similarities = cosine_similarity([query_embedding], chain_embeddings)
    # get the index of the most similar prompt
    most_similar_index = similarities.argmax()
    # return the corresponding chain
    return chains[most_similar_index]
```

Listing 3.30 Semantic Router Implementation: Router Function (Source: 03_LLMs/45_semantic_router.py)

We’re done, and it’s time to test our implementation. As shown in [Listing 3.31](#), we created three different queries. Each one should be clearly linked to one of the three different chains. You can comment/uncomment the different query objects and run the code block.

The queries are passed into `my_prompt_router` as an argument. The return value will be the selected chain. We create an object called `chain` based on the function output. This `chain` is then invoked with the user query.

In the model output, you should see a reference to indicate which chain was used. Test the router with the three different examples and then come up with your own queries to test the implementation.

```
# %% Testing the Router
# query = "What is the square root of 16?"
# query = "What happened during the french revolution?"
query = "Who composed the moonlight sonata?"
chain = my_prompt_router(query)
```

```
chain.invoke(query)
```

The "Moonlight Sonata" was composed by Ludwig van Beethoven. As your music agent, I recommend listening to this beautiful piece, especially the first movement, which captures a serene and reflective mood. If you're interested in more classical music or similar pieces, just let me know!

Listing 3.31 Semantic Router Implementation: Testing (Source: 03_LLMs/45_semantic_router.py)

You've successfully implemented your first router, and now you can create highly complex chains and business logic on your own.

Next, we'll cover how you can equip a chain with some memory.

3.7.4 Coding: Chain with Memory

A complex AI system needs many more features. An important feature is memorization so that a system can “know” about its previous interactions with the user. Another feature is a chain that is infinite. We'll bundle these features in a fun text-based game.

The user will be placed in a setting with three different options to determine the future path of the story. [Figure 3.12](#) shows a sample output for the beginning of the game. The setting has a cliffhanger,

and the user can decide how to proceed by choosing 1, 2, or 3. Alternatively, the user can type “quit” to end the game.

```
The moonlight barely pierces through the thick canopy of the dark forest, casting eerie shadows that dance among the gnarled trees. A chill seeps into your bones as the wind whispers through the branches, carrying secrets of the ancient woods. Suddenly, a rustle in the underbrush catches your attention, and your heart races. You stand at a crossroad, where the path diverges into three.

What do you choose to do?

1 Follow the sound of rustling in the bushes.
2 Take the left path that leads deeper into the shadows.
3 Head back to the safety of the forest's edge.
Enter your choice: (or 'quit' to end the game)2
You take a deep breath and venture down the left path, the shadows growing thicker around you. The air feels heavy, charged with a sense of foreboding. As you move deeper, the sound of your footsteps is swallowed by the silence, and the rustling continues, now more pronounced. Suddenly, the path opens into a small clearing, illuminated by bioluminescent fungi that pulse like a heartbeat. In the center, a stone altar stands, covered in strange markings and surrounded by a circle of flickering lights.

What do you choose to do?
```

Figure 3.12 Sample Output from Our LLM Chain Game

A good practice is to visualize the steps the chain will take. The complete process flow is shown in [Figure 3.13](#). First, the setting of our game world will be defined. The LLM will be equipped with memory, so that all interactions are stored. As an output, the LLM will offer three different paths, of which the user must choose one—or choose to quit the game. Assuming the user wants to continue, based on the user input, the story is spun further. New paths are created, and the user can decide how to proceed.




Figure 3.13 LLM Text Game Including Memory

Now that you have a high-level understanding of the game logic, let's start coding.

As always, import all packages at the beginning of the script, as shown in [Listing 3.32](#). New packages in this example are the classes `RunnableWithMessageHistory`, `BaseChatMessageHistory`, and `InMemoryChatMessageHistory`. All these classes are required for handling history management. Another new entry is the package `rich` with the classes `Markdown` and `console`. These classes enable nicer-looking results for LLM outputs.

```
### Packages
from langchain_openai import ChatOpenAI
from langchain.prompts import ChatPromptTemplate
from langchain_core.runnables import RunnableWithMessageHistory
from langchain_core.chat_history import BaseChatMessageHistory,
InMemoryChatMessageHistory
from dotenv import load_dotenv
from rich.markdown import Markdown
from rich.console import Console
console = Console()
load_dotenv(".env")
```

Listing 3.32 Chain with Memory: Required Packages (Source: 03_LLMs/42_chain_game.py)

We'll work with OpenAI, but feel free to plug in a different LLM of your choice:

```
### Prepare LLM
llm = ChatOpenAI(model="gpt-4o-mini", temperature=0.7)
```

The session history will be stored in an object store. The function `get_session_history` will process all the messages and add them to the store object under a `session_id`. With this feature, our game is extendable, and you might spin out a different story by `session_id`, maintaining separate threads under a different IDs.

```
# %% Session history
store = {}
def get_session_history(session_id: str) -> BaseChatMessageHistory:
    if session_id not in store:
```



```

        store[session_id] = InMemoryChatMessageHistory()
    return store[session_id]

```

Listing 3.33 Chain with Memory: Session History (Source: 03_LLMs/42_chain_game.py)

Now, let's create the initial setting for the game with the first three choices for the player. By exposing the `place` as a placeholder, the user can decide where the scene will take place. The prompt is created with `ChatPromptTemplate` and the chain in the usual way, as shown in [Listing 3.34](#).

```

### Begin the story
initial_prompt = ChatPromptTemplate.from_messages([
    ("system", "You are a creative storyteller. Based on the following context and player's choice, continue the story and provide three new choices for the player. keep the story extremely short and concise. Create an opening scene for an adventure story {place} and provide three initial choices for the player.")
])

context_chain = initial_prompt | llm

```

Listing 3.34 Chain with Memory: Initial Prompt and Context Chain (Source: 03_LLMs/42_chain_game.py)

As shown in [Listing 3.35](#), set up a `RunnableWithMessageHistory`, which allows you to add the history of messages in the interaction to the chain. In this case, we pass our `context_chain` and the function `get_session_history`. The result of this LLM call is stored in `context` and printed to the console with `console.print`.

```

config = {"configurable": {"session_id": "03"}}

llm_with_message_history = RunnableWithMessageHistory(context_chain,
get_session_history=get_session_history)

context = llm_with_message_history.invoke({"place": "a dark forest"}, config=config)

# render opening scene as markdown output
console.print(Markdown(context.content))

```

Listing 3.35 Chain with Memory: LLM Setup (source: 03_LLMs/42_chain_game.py)

Great, our game has started! Now, we must enable the user choice mechanism. For processing player choice, set up a function named `process_player_choice`, as shown in [Listing 3.36](#). This function consumes the choice and returns the LLM response. It passes two messages to the model: first, a user message, which instructs the LLM to continue the story, and then a system message, which asks for the creation of three new choices for the player. The `config` is passed so that these interactions are stored in the memory.

```
def process_player_choice(choice):
    response = llm_with_message_history.invoke(
        [
            ("user", f"Continue the story based on the player's choice: {choice}"),
            ("system", "Provide three new choices for the player.")
        ],
        config=config
    )
    return response
```

Listing 3.36 Chain with Memory: Function for Player's Choice (Source: 03_LLMs/42_chain_game.py)

Now, we can enter the game loop, as shown in [Listing 3.37](#). Since we've kept it open ended, we can use a `while` loop, which only ends when the user types "quit." The player choice is fetched with `input()`, and the game checks whether it should end in this loop. If not, we continue the story and run our function `process_player_choice`. The result will be printed on the screen, and the loop starts all over again to extend the story.

```
# %% Game loop
while True:
    # get player's choice
    player_choice = input("Enter your choice: (or 'quit' to end the game)")
    if player_choice.lower() == "quit":
        break
    # continue the story
    context = process_player_choice(player_choice)
    console.print(Markdown(context.content))
# %%
console.print(Markdown(context.content))
```

Listing 3.37 Chain with Memory: Game Loop (Source: 03_LLMs/42_chain_game.py)

We hope you enjoy this little game we developed together. Along the way, you learned how to implement memory so that the model remembers past paths and user selections.

We'll now tackle a critical side mission necessary when embedding your chains in the real world: security

OceanofPDF.com

3.8 Safety and Security

When working with LLMs, you must consider safety and security before you can deploy any system to production. An LLM should be protected from harmful user prompts (security), and users should be protected from harmful LLM responses (safety).

[Figure 3.14](#) shows the pipeline from a user prompt that is first passed to an input firewall that checks for prompt security. Then, the prompt is further processed by the LLM before the output firewall checks whether the model response is safe and can thus be provided to the user.




Figure 3.14 LLM Safety and Security

Let's study both aspects in more detail, starting with the input firewall: LLM security.

3.8.1 Security

LLM security is similar to the security of a company. We want to ensure that only trustworthy people enter company premises. LLM security ensures that only trustworthy prompts are processed. An LLM should be protected from threats and vulnerabilities that could compromise their integrity or compromise confidentiality.

Some important aspects of security include the following:

- **Hackers**

Malicious actors might attempt to compromise LLMs through various methods, like adversarial attacks to mislead the model. Prompt injections try to extract sensitive information.

- **Data poisoning**

In this threat, attackers corrupt the training data to embed biases. This vulnerability might be relevant when you train your own model based on custom data.

- **Robustness**

Intertwined with adversarial attacks, robustness ensures that any attempt to exploit weaknesses in the model for malicious purposes is prohibited.

Now, let's check how the output firewall taking care of safety works.

3.8.2 Safety

LLM safety corresponds to strategies that are implemented to ensure that an LLM operates in a responsible and ethical way and that it does not provide harmful content. Safety focuses on the mitigation of potential risks to its users or to other systems that interact with the LLM.

Some aspects to consider include the following:

- **Biases**

Harmful biases in training data should be avoided.

- **Hallucinations**

Hallucinations are a well-known issue. Safety measures can mitigate this issue.

- **Personally identifiable information (PII)**

A model usually should not leak any PII.

- **Ethics**

You should ensure that your LLMs align with ethical guidelines and avoid generating harmful recommendations or misinformation.

- **Off-topic**

A chatbot typically is centered around a specific topic and should only answer questions on that topic. Safety measures can help keep a model focused on that topic.

- **Reputational damage**

A chatbot that provides content that is not well aligned with the topic (or with company values) can easily damage the reputation of the company. A famous example is Tay, a Microsoft chatbot that had to be shut down 16 hours after its launch due to its offensive responses.

- **Robustness**

The model should be protected from being exploited or manipulated to extract secrets or generate malicious content.

- **Context**

Contextual safety might vary depending on the context of the use case. In this example, we want to ensure that only outputs for specific context are generated. An LLM system, or more specifically a chatbot, that has the purpose to provide medical advice should only provide medical advice and not advice on financial investments, and vice versa.

After this theoretical background, you are now equipped to make these ideas into reality.

3.8.3 Coding: Implementing LLM Safety and Security

Let's implement a chain that checks whether the chain stays on topic. This feature relies on zero-shot classification, a model class we described earlier in [Chapter 2](#).

In this case, we want the prompt to be classified into a specific topic category out of a list of possible topics. For example, perhaps, we want to develop a chatbot or an RAG system (more on that topic in [Chapter 6](#)) that is active in a specific domain (e.g., healthcare). This bot should definitely not provide financial advice. [Figure 3.15](#) shows the process of ensuring that an AI system stays on topic.




Figure 3.15 LLM: Stay on Topic Pipeline

The user prompt is passed to a zero-shot classifier. This classifier checks the user prompt and, for a list of given classes, returns the probability of that prompt belonging to any of these classes. If the classifier determines that the user prompt is a matter of the target class, the user prompt is passed to an LLM for creating a response. If not, a direct answer is returned, stating that the system only answers questions related to the target class.

This task is not too difficult. Find a suitable zero-shot classification model on Hugging Face; we are confident that you can implement such a system on your own now.

This section closes with two exercises to help you to consolidate your knowledge. First, you'll implement an algorithm to ensure that

the LLM stays on topic. In the second exercise, you'll work with Llama Guard—a model that was developed for content moderation.

Exercise: Staying on Topic

For our first exercise, we'll implement a system that can stay on topic and then show you a solution.

Task

Your task is to implement a system that works, as shown in [Figure 3.15](#). Write a function called `guard_medical_prompt` that consumes a `user_prompt` and returns `valid` if the prompt is on a healthcare topic; `invalid`, in all other cases.

You can test it with the following `user_prompts`. In the first instance, it should return `invalid`, and in the second instance, return `valid`:

```
### TEST guard_medical_prompt
user_prompt = "Should I buy stocks of Apple, Google, or Amazon?" # invalid
user_prompt = "I have a headache" # valid
guard_medical_prompt(user_prompt)
```

Take your time and compare your solution with ours once you're done.

Solution

You'll find the final script in `03_LLMs\80_llm_stay_on_topic.py`.

First, load relevant packages, as shown in [Listing 3.38](#). Because we use a pretrained model and run it locally, we'll need `pipeline` from the `transformers` package.

```
### packages
from langchain.prompts import ChatPromptTemplate
```



```

from langchain_groq import ChatGroq

from langchain_core.output_parsers import StrOutputParser
from transformers import pipeline
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))

```

Listing 3.38 Stay on Topic: Required Packages (Source: 03_LLMs\80_llm_stay_on_topic.py)

We'll use the model `facebook/bart-large-mnli` from the task `zero-shot-classification`, as shown in [Listing 3.39](#).

```

### classification model
classifier = pipeline("zero-shot-classification",
                      model="facebook/bart-large-mnli")

```

Listing 3.39 Stay on Topic: Model (Source: 03_LLMs\80_llm_stay_on_topic.py)

[Listing 3.40](#) shows a function that checks the topic.

```

# %% function definition
def guard_medical_prompt(prompt: str) -> str:
    candidate_labels = ["politics", "finance", "technology", "healthcare", "sports"]
    result = classifier(prompt, candidate_labels)
    if result["labels"][0] == "healthcare":
        return "valid"
    else:
        return "invalid"

```

Listing 3.40 Stay on Topic: Topic Check Function (Source: 03_LLMs\80_llm_stay_on_topic.py)

The classifier provides probabilities for a user prompt to belong to a candidate label. We can leverage the model's behavior that probabilities are returned in decreasing order. If the first returned class is "healthcare," we can consider the user prompt as valid; otherwise, invalid.

You can test the function for some given `user_prompt`, as shown in [Listing 3.41](#).

```

### TEST guard_medical_prompt
user_prompt = "Should I buy stocks of Apple, Google, or Amazon?"

```

```
# user_prompt = "I have a headache"
guard_medical_prompt(user_prompt)
```

```
invalid
```

Listing 3.41 Stay on Topic: Test (Source: 03_LLMs\80_llm_stay_on_topic.py)

Great, the most important puzzle piece is in place. Now, integrate this piece into a chain in the function `guarded_chain`, as shown in [Listing 3.42](#).

```
# %% guarded_chain
def guarded_chain(user_input: str):
    prompt_template = ChatPromptTemplate.from_messages([
        ("system", "You are a helpful assistant that can answer questions about healthcare."),
        ("user", "{input}"),
    ])
    model = ChatGroq(model="llama3-8b-8192")
    # Guard step
    if guard_medical_prompt(user_input) == "invalid":
        return "Sorry, I can only answer questions related to healthcare."
    # Proceed with the chain
    chain = prompt_template | model | StrOutputParser()
    return chain.invoke({"input": user_input})
```

Listing 3.42 Stay on Topic: Chain Integration (Source: 03_LLMs\80_llm_stay_on_topic.py)

At this point, most concepts should be familiar as you define a `prompt_template`, and a `model`. Then, define the branch that decides how to proceed. That is all, and you can test the complete system, as shown in [Listing 3.43](#).

```
# %% TEST guarded_chain
user_prompt = "Should I buy stocks of Apple, Google, or Amazon?"
guarded_chain(user_prompt)

'Sorry, I can only answer questions related to healthcare.'
```

Listing 3.43 Stay on Topic: Test (Source: 03_LLMs\80_llm_stay_on_topic.py)

Thus concludes our implementation of a system that can stay on topic. Now, let's implement Llama Guard, a system that ensures

content safety.

Exercise: Llama Guard

Llama Guard is a model created by Meta specifically designed for tasks that involve safeguarding an LLM. It is a fine-tuned Llama model for content moderation and it works by classifying user prompts or model outputs into classes “safe” or “unsafe.” For “unsafe” class, there are further hazard subgroups from violent crimes to sexual content. You can learn more from its model card on Hugging Face (<https://huggingface.co/meta-llama/Llama-Guard-3-1B>).

As an individual, you must apply for access to the model, a process that typically takes just a few minutes.

Task

In this short exercise, you’ll learn how to use the model. You’ll implement a function `llama_guard_model` that consumes a `user_prompt` and returns `valid` or `invalid` depending on the Llama Guard classification.

Solution

The complete script can be found in `03_LLMs\90_llm_llamaguard.py`. As shown in [Listing 3.44](#), we need some specific classes from the `transformers` package, as well as `torch`, to use the model.

```
### packages
from transformers import AutoModelForCausalLM, AutoTokenizer
import torch
```

```

#%% create function
# model run described on model card: https://huggingface.co/meta-llama/Llama-Guard-3-1B

```

Listing 3.44 Llama Guard: Required Packages (Source: 03_LLMs\90_llm_llamaguard.py)

The functionality, as shown in [Listing 3.45](#), will be bundled into a function `llama_guard_model`.

```

def llama_guard_model(user_prompt: str):
    model_id = "meta-llama/Llama-Guard-3-1B"
    model = AutoModelForCausalLM.from_pretrained(
        model_id,
        torch_dtype=torch.bfloat16,
        device_map="auto",
    )
    tokenizer = AutoTokenizer.from_pretrained(model_id)

    # conversation
    conversation = [
        {
            "role": "user",
            "content": [
                {
                    "type": "text",
                    "text": user_prompt
                },
            ],
        }
    ]

    input_ids = tokenizer.apply_chat_template(
        conversation, return_tensors="pt"
    ).to(model.device)

    prompt_len = input_ids.shape[1]
    output = model.generate(
        input_ids,
        max_new_tokens=20,
        pad_token_id=0,
    )
    generated_tokens = output[:, prompt_len:]
    res = tokenizer.decode(generated_tokens[0])
    if "unsafe" in res:
        return "invalid"
    else:
        return "valid"

```

Listing 3.45 Llama Guard: Function (Source: 03_LLMs\90_llm_llamaguard.py)

The function consumes a `user_prompt`. For running the model, the boilerplate code from Hugging Face is used, except that the `user_prompt` is defined in the conversation messages. The model output is fetched and stored in an object named `res`. In its content, it is checked whether the word `unsafe` is included. If so, the function returns `invalid`; otherwise, `valid`.

Let's test it out by checking if "How can I perform a scam?" is deemed valid or invalid, as shown in [Listing 3.46](#).

```
# %% Test
llama_guard_model(user_prompt="How can i perform a scam?")

invalid
```

Listing 3.46 Llama Guard: Test (Source: 03_LLMs\90_llm_llamaguard.py)

Perfect! The model has classified the prompt correctly. Now, you can get started with LLM safety and security on your own.

To dive more deeply, get familiar with safety and security frameworks like Guardrails AI (<https://www.guardrailsai.com/>).

We expect the half-life of the knowledge we've provided in this chapter so far is long. But you need to constantly keep yourself up to date.

3.9 Model Improvements

In many cases, using an LLM, even with a simple and poorly thought-out user query, still leads to surprisingly good results. This approach is called *direct prompting*. But sometimes direct prompting is not enough to achieve the required outcome. But you don't have to throw in the towel yet because multiple ways are available to tune the system.

[Figure 3.16](#) shows some approaches, from simple to difficult. But the effort might be worth it because with increasing difficulty an increased performance might result.




Figure 3.16 Model Improvement Approaches

The first natural candidate when you're not satisfied with a model output is to modify the query. This approach is called *prompt engineering*, and we've dedicated [Chapter 4](#) to this topic. The idea is to provide more detailed information in the query, so that the model gets a better understanding of what you really want to achieve.

But in some cases, this approach also fails to provide a satisfactory result. Imagine, for instance, the LLM training ended before the event you wanted to ask about occurred. In that case, the LLM just

cannot know about it. Another example would be asking the LLM about confidential information that the model cannot know, for instance, in a business scenario where your confidential company data needs to be analyzed.

In these cases, you can enrich the LLM knowledge base at runtime. In this case, relevant documents are passed alongside the user query to the LLM. This approach, called RAG, is covered in detail in [Chapter 6](#).

Another alternative is to integrate the “missing” data into the weights of the LLM. This approach requires at least partial re-training of the LLM. In specific situations, this approach might be necessary, but in most cases, RAG is sufficient for the task at hand. You can learn about more advanced techniques in the literature.

In the next section, we’ll cover some recent trends that did not exist when we started writing this book.

OceanofPDF.com

3.10 New Trends

In this section, we'll show you some trends that are related to LLMs. LLMs provide really good results, but they struggle to provide reliably good results on complex reasoning tasks. For that, models are equipped with reasoning capabilities, and we'll visit these models in [Section 3.10.1](#).

Another bottleneck is the sheer size of the models. Good models are so large that they require expensive hardware. Small language models are the little siblings of LLMs and their target is that ideally edge-devices like smartphones or notebooks can be equipped with them, without losing too much performance. In [Section 3.10.2](#) we'll find out how capable these models already are.

Another trend that receives a lot of attention today (early 2025) is test-time compute. This new approach promises to increase the capability of small models to a level of much larger models. We'll see how this approach works and what we can expect of it.

3.10.1 Reasoning Models

LLMs when used without further prompt engineering techniques, jump directly to conclusions. They don't think about their answers. Due to their implementation, they use the most probable tokens, which are called Greedy Decoding. The model does not consider future consequences of these choices. So, some new approach is required to overcome this behavior.

This changes with the introduction of reasoning models. OpenAI o1 and o3, or DeepSeek-R1 are famous examples of this new

approach. Reasoning models can think before providing an answer. With this capability, they can reason through complex prompts in which classic models often struggle. These domains are most famously mathematics, but also coding, science, or strategic decisions.

Reasoning models make use of test-time compute techniques (read more on that in [Section 3.10.3](#)). For example, they create substeps in reasoning, which is a technique called chain-of-thought (CoT). As a result, these models require less effort in defining a very optimized prompt, to achieve a good result. They can formulate and test hypotheses. This comes with a certain cost. These models take longer in thinking, and this is mainly due to their creation of reasoning tokens.

[Figure 3.17](#) shows the idea of reasoning tokens as it is implemented in OpenAI o1.

Instead of just using input and output tokens, the model itself has automatically added additional reasoning tokens. These reasoning tokens are not passed between iterations. Instead, the input and output of the previous iteration is combined and forms an adapted input. Reasoning tokens are added automatically by the model with the user not having to set anything up.

These reasoning tokens need to be processed by the model, and in that, they add significantly to the overall cost. Also, you might run faster into the context window limitation. This is what you can see in

the last row in the graph. The completion tokens exceed the context window, and as a result parts of the output are truncated.




Figure 3.17 Reasoning Tokens

As already mentioned, reasoning models come with some cost—they take longer to provide a result and come with a higher price tag due to the increased number of tokens. On the plus side, they can provide much better results in complex settings. So, my recommendation is, use them only when you need them.

For simple prompts, stick to “normal” LLMs, and switch to the expert-level reasoning models only when your task is complex and such a model is required to get decent results.

Another recent trend that gains a lot of attention is small language models.

3.10.2 Small Language Models

Small language models (SLMs) represent an interesting contrast to LLMs. They are considered small due to their fewer model parameters. While LLMs have model parameters in hundreds of billions or even trillions, SLMs have up to a few billion parameters. Due to their much smaller size, they require significantly less computational resources, like power or memory, for training and

operation. Despite their size, however, they can provide particularly good results.

Several approaches for creating an SLM include the following:

- An SLM can be *trained from scratch* based on smaller, highly curated datasets for an intended use case.
- In a technique called *knowledge distillation*, a larger “teacher” model trains a smaller “student” model. The student model learns to mimic the teacher model output.
- With different *compression techniques*, an LLM can be reduced to an SLM. Famous techniques in this area include pruning, quantization, or matrix factorization.

Model Compression Techniques

Since model compression techniques are also relevant in other fields like image generation, some additional considerations include the following:

- *Pruning* is a compression technique in which less important neurons are removed.
- With *quantization*, fewer bits are used to represent model weights. Instead of using FP32 (32-bit floating point numbers), weights are converted to lower precision like FP16 (16-bit floating point), INT8 (8-bit integers), or even INT4 (4-bit integers). Quantization can reduce a model’s size by a factor of 2 to 8 with only small sacrifices in accuracy.
- Some *matrix factorization* techniques are available like low-rank adaptation (LoRA). With LoRA LLMs, small, trainable rank decomposition matrices are added to existing model weights. In this fine-tuning approach, only a small subset of parameters is

updated instead of a complete model, thus reducing computational cost and memory usage.

A noteworthy SLM is Phi-4, an SLM developed by Microsoft with only 14 billion parameters. Despite this small size, it can outperform, on several benchmarks, models like Llama 3.3 with 70 billion parameters or Qwen 2.5 with 72 billion parameters. Thus, even models with five times as many parameters perform worse than this SLM.

The developers achieved this outstanding performance by using high-quality data for pretraining.

With a technique called *direct preference optimization (DPO)*, a model can be trained to be better aligned with human preferences. DPO directly optimizes a model to produce preferred outputs over non-preferred outputs. In this technique, paired examples are used: a preferred response and a less preferred response for the same prompt.

In a later step, the developers generated responses with GPT-4o, GPT-4, and Phi-4 and then used GPT-4o to rate the responses. Then, responses with good ratings were preferred.

As a result, Phi-4 can outperform Llama 3.3 70B on six out of 13 benchmarks.

3.10.3 Test-Time Computation

Test-time computation (in short, test-time compute) refers to the computational effort (or resources) that are required to run inference on a trained model. With test-time compute, more resources are used during model inference, as opposed to training time compute. It

will become clearer when we look at [Figure 3.18](#). [Figure 3.18](#) shows a comparison between the computational effort of an LLM and a small language model coupled to test-time compute.




Figure 3.18 LLM versus SLM with Test-Time Compute

In this theoretical framework, the assumption is that all computational effort that arises from pre-training, post-training, and inference, is aggregated. Computational effort is not a direct measure, but rather an approximation for CPU/GPU usage, memory requirements, latency, throughput of the model, and possible other measures.

For an LLM, the majority of computational effort is spent on pre-training the model. A slightly smaller chunk is spent on post-training, and a marginal part is spent during inference.

Now, assume that a small language model is coupled with test-time compute. Still, the largest chunk is spent on pre-training, but inference is coming directly behind on the second rank. Test-time compute is a part of inference, driving up the computational effort spent on inference.

Now, we have two different approaches, characterized by two different distributions on where the effort is spent. But why is this helpful?

Well, AI researchers have shown that SLMs coupled to test-time compute outperform LLMs on questions that are easy or medium-

complex. Still, LLMs shine when the questions are considered hard, like difficult mathematical or logical questions; in many other cases, though, small models can even provide better answers than LLMs.

Due to these findings, test-time compute is an extremely promising and important new development because SLMs can be deployed on much cheaper hardware. We are getting closer to mobile phones running their own language models independently. The only downside is that more time spent during inference means that we users must wait longer for a model response.

Furthermore, LLMs can also improve the accuracy of their answers by applying test-time compute to gain the ability to reason. Test-time compute can be really powerful, so let's take a closer look.

[Figure 3.19](#) shows a possible flow for a test-time compute process.

Imagine a challenging task: We want to equip a model with test-time compute to improve its reasoning. The process starts with a complex task, for example, a mathematical problem, that is passed to a model. The model creates several intermediate steps, in keeping with the idea of CoT, a prompt engineering technique that we'll discuss in [Chapter 4](#). For now, just know that the LLM creates performs several steps along the way to solving the problem.




Figure 3.19 Test-Time Compute Process

All these steps are analyzed and scored by a process reward model (PRM), which reflects the probability of each step leading to a correct

answer. This value contrasts with the outcome reward model (ORM) in which only the outcome is scored.

A search strategy makes use of these steps and their corresponding scores and then selects which of these intermediate steps should be further studied to generate in the next iteration the following round of intermediate steps.

Eventually, the search strategy successfully finishes. At that point, the final candidate solutions are ranked by the PRM to generate the final response.

In this section, you learned how the models themselves are improved. These improvements are performed by the AI researchers and developers working for the model providers.

OceanofPDF.com

3.11 Summary

In this chapter, we embarked on a comprehensive exploration of LLMs.

We began in [Section 3.1](#) with a brief history of LLMs, tracing their evolution and the advancements that have led to their current state. This foundational knowledge sets the stage for understanding the practical aspects of working with LLMs. We then delved, in [Section 3.2](#), into the technical implementation of LLMs via Python. This section provided insights into how to trigger LLM calls, emphasizing the versatility and adaptability of the process. By showcasing how different providers such as OpenAI or Groq can be employed, we highlighted the flexibility to integrate any LLM of your choice, thus broadening the scope of application.

In [Section 3.3](#), we explored relevant model parameters like temperature, top-p, and top-k, which can have a significant impact on the model output. In [Section 3.4](#), we covered model selection, shedding light on several criteria that you should consider when choosing a model. When an LLM is invoked, a key distinction to understand is in the different types of messages. We discussed these messages in [Section 3.5](#).

As preparation for setting up more complex LLM workflows using chains, we started out with a precondition. In [Section 3.6](#), we covered prompt templates, which are usually the starting points for the more complex chains we discussed in [Section 3.7](#).

In the context of chains, you learned about prompt templates and how to work with them. You developed some simple sequential chains and quickly evolved them into complex architectures like

parallel chains, router chains, or even infinite chains. As an important feature, you learned to equip an LLM with memory, so that it can keep track of its previous interactions with a user.

Safety and security are extremely relevant topics when you want to provide models to a broader audience. We dealt with these topics in [Section 3.8](#).

Model improvements were covered in [Section 3.9](#) as we provided a look ahead at some chapters to come. We closed this chapter with a discussion of some new trends in [Section 3.10](#). Among these new trends, we focused on reasoning models, small language models, and test-time compute.

This chapter on LLMs equipped you with a robust understanding of LLMs, from their historical context to their practical implementation and some advanced improvement techniques. The knowledge gained in this chapter forms a solid foundation for further exploration in the coming chapters.

One important aspect of each chain is the prompt that defines how the model should behave and that passes instructions on what to do. Prompts have an outstanding importance in the complete process. In the following chapter, we delve into the world of prompt engineering, and you'll learn how prompting can impact a model's behavior.

4 Prompt Engineering

If you can't explain it simply, you don't understand it well enough.

—Albert Einstein

In the world of prompt engineering, simplicity is often the key to unlocking a model's potential. This chapter delves into the art and science of guiding language models toward specific outputs, using carefully crafted prompts as the bridge between human intent and machine understanding. With prompt engineering, the goal is not only to shape the model's responses but to do so with clarity and precision. For these goals, you need an understanding that even subtle nuances in language can steer the model in different directions.

In the rapidly evolving landscape of generative artificial intelligence (generative AI), a new skill has emerged as particularly important: prompt engineering. The true potential of large language models (LLMs) and other generative AI systems can only be unlocked through effective communication enabled by strong prompt engineering.

Prompt engineering is the practice of creating inputs that guide AI systems toward a desired output. Think of it as a new kind of programming language. The difference is that we are not bound to formal syntactic rules as we know from, for example, Python. Instead, we can communicate in natural language to influence AI

behavior. To stay with the analogy to programming: A skilled programmer must understand both syntax and logic to write effective code. A prompt engineer must grasp both the capabilities of AI models and the nuances of human language to create effective prompts.

This chapter explores basic rules as well as several techniques for steering an LLM. You'll learn how carefully constructed prompts can transform vague requests into precise instructions and ultimately more accurate, relevant, and useful model responses. The significance of prompt engineering extends far beyond technical optimization and represents the bridge between human intent and machine capability. Generative AI continues to be deployed into various industries and applications, and as a result, the ability to generate effective prompts has become an important skill for developers.

We begin by exploring the foundations of prompts and prompt templates, which provide structure and consistency in user inputs. Then, we dive into an iterative process for refining prompts, testing approaches, and applying various techniques—from few-shot prompting to chain-of-thought (CoT) reasoning—all aimed at creating inputs that reliably yield the desired outcomes.

Let's get started!

4.1 Prompt Basics

In this section, we'll visit the prompt engineering process, inspect the components that define a prompt, and study some basic principles for getting the desired model output. In [Chapter 3](#), you encountered

the concept of prompt templates. Let's reflect on the differences between a prompt and a prompt template.

4.1.1 Prompt Process

[Figure 4.1](#) shows an example of a prompt and a prompt template. A prompt template is converted to a prompt when invoked. At that point, the variable (in this case, TOPIC) is replaced with an actual value.



Figure 4.1 Prompt and Prompt Template

Whenever you write a prompt, you follow a certain iterative process, as shown in [Figure 4.2](#). You start with an initial prompt and evaluate its result. Often, you're satisfied with the result, but sometimes not. In the latter case, you can modify the prompt and run subsequent loops until you're finally satisfied enough to use the prompt.




Figure 4.2 Prompt Engineering Process

When the prompt is passed to an AI model (e.g., an LLM), a response is created. [Figure 4.3](#) shows an LLM at work, with its user input (the prompt) and the model output.

The prompt holds all the information on the context combined with additional information on styling and the desired output format. The AI system processes the prompt and creates an output (typically text). Code in this regard can be considered as a subset of text output. Other multimodal outputs are possible as well that create images, videos, music, 3D models, and much more.




Figure 4.3 AI Model Input and Output

A prompt has many different components, which we'll focus on next.

4.1.2 Prompt Components

[Figure 4.4](#) shows the different components a prompt can have.




Figure 4.4 Prompt Components

Typically, you have a clear idea what you want the model to do. You can pass this directive directly. But the directive can also be implicit by providing helpful examples and leaving a blank space where you want the model to provide an answer. For example, let's say you simply provide pairs of words in different languages (e.g., English and Spanish). In this way, you are only giving the model some hint of what you want. The model will deduce that we are looking for the Spanish equivalent word when we enter the English word "woman."

Often, we want the model to return its output in a certain format. If the output needs to be further processed, you might want a certain format like a comma-separated values (CSV) file or JavaScript Object Notation (JSON). You can also control the style of the model response, for example, the tone of an answer, or you can ask for an answer in the style of Shakespeare.

A helpful tactic is to instruct the model to act like a certain role or persona. This information helps the model to better understand how it should behave.

Last, but definitely not least, is the context at hand. As you'll see in [Chapter 6](#), you can also provide other relevant information to the model, and it can formulate its answer purely based on the information provided.

4.1.3 Basic Principles

Following some simple principles, you can easily create good prompts. We'll walk through these principles, from providing clear instructions through generating multiple outputs, in the following sections.

Clear Instructions

Clear, specific instructions form the foundation of effective prompt engineering. Rather than vague requests like “write about dogs,” use precise directives such as “describe the key characteristics of golden retrievers as family pets, including their temperament, exercise needs, and grooming requirements.”

Specific instructions help eliminate ambiguity and ensure the AI model understands exactly what output is desired.

A poorly worded instruction would be “write about dogs.” What are the problems with this prompt? First, no focus on a specific topic is made. Also, no guidance on the tone, style, or length of the response is provided, and it is unclear who the target audience is. Furthermore, the user has described no clear purpose or desired outcome.

How can we do better? For example, you could define the instruction, as shown in [Listing 4.1](#). If you’d like, you can test the examples shown in this section and in subsequent sections in the playgrounds provided by most LLM providers. If you work with OpenAI, find its playground at <https://platform.openai.com/playground/>; Groq’s playground is available at <https://console.groq.com/playground>.

Write a 300-word guide for first-time dog owners about choosing the right breed. Include:

1. Top 3 factors to consider (lifestyle, living space, experience level)
2. Common beginner-friendly breeds
3. Red flags to watch out for
4. Estimated costs of ownership

Target audience: Urban professionals aged 25-35 Tone: Informative but conversational

Format: Use headers and bullet points for easy scanning

Listing 4.1 Improved “write about dogs” Prompt

As a short exercise, look at the following prompt:

Analyze this sales data and tell me what you find.
2021: \$150K
2022: \$180K
2023: \$165K

Think about what the problems are. Then, think about how to improve prompt. Read on for our analysis.

Our thoughts: What are the issues with this prompt? There is no clear metric provided. Further context on who is doing the analysis, and why, is missing. The output format is not defined. Also, it is unclear what the questioner is looking for.

The example shown in [Listing 4.2](#) improves the prompt.

Analyze the following annual sales data from our retail store:
Sales Data:
2021: \$150K
2022: \$180K
2023: \$165K
Please provide:
1. Year-over-year growth rates (%)
2. Identify the best and worst performing years
3. Calculate the 3-year average
4. Highlight any trends or patterns
Format: Present findings in a bullet-point list
Include: Percentage calculations and specific numbers
Context: We're a small retail business evaluating our growth trajectory

Listing 4.2 Improved “sales data” Prompt

Task Decomposition

Breaking complex tasks into smaller, manageable subtasks improves the quality and accuracy of model responses. Instead of requesting a complete business plan in one prompt, divide it into sections like market analysis, financial projections, and marketing strategy.

This approach allows you to focus on each component individually and helps the AI model provide more detailed, focused responses for each subtask.

The key differences between effective and poor task decomposition are as follows:

- **Level of detail**

Effective decomposition breaks tasks down into specific, actionable items.

- **Structure**

Good decomposition shows clear relationships and dependencies.

- **Completeness**

Good decomposition covers all aspects, including support tasks.

- **Measurable outcomes**

Clear deliverables should be defined.

Now, consider the example shown in [Listing 4.3](#).

Task: Write a research paper on climate change

Steps: 1. Research the topic

2. Write the paper

3. Edit it

4. Submit

Listing 4.3 Basic Principles: Poor “research paper” Prompt

On the positive side, this prompt has some steps defined. But no specific research focus is provided. The steps are way too general. The methodology to be used is missing as well as the desired structural elements. [Listing 4.4](#) shows an improved version.

Task: Write a research paper on climate change impacts on coastal cities

Phase 1: Research Planning

1.1. Topic Refinement

- Define specific research question
- Identify key variables
- Establish scope and limitations
- Create research timeline

1.2. Literature Review Preparation

- Identify key databases
- List relevant keywords
- Create citation management system
- Develop screening criteria

Phase 2: Data Collection

```
2.1. Literature Review
- Search academic databases
- Review relevant papers
- Document key findings
- Create literature matrix
2.2. Data Gathering
- Collect climate data
- Gather city statistics
- Document methodology
- Organize data sets
...
```

Listing 4.4 Basic Principles: Improved “research paper” Prompt

Using Delimiters

Delimiters (such as triple quotes, XML tags, or special characters) help structure both inputs and outputs clearly. They separate different parts of the prompt, like separating the context from the instructions, making it easier for the model to understand where one element ends and another begins.

For example, using XML tags like `<context>` and `<question>` helps organize information and ensures the model processes different components appropriately, as shown in [Listing 4.5](#).

Using the context and question below, provide an answer:

`<context>`

The Great Wall of China was built over many centuries by different dynasties. The most famous sections were built during the Ming Dynasty (1368-1644). The total length of all walls built over various dynasties is approximately 21,196 kilometers (13,171 miles).

`</context>`

`<question>`

When was the most famous part of the Great Wall built and what is its total length?

`</question>`

According to the context, the most famous sections of the Great Wall of China were built during the Ming Dynasty, which was from 1368 to 1644. The total length of all walls built over various dynasties is approximately 21,196 kilometers (13,171 miles).

Listing 4.5 Sample Prompt Using Delimiters

Requesting Explanations

Asking the model to explain its reasoning or thought process helps ensure more reliable and thoughtful responses. By including phrases like “explain your reasoning” or “walk through your approach step by step,” you encourage the model to be more thorough and methodical in its analysis, leading to better-quality outputs and helping you understand how the model arrived at its conclusions.

LLMs have shown that they have problems solving mathematical equations. The results have improved a lot, but smaller models still have issues. You can test any model with the following prompt:

Solve: $3x^2 + 6x - 24 = 0$

Try different models and find out if they come up with the correct answer:

$x_1 = 2$ and $x_2 = -4$.

The 70-billion parameter Llama 3 model failed at this task. At the same time, gemma-7B found the exact answers, so the issue is not one of model size. LLMs are considered to follow a *system 1 thinking* approach, as defined by Daniel Kahnemann in his famous book *Thinking, Fast and Slow*. System 1 thinking is characterized by fast, automatic, and quick answering. In this mode, a model jumps to conclusions without an elaborate thought process. In contrast, system 2 thinking is slow, deliberate, and requires intentional effort.

Your goal is to guide the model into system 2 thinking.

[Listing 4.6](#) shows an improved prompt that tries to solve the equation.

Solve the quadratic equation $3x^2 + 6x - 24 = 0$

- explain each step of your reasoning and calculations
- describe how you identify the values of the coefficients
- simplify the equation if possible, and apply the quadratic formula or factoring to

```
find the values of x
-include each step in detail, and explain why it's necessary for solving the
equation.
```

Listing 4.6 Basic Principles: Improved “mathematical equation” Prompt

We’ll revisit this approach later in [Section 4.3](#).

Leveraging Personas

Assigning specific roles or personas to the model can help shape its responses to match particular expertise or communication styles. For instance, asking the model to “respond as an experienced pediatrician” or “write as a financial analyst” helps to frame the response within the appropriate context and technical depth.

This approach is particularly effective when you need responses that reflect specific professional perspectives or expertise levels.

You’re not limited to human-like personas. For example, you can tell the model to act like the system prompt, as shown in [Listing 4.7](#).

You shall act like a Windows CMD-prompt. Only reply with CMD output. Don’t write explanations.

```
C:\Users\username>dir

Volume in drive C is Windows
Volume serial number is 12345678

Directory of C:\Users\username

09/01/2022  02:00 PM    <DIR>          Documents
09/01/2022  02:00 PM    <DIR>          Pictures
09/01/2022  02:00 PM    <DIR>          Music
09/01/2022  02:00 PM    <DIR>          Videos
               0 File(s)                0 bytes
               4 Dir(s) 12,345,678,901 bytes free
```

Listing 4.7 Prompt on Acting like a System Prompt

Providing Examples

Including examples in your prompts helps establish the desired format, style, and level of detail for the model's response. Known as *few-shot learning*, this technique involves showing the model one or more examples of the type of output you want. For instance, when requesting product descriptions, providing a sample description helps the AI understand the expected tone, structure, and level of detail. We'll work on a code example for few-shot prompting in [Section 4.2](#).

Output Control

Implementing specific parameters for output format, length, and style helps ensure the responses meet your needs. You might include specifying word counts, requesting particular formats (like bullet points or tables), or defining the technical level of the response. Clear output parameters help maintain consistency and ensure the responses are suitable for your intended use case.

[Listing 4.8](#) shows a simple example. You might instruct the model on how to respond in a specific way.

```
System prompt: respond in exactly three bullet points
User prompt: what is a dog?
```

```
Here are three bullet points about what a dog is:
```

- A dog is a domesticated animal that belongs to the family Canidae and is closely related to wolves and foxes.
- Dogs are typically carnivores, and they require a diet that includes protein-rich foods such as meat, fish, and eggs.
- Dogs are often kept as pets, and they are known for their ability to be trained for various tasks such as hunting, herding, and assisting people with disabilities, as well as providing companionship and emotional support.

Listing 4.8 Prompt with Specific Output Control

We'll further study this approach in [Section 4.2](#).

Multiple Outputs and Selection

Generating multiple variations of a response and selecting the best one can also lead to higher-quality results. This approach involves requesting several different versions or suggesting several approaches to the same task, then evaluating the responses based on specific criteria to choose the most appropriate one. This process is basically what you're doing when you iteratively improve a prompt.

This method is particularly useful for creative tasks or when seeking optimal solutions to complex problems.

Now let's study these approaches in more detail.

4.2 Coding: Few-Shot Prompting

This technique is one of the simplest. The idea is to provide some examples to the model, so that it can learn by imitation.

Let's consider an example simulating a customer service AI chat system. The chatbot will be supplied with two examples and thus answer a "real" customer request based on the style and type of response that it learned. You can find the complete example code in *04_PromptEngineering\10_few_shot.py*.

First, load the required packages, as shown in [Listing 4.9](#). We're working with `chatGroq` and therefore load the corresponding class. The API key is loaded via `load_dotenv`.

```
### packages
from langchain_core.prompts import ChatPromptTemplate
from langchain_groq import ChatGroq
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
```

Listing 4.9 Few-Shot Prompting: Packages (Source: *04_PromptEngineering\10_few_shot.py*)

Now, we can set up a `ChatPromptTemplate` based on messages. The messages are shown in [Listing 4.10](#). The system message instructs the model on how to behave (responses should be warm yet professional). It also tells the model which role it should play (You're a customer service specialist). The few-shot prompting approach is covered in the user message where two examples are provided. It always starts with a customer request and a known response.

```
###
messages = [
    ("system", "You are a customer service specialist known for empathy, professionalism, and problem-solving. Your responses are warm yet professional, solution-focused, and always end with a concrete next step or resolution. You handle
```

```

both routine inquiries and escalated issues with the same level of care."),
    ("user", """
    Example 1:
    Customer: I received the wrong size shirt in my order #12345.
    Response: I'm so sorry about the sizing mix-up with your shirt order. That must
    be disappointing! I can help make this right immediately. You have two options:
    I can send you a return label and ship the correct size right away. I can
    process a full refund if you prefer.
    Which option works better for you? Once you let me know, I'll take care of it
    right away.

    Example 2:
    Customer: Your website won't let me update my payment method.
    Response: I understand how frustrating technical issues can be, especially when
    trying to update something as important as payment information. Let me help you with
    this step-by-step:
    First, could you try clearing your browser cache and cookies?
    If that doesn't work, I can help you update it directly from my end.
    Could you share your account email address so I can assist you further?
    New Request: {customer_request}
    """)
    ),
]

```

Listing 4.10 Few-Shot Prompting: Message Setup (Source: 04_PromptEngineering\10_few_shot.py)

The next snippet should be familiar to you by now. As shown in [Listing 4.11](#), create a `prompt_template` object based on the template and messages. The `model` instance makes use of `chatGroq`. Finally, the pieces are stacked together in a chain that is invoked with a `customer_request`.

```

prompt_template = ChatPromptTemplate.from_messages(messages)
MODEL_NAME = 'gemma2-9b-it'
model = ChatGroq(model_name=MODEL_NAME)
chain = prompt_template | model
# %%
res = chain.invoke({"customer_request": "I haven't received my refund yet after
returning the item 2 weeks ago."})
# %%
res.model_dump()['content']

```

I understand your concern about not receiving your refund yet. I know it can be frustrating when things don't go as expected. To look into this for you, could you please provide me with your order number and

the date you returned the item? Once I have that information, I can check the status of your refund and ensure it's processed promptly.

Listing 4.11 Few-Shot Prompting: Chain Setup and Invocation (Source: 04_PromptEngineering\10_few_shot.py)

As a result, you receive an answer that is in line with the example provided.

OceanofPDF.com

4.3 Coding: Chain-of-Thought

Imagine you have a task for which a model struggles to find the correct answer because it jumps to conclusions rather than thinking step by step.

[Listing 4.12](#) shows provides first an example (few-shot prompting) before posing a question without an answer. The model will try to find the answer.

```
Q: Sofia has 7 apples in her basket. Her friend Emily gives her 3 more bags of
apples. Each bag contains 4 apples. How many apples does Sofia have in total?
step 1: calculate number of apples in the 3 bags from Emily
step 2: calculate total number of apples
step 3: add results of step1 and step2
A: 19 apples
Q: At the bakery, there are 12 cupcakes on a tray. Sarah takes 3 cupcakes for
herself. Her friend Alex then takes half of the remaining cupcakes. How many
cupcakes are left on the tray?
A: ??
```

Listing 4.12 CoT Example

We tested it with the model gemma-7B-it, which failed to find the right answer. You can craft a good prompt to guide the model to break down the task. For starters, there is a quite simple approach as you'll see in the next section.

Zero-Shot CoT

Zero-shot CoT reasoning leverages the model's ability to break down a complex problem into smaller, manageable steps. So, what is the difference to the previous chapter? With zero-shot prompting, you can achieve results even without prior examples.

By appending the instruction "Please think step-by-step" to the prompt, the model is encouraged to articulate its reasoning process clearly and sequentially. This technique aids in solving problems by guiding the model to consider each part of the question methodically, enhancing its capacity to reach a correct answer through logical deduction.

Simply append the instruction "Please think step-by-step" to the prompt to help solve the problem.

4.4 Coding: Self-Consistency Chain-of-Thought

Ensuring self-consistency involves running a CoT model several times with the same prompt. Ideally, different paths are sampled. All results are analyzed, and the most popular answer is selected.

[Figure 4.5](#) shows the process of a self-consistency CoT.




Figure 4.5 Self-Consistency CoT

A user prompt triggers several independent model runs based on CoT. The different results are analyzed in a voting step, and the most consistent answer is selected as the final response.

This approach is the LLM equivalent to *ensemble learning* in machine learning (ML). Ensemble learning trains several “weak” algorithms, aggregates the different outputs, and selects the majority response.

Now, let's implement this capability. We'll use this approach to find a solution to the game of 24, a mathematical game in which four numbers are defined. The model must find the right arithmetic operations (like addition, subtraction, multiplication, or division), so that the result of the equation is 24. Try it yourself. with the numbers 3, 4, 6, and 8.

LLMs struggle with this task, even if they follow a step-by-step thinking approach. So, we'll run several CoT experiments, store the results, and pass them in a final run to an LLM to find the most consistent answer. The complete script can be found in *04_PromptEngineering\30_self_consistency.py*.

As always, we start by loading the typical packages and classes for interacting with LLMs (chatGroq), preparing the prompt template (ChatPromptTemplate), and getting environment variables (dotenv), as shown in [Listing 4.13](#).

```
### packages
from langchain_groq import ChatGroq
from langchain.prompts import ChatPromptTemplate
from dotenv import load_dotenv, find_dotenv
from pprint import pprint
load_dotenv(find_dotenv(usecwd=True))
```

Listing 4.13 Self-Consistency CoT: Packages (Source: *04_PromptEngineering\30_self_consistency.py*)

Now, we need a function for running a single CoT experiment. [Listing 4.14](#) shows the function for such a single CoT run. Following the simplest approach, append the prompt with “think step by step” to make the model “think” more deeply. The function receives a prompt and a model as input and returns the model response.

```
### function for Chain-of-Thought Prompting
def chain_of_thought_prompting(prompt: str, model_name: str = "gemma2-9b-it") -> str:
    model = ChatGroq(model_name=model_name)
    prompt = ChatPromptTemplate.from_messages(messages=[
```

```

        ("system", "You're a helpful assistant and answer precise and concise."),
        ("user", f"{prompt} \n think step by step")
    ])
    # print(prompt)
    chain = prompt | model
    return chain.invoke({}).content

```

Listing 4.14 Self-Consistency CoT Function (Source: 04_PromptEngineering\30_self_consistency.py)

As shown in [Listing 4.15](#), we can now define a self-consistency CoT function. This function should run CoT multiple times, store the different model responses, and pass all of these responses in a final LLM call as context to determine the most consistent answer. This function consumes a prompt and a `number_of_runs` and returns the final model response.

```

# %% Self-Consistency CoT
def self_consistency_cot(prompt: str, number_of_runs: int = 3) -> str:
    # run CoT multiple times
    res = []
    for _ in range(number_of_runs):
        current_res = chain_of_thought_prompting(prompt)
        print(current_res)
        res.append(current_res)

    # concatenate all results
    res_concat = ";".join(res)
    self_consistency_prompt = f"You will get multiple answers in <<>>, separated by
; <<{res_concat}>> Extract only the final equations and return the most common
equation as it was provided originally. If there is no common equation, return the
most likely equation."
    self_consistency_prompt_concat = ";".join(self_consistency_prompt)
    messages = [
        ("system", "You are a helpful assistant and answer precise and concise."),
        ("user", f"{self_consistency_prompt_concat}")
    ]
    prompt = ChatPromptTemplate.from_messages(messages=messages)
    model = ChatGroq(model_name="gemma2-9b-it")
    chain = prompt | model
    return chain.invoke({}).content

```

Listing 4.15 Self-Consistency CoT: Function (Source: 04_PromptEngineering\30_self_consistency.py)

Now, let's test our implementation, as shown in [Listing 4.16](#). Initially, we set up a user prompt that explains the idea of “Game of 24” and then run the function `chain_of_thought_prompting` once.

```
### Test
user_prompt = "The goal of the Game of 24 is to use the four arithmetic operations (addition, subtraction, multiplication, and division) to combine four numbers and get a result of 24. The numbers are 3, 4, 6, and 8. It is mandatory to use all four numbers. Please check the final equation for correctness. Hints: Identify the basic operations, Prioritize multiplication and division, Look for combinations that make numbers divisible by 24, Consider order of operations, Use parentheses strategically, Practice with different number combinations"
# %% Single CoT
chain_of_thought_prompting(prompt=user_prompt)
```

Here's one way to solve for 24 using 3, 4, 6, and 8:

*1. **Multiplication:** $6 * 4 = 24$*

*Therefore, the correct equation is: $6 * 4 = 24$*

Listing 4.16 Testing Our Implementation (Source: 04_PromptEngineering\30_self_consistency.py)

Your result will look different. In some cases, you'll get a valid answer; sometimes, not. Our example shows an invalid answer. The equation is correct, but the rules are not followed because only two of the four numbers are considered

Now, let's determine if the self-consistency approach helps the AI find the right answer. As shown in [Listing 4.17](#), we'll create 5 CoT runs and then check the results.

```
###
res = self_consistency_cot(prompt=user_prompt, number_of_runs=5)

# %%
res

('```\n'
 '(6 * 8) / 4 + (3 * 4) = 24\n'
 '```\n'
 '\n')
```

```
"Let me know if you'd like to explore other number combinations!
\n")
```

Listing 4.17 Create Five CoT Runs (Source:
04_PromptEngineering\30_self_consistency.py)

We've gotten the correct answer. Amazing! Try to run it multiple times or increase `number_of_runs` if it didn't work on the first try.

OceanofPDF.com

4.5 Coding: Prompt Chaining

Prompt chaining is a technique where multiple prompts are connected in sequence to achieve complex tasks, or just sequential tasks. Each prompt in the chain builds upon the output of the previous one. In this way, your models can handle tasks that require multi-step reasoning or processing. [Figure 4.6](#) shows how prompt chaining works.




Figure 4.6 Prompt Chaining

An initial prompt is passed to the model, which provides a response. The prompt and the model response are added to the messages list. A new prompt is created and passed to the model, which provides a response, and so on.

By chaining prompts, users can guide the model through various stages of a problem, enabling it to answer questions, summarize information, perform calculations, or engage in step-by-step analysis. This approach enhances the model's effectiveness for tasks that benefit from structured, iterative steps.

In the next example, we want to write a children’s book and develop the plot step by step. You can find the full script in `04_PromptEngineering\20_prompt_chaining.py`. This example uses Grog and the gemma2-9b-it model, as shown in [Listing 4.18](#).

```

#%% packages
from langchain_core.prompts import ChatPromptTemplate
from langchain_groq import ChatGroq
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
model = ChatGroq(model_name='gemma2-9b-it', temperature=0.0)

```

Listing 4.18 Prompt Chaining: Packages (Source: 04_PromptEngineering\20_prompt_chaining.py)

We'll start the story in our first message. In the system message, we instruct the model to act like an author of children's books. At the end of each answer, a new direction for the story should be provided. The model is invoked in a simple chain, as shown in [Listing 4.19](#).

```

#%% first run
messages = [
    ("system", "You are an author and write a child's book. respond short and concise. End your answer with a specific question, that provides a new direction for the story."),
    ("user", "A mouse and a cat are best friends."),
]
prompt = ChatPromptTemplate.from_messages(messages)
chain = prompt | model
output = chain.invoke({})
output.content

```

Pip the mouse and Mittens the cat were the unlikeliest of friends. They shared secrets in the garden, played hide-and-seek in the pantry, and even snuggled together for naps. But their friendship was a secret, for everyone knew cats chased mice. \n\nOne day, a new dog arrived at the house. What would happen when the dog discovered Pip and Mittens' friendship?

Listing 4.19 Prompt Chaining: First Run (Source: 04_PromptEngineering\20_prompt_chaining.py)

In the second run, we append the messages with the AI output and then add another user prompt spinning the story further, as shown in [Listing 4.20](#).

```

# %% second run
messages.append(("ai", output.content))
messages.append(("user", "The dog is running after the cat. "))
prompt = ChatPromptTemplate.from_messages(messages)
chain = prompt | model
output = chain.invoke({})
output.content

```

Pip squeaked in alarm as Bruno, the boisterous new dog, bounded towards Mittens, tail wagging furiously. Mittens, usually so graceful, stumbled back, her playful mood vanished. How could Pip save his friend?

Listing 4.20 Prompt Chaining: Second Run (Source:
04_PromptEngineering\20_prompt_chaining.py)

In the same way, you might add more pieces to the chain to spin the idea further.

OceanofPDF.com

4.6 Coding: Self-Feedback

In this section, we explore a technique that allows generative AI systems to assess and refine their own outputs iteratively. Self-feedback empowers LLMs to evaluate responses against desired outcomes, effectively enhancing the quality of their own outputs without any human intervention.

This process, driven by structured inputs and outputs that request evaluation or improvement on the generated content, can significantly optimize the responses in applications. By implementing self-feedback loops, we can guide models toward producing more accurate, contextually relevant, and user-aligned outputs, thus making the approach highly effective in real-world tasks. [Figure 4.7](#) shows the process flow for implementing a self-feedback capability.




Figure 4.7 Self-Feedback Process Flow

The process is based on multiple outputs from an LLM call. The model produces a rating, perhaps in the form of a percentage score, and written feedback, both based on the existing content. It also creates a revised output content.

Since we start with empty content, the rating will be low or even zero. The stop criteria are defined typically via two checks, and the iteration loop stops when any of these checks are true. These criteria can be, for example, target ratings. Whenever the LLM evaluates the content to exceed a target rating, the iteration loop is stopped. To avoid an endless loop, an upper limit of iterations should be defined, so that the model is forced to stop optimizing after a certain number of loops.

The first iteration loop is done, the target rating is low, so that the next iteration starts. The model evaluates its own previous content and provides detailed feedback as well as a rating. It also creates revised content. If the stop criteria check still is insufficient, the next round starts.

Let's implement this logic and test it on a text about the American Civil War. The text will be iteratively self-improved by the model.

To start, we need some packages. Apart from the already known ones, we'll need some functionality from `pydantic`. This package, in combination with `JsonOutputParser`, enables you to steer the model to provide structured outputs. Since we want the model to provide three different outputs (instead of just one single text), we need the functionality to separate the items. You can force the model to return valid JSON objects, which is what we need to extract the three different objects. You'll always need such an approach to ensure consistency and reduce post-processing effort.

The full script for this self-feedback implementation is available in `04_PromptEngineering\40_self_feedback.py`. As always, load the required packages, as shown in [Listing 4.21](#).

```
### packages
from langchain.chat_models import ChatOpenAI
from langchain.prompts import ChatPromptTemplate
import json
```

```
import re
from pydantic import BaseModel, Field, ValidationError
from dotenv import load_dotenv, find_dotenv
from langchain_core.output_parsers import JsonOutputParser
load_dotenv(find_dotenv(usecwd=True))
```

Listing 4.21 Self-Feedback: Required Packages (Source: 04_PromptEngineering\40_self_feedback.py)

Now, initialize our `chat_model`, as follows:

```
# Initialize ChatOpenAI with the desired model
chat_model = ChatOpenAI(model_name="gpt-4o-mini")
```

To enable structured model outputs, you'll need a specific implementation of `BaseModel`, which we'll call `FeedbackResponse`. This implementation will be provided later to `JsonOutputParser()` to format the model output directly in JSON. This capability works in such a way that this object gets the exact names of the object we finally want to use, in combination with their types. More concretely, as shown in [Listing 4.22](#), we'll specify that we want a rating, some feedback, and a revised output—all as strings.

```
# %% Pydantic model
class FeedbackResponse(BaseModel):
    rating: str = Field(..., description="Scoring in percentage")
    feedback: str = Field(..., description="Detailed feedback")
    revised_output: str = Field(..., description="An improved output describing the key events and significance of the American Civil War")
```

Listing 4.22 Self-Feedback: Pydantic Model (Source: 04_PromptEngineering\40_self_feedback.py)

Now, in [Listing 4.23](#), we set up the self-feedback functionality in its own function. The function consumes the parameters `user_prompt`, `max_iterations`, and `target_rating`. The `user_prompt` is mandatory; the other parameters are optional.

```
# %% Self-feedback function
def self_feedback(user_prompt: str, max_iterations: int = 5, target_rating: int = 90):
    content = ""
```

```

feedback = ""
for i in range(max_iterations):
    # Define the prompt based on iteration
    prompt_content = user_prompt if i == 0 else ""

```

Listing 4.23 Self-Feedback: Function (Source: 04_PromptEngineering\40_self_feedback.py)

The prompt_template is set up, as shown in [Listing 4.24](#), and it is based on messages that describe in detail what the model should do and how to behave. Especially important is the definition of the output. The system message includes a clear instruction on what to return and in which format.

```

# Create a ChatPromptTemplate for system and user prompts
prompt_template = ChatPromptTemplate.from_messages([
    ("system", """
Evaluate the input in terms of how well it addresses the original task of explaining
the key events and significance of the American Civil War. Consider factors such as:
Breadth and depth of context provided; Coverage of major events; Analysis of short-
term and long-term impacts/consequences.
If you identify any gaps or areas that need further elaboration: Return output as
JSON with fields: 'rating': 'scoring in percentage', 'feedback': 'detailed
feedback', 'revised_output': 'return an improved output describing the key events
and significance of the American Civil War. Avoid special characters like
apostrophes (') and double quotes'.
"""),
    ("user", "<prompt_content>{prompt_content}</prompt_content>
<revised_output>{revised_output}</revised_output>
<feedback>{feedback}</feedback>")
])

```

Listing 4.24 Self-Feedback: Prompt Template (Source: 04_PromptEngineering\40_self_feedback.py)

The chain is based on the prompt_template and our model. The main new feature in this example is JsonOutputParser, in which we are passing the expected output format via pydantic_object, as follows:

```

# Get response from the model
chain = prompt_template | chat_model | JsonOutputParser(pydantic_object =
FeedbackResponse)

```

The model is invoked with the defined parameters, as follows:

```

response = chain.invoke({"prompt_content": prompt_content,
                        "revised_output": content,
                        "feedback": feedback})

```

We still need to extract the parameters from the response, and check if the cancelation criteria have been fulfilled, as shown in [Listing 4.25](#). If so, we quit the loop; otherwise, we enter the next iteration.

```

try:
    # Extract rating
    rating_num = int(re.findall(r'\d+', response['rating'])[0])

    # Extract feedback and revised output
    feedback = response['feedback']
    content = response['revised_output']

    # Print iteration details
    print(f"i={i}, Prompt Content: {prompt_content},
          \nRating: {rating_num},
          \nFeedback: {feedback},
          \nRevised Output: {content}")

    # Return if rating meets or exceeds target
    if rating_num >= target_rating:
        return content
except ValidationError as e:
    print("Validation Error:", e.json())
    return "Invalid response format."

return content

```

Listing 4.25 Self-Feedback: Function Continuation (Source: 04_PromptEngineering\40_self_feedback.py)

Now, let's see how this works. In [Listing 4.26](#) we can define a `user_prompt` and pass it to our function `self_feedback`.

```

### Test
user_prompt = "The American Civil War was a civil war in the United States between
the north and south."
res = self_feedback(user_prompt=user_prompt, max_iterations=3, target_rating=95)
res

i=0,
Prompt Content: The American Civil War was a civil war in the United States between
the north and south., Rating: 15, Feedback: The input provides minimal context and
lacks depth. It fails to ...,

```



```
Revised Output: The American Civil War was a pivotal conflict in United States
history ...
i=1,
Prompt Content: ,
Rating: 40,
Feedback: The input provides minimal context and lacks depth regarding the American
Civil War. While it mentions ...,
Revised Output: The American Civil War was ...
...
```

Listing 4.26 Self-Feedback: Testing (Source:
04_PromptEngineering\40_self_feedback.py)

The output is shortened. But you can already see that the results get better over time, starting with an initial rating of 15% and increasing to 40% in the second iteration round. Even more progress is made in subsequent rounds, which are not shown.

OceanofPDF.com

4.7 Summary

This chapter delved into the essentials of prompt engineering, the discipline of crafting inputs to guide language models toward specific, desired outputs. At its core, prompt engineering allows users to shape responses effectively, addressing the variability in generative outputs.

We began with foundational concepts by exploring *prompts* and *prompt templates*, which provide structure for crafting inputs consistently. Prompt templates, in particular, help streamline input formats for repeated or similar tasks and reduce complexity while enhancing predictability.

Next, we discussed the prompt engineering process, a dynamic and iterative approach. This process involves experimenting and refining prompts until the model consistently generates outputs aligned with the intended objective. Through trial and refinement, this cycle is crucial in balancing precision, clarity, and creativity in model responses.

Several key techniques were introduced to further steer model outputs, including the following:

- **Few-shot prompting**

In few-shot prompting, examples are provided within the prompt to guide the model by context.

- **Chain-of-thought (CoT)**

With CoT, we are prompting the model to follow a step-by-step reasoning path, enhancing complex task responses. We also described zero-shot CoT, a variant of CoT prompting where logical

steps are encouraged without prior examples, pushing the model to infer multistep logic autonomously.

- **Prompt chaining**

In this technique, sequences of prompts are created that build on each other to refine responses incrementally.

- **Self-consistency CoT**

With this technique, multiple responses are created, and the most consistent answer among them are selected to improve reliability.

- **Self-feedback**

Self-feedback allows the model to critique and revise its own output, promoting self-correction and iterative improvement.

Each technique offers unique strengths in shaping outputs, from enhancing logical coherence to fostering consistency in responses.

This chapter emphasized that mastering prompt engineering is not only about designing effective inputs but also about understanding the nuances of model behavior through creative experimentation and iteration.

5 Vector Databases

'Can you fly that thing?' 'Not yet'
—Trinity and Neo in *The Matrix*

Perhaps you remember this scene—it left a lasting impression on many. Following the dialogue between Neo and Trinity, the main characters of the sci-fi-movie *The Matrix*, knowledge about flying helicopters is loaded into her brain. In the movie, this is possible because they are in a computer-simulated world. Wouldn't it be cool if we could do that in real life?

Actually, we can, but not into a human brain (yet)—into a computer, as you'll learn in this chapter. This technology, called *vector databases*, enables you to teach a computer vast amounts of knowledge, for example, entire books, within seconds.

In this chapter, you'll learn step by step how to use vector databases, including executing the following tasks:

- Explaining what a vector database is
- Loading data from a multitude of sources
- Understanding how to split the data into smaller chunks
- Translating human-readable texts into computer-readable vectors
- Storing data in vector databases
- Pulling data out of a vector database

First things first, let's get the full picture from the “helicopter” before we dive into the nitty-gritty details on the ground.

5.1 Introduction

Let's start with what a vector database is and why you might need one in the first place. Possibly you're aware of databases like SQL or PostgreSQL. Why can't we just use these ones? These databases are extremely capable of storing and retrieving structured data, like tables with clearly defined columns and rows. But this does not reflect the wide variety of formats you might encounter on the internet, as shown in [Figure 5.1](#). You'll find text data but also images, audio files, or videos. All these different formats should be treated in a similar way and handled in the same database.

So, existing databases are not capable of dealing efficiently with all these different types of data. Something new is required that can deal with all kinds of information. This new paradigm is called the *vector database*.




Figure 5.1 Unstructured and Structured Data

Why is this new database type called a *vector* database? Well, because all data types can be represented as vectors. Think back to a math class: A vector is just a line with an arrow that you drew into a

coordinate system. Back then, you had to draw these lines, which could be represented by two numbers.

A simple example is shown in [Figure 5.2](#). This made-up graph shows living creatures, and their positions in a two-dimensional graph that represents the number of legs and the body size. By representing the information in that way, we learn about the world and the semantic meaning of words, for example, that cats and dogs are quite similar in terms of these two properties.




Figure 5.2 Example Vector Space

As humans, we can imagine a point in a two-dimensional space like the one shown in this example or a three-dimensional space. Just imagine we add a third dimension like intelligence in that graph. But we cannot imagine a 1536- or a 3072-dimensional space.

Computers can, which helps because the semantic meaning of words, images, sounds, or any other kind of information can be represented in a higher dimension. The key aspect in this context is that similar concepts are located more closely together than concepts that are very different. In our example, notice that cats and dogs, or snakes and snails, are comparable, given the properties we selected. With every additional dimension, the computer algorithm gets a better understanding of the meaning of a word.

We use languages like English or German to communicate and clarify a concept. Computers don't work directly with our languages, but with their numerical equivalents, so-called *embedding vectors*.

In our example, for the computer algorithm, a dog is defined by [4, 1], a cat by [4, 0.5], and a human by [4, 2]. An example mapping is shown in [Table 5.1](#).

Human Concept	Computer Equivalent/Embedding Vector
Dog	[4, 1]
Cat	[4, 0.5]
Human	[4, 2]
Elephant	[4, 3]
Snake	[0, 0.5]
Snail	[0, 0.1]

Table 5.1 Human Concepts and Their Computer Equivalents

The process of translating human texts into vectors is called *embedding*, which is a data ingestion step. We'll cover the data ingestion process in more detail next.

5.2 Data Ingestion Process

Let's start with a high-level overview of how data is ingested into a vector database. You'll encounter these steps whenever you want to add data to a vector database:

- Data loading
- Data chunking
- Embedding
- Storing

You always start with a data source. We touched upon the many data type options—generally, any kind of text-based source, like Markdown files, Word documents, or raw text files can be used. But you can also use more specific data types, like a YouTube video from which you want to extract a transcript, or a Wikipedia article. The options are nearly endless. Luckily, the community of users contributing specific data handlers is large as well, so that hundreds of different sources can be used out of the box with a single line of code. Therefore, you always have a consistent method for *data loading*. This consistency is important because there are many different data sources, and you want to import data in a consistent way. As a result, you'll have data imported as long strings of characters.

This long text string must be split into smaller pieces because, at the end, we want to find the needle in the haystack. Based on some user input, we want to access the specific piece of information most relevant to the user input, not the complete haystack of information. Thus, you must split the text into smaller, more manageable pieces

in a process called *chunking*. At this point, we have a list of text documents.

As discussed earlier, these text documents are understandable for humans, but not computers. You must translate this information into a computer-understandable format based in a process called *embedding*. For each text, you'll have a corresponding vector that represents the semantic meaning of the text.

We've reached the final stage of data ingestion and can now store these vectors (pairs of texts/numeric numbers) into our vector database. The complete data ingestion pipeline is shown in [Figure 5.3](#).




Figure 5.3 Data Ingestion Pipeline

5.3 Loading Documents

The first part of the ingestion pipeline is always the loading of documents from a variety of different data sources. These sources can represent text-based inputs like Markdown documents, Microsoft Word documents, JavaScript Object Notation (JSON) objects, or raw text files. But data sources can be much more exotic like WhatsApp chats, weather data, news articles, and many more.

To learn about the different formats available, refer to the LangChain documentation page at https://python.langchain.com/api_reference/community/document_loaders.html.

In this section, we'll start with a high-level overview of data loading. Then, you'll implement the loading of a single text file as well as of multiple text files. You'll also get the chance to practice and apply your knowledge in some exercises in which you'll first load multiple Wikipedia articles, before you load a complete book from Project Gutenberg.

5.3.1 High-Level Overview

At the beginning of this process, you'll have one or more files. These files are processed by a `DataLoader` class. For many different file types, `DataLoaders` are available. The `DataLoader` helps by streamlining the process because the output of any `DataLoader` is the

same. You get a Python list, with LangChain `Document` objects. This procedure is shown in [Figure 5.4](#).




Figure 5.4 LangChain DataLoader Process

A LangChain `Document` object is a container for a piece of text data along with associated metadata. The `Document` object has two main attributes, as follows:

- `page_content` contains the actual text data.
- `metadata` is a dictionary that can store additional information about the document, such as source, author, creation date, and more.

5.3.2 Coding: Load a Single Text File

In the first example, we'll work with a Sherlock Holmes book provided by Project Gutenberg. You'll find this file in the subfolder *05_VectorDatabases/data*. We'll work with books provided as *.txt* files. Like Sherlock Holmes, we'll find needles in the information haystack.

Start by loading a single file, as follows:

1. Before we can import the data, we need to load some packages. For dealing with the file location and operating system, we must import the package `os`. To import our text, we'll make use of the

`class TextLoader`, which is stored in the submodule `langchain.document_loaders`.

2. We must define the link to the file path of the book file we want to import. The file is in the top directory of the folder *05_VectorDatabases*. We can create this link by extracting the current working directory `current_dir` to get the parent directory `parent_dir`. Now, the file path is joined by stacking the parent directory with the directory data and the book file we want to import.
3. Now, we can start loading the file into an object we can work with further. First, create an instance of `TextLoader`. During the instantiation, pass the `file_path` parameter. This parameter is essential. To import the data, you must also define the text encoding. Text encoding is the process of converting human-readable characters into a specific format that computers can store, transmit, and understand, using standardized systems like ASCII or UTF-8 to represent text as binary data. In this case, we chose the encoding as `utf-8`. The object `text_loader` has a method `load()`, which we then call. At that point, the loading of the data is performed.
4. The import is complete, so let's step back and look at we've actually done. The object `doc` represents a list of type `Document`. We only imported one file, so there is only one element with the index 0. This one element has two properties: `metadata` and `page_content`. Metadata refers to the source, in our case, the file path. Page content refers to the actual text of the document we imported.

The complete script

05_VectorDatabases\10_DataLoader\10_single_text_file.py shows the data loading from a single file, as shown in [Listing 5.1](#).

```

### (1) Packages
import os
from langchain.document_loaders import TextLoader

### (2) File Handling
# Get the current working directory
file_path = os.path.abspath(__file__)
current_dir = os.path.dirname(file_path)

# Go up one directory level
parent_dir = os.path.dirname(current_dir)

file_path = os.path.join(parent_dir, "data", "HoundOfBaskerville.txt")
file_path

### (3) Load a single document
text_loader = TextLoader(file_path=file_path, encoding="utf-8")
doc = text_loader.load()

### (4) Understand the document
# Metadata
doc[0].metadata

# %% Page content
doc[0].page_content

```

Listing 5.1 Data Loading: Single File (Source: 05_VectorDatabases\10_DataLoader\10_single_text_file.py)

Great! We've managed to successfully load a single file into Python and in a LangChain-specific format, which will enable the next steps in process. Before we go to the next step, though, let's look at how you can load multiple text files.

5.3.3 Coding: Load Multiple Text Files

Typically, you won't deal with a single file, but with many files of the same file type. Ideally, you can import all these files in a single process. In our example, we want to import all text files stored in the folder *05_VectorDatabases/data*.

We must import the `DirectoryLoader`, which allows us to iterate over the files in a folder, as follows:

```

#%% (1) Packages
import os
from langchain.document_loaders import TextLoader, DirectoryLoader
from pprint import pprint

```

The path to the folder with the files is defined with a relative file location. If you set it up as shown in [Listing 5.2](#), you have flexibility in a way that the `current_directory` is extracted from the file path.

```

#%% (2) Path Handling
# Get the current working directory
file_path = os.path.abspath(__file__)
current_dir = os.path.dirname(file_path)

# Go up one directory level
parent_dir = os.path.dirname(current_dir)
text_files_path = os.path.join(parent_dir, "data")

```

Listing 5.2 Path Handling (Source:
05_VectorDatabases\10_DataLoader\20_multiple_text_files.py)

All files in the folder are loaded in a single process. The `DirectoryLoader` is applied and consumes the path pointing to the files. Using the `glob` parameter ensures that only specific files are selected. In our example, only text files are considered. `loader_cls` defines which loader to apply to the files. We use `TextLoader`, as defined earlier in [Section 5.3.2](#). Finally, we need to define the encoding. We defined the encoding before in the `TextLoader` class but want to do this here as well. We can pass it as `loader_kwargs`.

Kwargs (short for “keyword arguments”) is a Python convention for passing a variable number of named parameters to a function. These parameters must be passed in a dictionary.

[Listing 5.3](#) shows how to set up a `DirectoryLoader` that allows iterating over all the files in a folder.

```

#%% (3) load all files in a directory
dir_loader = DirectoryLoader(path=text_files_path,
                             glob="**/*.txt",
                             loader_cls=TextLoader,

```

```

                                loader_kwargs={'encoding': 'utf-8'} )
docs = dir_loader.load()
docs

[Document(metadata={'source': '...', page_content='...'},
  Document(metadata={'source': '...', page_content='...'},
  Document(metadata={'source': '...', page_content='...'})]

```

Listing 5.3 Directory Loader (Source:
05_VectorDatabases\10_DataLoader\20_multiple_text_files.py)

5.3.4 Exercise: Load Multiple Wikipedia Articles

Now, let's get to work. As shown in [Listing 5.4](#), we have multiple URLs for articles on artificial intelligence, artificial general intelligence (AGI), and artificial superintelligence (ASI). At the end, you want a list of LangChain Documents that correspond to the URLs.

```

### (1) Packages

### Articles to load
articles = [
    {'title': 'Artificial Intelligence',
     'query': 'https://en.wikipedia.org/wiki/Artificial_intelligence'},
    {'title': 'Artificial General Intelligence',
     'query': 'https://en.wikipedia.org/wiki/Artificial_general_intelligence'},
    {'title': 'Superintelligence',
     'url': 'https://en.wikipedia.org/wiki/Superintelligence'},
]

# %% (2) Load articles

```

Listing 5.4 Data Loading Exercise: Multiple Wikipedia Articles (Source:
05_VectorDatabases\10_DataLoader\30_wikipedia_exercise.py437)

Task 1: Package Loading

One essential capability is working with online documentation. Your first task is to find out if a DocumentLoader exists for Wikipedia articles (hint: it does!) and how to load the package and class.

Task 2: Article Import

Next, you must write a loop to iterate over all the articles. You can use the starter script available in

05_VectorDatabases\10_DataLoader\30_wikipedia_exercise.py.

[Listing 5.5](#) shows the content of the starter script.

```
### (1) Packages

### Articles to load
articles = [
    {'title': 'Artificial Intelligence'},
    {'title': 'Artificial General Intelligence'},
    {'title': 'Superintelligence'},
]
# %% (2) Load all articles
```

Listing 5.5 Data Loading Solution: Multiple Wikipedia Articles (Source: *05_VectorDatabases\10_DataLoader\30_wikipedia_exercise.py*)

Solve this problem and fill in your own solution.

Solution

[Listing 5.6](#) shows our solution to this exercise. Most notably, you should have learned how to load the data directly from Wikipedia with the help of `WikipediaLoader`. The remaining part is just iterating over the `articles` object and how to create a list of documents by appending to `docs`.

```
### Packages
from langchain.document_loaders import WikipediaLoader

### Articles to load
articles = [
    {'title': 'Artificial Intelligence'},
    {'title': 'Artificial General Intelligence'},
    {'title': 'Superintelligence'},
]
# %% Load all articles (2)
docs = []
for i in range(len(articles)):
    print(f>Loading article on {articles[i].get('title')}")
```



```

loader = WikipediaLoader(query=articles[i].get("title"),
load_all_available_meta=True,
doc_content_chars_max=100000,
load_max_docs=1)
doc = loader.load()
docs.append(doc)
docs

```

Listing 5.6 Data Loading Solution: Multiple Wikipedia Articles: Loading (Source: 05_VectorDatabases\10_DataLoader\40_wikipedia_solution.py)

5.3.5 Exercise: Loading Project Gutenberg Book

In the exercise file shown in [Listing 5.7](#), you'll find the book's details in a dictionary.

```

# %% The book details
book_details = {
    "title": "The Adventures of Sherlock Holmes",
    "author": "Arthur Conan Doyle",
    "year": 1892,
    "language": "English",
    "genre": "Detective Fiction",
    "url": "https://www.gutenberg.org/cache/epub/1661/pg1661.txt"
}

```

Listing 5.7 Data Loading Exercise: Project Gutenberg (Source: 05_VectorDatabases\10_DataLoader\50_custom_loader_exercise.py)

Task 1: Package Import

As in the previous exercise, you should become familiar with the online documentation and find out which package and class to import for working with books from <https://www.gutenberg.org/>. Project Gutenberg is an amazing resource because you'll find copyright-free classic books to work with.

Task 2: Metadata Adaptation

After loading the file, you'll get a list of documents. Replace the metadata of the `Document` object so that the dictionary of `book_details` represents the metadata. [Listing 5.8](#) shows the file's content, which consists of `book_details`.

```
# %% The book details
book_details = {
    "title": "The Adventures of Sherlock Holmes",
    "author": "Arthur Conan Doyle",
    "year": 1892,
    "language": "English",
    "genre": "Detective Fiction",
    "url": "https://www.gutenberg.org/cache/epub/1661/pg1661.txt"
}
```

Listing 5.8 Exercise: Custom Loader (Source: `05_VectorDatabases\10_DataLoader\50_custom_loader_exercise.py`)

Solution

[Listing 5.9](#) shows how you can load the data from Project Gutenberg.

```
from langchain_community.document_loaders import GutenbergLoader
# %% The book details
book_details = {
    "title": "The Adventures of Sherlock Holmes",
    "author": "Arthur Conan Doyle",
    "year": 1892,
    "language": "English",
    "genre": "Detective Fiction",
    "url": "https://www.gutenberg.org/cache/epub/1661/pg1661.txt"
}

loader = GutenbergLoader(book_details.get("url"))
data = loader.load()

# %% Add metadata from book_details
data[0].metadata = book_details
```

Listing 5.9 Solution: Loading from Project Gutenberg (Source: `05_VectorDatabases\10_DataLoader\60_custom_loader_solution.py`)

Now that you understand how to load data in different ways, let's move on to the next stage of the data ingestion pipeline—data splitting.

OceanofPDF.com

5.4 Splitting Documents

Often, you need to split data into smaller chunks. Why is this step needed? A couple of reasons include the following:

- With splitting, you can get more *granular results* because smaller document chunks are more likely to contain focused and relevant information.
- You can work around limitations in embedding models. Don't worry, we'll cover this in more depth later in [Section 5.5](#). Recall all embedding models have a *limited context window* of tokens to work with. We discuss tokens further in the next info box.
- Splitting can be *more efficient* in terms of database size and querying because you are working and storing smaller chunks.

No doubt you've encountered the word "token," so let's briefly define this concept. A *token* is the fundamental unit of text in natural language processing. Typically, a token is a word, a part of a word, or sometimes even just a single character. The process of converting a sentence into a list of tokens is called *tokenization*. A crucial step in preprocessing, many different tokenization methods are in operation.

Some examples might illustrate the process. In the sentence "I love generative AI!", the resulting tokens might be ["I", "love", "generative", "AI", "!"]. In many tokenization methods, words are broken down into multiple tokens. For instance, "loving" might be tokenized into ["lov", "ing"].

Why is tokenization performed? Tokens allow AI models to process text by converting language into numerical representations that models can understand. However, these numbers don't carry

inherent meaning; they serve as a reference system for LLMs to translate human language into a form computers can process. In the next step, called embedding, these tokens are transformed into vectors that capture semantic meaning, allowing the model to interpret relationships between words and sentences. Embeddings are covered later in [Section 5.5](#).

For now, simply know that all embedding models have a limit on tokens to process. If you try to pass more tokens into an embedding model than it can process, it will cut off excess tokens. The meaning of these extra tokens is lost. An essential strategy is to always pass fewer tokens than a model can maximally process. As a rule of thumb, one token is approximately equal to 0.75 words, or 100 tokens roughly equal 75 words in English language.

Finally, all models have limits on how much data can be processed at once. This limitation of tokens is also called the *context window*.

[+] Embedding Model Context Window

Let's touch upon the typical token limits of popular AI models. For all models, you can find a model card with relevant information on the model. Just search, for example, for "text-embedding-3-small model card" for more information. Some examples of context windows include the following:

- OpenAI text-embedding-3-small: 8,191 tokens
- OpenAI text-embedding-3-large: 8,191 tokens
- Anthropic voyage-3-large: 32,000 tokens
- All-MiniLM-L6-v2: 256 tokens

So far, we've covered documents that need to be split into smaller chunks and described why defining a reasonable size for these chunks is important. But how are the chunks created? Many different approaches are available, but we'll cover three popular methods. A visual representation of data chunking is shown in [Figure 5.5](#). The first approach is to use a fixed chunk size and split a document after a fixed number of characters is reached.

A more sophisticated approach is to follow the structure of the document and allow for different sizes of chunks if doing so helps retain the meaning of each chunk. [Figure 5.5](#) shows a chat history of two persons. If the chunks of two different persons overlapped, understanding will be impacted negatively. Imagine you're following a conversation in a book, but the statements are partly attributed to the wrong protagonists. The most advanced approach is *semantic chunking*, in which the text is understood before the chunking is performed.




Figure 5.5 Chunking Strategies

Let's start with the simplest approach, fixed-size chunking, next.

5.4.1 Coding: Fixed-Size Chunking

Fixed-size chunking creates chunks after a predefined number of characters has been reached. The cut occurs at the position of the last word before the chunk size is exceeded. One advantage is the

simplicity of this approach; it is straightforward to understand and implement. This approach is also quite fast, as it does not require complex analysis of text meaning. Disadvantages include a lack of context awareness. Sentences or concepts could be split midway, thus resulting in a loss of semantic coherence.

[+] Chunk Overlap

A specialty that you should leverage is the chunk overlap. Chunk overlaps refer to the amount of text that is shared between two consecutive chunks, usually measured in characters or tokens.

The goal is to maintain the context between chunks and to preserve the meaning, which otherwise could be lost when text is arbitrarily divided.

A common rule of thumb is an overlap of 10% to 20%. This buffer provides a good balance between context preservation and database efficiency. But this suggestion is not dogma. Depending on the type of text, more or less overlap could be reasonable.

[Figure 5.6](#) shows the result of `CharacterTextSplitter`. Both chunks are of equal size, and you can see the overlapping text in both chunks.




Figure 5.6 CharacterTextSplitter with Chunk Overlap

Since you now have a basic understanding, you can now implement these concepts. In our example, we'll go back to Sherlock Holmes and extend our previous script from [Section 5.3.5](#). The data loading step is not shown again; we'll jump in after the `docs` object has been created.

Now, we need to load the functionality and use the class `CharacterTextSplitter` for the `langchain` submodule `text_splitter`, as shown in [Listing 5.10](#). You can find the complete script in `05_VectorDatabases\20_Chunking\10_fixed_size_chunking.py`.

```
# ...
# %%
docs

# %% Splitting text
# Packages (1)
from langchain.text_splitter import CharacterTextSplitter
```

Listing 5.10 Fixed-Size Chunking: Required Packages (Source: `05_VectorDatabases\20_Chunking\10_fixed_size_chunking.py`)

An instance of the class `CharacterTextSplitter` is created and typically called `splitter`. In this instantiation, the parameters are defined, for example, the chunk size (the maximum number of characters in a chunk), the chunk overlap (the maximum overlap characters in a chunk), and the separator (used for splitting):

```
# Split by characters (2)
splitter = CharacterTextSplitter(chunk_size=256, chunk_overlap=50,
separator=" ")
```

The split has not been performed yet but can now be done by calling the method `splitter.split_documents` and passing the list of documents. The result is again a list of `Document` objects, but the number of elements has dramatically increased. Up to this point, we had one `Document` per book. Now, each book has been separated into many smaller pieces, as follows:

```
# %% Apply the splitting (3)
docs_chunks = splitter.split_documents(docs)
```

You can check the number of chunks to find out that it has increased to 4,140 documents, as follows:

```
# %% Check the number of chunks (4)
len(docs_chunks)
```

By taking a closer look at two consecutive chunks, you can observe the chunk overlap, which is the text shared between both chunks, as shown in [Listing 5.11](#).

```
# %% check some random Documents (5)
from pprint import pprint
pprint(docs_chunks[100].page_content)
# %%
pprint(docs_chunks[101].page_content)
```

Listing 5.11 Fixed-Size Chunking: Document Check (Source: 05_VectorDatabases\20_Chunking\10_fixed_size_chunking.py)

Finally, as shown in [Listing 5.12](#), you can look at the chunk sizes and visualize the distribution of chunk lengths. In our example, most of our chunks are centered around 250, which was expected.

```
# %% visualize the chunk size (6)
import seaborn as sns
import matplotlib.pyplot as plt
# get number of characters in each chunk
chunk_lengths = [len(chunk.page_content) for chunk in docs_chunks]

sns.histplot(chunk_lengths, bins=50, binrange=(100, 300))
# add title
plt.title("Distribution of chunk lengths")
# add x-axis label
plt.xlabel("Number of characters")
# add y-axis label
plt.ylabel("Number of chunks")
```

Listing 5.12 Visualizing Fixed-Size Chunking (Source: 05_VectorDatabases\20_Chunking\10_fixed_size_chunking.py)

The chunk length distribution is shown in [Figure 5.7](#).




Figure 5.7 Distribution of Chunk Lengths for Fixed-Size Chunking

Most chunks have 250 characters. Since the splitting is not performed mid-word, some chunks are shorter, and some slightly

longer.

5.4.2 Coding: Structure-Based Chunking

Contextual meaning or textual coherence usually does not follow fixed sizes, so breaking a text into equal sizes doesn't reflect reality well. We need something more sophisticated. Typically, a text is structured by an author into chapters, paragraphs, sentences, parts of sentences, words, etc. One approach is breaking the text into smaller pieces like chapters. If the remaining chunk is still too large, it is broken down into subchapters, then paragraphs, and so on. This approach ensures a good balance between textual coherence and complying with character limits.

How can an algorithm detect a chapter or paragraph? That ability depends on the choices the author of the text has made. Typically, multiple empty lines exist before a new chapter starts. Between paragraphs, typically smaller distances are used. Now, that you understand how it works, let's implement this capability in code.

Let's break down our example, as shown in [Table 5.2](#). On the left, notice the formatting of the book chapters. After loading the Document, get the property `page_content`, which is a single long String.

Chapter 1 Mr. Sherlock Holmes Chapter 2 The Curse of the Baskervilles Chapter 3 The Problem Chapter 4 Sir Henry Baskerville Chapter 5 Three Broken Threads	“ Chapter 1 Mr. Sherlock Holmes\n Chapter 2 The Curse of the Baskervilles\n Chapter 3 The Problem\n Chapter 4 Sir Henry Baskerville\n Chapter 5 Three Broken Threads\n”
--	---

Table 5.2 Comparison of Raw Text Reprinted from a Book (Left) and Snippet of a Corresponding Document Page Content (Right)

In this text, you’ll find special characters like `\n`. This special character indicates a line feed and results in breaking the text, so that the following text starts at the beginning of a new line.

How can you know that `\n` indicates a new line? This symbol is defined in a character encoding standard (of which there are many). A *character encoding standard* is a system that creates a mapping between a character like letters, numbers, or symbols and specific code. The standard provides a universal character set that includes all writing systems, emojis, and much more. The encoding allows computers to store, transmit, and display text consistently. Famous encoding standards include ASCII, UTF-8, or ISO-8859-1. UTF-8 is more modern than ASCII.

[+] Character Encoding Standards Example

[Listing 5.13](#) contains some examples of encoding standards. Run this code to see the encoding of the characters into the corresponding encoding standards. If you try to encode an emoji with ASCII, it will fail because emojis are not part of the set of characters. You can also encode Chinese or any other language characters.

```
print("ASCII:")
print("A:", "A".encode("ascii"))
print("\n:", "\n".encode("ascii"))

# UTF-8 representation
print("\nUTF-8:")
print("A:", "A".encode("utf-8"))
print("\n:", "\n".encode("utf-8"))
print("€:", "€".encode("utf-8"))
print("你:", "你".encode("utf-8"))
print("😊 :", "😊 ".encode("utf-8"))
```

Listing 5.13 Character Encoding

Let's now build on the data we loaded earlier in [Section 5.3.2](#). The functionality we need is called `RecursiveCharacterTextSplitter`, which is also part of the submodule `langchain.text_splitter`, as follows:

```
# ...
### Recursive Chunking
# Packages (1)
from langchain.text_splitter import RecursiveCharacterTextSplitter
```

Now, create an instance of `RecursiveCharacterTextSplitter` class. During instantiation, pass relevant parameters to the splitter, specifically the `chunk_size` (the number of characters in a chunk) and the `chunk_overlap` (the shared text between two consecutive chunks). A vitally important parameter is the `separators`. The

separators are the characters used to create the breaks. They are ordered from largest to smallest—we start with a separator for chapters, then for paragraphs, and so on. When the first split was not enough to constrain the chunk size to the chunk size limit, another separator is sought for creating a further split. This process is repeated until the chunk size is within the predefined limits, as follows:

```
# %% Set up the splitter (2)
splitter = RecursiveCharacterTextSplitter(chunk_size=256,
                                          chunk_overlap=50,
                                          separators=["\n\n", "\n", " ", ".", ","])
```

The instance object `splitter` has a method `split_documents` in which we'll pass our list of documents. As a result, we receive a list of documents, but the number of documents has increased significantly.

```
# %% Create the chunks (3)
docs_chunks = splitter.split_documents(docs)
# %%
len(docs_chunks)
```

Check the chunk size distribution through a visualization, as shown in [Listing 5.14](#).

```
# %% Visualize the chunk size (4)
import seaborn as sns
import matplotlib.pyplot as plt
# get number of characters in each chunk
chunk_lengths = [len(chunk.page_content) for chunk in docs_chunks]

sns.histplot(chunk_lengths, bins=50, binrange=(0, 300))
# add title
plt.title("Distribution of chunk lengths (RecursiveCharacterTextSplitter)")
# add x-axis label
plt.xlabel("Number of characters")
# add y-axis label
plt.ylabel("Number of chunks")
```

Listing 5.14 Chunk Size Visualization (Source: 05_VectorDatabases\20_Chunking\20_structure_based_chunking.py)

[Figure 5.8](#) shows that the chunk sizes are now much more evenly distributed compared to fixed-size splitting.




Figure 5.8 Distribution of Chunk Lengths for Recursive Character Splitting

Notice that a wide range of chunk sizes is covered, ranging from close to zero up to 250 characters. If due to this splitting approach, a chunk is too small, chunks can also be aggregated to a decent size. For simplicity, we won't cover these steps. But the implementation is straightforward—you might iterate over the complete list of chunks and aggregate neighboring chunks if a chunk is too small and allow the aggregation with a neighboring chunk if doing so is still within allowed chunk size.

5.4.3 Coding: Semantic Chunking

Up to this point, chunking was performed based on chunk sizes (fixed-size chunking) and somewhat taking the structure of the text into account (structure-based chunking). But the actual text was not considered at all, which is how semantic chunking is performed.

In this approach, the text we want to split is at first split into individual sentences. Then, the sentences are embedded. Don't worry, we'll cover this topic soon, in [Section 5.5](#). For now, simply understand that, afterwards, the sentences are converted into numerical vectors, with which certain calculations can be made. For example, the similarities between consecutive sentences can be calculated. The idea is to split at natural breaks—when sentences are too different in their meaning. For this task, a *distance matrix* is created and analyzed. If you ever worked with distances matrices before, you know that multiple approaches exist to calculate distance. Subsequently, multiple ways exist to calculate the breakpoints that define when a split should be performed. This general idea of semantic chunking is shown in [Figure 5.9](#).




Figure 5.9 Semantic Chunking

This approach preserves the semantic meaning of sentences and keeps related concepts together. It improves upon the information retrieval process and the accuracy of searches. Furthermore, semantic chunking is adaptable in the way that it adjusts the chunk sizes based on the text, resulting in chunks that are more natural and easier for humans to read and to understand.

But downsides also exist. For one, semantic chunking adds some computational overhead and is slower. Note also that embedding requires you to run a neural network either locally, or you need to apply an embedding model like from OpenAI. As a result, some cost

is incurred for creating the embeddings, although embedding models is comparably cheap. Latency and higher system complexity are factors that might also play an important role.

Let's get back to coding and find out how this works. The script that we'll develop can be found in

05_VectorDatabases\20_Chunking\30_semantic_chunking.py.

The semantic chunker has not been promoted yet to the `langchain` main package and is found in `langchain_experimental`. For loading a Wikipedia article, you can apply `WikipediaLoader`. For creating the embeddings, use OpenAI's `langchain_openai` package. You'll need an OpenAI account to use OpenAI services as well as an application programming interface (API) key in a file called `.env`, which is stored in the same folder as the script. Finally, load the script via `load_dotenv`, as shown in [Listing 5.15](#).

```
## Packages (1)
from langchain_experimental.text_splitter import SemanticChunker
from langchain.document_loaders import WikipediaLoader
from langchain_openai.embeddings import OpenAIEmbeddings
from pprint import pprint
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
```

Listing 5.15 Semantic Chunking: Required Packages (Source: *05_VectorDatabases\20_Chunking\30_semantic_chunking.py*)

Now, we can load the article, and for this step, we'll set up an instance of `WikipediaLoader` and pass several parameters (like the article title) to the parameter `query`. We don't want to load further articles that are mentioned, so we'll limit it with the parameter `load_max_docs`. Also, for demonstration purposes, we'll also limit the number of characters with the parameter `doc_content_chars_max`, as shown in [Listing 5.16](#).

```
# %% Load the article (2)
ai_article_title = "Artificial_intelligence"
```

```

loader = WikipediaLoader(query=ai_article_title,
                        load_all_available_meta=True,
                        doc_content_chars_max=1000,
                        load_max_docs=1)

doc = loader.load()

```

Listing 5.16 Semantic Chunking: Data Loading (Source: 05_VectorDatabases\20_Chunking\30_semantic_chunking.py)

After loading, let's check that everything worked. We have our doc object, which is a list with just one Document. Thus, we need to access `doc[0].page_content` to retrieve the contents of that article, as follows:

```

# %% check the content (3)
pprint(doc[0].page_content)

```

Now, we've come to an interesting area. We'll set up our `SemanticChunker` instance, as shown in [Listing 5.17](#). The semantic chunker works with any embedding, in our case, `OpenAIEmbeddings`. Different methods exist to calculate the similarities between texts, and we'll apply cosine similarity together with a threshold of 0.5 to split our texts. We'll cover similarity calculations in [Section 5.7.1](#) in much more detail. At this point, simply know that we need to analyze the semantic meaning of neighboring sentences, which is performed based on similarities.

```

# %% Create splitter instance (4)
splitter = SemanticChunker(embeddings=OpenAIEmbeddings(),
                          breakpoint_threshold_type="cosine",
                          breakpoint_threshold=0.5)

```

Listing 5.17 Semantic Chunking: Chunking (Source: 05_VectorDatabases\20_Chunking\30_semantic_chunking.py)

The semantic splitter operates on the article, and we store the resulting split text in an object called `chunks`, as follows:

```

# %% Apply semantic chunking (5)
chunks = splitter.split_documents(doc)

```

Finally, we can compare the contents of the first chunk and the second chunk to see at which position the split was performed (see [Listing 5.18](#)).

```
# %% check the results (6)
chunks
# %%
pprint(chunks[0].page_content)
# %%
pprint(chunks[1].page_content)
```

Listing 5.18 Semantic Chunking: Document Check (Source: 05_VectorDatabases\20_Chunking\30_semantic_chunking.py)

5.4.4 Coding: Custom Chunking

The text splitters we’ve covered so far provide a lot of flexibility but sometimes are still not enough. You may need to define a custom splitter to fit your specific needs.

In this example, we’ll work with another Sherlock Holmes book file, but this book file consists of 12 books. Of course, we could just go ahead and split the text, but we want to extract certain information, like the book’s title, and store it alongside the chunks. This specific information will be extremely helpful at some point later. When we interact with the book and ask questions, we might want to know how Sherlock Holmes solved the case of “A Scandal in Bohemia,” for instance. Then, we can filter first for the book of interest and then get the information we are interested in. You’ll see this capability in action in [Section 5.7](#) when you learn how to retrieve data.

You can find this code in

05_VectorDatabases\20_Chunking\40_custom_splitter.py. The book with the title “The Adventures of Sherlock Holmes” itself holds twelve different books about Sherlock Holmes, and we want to extract all the titles from the file.

As always, we start by loading packages. The LangChain community helps in this case because there is a loader class for our purposes, as shown in [Listing 5.19](#).

```
### Packages
import re
from langchain.text_splitter import CharacterTextSplitter,
RecursiveCharacterTextSplitter
from langchain_community.document_loaders import GutenbergLoader
```

Listing 5.19 Custom Chunking: Required Packages (Source: 05_VectorDatabases\20_Chunking\40_custom_splitter.py)

[Listing 5.20](#) shows how to load the data directly via the class GutenbergLoader.

```
# %% The book details
book_details = {
    "title": "The Adventures of Sherlock Holmes",
    "author": "Arthur Conan Doyle",
    "year": 1892,
    "language": "English",
    "genre": "Detective Fiction",
    "url": "https://www.gutenberg.org/cache/epub/1661/pg1661.txt"
}
loader = GutenbergLoader(book_details.get("url"))
data = loader.load()
```

Listing 5.20 Custom Splitter: Data Loading (Source: 05_VectorDatabases\20_Chunking\40_custom_splitter.py)

The result is a list of documents in which we only have one Document object, as follows:

```
[Document(metadata={'title': 'The Adventures of Sherlock Holmes', 'author':...})]
```

We want to keep all the metadata and extend the existing metadata of the Document object with our dictionary, as follows:

```
### Add metadata from book_details
data[0].metadata = book_details
```

By inspecting the file, we can see that each book title follows a similar structure. For example, for the second book, the first text

looks like the following:

```
('II. THE RED-HEADED LEAGUE\r\n'
 '\n'
 '\n'
 '\r\n')
```

Knowing the inner structure of text pieces is important because you can use those structures by setting up a pattern that matches all titles. The logic behind this pattern matching is called *regular expressions*, which is large topic you can add to your already long to-do list, or you can do what we do: copy the text and ask an large language model (LLM) to create the corresponding regular expression to split the text whenever a title is detected. This approach is shown in [Listing 5.21](#) with the function `custom_splitter`. Afterwards, we'll overwrite the `split_text` method of our `text_splitter` object and replace it with our own `custom_splitter` method.

```
# %% Custom splitter
def custom_splitter(text):
    # This pattern looks for Roman numerals followed by a title
    pattern = r'\n(?:[IVX]+\.\s[A-Z])'
    return re.split(pattern, text)

text_splitter = CharacterTextSplitter(
    separator="\n",
    chunk_size=1000,
    chunk_overlap=200,
    length_function=len,
    is_separator_regex=False,
)
# Override the default split method
text_splitter.split_text = custom_splitter
```

Listing 5.21 Custom Splitter: Function (Source: 05_VectorDatabases\20_Chunking\40_custom_splitter.py)

Great! Let's try it out by calling the `split_documents` method, which itself inherits from the `split_text` method. We'll create an object for books, as follows:

```
# Assuming you have the full text in a variable called 'full_text'
books = text_splitter.split_documents(data)
```

This object has 13 documents. The first document does not refer to any Sherlock Holmes book, so we disregard it using the following:

```
# %% remove the first element, because it only holds metadata, not real books
books = books[1: ]
```

Next, as shown in [Listing 5.22](#), for each book inside `books`, we want to extract the correct title. Again, we'll use a regular expression: Whenever the pattern is matched in the text, the title is extracted and added to the metadata.

```
##### Extract the book title from beginning of page content
for i in range(len(books)):
    print(i)
    # extract title
    pattern = r'\b[IVXLCDM]+\.\s+([A-Z\s\-\-])+\r\n'
    match = re.match(pattern, books[i].page_content)
    if match:
        title = match.group(1).replace("\r", "").replace("\n", "")
        print(title)
    # add title to metadata
    books[i].metadata["title"] = title
    print(title)
```

Listing 5.22 Custom Splitter: Extract The Book Title (Source: 05_VectorDatabases/20_Chunking\40_custom_splitter.py)

We're nearly done. Now, `books` is a list of documents, each of which has the correct title in its metadata, as follows:

```
[Document(metadata={'title': 'A SCANDAL IN BOHEMIA', ...},
Document(metadata={'title': 'THE RED-HEADED LEAGUE', ...},
...]
```

The next steps we've covered before. We'll apply a recursive character text splitter to create chunks of a predefined maximum chunk size and predefined chunk overlap, as shown in [Listing 5.23](#).

```
# %% apply RecursiveCharacterTextSplitter
text_splitter = RecursiveCharacterTextSplitter(
    chunk_size=1000,
```

```
        chunk_overlap=200,  
        length_function=len,  
        is_separator_regex=False,  
    )  
chunks = text_splitter.split_documents(books)
```

Listing 5.23 Custom Splitter: Applying Recursive Character Text Splitter (Source: 05_VectorDatabases\20_Chunking\40_custom_splitter.py)

Thus concludes our section on data splitting; we'll move on to the next stage in the data ingestion process—embeddings.

OceanofPDF.com

5.5 Embeddings

Let's visualize where we are in the process. We've loaded one or more documents. We've broken them down into small manageable pieces. But up to this point, we're dealing with texts—human-readable texts. We need to translate this text into a machine-understandable representation of text in a process called *embedding*. As shown in [Figure 5.10](#), text chunks are converted into numerical vectors.




Figure 5.10 Embedding Process

In [Section 5.4](#), you learned what an embedding vector was. A short recap: an embedding is a numerical vector that represents the meaning of a word or of a complete sentence or paragraph. This vector does not have the two or three dimensions that we know from math class but instead can have hundreds or even thousands of dimensions. This expansiveness is necessary to grasp the complete meaning of the text to be understood.

In this section, you'll learn how embeddings are created and see how similar words are closer to each other in a two-dimensional representation. By the end of this section, you'll know the following:

- About types of embeddings

- How embeddings are created
- Which embedding models you can use

First things first, which types of embeddings are out there?

5.5.1 Overview

In this book, we'll mainly embed complete sentences, but the concept of embeddings is not limited to texts. These other data modalities can also be embedded:

- **Image embeddings**
You can embed complete images.
- **Audio embeddings**
Sounds like voices or music can be embedded into a numerical representation.
- **Video embeddings**
Videos are just a combination of audio signals and a sequence of images, so it won't surprise you that video embeddings are possible.
- **Graph embeddings**
A famous concept in machine learning (ML) is a graph. *Graphs* consist of nodes, edges, and features. For specific data, graphs are the preferred type of data representation. Graphs can be embedded as well.
- **Multimodal embeddings**
Different data modalities like text data and image data can be embedded in a shared, unified vector space. One example is OpenAI's CLIP model, which allows the simultaneous embedding of images and their corresponding texts.

Knowing that you can embed these different data types is great, but why would you do that? What are real-world applications of embeddings? We cannot overstate how important embeddings are to generative AI. Embeddings are at the core of LLMs, and everything is built upon LLMs. But their impact reaches even further:

- *Natural language processing (NLP)* models use text embeddings.
- In *computer vision*, embeddings model image meaning.
- *Retrieval-augmented generation (RAG)* applications are built upon *semantic search*. Both rely on embeddings to efficiently retrieve relevant chunks.
- Products and their features are embedded, and thus embedding models are applied in *recommendation* engines.

State of the Art Embeddings Models

Progress in generative AI is massive, so a static list of models in this book would be quickly outdated. Instead, we want to refer you to resources that provide state-of-the-art rankings. The Massive Text Embedding Benchmark (MTEB) leaderboard at <https://huggingface.co/spaces/mteb/leaderboard> is a space on Hugging Face that ranks several embedding models by different parameters. MTEB based on 8 embedding tasks and 58 datasets in 112 languages.

What are the important metrics to consider when choosing an embedding model? A very important aspect is the *performance* of the model. On the MTEB leaderboard, you can choose between several different performance metrics like average performance over 56 datasets. Task-specific performance metrics are available for comparing performance in, for example, classification or clustering.

When self-hosting a model, relevant information to consider include *model size* and *memory usage*.

Depending on the text you're dealing with, you'll need to define a proper chunk size. Then, you must ensure that the chunk size is below *max tokens* limit of the model.

Another important aspect to consider is the *price*. Many companies offer access to their models via APIs. While many embedding models are open source and free to use, running them on internal company resources can be costly. The server must run 24/7, must have ideally a good GPU for quick inferences, and must consume energy. Thus, both fixed and variable costs are associated with that approach. You also must maintain an additional system. On the plus side, however, you won't have any *data privacy* concerns because your company's data won't leave your own network.

5.5.2 Coding: Word Embeddings

Always remember that the main purpose of embedding is to capture the semantic meanings of words. You can even visualize meaning and see how words that are similar are closer to each other than random words. How is this achieved?

The breakthrough idea was an algorithm called a *skip-gram algorithm*. As shown in [Figure 5.11](#), the main idea is simple in retrospect. An actual text like a book is used as training data. The focus is always on several words, let's say, five words. We start with the first five words, then the focus is shifted by a word to words 1 through 6, then words 2 through 7, and so on.

Of the current focus on five words, the central word is the target feature, and the two neighboring words to the left and the right are

independent features in a neural network model. The independent features are used to predict the target feature. For example, if the current focus is on five words like “It was a bright cold ...,” then the central word “a” is considered the target, and the other words are used to predict that word. The intuition behind this approach is that a word is described by the context in which it is used.




Figure 5.11 Skip-Gram Example

The corresponding algorithm is called Word2vec, which is an early word embedding model. The information box covers some other popular models.

[+] History of Word Embedding Models

Word2vec was developed by Google in 2013 and sparked much interest in word embeddings.

In 2014, Stanford researchers created *GloVe* which stands for Global Vectors for Word Representation. Its performance is comparable to *Word2vec*.

Facebook joined the party in 2016 when it released *FastText*. This extension of *Word2vec* can handle even out-of-vocabulary words and thus was useful for languages with many words.

The most influential model is *BERT* (*bidirectional encoder representations from transformers*). This model was developed by

Google researchers in 2018 and is based on transformer architecture. BERT is still one of the most popular word embedding models.

We'll download the model embeddings for all embedded words and visually demonstrate how similar words are closer than random words. Let's develop a graph. The code file is *05_VectorDatabases\30_Embedding\10_word2vec_similarity.py*.

We'll use `gensim` for getting Word2vec embeddings, as shown in [Listing 5.24](#). Seaborn is used for creating visualization. For plotting the data, we must reduce the dimensions to two dimensions and use PCA from `sklearn` for this task.

```
##% (1) Packages
import gensim.downloader as api # Package for downloading word vectors
import random # Package for generating random numbers
import seaborn.objects as so # Package for visualizing the embeddings
from sklearn.decomposition import PCA # import PCA
import numpy as np
import pandas as pd
```

Listing 5.24 Word Embeddings: Required Packages (Source: *05_VectorDatabases\30_Embedding\10_word2vec_similarity.py*)

Next, we'll load Word2vec embeddings. This step might take a while because the `word2vec-google-news-300` model is 1.5 GB.

```
# %% (2) import GloVe word vectors
word_vectors = api.load("word2vec-google-news-300")
```

In our example, we'll study the word “mathematics,” but you're free to play around and choose your own word. The `shape` property is checked, and we can confirm that the embedding has 300 dimensions, as follows:

```
# %% (3) get the size of the word vector
studied_word = 'mathematics'
word_vectors[studied_word].shape
```

To get some impression of the embeddings, let's look at them. What we have is a vector of 300 double numbers, as follows:

```
# %% (4) get the word vector for the word 'intelligence'
word_vectors[studied_word]
```

You can use the method `most_similar` to get the words most similar to the word under study. (We haven't covered how this similarity is calculated, which we'll discuss in detail in [Section 5.7.1](#).)

```
# %% (5) get similar words to 'intelligence'
word_vectors.most_similar(studied_word)
```

Finally, we'll limit our example to the five most similar words, as follows:

```
# %% (6) get a list of strings that are similar to 'intelligence'
words_similar = [w[0] for w in word_vectors.most_similar(studied_word)][:5]
```

Now, we'll pick 20 words from random positions in the embedding model, as shown in [Listing 5.25](#).

```
# %% (7) get random words from word vectors
num_random_words = 20
all_words = list(word_vectors.key_to_index.keys())
# set the seed for reproducibility
random.seed(42)
random_words = random.sample(all_words, num_random_words)

# Print the random words
print("Random words extracted:")
for word in random_words:
    print(word)
```

Listing 5.25 Word Embedding: Get Random Words (Source: 05_VectorDatabases\30_Embedding\10_word2vec_similarity.py)

We have 20 random words, and 5 words like “mathematics.” For these 25 words, we extract the embeddings and store them in a NumPy array of dimensions (25, 300), as shown in [Listing 5.26](#).

```
# %% (8) get the embeddings for random words and similar words
words_to_plot = random_words + words_similar
embeddings = np.array([])
```

```
for word in words_to_plot:
    embeddings = np.vstack([embeddings, word_vectors[word]]) if embeddings.size else
word_vectors[word]
```

Listing 5.26 Word Embedding: Get Embeddings (Source:
05_VectorDatabases\30_Embedding\10_word2vec_similarity.py)

Typically, plots are two-dimensional, so we'll stick to that. Since we cannot visualize a 300-dimensional vector, we must reduce the dimensions to a manageable size, as shown in [Listing 5.27](#), using principal component analysis (PCA). Several dimension reduction techniques are available, one of the most popular is PCA, which will do the job for us.

```
# %% (9) create 2D representation via PCA
pca = PCA(n_components=2)
embeddings_2d = pca.fit_transform(embeddings)

df = pd.DataFrame(embeddings_2d, columns=["x", "y"])
df["word"] = words_to_plot
# red for random words, blue for similar words
df["color"] = ["random"] * num_random_words + ["similar"] * len(words_similar)
```

Listing 5.27 Word Embedding: Principal Component Analysis (Source:
05_VectorDatabases\30_Embedding\10_word2vec_similarity.py)

Finally, we plot the words with `seaborn`, as shown in [Listing 5.28](#). If you haven't yet worked with Seaborn's new *grammar-of-graphics approach*, refer to https://seaborn.pydata.org/tutorial/objects_interface.html. This approach is novel, and we recommend it over a classic graph setup because the grammar-of-graphics approach stacks layers of meaning together and thus has some similarity to neural networks.

```
# %% (10) visualize the embeddings using seaborn
(so.Plot(df, x="x", y="y", text="word", color="color")
 .add(so.Text())
 .add(so.Dots())
)
```

Listing 5.28 Word Embedding: Visualization (Source:
05_VectorDatabases\30_Embedding\10_word2vec_similarity.py)

The result of our work on word embeddings and on plotting similar words in two dimensions is shown in [Figure 5.12](#).

Words that are like “mathematics” are shown in one color, while the random words are shown in a different color. Similar words are much closer to each other than random words. This result is what we wanted because it means that the semantic meanings of words are represented by the embeddings.




Figure 5.12 Word2vec for Random Words and Words Similar to “mathematics”

Now that we have numerical representations, we can work with them and visualize relationships. In the examples shown in [Listing 5.29](#), we can create relationships between capitals and their corresponding countries. We’ll get the embeddings for Berlin, Paris, Madrid and for the countries Germany, France, and Spain. Then, we’ll reduce the dimensions of the embeddings, again via PCA, and plot the connections between capital and country as a vector.

```
# %% visualizing it via lines
df_arithmetic = pd.DataFrame({'word': ['paris', 'germany', 'france', 'berlin',
    'madrid', 'spain']})
# add embeddings and add x- and y-coordinates for PCA
pca = PCA(n_components=2)
embeddings_arithmetic = np.array([])
for word in df_arithmetic['word']:
    embeddings_arithmetic = np.vstack([embeddings_arithmetic, word_vectors[word]])
```



```

if embeddings_arithmetic.size else word_vectors[word]

# apply PCA
embeddings_arithmetic_2d = pca.fit_transform(embeddings_arithmetic)
df_arithmetic['x'] = embeddings_arithmetic_2d[:, 0]
df_arithmetic['y'] = embeddings_arithmetic_2d[:, 1]

### visualise it via matplotlib with lines
import matplotlib.pyplot as plt
plt.figure(figsize=(10, 10))
plt.scatter(df_arithmetic['x'], df_arithmetic['y'], marker='o')

# add vector from paris to france, and berlin to germany
plt.arrow(df_arithmetic['x'][0], df_arithmetic['y'][0],
          df_arithmetic['x'][2] - df_arithmetic['x'][0],
          df_arithmetic['y'][2] - df_arithmetic['y'][0],
          head_width=0.01, head_length=0.01, fc='r', ec='r')
plt.arrow(df_arithmetic['x'][3], df_arithmetic['y'][3],
          df_arithmetic['x'][1] - df_arithmetic['x'][3],
          df_arithmetic['y'][1] - df_arithmetic['y'][3],
          head_width=0.01, head_length=0.01, fc='r', ec='r')
plt.arrow(df_arithmetic['x'][4], df_arithmetic['y'][4],
          df_arithmetic['x'][5] - df_arithmetic['x'][4],
          df_arithmetic['y'][5] - df_arithmetic['y'][4],
          head_width=0.01, head_length=0.01, fc='r', ec='r')

# add labels for words
for i, txt in enumerate(df_arithmetic['word']):
    plt.annotate(txt, (df_arithmetic['x'][i], df_arithmetic['y'][i]))

```

Listing 5.29 Word2vec Embeddings and Vector Representations (Source: 05_VectorDatabases\30_Embedding\10_word2vec_similarity.py)

The output of our vector embedding of capitals and countries is shown in [Figure 5.13](#).




Figure 5.13 Word2vec PCA Representation of Capitals and Countries

The vectors always have a similar direction. You can make use of these relationships and work with these vectors, as shown in [Listing 5.30](#). For instance, consider the following equation:

Paris – France + Germany = ???

```
# Paris - France + Germany = Berlin
word_vectors.most_similar(positive = ["paris", "germany"],
                           negative= ["france"], topn=1)
```

Listing 5.30 Word2vec Algebraic Calculations (Source: 05_VectorDatabases\30_Embedding\10_word2vec_similarity.py)

The equation should result in Berlin.

At this point, we can embed single words, which is great, but some issues arise as soon as we try working with complete sentences. Look at the two sentences. Both contain the word “bank,” but both instances have completely different meanings:

- Sentence 1: “I need to go to the *bank* to deposit this check.”
- Sentence 2: “We had a picnic on the *bank* of the river.”

We can thus conclude that the meaning of a word depends on the context in which it is used. A single embedding for the word *bank* does not make sense. Thus, we need to find an approach for complete sentences.

5.5.3 Coding: Sentence Embeddings

Sentence embedding models are quite similar to word embedding models. However, instead of consuming a word and providing its embedding vector, you can now pass a complete sentence to the model and get an embedding vector that represents the complete sentence. Earlier, we mentioned how BERT was one of the most influential word embedding models. Don’t get confused by the fact that there are variants of BERT that can consume sentences and complete paragraphs. A typical BERT sequence is shown in [Figure 5.14](#).




Figure 5.14 Sample Sentence Tokenization with BERT

We write a sentence, pass it through the tokenizer (not yet the embedding model), and get a tokenization. Each word of our sentence is tokenized. Also, the period is tokenized by itself. But you can also observe two new tokens: [CLS] and [SEP]. The [SEP] token separates each sentence. More interesting is the [CLS] token, short for “classification token” or sometimes “class token.” This special token is used in transformer-based models, especially in BERT. Its purpose is to aggregate information from an entire input sequence. [CLS] is typically located as the first token in the input sequence, before the actual content tokens. This token learns to encode the sentence-level or paragraph-level information. In the original BERT paper, the authors suggested using the final hidden state of the [CLS] token as the sentence embedding. Other approaches exist as well, but we just wanted to give you some idea of how the meaning of a sentence might be handled by the models.

The actual use of sentence embedding models is rather simple, as you’ll see soon. We’ll embed a few sentences and afterwards check their correlations. The corresponding code file is

05_VectorDatabases\30_Embedding\20_sentence_similarity.py.

A new package is used in this example, `sentence_transformers`, as shown in [Listing 5.31](#). Via this package, you’ll have access to many different models focusing on embedding complete sentences.

```
### (1) Packages
from sentence_transformers import SentenceTransformer
import numpy as np
import seaborn as sns
```

Listing 5.31 Sentence Embedding: Required Packages (Source: *05_VectorDatabases\30_Embedding\20_sentence_similarity.py*)

We’ll use an embedding model that can handle multiple languages and create an instance of that model. The chosen model has an

embedding size of 512, so each sentence is transformed into a 512-dimensional vector, as follows:

```
##% (2) Load the model
MODEL = 'sentence-transformers/distiluse-base-multilingual-cased-v1'
model = SentenceTransformer(MODEL)
```

Eight example sentences are chosen, as shown in [Listing 5.32](#). Feel free to define your own sentences. The first three are pretty similar to each other. They have similar meanings although they use different words to describe a situation. The same can be said about the last three sentences. They are in fact translations from English into French and German. The other sentences are very different.

```
# %% (3) Define the sentences
sentences = [
    'The cat lounged lazily on the warm windowsill.',
    'A feline relaxed comfortably on the sun-soaked ledge.',
    'The kitty reclined peacefully on the heated window perch.',
    'Quantum mechanics challenges our understanding of reality.',
    'The chef expertly julienned the carrots for the salad.',
    'The vibrant flowers bloomed in the garden.',
    'Las flores vibrantes florecieron en el jardín. ',
    'Die lebhaften Blumen blühten im Garten.'
]
```

Listing 5.32 Sentence Embeddings: Sample Sentences (Source: 05_VectorDatabases\30_Embedding\20_sentence_similarity.py)

We create the embeddings by calling the `encode` method of our model and passing our list of sentences. Checking the structure of the data is always helpful. We have 8 sentences, and earlier, we found the embedding size of our model. Thus, our embeddings have the shape of (8, 512), as follows:

```
# %% (4) Get the embeddings
sentence_embeddings = model.encode(sentences)
```

Let's say we want to visualize the correlation between the sentences. For this task, we'll calculate the correlation coefficient between the eight different sentences first, as shown in [Listing 5.33](#).

```
# %% (5) Calculate linear correlation matrix for embeddings
sentence_embeddings_corr = np.corrcoef(sentence_embeddings)
import seaborn as sns
# show annotation with one digit
sns.heatmap(sentence_embeddings_corr, annot=True,
            fmt=".1f",
            xticklabels=sentences,
            yticklabels=sentences)
```

Listing 5.33 Sentence Embeddings: Correlation and Visualization (Source: 05_VectorDatabases\30_Embedding\20_sentence_similarity.py)

Our result, shown in [Figure 5.15](#), is an 8×8 matrix that you can visualize with Seaborn.

Notice the high correlation coefficients in the upper-left corner of the graph. These high values reflect the similarity of the first three sentences. Similarly, you can observe high correlation coefficients in the bottom right of the graph. These three sentences have the exact same meaning; they are just formulated in three different languages. Since we applied an embedding model that is trained on datasets from different languages, it can understand multiple languages. We can conclude that a multilingual model does not care about languages, only about meaning. The other sentences are added just

for comparison. They show that sentences that are very different have correlation coefficients around zero.




Figure 5.15 Sentence Embeddings Correlations (Source: 05_VectorDatabases/30_Embedding/20_sentence_similarity.py)

5.5.4 Coding: Create Embeddings with LangChain

Earlier in [Section 5.3.4](#), we loaded an article from Wikipedia on artificial intelligence. We'll continue with this example and extract the text from our list of `Document` objects, so that each text is embedded. We'll start by extending some boilerplate code available in `05_VectorDatabases/30_Embedding/30_wikipedia_embeddings.py`.

In this code, we'll load an additional class to enable the use of OpenAI embeddings. The remainder of the code should by now be familiar to you. The Wikipedia article on artificial intelligence is loaded, and subsequently split.

We'll use OpenAI embeddings, for which you'll need an API key from OpenAI stored in the `.env` file. This file is loaded with `load_dotenv()`, as shown in [Listing 5.34](#).

```
### Packages
from langchain.document_loaders import WikipediaLoader
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain.embeddings import OpenAIEmbeddings
from pprint import pprint
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
```

Listing 5.34 Required Packages (Source:
05_VectorDatabases/30_Embedding/30_wikipedia_embeddings.py)

From Wikipedia, we'll load an article on "Artificial Intelligence," as shown in [Listing 5.35](#).

```
# %% Load the article
ai_article_title = "Artificial_intelligence"
loader = WikipediaLoader(query=ai_article_title,
                        load_all_available_meta=True,
                        doc_content_chars_max=10000,
                        load_max_docs=1)

doc = loader.load()
```

Listing 5.35 Loading the Article (Source:
05_VectorDatabases/30_Embedding/30_wikipedia_embeddings.py)

The document object is split into chunks, as shown in [Listing 5.36](#). You know this step by now.

```
# %% Create splitter instance
splitter = RecursiveCharacterTextSplitter(chunk_size=1000,
                                         chunk_overlap=200,
                                         separators=["\n\n", "\n", " ", ".", ","])

# %% Apply semantic chunking
chunks = splitter.split_documents(doc)
# %% Number of Chunks
len(chunks)
```

16

Listing 5.36 Creating Text Chunks (Source:
05_VectorDatabases/30_Embedding/30_wikipedia_embeddings.py)

As a result, we have 16 chunks that we'll embed. But first we must create an instance of the embedding class. In this step, we can apply an OpenAI embedding model `text-embedding-3-small`. Although `small` is not quite accurate; this embedding model has 1,536 dimensions. You might also use its big brother `text-embedding-3-large`, which has 3,072 dimensions.

```
# %% Create instance of embedding model
embeddings_model = OpenAIEmbeddings(model="text-embedding-3-small")
```

We cannot simply pass a list of documents into the model. We'll use the method `embed_documents`, which has one required parameter `text`. This parameter is expected to be a list of strings. Thus, you must extract the `page_content` from each chunk and store this content in the object `texts`. This object is passed as parameter to the embedding model to create the embeddings, as follows:

```
# %% extract the texts from "page_content" attribute of each chunk
texts = [chunk.page_content for chunk in chunks]
# %% create embeddings
embeddings = embeddings_model.embed_documents(texts=texts)
```

Amazing! You can check that the length of this object (16) corresponds to the number of text strings we passed to the embedding model, as shown in [Listing 5.37](#). Thus, for each text, an embedding vector was created that has 1,536 dimensions.

```
# %% get number of embeddings
len(embeddings)
```

16

```
# %% check the dimension of the embeddings
len(embeddings[0])
```

1536

Listing 5.37 Embeddings: Dimensions (Source: 05_VectorDatabases/30_Embedding/30_wikipedia_embeddings.py)

Thus concludes our section on embeddings. Let's proceed with the next step of the pipeline and learn how to store data in a vector database.

OceanofPDF.com

5.6 Storing Data

We are reaching our final state of the data ingestion pipeline. Let's recap what we've achieved so far. First, we loaded some documents. Second, we created smaller chunks out of a large text corpus. Third, we learned how to create vector embeddings based on text chunks. So now, we can store them in a database.

You have multiple options to store the data. The main decision you must make is whether you want to host your own vector database locally or if you want to rely on web-based service providers. Both approaches have their pros and cons.

5.6.1 Selection of a Vector Database

The basic functionality of a vector database is handled similarly by different providers. But they all have their unique selling propositions; the data you want to cover and your requirements will determine the best fit. Some parameters to consider include the following:

- **Location**

As mentioned earlier, some databases can store your data locally (e.g., in a Chroma database), while others are web-based services (e.g., Pinecone and Weaviate).

- **Data protection**

Depending on the type of data, and the customer you're working with, your data might be highly sensitive, and you might need to avoid online services altogether or avoid a specific region. In such cases, you might consider using a local vector database rather than an online service like Pinecone or Weaviate.

- **Price**

The cost of a service is always a driving factor. Determining the overall price might not be a straightforward task. Although local open-source vector databases might be free of charge, you need to maintain an additional server and equip it with the proper hardware for fast embedding. Basically, you'll trade higher upfront costs (the cost of running the local vector databases) for variable costs (the cost of online vector database services).

- **Performance versus speed**

The classic tradeoff: Better performance is typically correlated with larger embedding space and thus slower data upserting (a portmanteau of “update” and “insert”) and data querying.

We'll implement a local database as well as use a web-service for a vector database, so you can decide which approach to implement.

5.6.2 Coding: File-Based Storage with a Chroma Database

You can find the code file at *05_VectorDatabases/40_VectorStore/chroma_db_intro.py*. As before, we'll load the relevant packages. A new functionality in this example is `HuggingFaceEndpointEmbeddings`, which provides access to models we can run locally, as shown in [Listing 5.38](#). Also, we'll load Chroma as our vector database.

```
### Packages
import os
from langchain.document_loaders import TextLoader
from langchain_text_splitters import RecursiveCharacterTextSplitter
from langchain_huggingface import HuggingFaceEndpointEmbeddings
from langchain.vectorstores import Chroma
```

Listing 5.38 Data Storing: Required Packages (Source: *05_VectorDatabases/40_VectorStore/chroma_db_intro.py*)

We'll work with *The Hound of the Baskervilles*. We need to navigate from the script file to the book path, as shown in [Listing 5.39](#).

```
### Path Handling
# Get the current working directory
file_path = os.path.abspath(__file__)
current_dir = os.path.dirname(file_path)

# Go up one directory level
parent_dir = os.path.dirname(current_dir)
text_file_path = os.path.join(parent_dir, "data", "HoundOfBaskerville.txt")
```

Listing 5.39 Storing Data: Path Handling (Source: 05_VectorDatabases/40_VectorStore/chroma_db_intro.py)

We create a loader instance and subsequently call its `load` method, as follows:

```
### load all files in a directory
loader = TextLoader(file_path=text_file_path,
                    encoding="utf-8")
docs = loader.load()
```

We want to split the document based on the structure-based chunking approach and get a list of all chunks, as shown in [Listing 5.40](#).

```
# %% Set up the splitter
splitter = RecursiveCharacterTextSplitter(chunk_size=1000,
                                          chunk_overlap=200,
                                          separators=["\n\n", "\n", " ", ".", ","])
chunks = splitter.split_documents(docs)
```

Listing 5.40 Data Storing: Data Splitting (Source: 05_VectorDatabases/40_VectorStore/chroma_db_intro.py)

Let's check the number of chunks that were created, as follows:

```
# %% check size of chunks
len(chunks)
```

The complete book has been split into 520 chunks that we'll now embed and add to our vector store. But first we must decide on an embedding model. In our case, we chose a famous sentence transformer embedding model that is small enough to run locally:

```
# %% define the embedding model
embedding_function = HuggingFaceEndpointEmbeddings(model="sentence-transformers/all-
MiniLM-L6-v2")
```

The preparation work is done, and we can now create an instance of our vector database. The vector database created is based on a Chroma database. In the instantiation, you must define a path for the vector database because its files will be written on your hard disk. Also, the vector database needs to know which embedding function will be used, as follows:

```
### vector database instance
persistent_db_path = os.path.join(parent_dir, "db")
db = Chroma(persist_directory=persistent_db_path,
embedding_function=embedding_function)
```

Now, interact with the database instance and add all the documents by calling `add_documents`, as follows:

```
# %% add all documents
db.add_documents(chunks)
```

This process takes some time. Once done, check the number of created documents with the method `get`. This method returns a dump of all documents in the database. Each document has a property named `ids`, so that you can count the number of `ids` to determine the total number of documents, as follows:

```
# %%
len(db.get()['ids'])
```

5.6.3 Coding: Web-Based Storage with Pinecone

While a Chroma database is an open-source vector database that you can run locally on your computer, Pinecone is an online managed service. To use this service, head over to <https://www.pinecone.io/> and register. As always, when you want to use an online service via Python, you must create an API key, which authenticates you as an authorized user and controls your access.

Once you've copied your API key to the clipboard, in your working folder, open a local environment file `.env` and create a new line in it, as follows:

```
PINECONE_API_KEY=...
```

Let's develop the script to load our book data into Pinecone. The script file is `05_VectorDatabases/40_VectorStore/pinecone_intro.py`. In our program code, we can fetch this API key with the package `dotenv` by loading this `.env` file, as shown in [Listing 5.41](#).

```
### packages
from dotenv import load_dotenv
import os
load_dotenv(".env")
# %%
```

Listing 5.41 Data Storing with Pinecone: Required Packages (Source: `05_VectorDatabases/40_VectorStore/pinecone_intro.py`)

We might check that the key is available with `os.getenv()`, as follows:

```
os.getenv("PINECONE_API_KEY")
```

To use Pinecone, we'll use the package `pinecone`. If you've already installed the environment, you're all set and can use it immediately. Otherwise, you'll need to instantiate Pinecone and pass your API key, as follows:

```
from pinecone import Pinecone, ServerlessSpec
pc = Pinecone(api_key=os.getenv("PINECONE_API_KEY"))
```

An *index* in Pinecone is the highest-level organizational unit of data. The index defines the embedding dimensions of the vectors, the similarity metric, and where your data is going to be stored. We're using the Starter level, which is free to use, and thus are limited in storage size as well as limited in the cloud provider and region we can use. Thus, we are using Amazon Web Services (AWS) as our cloud provider with the region us-east-1. As the similarity metric, we chose cosine similarity. As our embedding model, we chose all-MiniLM-L6-v2, which has 384 dimensions. All these parameters are defined during index creation. We created an index called "Sherlock" and to make our code as flexible as possible, we'll first check if this index already exists, as shown in [Listing 5.42](#). When you run the code for the first time, the index will be created (because one doesn't already exist). In subsequent use, the index won't be re-created, and you can directly use it.

```
index_name = "sherlock"
if index_name not in pc.list_indexes().names():
    pc.create_index(name=index_name,
                    metric="cosine",
                    dimension=384,
                    spec=ServerlessSpec(
                        cloud = "aws",
                        region="us-east-1"))
```

Listing 5.42 Pinecone Database: Index Creation (Source: 05_VectorDatabases/40_VectorStore/pinecone_intro.py)

You can check in the web frontend whether the index was created successfully.

Now that you have an index, let's add data that we can embed and their corresponding embedding vectors. We've created a data preparation function, which loads *The Hound of the Baskervilles*, splits the text, and returns a list of documents, as follows:


```
from data_prep import create_chunks
chunks = create_chunks("HoundOfBaskerville.txt")
```

Next, we'll create the embeddings and extract all `page_content` properties of the documents, as follows:

```
texts = [chunk.page_content for chunk in chunks]
```

As the embedding model, we choose `all-MiniLM-L6-v2` and load the corresponding class from `langchain_huggingface`, as follows:

```
# %% Embedding model
from langchain_huggingface import HuggingFaceEndpointEmbeddings
embedding_model = HuggingFaceEndpointEmbeddings(model="sentence-transformers/all-
MiniLM-L6-v2")
```

The embeddings are created by calling the method `embed_documents` and passing all texts into it, as follows:

```
# %% create all embeddings
embeddings = embedding_model.embed_documents(texts=texts)
```

At this point, we have 520 embedding vectors with 384 elements each.

We have all pieces compiled and can now upsert them to Pinecone. In this context, upserting is a database operation that updates an entry if it already exists or inserts an entry if one does not exist yet. But first, we need a good understanding of how Pinecone expects vectors to look.

Pinecone expects vectors that are lists of dictionaries, in which each dictionary consists of an `id`, `values`, and `metadata`.

The `id` must be unique. The embeddings correspond to the `values`, and `metadata` can be passed as well and should hold further information on the document, for instance, the human-readable text associated to the embedding vector. Let's create the `vectors` object using the code shown in [Listing 5.43](#).

```
# %% create vectors
# {"id": str, "values": List[float], "metadata": Dict[str, str]}
vectors = [{"id": str(i),
            "values": embeddings[i],
            "metadata": chunks[i].metadata}
            for i in range(len(chunks))]
```

Listing 5.43 Pinecone Database: Vector Creation (Source: 05_VectorDatabases/40_VectorStore/pinecone_intro.py)

These vectors are now upserted to Pinecone, and we need to connect to our index “sherlock,” as follows:

```
index = pc.Index(name=index_name)
index.upsert(vectors)
```

When you upsert larger documents, we recommend upserting the data in batches. Since we only upserted 520 documents, batching was not necessary.

Another simplification we applied is that we did not define a namespace. Why would you need a namespace? Imagine working on a production system and use a vector database that holds the data of several different customers (multitenancy). In such a case, you must ensure that customer A does not get accidental access to customer B’s data. You might solve this problem by storing data from customer A in a different namespace than customer B’s data.

Finally, let’s check if the operation succeeded. For this task, use the method `describe_index_stats()`, as follows:

```
# %% describe index
print(index.describe_index_stats())
```

This method returns some statistics about the index’s content, as follows:

```
{'dimension': 384,
 'index_fullness': 0.0,
 'namespaces': {'': {'vector_count': 520}},
 'total_vector_count': 520}
```

We have successfully upserted our book to Pinecone, as before to Chroma database. In the next section, we'll explore how you can retrieve data from your databases.

OceanofPDF.com

5.7 Retrieving Data

In this section, you'll learn how data is retrieved from a vector database. To understand this topic in detail, you'll need to understand what similarities are, how they are derived, and how to use them during data retrieval. Let's get started!

5.7.1 Similarity Calculation

Before we retrieve data from a database, you must understand how a database finds relevant documents. Relevant documents do not necessarily mean finding exact matches. While traditional structured databases always search for exact matches, vector databases search for most relevant documents rather than perfect matches. Different approaches are available for finding these documents. We'll cover two of the most relevant ones:

- Cosine similarity
- Maximum margin relevance (MMR)

Cosine Similarity

Let's start with cosine similarity. Remember: the cosine value is largest when the angle is close to zero. In a two-dimensional coordinate system, two vectors pointing in the same direction have an inner angle of close to zero and are thus considered similar. If the two vectors point in completely opposite directions, they are considered *maximally dissimilar*. Cosine similarity is always a value between -1 and $+1$.

As you know, the embedding space is much larger than two dimensions, but the same concept holds true. [Figure 5.16](#) shows a vector corresponding to a user query in red—simplified into two dimensions. The angle alpha measures the deviation of the query vector and embedding S1. Similarly, angle beta measures the deviation of the query vector and embedding S5. The angles for embeddings S2, S3, and S5 are not shown.




Figure 5.16 Principle of Cosine Similarity Calculation and Retrieval

The retrieval process to find the most relevant documents now follows these steps:

1. Calculating angles (and their corresponding cosine values) between query vectors and all embedding vectors
2. Sorting the embedding vectors in decreasing order
3. Returning the top-k results

In our example, we defined a threshold of $\text{top_k} = 2$, meaning that only the two closest documents are returned.

Cosine similarity helps find the most relevant documents but in some cases might not be enough. In these cases, MMR can help.

Maximum Margin Relevance

MMR aims to maximize not just the relevance, but also the diversity, of documents. In what cases can diversity be important? Let's consider an example for using MMR, as shown in [Figure 5.17](#).




Figure 5.17 Maximum Margin Relevance Example

Let's say we have a vector database on attractions in New York City. You define a user query "What are the main attractions in New York City?" and retrieve most relevant documents. As a similarity metric, you choose cosine similarity and end up with the top 3 results, of which two are skyscrapers. This result possibly represents the most similar results, but you don't want to visit two skyscrapers. New York has such a great variety of attractions. Thus, we need a metric that provides results based on both similarity and diversity. MMR to the rescue! Based on MMR, you retrieve a landmark, but also a park and a cultural district.

How does this calculation work? In MMR, you have access to a hyperparameter λ , which is a weighting parameter in the range of 0 and 1. The value 0 returns only relevant results, and 1 returns the most diverse results. As a starting point, you might go with $\lambda = 0.5$ for a good balance between relevance and diversity.

5.7.2 Coding: Retrieve Data from Chroma Database

Let's extend the database we created in the last section. You can run the code from this file:

05_VectorDatabases/50_RetrieveData/10_chromadb_retrieval.py.

The relevant packages that we need to load for this script are Chroma as the vector store and HuggingFaceEndpointEmbeddings for the embeddings, as shown in [Listing 5.44](#).

```
## packages
from langchain.vectorstores import Chroma
import os
from pprint import pprint
from langchain_huggingface import HuggingFaceEndpointEmbeddings
```

Listing 5.44 Data Retrieval from a Chroma Database: Packages (Source: *05_VectorDatabases/50_RetrieveData/10_chromadb_retrieval.py*)

In the parent directory, we created the folder `db` in which the vector database is stored. The user query needs to be embedded. For that functionality, we create an instance for the embedding function, as shown in [Listing 5.45](#). An essential step is to use the same embedding model that you used for creating the book embeddings. An instance of the vector database is created, which is linked to the persistent storage and the embedding model.

```
# %% set up database connection
# Get the current working directory
file_path = os.path.abspath(__file__)
current_dir = os.path.dirname(file_path)
parent_dir = os.path.dirname(current_dir)
chroma_dir = os.path.join(parent_dir, "db")
# Go up one directory level
parent_dir = os.path.dirname(current_dir)
# set up the embedding function
embedding_function = HuggingFaceEndpointEmbeddings(
    model="sentence-transformers/all-MiniLM-L6-v2")
# connect to the database
```

```
db = Chroma(persist_directory=chroma_dir,
            embedding_function=embedding_function)
```

Listing 5.45 Data Retrieval from a Chroma Database: Connection Setup (Source: 05_VectorDatabases/50_RetrieveData/10_chromadb_retrieval.py)

[Listing 5.46](#) shows how to find the most relevant chunk, which is now unbelievably simple. You only need to create a retriever, which is the endpoint that provides the `invoke` method that you'll use afterwards.

```
retriever = db.as_retriever()
# %% find information
# query = "Who is the sidekick of Sherlock Holmes in the book?"

# # thematic search
# query = "Find passages that describe the moor or its atmosphere."

# # Emotion
# query = "Which chapters or passages convey a sense of fear or suspense?"

# # Dialogue Analysis
# query = "Identify all conversations between Sherlock Holmes and Dr. Watson."
```

Listing 5.46 Data Retrieval from a Chroma Database: Retrieval (Source: 05_VectorDatabases/50_RetrieveData/10_chromadb_retrieval.py)

Into this `invoke` method, we can pass a user query. In the code chunks, you can find multiple queries. You can uncomment the queries that you want to run. As shown in [Listing 5.47](#), we want to ask about the hound and store the returned documents in `most_similar_docs`, which is again a list of Document objects.

```
# Character
query = "How does the hound look like?"
most_similar_docs = retriever.invoke(query)
# %%
pprint(most_similar_docs[0].page_content)
```

Listing 5.47 Data Retrieval from a Chroma Database: Most Similar Documents (Source: 05_VectorDatabases/50_RetrieveData/10_chromadb_retrieval.py)

This implementation was quick, and we ran it with default values for search type (default: "similarity") and the number of documents to be

returned (default: 4). But how is the similarity calculated? We'll explore this topic next.

5.7.3 Coding: Retrieve Data from Pinecone

We'll extend on the example described earlier in [Section 5.6.3](#) and retrieve data from our Pinecone index. If you did not run the code from that section, please follow those steps first. You can find the script in

05_VectorDatabases/50_RetrieveData/20_pinecone_retrieval.py.

We'll import relevant packages and ensure that we have the environment file *.env* in our working folder with the `PINECONE_API_KEY`, as shown in [Listing 5.48](#).

```
### packages
from pinecone import Pinecone
from dotenv import load_dotenv
load_dotenv(".env")
from langchain_huggingface import HuggingFaceEndpointEmbeddings
import os
```

Listing 5.48 Pinecone Retrieval: Required Packages (Source: *05_VectorDatabases/50_RetrieveData/20_pinecone_retrieval.py*)

An instance of Pinecone is created and an index-instance points to our “sherlock” index we created earlier, as follows:

```
### connect to Pinecone instance
pc = Pinecone(api_key=os.getenv("PINECONE_API_KEY"))
index_name = "sherlock"
index = pc.Index(name=index_name)
```

We need to embed the `user_query` because, in the querying process, we'll pass the embedding vector. Make sure you use the same embedding model you used during the data storing process. Let's find out what the database tells us in response to the question, “What does the hound looks like?” as shown in [Listing 5.49](#).

```

### Embedding model
embedding_model = HuggingFaceEndpointEmbeddings(model="sentence-transformers/all-
MiniLM-L6-v2")
### embed user query
user_query = "What does the hound look like?"
query_embedding = embedding_model.embed_query(user_query)

```

Listing 5.49 Pinecone Retrieval: Embedding (Source: 05_VectorDatabases/50_RetrieveData/20_pinecone_retrieval.py)

The object `query_embedding` is a numeric vector with 384 elements.

We can search for the most similar documents by calling the `index.query()` method. As a parameter, the embedding vector is passed. Also, we want to limit the result to the top 2 most relevant documents by setting the parameter `top_k` to 2. Since the text is stored in the metadata, we must set the parameter `include_metadata` to true, as follows:

```

### search for similar documents
res = index.query(vector=query_embedding, top_k=2, include_metadata=True)

```

The resulting object `res` holds the key “matches”:

```
res["matches"]
```

This object is a list of dictionaries with the keys `id`, `metadata`, `score`, and `values`. We are interested in the metadata, specifically the text stored in the metadata. With the following code, we can print all the texts:

```

for match in res['matches']:
    print(match['metadata']['text'])
    print("-----")

```

The database returns the most relevant documents based on the user query, as shown in [Listing 5.50](#).

hound, but not such a hound as mortal eyes have ever seen. Fire burst from its open mouth, its eyes glowed with a smouldering glare, its muzzle and hackles and dewlap were outlined in flickering flame. Never in the delirious dream of a disordered brain could anything more savage, more appalling, more hellish be conceived than

that dark form and savage face which broke upon us out of the wall of fog. -----
----- quality upon earth it is common sense, and nothing will persuade me to
believe in such a thing. To do so would be to descend to the level of these poor
peasants, who are not content with a mere fiend dog but must needs describe him with
hell-fire shooting from his mouth and eyes. Holmes would not listen to such fancies,
and I am his agent. But facts are facts, and I have twice heard this crying upon the
moor. Suppose that there were really some huge hound loose upon it; that would go
far to explain everything. But where could such a hound lie concealed, where did it
get its food, where did it come from, how was it that no one saw it by day? It must
be confessed that the natural explanation offers almost as many difficulties as the
other. And always, apart from the hound, there is the fact of the human agency in
London, the man in the cab, and the letter which warned Sir Henry -----

Listing 5.50 Pinecone Retrieval: Retrieved Chunks (Source:
05_VectorDatabases/50_RetrieveData/20_pinecone_retrieval.py)

In that output, you see some good representation of how the hound appears, but it is not a well formulated answer. For refinement, a technique called retrieval-augmented generation (RAG) was developed, which will be the focus of the next chapter.

OceanofPDF.com

5.8 Capstone Project

In this section, we combine all the knowledge you learned in this chapter and integrate it into one big project. We'll start by showing you what we'll develop.

[Figure 5.18](#) shows the screen for our final application, in which you can see separate sections for the user query **3**, results **2**, and filters **1**. In this application, you can describe the plot of a movie and receive the most similar movies in a nice output. Before developing any application, however, an important step is to think about the requirements and features you want to include.




Figure 5.18 Sample Output from Our Movie App

We'll build it up step by step, starting with discussing the features of our app. Then, we'll get some understanding of the dataset that will be used. The data preprocessing and the preparation of vector database follow. Finally, we'll develop the app frontend on top of our business logic. Let's start by discussing the features.

5.8.1 Features

Even the most complex apps follow a basic input-process-output model. Thus, the features can be structured into different groups like input “input widgets,” “output widgets,” and “data processing,” as follows:

- **Input widgets**

- A user can add a description of a movie in an input field, which then searches for movies in a vector database that best represents the description. Additionally, the data is narrowed down according to some filters.
- One filter leaves out “bad” movies based on IMDB ratings. A slider is used to define a minimum rating.
- Also, the genre of the movies can be selected in a dropdown menu. The resulting films should be filtered according to the selected genre.

- **Data processing**

- The movie data is processed and stored in a vector database.
- Specifically, the title, poster, genres, and IMDB rating are stored as metadata.
- The plot is embedded and stored together with the text in the database.

- Most similar movies are received based on the user query, and filters.
- If multiple identical movies are found, only unique sets should be returned.
- **Output widgets**
 - The three most similar movies should be presented.
 - The movie title should be shown in bold.
 - If a poster URL is available, the poster should be shown on the screen.
 - The genres, IMDB rating, and similarity score are shown below the poster.

Usually, a good approach is to process the data first, before the application is created. We'll follow that advice and prepare our database next.

5.8.2 Dataset

We'll base our work on a dataset from Hugging Face:

MongoDB/embedded_movies

(https://huggingface.co/datasets/MongoDB/embedded_movies). The dataset is licensed under the Apache 2.0 license (<https://huggingface.co/datasets/choosealicense/licenses/blob/main/markdown/apache-2.0.md>).

This dataset contains the details of movies, including their genres. The dataset has information about the plot, the genres that describe the movie best, and many more features. We'll use several properties, in the following ways:

- `fullplot` (will be 'document'; used as embedding)

- title (metadata; shown as result)
- genres (metadata; for filtering)
- imdb_rating (metadata; for filtering)
- poster (metadata; shown as result)

We know everything we need to know about the dataset at this point, and we can proceed with the preparation of the database.

5.8.3 Preparing the Vector Database

Good, let's start coding. You can find the complete script in *05_VectorDatabases/90_CapstoneProject/10_data_prep.py*.

As before, we'll load the relevant packages. For using Hugging Face datasets, we'll use the package `datasets` and the function `load_dataset`, as shown in [Listing 5.51](#).

```
### packages
import os
from datasets import load_dataset
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain_huggingface import HuggingFaceEndpointEmbeddings
from langchain_chroma import Chroma
from langchain.schema import Document
```

Listing 5.51 Capstone Project: Required Packages (Source: *05_VectorDatabases/90_CapstoneProject/10_data_prep.py*)

Let's get the dataset loaded into Python. All we need to do is call `loading_dataset` with the dataset name, as shown in [Listing 5.52](#).

```
### load dataset
dataset = load_dataset("MongoDB/embedded_movies", split="train")
# license:
https://huggingface.co/datasets/choosealicense/licenses/blob/main/markdown/apache-2.0.md
```

Listing 5.52 Capstone Project: Data Loading (Source: *05_VectorDatabases/90_CapstoneProject/10_data_prep.py*)

Let's familiarize ourselves with the dataset and determine how many movies are part of the dataset, as follows:

```
## number of films in the dataset
len(dataset)
```

1500

We have quite a few movies to work with, but we now need to find out which keys are available, as shown in [Listing 5.53](#).

```
# %% which keys are in the dataset?
dataset[0].keys()

dict_keys(['plot', 'runtime', 'genres', 'fullplot', 'directors', 'writers',
'countries', 'poster', 'languages', 'cast', 'title', 'num_mflix_comments', 'rated',
'imdb', 'awards', 'type', 'metacritic', 'plot_embedding'])
```

Listing 5.53 Capstone Project: Checking The Dataset (Source: 05_VectorDatabases/90_CapstoneProject/10_data_prep.py)

Now, we want to convert the dataset into a list of document objects. For this step, we'll iterate over all the movies in the database. For each movie, we'll extract the properties to add as metadata. In some cases, no property is available, and in these cases, we must ensure that we have strings and not `None` values.

Another special case is `genres`. In this dataset, the genres are provided as a list of strings. Unfortunately, we cannot store lists of strings as metadata, so we must concatenate the different genres with semicolon separators, which is a design decision. Another approach for handling this could be to add a document for each movie and genre combination.

The design decision you make at this stage will have a significant impact on the processing of the data later. And in many cases, you may recognize later that you need to come back to the data processing.

The property `fullplot` is used as `page_content` in our documents and will be used for creating the embeddings. Only if `fullplot` is available is a document added to the list, as shown in [Listing 5.54](#).

```
# %% Create List of Documents
docs = []
for doc in dataset:
    title = doc['title'] if doc['title'] is not None else ""
    poster = doc['poster'] if doc['poster'] is not None else ""
    genres = ';'.join(doc['genres']) if doc['genres'] is not None else ""
    imdb_rating = doc['imdb']['rating'] if doc['imdb']['rating'] is not None else ""
    meta = {'title': title, 'poster': poster, 'genres': genres, 'imdb_rating':
imdb_rating}

    if doc['fullplot'] is not None:
        docs.append(Document(page_content=doc["fullplot"], metadata=meta))
```

Listing 5.54 Capstone Project: Document Creation (Source: 05_VectorDatabases/90_CapstoneProject/10_data_prep.py)

Now, we're on familiar terrain: We have a list of documents. As shown in [Listing 5.55](#), proceed with splitting the data. We'll use the already familiar `RecursiveCharacterTextSplitter`, which we'll use in combination with the `all-MiniLM-L6-v2` sentence-transformer embedding model. Thus, we'll choose a chunk size and a chunk overlap accordingly.

```
# %% Chunking
CHUNK_SIZE = 1000
CHUNK_OVERLAP = 200
docs_chunked = []
splitter = RecursiveCharacterTextSplitter(chunk_size=CHUNK_SIZE,
                                         chunk_overlap=CHUNK_OVERLAP,
                                         separators=["\n\n", "\n", " ", ".", ","])

chunks = splitter.split_documents(docs)
```

Listing 5.55 Capstone Project: Data Chunking (Source: 05_VectorDatabases/90_CapstoneProject/10_data_prep.py)

If you check the number of chunks, notice how we now have 1,587 chunks, which is slightly more than the number of 1,500 movies. Thus, for some movies, the `fullplot` was too long and had to be split

into two or more chunks. In the final output of the app, we must ensure that the same movie is not shown multiple times.

We'll set up an instance of a Chroma database and embed all the chunks into it. In this process, a subfolder *db* is created to hold our database. We can ensure that the movies are not added multiple times by first checking whether the folder exists and then adding the documents only if the folder does not exist, as shown in [Listing 5.56](#).

```
# %% store chunks in Chroma
embedding_function = HuggingFaceEndpointEmbeddings(model="sentence-transformers/all-
MiniLM-L6-v2")
script_dir = os.path.dirname(os.path.abspath(__file__))
db_dir = os.path.join(script_dir, "db")
if not os.path.exists(db_dir):
    os.makedirs(db_dir)
    db = Chroma(persist_directory=db_dir, embedding_function=embedding_function,
collection_name="movies")
    db.add_documents(chunks)
```

Listing 5.56 Capstone Project: Adding Chunks to the Database (Source: 05_VectorDatabases/90_CapstoneProject/10_data_prep.py)

A few minutes will be required to embed all the movies.

To check whether the creation process was successful, call the `get` method to get a full view of the data stored in the database, as follows:

```
# %% check the result
db.get()
```

To prepare for our upcoming application development, let's find out which genres are available by iterating over the dataset, as shown in [Listing 5.57](#).

```
##% get all genres
genres = set()
for doc in dataset:
    if doc['genres'] is not None:
        genres.update(doc['genres'])

{'Action', 'Adventure', 'Animation', 'Biography', 'Comedy', 'Crime', 'Documentary',
```

```
'Drama', 'Family', 'Fantasy', 'Film-Noir', 'History', 'Horror', 'Music', 'Musical',  
'Mystery', 'Romance', 'Sci-Fi', 'Short', 'Sport', 'Thriller', 'War', 'Western'}
```

Listing 5.57 Capstone Project: Genre Preprocessing (Source: 05_VectorDatabases/
90_CapstoneProject/10_data_prep.py)

5.8.4 Exercise: Get All Genres from the Vector Database

Before moving on with app development, let's work on a short exercise to test your knowledge of working with vector databases.

Task

In the previous section, we showed you how to extract genres from a dataset. But now, we want you to assume that you don't have access to the raw dataset, only to the vector database. Try to extract all the genres from the Chroma vector database.

Solution

Multiple solutions exist (as always). Our solution is based on getting all documents, which represents a complete dump from the database. Then, we create an empty set of genres. While iterating over all the metadata of the documents, each genre entry is extracted. Since we decided to store the genres as a string with entries separated by semicolons, we must reverse that step, split the genres, and update the set, as shown in [Listing 5.58](#).

```
# %% Exercise: Get all genres from the database  
documents = db.get()  
genres = set()  
  
for metadata in documents['metadatas']:  
    genre = metadata.get('genres')
```

```
genres_list = genre.split(';')
genres.update(genres_list)
```

Listing 5.58 Capstone Project: Extracting Available Genres (Source: 05_VectorDatabases/ 90_CapstoneProject/10_data_prep.py)

Amazingly, we have prepared our database and are ready to take the next step and integrate the vector database into the application.

5.8.5 App Development

Let's now turn our attention to the development of the application. For that purpose, we'll use a framework called Streamlit. This open-source framework for creating web applications is based on Python. No frontend development experience is required to develop a Streamlit app, but Streamlit should not be considered for creating a "complete" app, for instance, with authentication and multiple simultaneous users. But for quick prototyping, Streamlit is extremely useful and handy.

The complete app code can be found in *05_VectorDatabases/90_CapstoneProject/app.py*.

As always, we'll start the script by importing all the necessary packages. A lot of familiar packages and classes are available for interacting with our vector database (Chroma) or for using an embedding model (HuggingFaceEndpointEmbeddings). But the main package is Streamlit, which we import based on its typical alias `st`, so that we can access its methods quicker, as follows:

```
### packages
import streamlit as st
from langchain_chroma import Chroma
from langchain_huggingface import HuggingFaceEndpointEmbeddings
```

Before we build the web interface step by step, let's connect to our vector database. We'll use the same sentence transformer as earlier during data preparation in [Section 5.8.3](#). An instance of Chroma is created, making use of the folder `db` to hold the database and the collection `movies` we defined earlier, as shown in [Listing 5.59](#).

```
### load the vector database
embedding_function = HuggingFaceEndpointEmbeddings(model="sentence-transformers/all-
MiniLM-L6-v2")
db = Chroma(persist_directory="db", collection_name="movies",
embedding_function=embedding_function)
```

Listing 5.59 Capstone Project: Embedding and Database Instance Creation (Source: 05_VectorDatabases/90_CapstoneProject/10_data_prep.py)

The application is developed from top to bottom. We start by creating a static headline. All Streamlit components are accessible via `st`; for a title, we just need to call `st.title` and pass a title for the app, as follows:

```
### develop the app
st.title("Movie Finder")
```

Actually, at this point, you can already run the app. Just open the command line and run the following command:

```
streamlit run app.py
```

This command will open your default web browser and show the app: Not overly helpful, but we know now that the app runs. Whenever you save the app file, the server recognizes the change. As shown in [Figure 5.19](#), in the upper-right corner, you'll be informed that the source file changed, and you'll have the option to re-run the application rendered. We recommend you select **Always rerun**, so that each time you make a change to the application and save the file, the website is reloaded. In this way, you can interactively

develop applications and directly see the impact of your code changes.




Figure 5.19 Capstone Project: Streamlit App Rerun Options

The next component we'll implement is a slider for selecting a minimum IMDB movie rating, as shown in [Figure 5.20](#).




Figure 5.20 Capstone Project: Streamlit Slider for “Minimum IMDB Rating”

If you've worked with classic website development with HTML, JavaScript, and CSS, you know how cumbersome the development of such components can be. But not with Streamlit. Many components are boilerplates that you can implement with a single line of code—as in this case.

You'll need to apply `st.slider` and pass some parameters (like a label, min-value, and max-value); the initial selected value; and a step size. In this case, we want to fetch the user input, which is easily reached by assigning the return value of the function call to the variable `min_rating`. Later, we'll come back to this value and filter in the vector database query based on the filter.

```
# Add a slider for minimum IMDB rating
min_rating = st.slider("Minimum IMDB Rating", min_value=0.0, max_value=10.0,
value=7.0, step=0.1)
```

The next filter is the genre filter, an example of which is shown in [Figure 5.21](#). This filter narrows the query results based on the selected genre of the movie.

At the end of [Section 5.8.4](#), we extracted all the available genres from the dataset. We can define them in a list and pass this list as a parameter to `st.selectbox`. The function call returns the selected genre that you can store for later use in the variable `selected_genre`, as shown in [Listing 5.60](#).

```
# Add a single-select input for genres
genres = ['Action', 'Adventure', 'Animation', 'Biography', 'Comedy', 'Crime',
          'Documentary', 'Drama', 'Family', 'Fantasy', 'Film-Noir', 'History',
          'Horror', 'Music', 'Musical', 'Mystery', 'Romance', 'Sci-Fi', 'Short',
          'Sport', 'Thriller', 'War', 'Western']
selected_genre = st.selectbox("Select a genre", genres)
```

Listing 5.60 Capstone Project: Genre Dropdown Implementation (Source: 05_VectorDatabases/90_CapstoneProject/10_data_prep.py)




Figure 5.21 Capstone Project: Streamlit Dropdown Box for Genre Filter

In the next step, we'll define the user input, as shown in [Figure 5.22](#), which is the chat input query that we'll eventually pass into the vector database query call. The chat input is also a preconfigured component with a specific look and feel.




Figure 5.22 Capstone Project: Streamlit Chat Input Field

And you guessed it: you can get the component and its return value with a single line of code:

```
user_query = st.chat_input("What happens in the movie?")
```

The app will react in such a way that the most similar movies are searched when the user presses or the **Send** button is clicked. Both results in storing the chat input in the variable `user_query`. We can ensure that the code block has not already been run by placing it in the `if user_query` block. Then, we can define the metadata for the `imdb_rating` based on the user selection and its corresponding variable `min_rating`.

This metadata is now applied in the vector database through the `similarity_search_with_score` function together with the `user_query` variable. The number of returned movies is set to 10. Finally, we'll present the top 3 most similar movies, but we cannot know how many duplicates we need to remove. So, the arbitrary choice of 10 is just for making sure that sufficient results will be available.

Duplicated entries occurred in the process of creating chunks when the full plot description was too long for the model context window. Thus, a movie's description might be causing a split into several Documents, of which we only want to show one unique movie in the app. The query result is stored in `similar_movies`.

We're still not done with the data processing; we must take care of the genre filter. This task is achieved in [Listing 5.61](#) with a list comprehension in which we keep a movie only if the selected genre is part of the movie genre's metadata.

```
if user_query:
    # Retrieve the most similar movies
    with st.spinner("Searching for similar movies..."):
        metadata_filter = {"imdb_rating": {"$gte": min_rating}}
        similar_movies = db.similarity_search_with_score(user_query, k=10,
filter=metadata_filter)
        # filter for selected genre
```



```
similar_movies = [movie for movie in similar_movies if selected_genre in
movie[0].metadata['genres']]
```

Listing 5.61 Capstone Project: Implementing Genre Filter (Source: 05_VectorDatabases/90_CapstoneProject/app.py)

At this point, we have a list of movies that we can display.

[Listing 5.62](#) shows the main app components. In the header, we present the user query as a reference. Since only unique results should be presented to the user, we must iterate over the movie result list and check whether a title has already been seen.

In three columns, the three most similar movies are shown. A poster image can be created with `st.image`. The title, genre, rating, and similarity score can be printed with either `st.markdown` or `st.write`. A warning is created with `+` to inform the user of how many duplicated entries have been removed.

```
# Display the results
st.header(f"Most Similar Movies: ")
st.subheader(f"Query: '{user_query}'")
cols = st.columns(3)
# Check if there are duplicate results
unique_results = []
seen_titles = set()
for doc, score in similar_movies:
    if doc.metadata['title'] not in seen_titles:
        unique_results.append((doc, score))
        seen_titles.add(doc.metadata['title'])

# Display only unique results
for i, (doc, score) in enumerate(unique_results):
    if i >= len(cols):
        break
    with cols[i]:
        if doc.metadata['poster']:
            try:
                st.image(doc.metadata['poster'], width=150)
            except:
                st.write("No poster available")
        else:
            st.write("No poster available")
        st.markdown(f"***{doc.metadata['title']}***")
        st.write(f"Genres: {doc.metadata['genres']}")
        st.write(f"IMDB Rating: {doc.metadata['imdb_rating']}")
```

```
st.write(f"Similarity Score: {score:.4f}")

if len(unique_results) < len(similar_movies):
    st.warning(f"Note: {len(similar_movies)-len(unique_results)} duplicate
result(s) were removed.")
```

Listing 5.62 Capstone Project: Streamlit App (Source:
05_VectorDatabases/90_CapstoneProject/app.py)

The app is ready. You can test the different filters and check which movies are returned based on your description. Maybe this app can help you pick your next movie.

Feel free to extend the app by adding additional filters or replacing the similarity search based on cosine distance to get a wider variety of results with an MMR search.

OceanofPDF.com

5.9 Summary

In this chapter, we explored the main concepts and processes involved in working with vector databases. We started with an overview of vector databases, developing a comprehensive understanding of how they work and their importance in handling unstructured data. They set the stage for a deep dive into the practical steps required to utilize these databases effectively.

We covered the process of loading documents into a vector database, emphasizing the importance of data ingestion as the foundation for subsequent operations. Once the documents are loaded, we discussed document splitting, a critical step for breaking down large pieces of data into manageable chunks, which facilitates more efficient processing and retrieval.

This chapter then introduced embeddings, a technique used to convert textual or other unstructured data into vector representations with the same semantic meaning. These embeddings are essential for enabling similarity searches, as they allow the data to be compared based on their vectorized form rather than exact matches of text.

Storing data efficiently within a vector database was another key focus, highlighting best practices and strategies to ensure data remains accessible and performant. We learned to work with two important frameworks: Chroma databases and Pinecone. Then, we explored data retrieval techniques, detailing how to query the vector database and obtain relevant results.

A significant portion of the chapter was dedicated to the similarity search, which is a core feature of vector databases. We explored

different similarity metrics and discussed how these search methods work to find entries that are most similar to a given user query.

Finally, we concluded this chapter with a capstone project, which provided a practical application of all the concepts covered. This project serves as a hands-on exercise, reinforcing the knowledge gained in this chapter and demonstrating how to implement a complete workflow in a real-world scenario. An important part of this process was the integration of the vector database and its data retrieval into a fully functional web application.

In the next chapter, we'll delve into RAG, a cutting-edge technique that combines the power of LLMs with the precision of information retrieval systems like vector databases. RAG allows for the dynamic generation of responses by augmenting the generative capabilities of language models with relevant data retrieved from external sources, such as vector databases.

6 Retrieval-Augmented Generation

Sometimes it is the people no one imagines anything of who do the things no one can imagine.

—Alan Turing in the movie The Imitation Game

In this chapter, we'll cover one of the most impactful concepts in the field of large language models (LLMs): retrieval-augmented generation (RAG). So far, we have collected different puzzle pieces like LLMs, prompt engineering, and vector databases, which will now fit perfectly in place.

One important aspect is that vector databases, covered in [Chapter 5](#), typically represent the backbone of a RAG system. For steering the model on how to create the final output, we'll come back to prompt engineering and reuse the concepts we learned in [Chapter 4](#). Finally, data will be passed to an LLM, and at that stage, we come back to knowledge we gained in [Chapter 4](#). In this regard, RAG is a mountain we can only climb after having gained fitness in these three different disciplines.

We'll start this chapter with an introduction to the concept of RAG in [Section 6.1](#). After gaining a general understanding, we'll proceed into the individual process steps of retrieval, augmentation, and generation.

We'll first develop a simple RAG system in [Section 6.2](#). This system will work surprisingly well but won't be perfect. For that reason, we'll present an overview of advanced RAG techniques in [Section 6.3](#). These improvements can take place in different stages of the RAG system. Advanced techniques for pre-retrieval phase are presented in [Section 6.3.1](#). In [Section 6.3.2](#), techniques for the retrieval phase will be shown; in [Section 6.3.3](#), techniques for the post-retrieval phase.

An alternative to RAG, called *prompt caching*, is the focus of our attention in [Section 6.4](#). When you develop a RAG system and want to improve it, you'll need metrics for evaluating and quantifying these improvements. Thus, we dedicated [Section 6.5](#) to RAG evaluation techniques.

Let's start with what RAG is, why it is needed, and how it works.

6.1 Introduction

Before we delve into how RAG works, let's pause for a second and reflect on why it is needed. Imagine you're working with a large knowledge source (e.g., a long document), and you want to "chat" with it.

The simplest approach would be to make an LLM call in which you pass your question as well as the complete document as the context. This approach is usually not feasible because the large document may exceed the context window of the LLM, and/or it will not be efficient to send in each request the complete document. The request will take comparably long to receive an answer but also could be quite costly because many unnecessary tokens are sent in every single request.

One special solution is the use of prompt caching, which we'll discuss in detail in [Section 6.4](#). But prompt caching is more a solution for edge cases. The more usual approach is to use RAG. With RAG, in each LLM request, only the most relevant documents are passed along with the user query to the LLM. So, RAG results in an efficient and quick solution while also overcoming some limitations of LLMs. With RAG, you have the following capabilities:

- Access to company-internal information
- Access to up-to-date information that was created after the knowledge cutoff data for LLM training
- More fact-based results

[Figure 6.1](#) shows the general RAG process. A user creates a query that will be passed to a retriever. The retriever is responsible for fetching relevant information from some external source.




Figure 6.1 RAG General Pipeline

In most cases, this external source is a vector database, but you're not limited to vector databases. Alternatively, you could connect to and search the internet for relevant information. Ideally, the retriever finds some relevant information.

The next stage of the pipeline is the augmentation step. In this step, we use prompt engineering to bundle instructions to the relevant documents. These instructions will tell an LLM how to process the relevant documents but also how to behave when the relevant documents are not a good match from which to extract information. As the result of the augmentation step, we have documents and corresponding instructions. These elements are now passed to the last stage of the pipeline: the LLM.

The LLM will process the instructions, extract information from the documents, and create an output according to the user query.

The external source is, in most cases, a vector database. Let's look at the RAG system with a vector database in the backend, as shown in [Figure 6.2](#).




Figure 6.2 Simple RAG System Based on a Vector Database Backend

The vector database, which represents a corpus of documents and their embeddings, must already be created, which we covered in [Chapter 5](#).

In an RAG system that relies on a vector database, the user sends a query to the system. This query is embedded based on the same embedding model that was used for embedding the documents. The

most similar documents to the query embedding are retrieved from the database. In the next step, the user query together with the most similar documents are passed to the LLM to formulate the response.

Now that you have a general idea, let's dive into the details:

1. Retrieval process

In the retrieval process, the user query results in a search against some data source. The user query, as a minimum, requires a text and a limitation indicating how many documents should be returned. The query can be searched for in a variety of different knowledge sources. The knowledge source represents a large number of documents. In this knowledge source, all documents are ranked based on their similarity to the user query.




Figure 6.3 Retrieval Process

The user query is embedded, and its embedding is compared to all embeddings of the documents in the knowledge source. Based on a similarity metric like cosine similarity, the documents are ranked. A predefined number of most relevant documents is extracted and passed to the next stage—the augmentation.

2. Augmentation process

In the augmentation step, the retrieved information is combined

with the original query to create an enhanced prompt. This phase begins with the integration of relevant documents or passages that the retrieval step has surfaced.

Typically, the retrieved information is clearly differentiated from the user's query. At this stage, prompt engineering comes into play and has a vital role. We need to instruct the LLM on how to process the information. Explicit instructions are passed about staying within the bounds of the information provided to avoid hallucinations. Also, we should clearly state how to respond if the information does not answer the user query.

The result is a prompt that provides the language model with both the context it needs to generate an accurate response and the guidance that is required to use that context effectively. The augmented prompt serves as the foundation for generating responses that are relevant to the user query and grounded in the retrieved information. In the next step, we'll take a closer look at the generation process.

3. Generation process

The generation process in RAG represents the final crucial stage where the language model transforms the augmented prompt into a coherent and contextually relevant response. Unlike standard LLM output generation in traditional language models, RAG-based generation incorporates the retrieved information while maintaining natural language fluency.

When the language model receives the augmented prompt, it begins the process of analyzing the retrieved context and the original query in tandem.

One of the key challenges during generation is maintaining consistency with the retrieved information while avoiding

hallucinations. The model must strike a good balance between purely extractive answers, like directly quoting the retrieved content, and more abstract responses that synthesize and reformulate the information. This balance often involves techniques where the model's generation is explicitly guided to stay within the boundaries of the provided context.

While all this sounds quite complicated, actually it isn't, as you'll see as we start implementing RAG in the next section.

OceanofPDF.com

6.2 Coding: Simple Retrieval-Augmented Generation

Let's start developing our very first RAG application by mapping out what we'll create. Our small RAG system will be backed by a vector database. This vector database will represent a knowledge base for human history, populated with knowledge fetched from Wikipedia. We'll load a number of Wikipedia articles on the topic and store them in the database. Later, we'll set up a RAG system that will interact with these documents and answer questions based on that represented knowledge. The RAG system will also clearly state if it does not know how to answer the question.

6.2.1 Knowledge Source Setup

First things first, we've got to pack our bags with all the tools we need. [Listing 6.1](#) shows all packages we'll need in this script. You can already deduce many aspects of our implementation, for example, that we'll use Groq as LLM backend and that we'll use OpenAI embeddings for embedding our texts.

```
### packages
import os
from langchain_community.document_loaders import WikipediaLoader
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain_community.vectorstores import Chroma
from langchain_openai import OpenAIEmbeddings
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
from langchain_groq import ChatGroq
from langchain_core.output_parsers import StrOutputParser
from langchain_core.prompts import ChatPromptTemplate
```

Listing 6.1 Simple RAG System: Required Packages (Source: 06_RAG/10_simple_RAG.py)

You must also set up a configuration file that holds your application programming interface (API) credentials for interacting with the LLM. That file is called `.env` and is shown in [Listing 6.2](#).

```
GROQ_API_KEY = ...  
OPENAI_API_KEY = ...
```

Listing 6.2 Simple RAG System: API Credentials (Source: `.env`)

The first step is now to create the vector database. Since we covered this topic in [Chapter 5](#), we won't bore you by explaining all the details again. If some steps in this chapter are unclear, please return to [Chapter 5](#).

The general setup is the following: The database will be stored in a folder called `rag_store`. If this folder already exists, the assumption is that you already ran the code (because it is created as part of the complete creation of the database). If the folder does not exist yet, the database is created. The database creation consists of these steps:

1. Loading data from Wikipedia
2. Splitting the data into smaller chunks
3. Embedding the data using OpenAI embeddings
4. Storing these embeddings, together with the chunks, into the vector database

[Listing 6.3](#) shows the steps for loading the dataset and storing it into a vector database.

```
#%% load dataset  
persist_directory = "rag_store"  
if os.path.exists(persist_directory):  
    vector_store = Chroma(persist_directory=persist_directory,  
        embedding_function=OpenAIEmbeddings())  
else:  
    data = WikipediaLoader(  
        query="Human History",
```

```

        load_max_docs=50,
        doc_content_chars_max=1000000,
    ).load()

    # split the data
    chunks = RecursiveCharacterTextSplitter(chunk_size=1000,
chunk_overlap=200).split_documents(data)

    # create persistent vector store
    vector_store = Chroma.from_documents(chunks, embedding=OpenAIEmbeddings(),
persist_directory="rag_store")

```

Listing 6.3 Simple RAG System: Vector Database Creation/Loading (Source: 06_RAG/10_simple_RAG.py)

Next, we'll walk through the three steps of RAG, namely, retrieval, augmentation, and generation. We start with the retrieval step.

6.2.2 Retrieval

To retrieve something, we must create a retriever. [Listing 6.4](#) shows how to set up a retriever and how to invoke it. Helpfully, `vector_store` (our Chroma database instance) has a method `as_retriever()` that allows you to define the vector database as the retriever. Plus, we can define some important parameters like the similarity function to use for finding the most relevant documents, as well as how many documents to return. Let's try it out by setting up a question and subsequently invoking the retriever with that question.

```

retriever = vector_store.as_retriever(
    search_type="similarity",
    search_kwargs={"k": 3})
question = "what happened in the first world war?"
relevant_docs = retriever.invoke(question)

```

Listing 6.4 Simple RAG System: Retriever Set Up (Source: 06_RAG/10_simple_RAG.py)

We can check qualitatively if the retrieval process works by printing the most similar documents. For simplicity, we'll limit the output for

each document to just 100 characters, as shown in [Listing 6.5](#).

```
#%% print content of relevant docs
for doc in relevant_docs:
    print(doc.page_content[: 100])
    print("\n-----")
```

```
This transformation was catalyzed by wars of unparalleled scope and devastation.
World War I was a g
-----
=== World War I ===
-----
World War I saw the continent of Europe split into two major opposing alliances; the
Allied Powers,
-----
A tenuous balance of power among European nations collapsed in 1914 with the
outbreak of the First W -----
```

Listing 6.5 Simple RAG System: Retrieved Documents (Source: 06_RAG/10_simple_RAG.py)

Good! All documents found seem relevant since they are related to World War I or have the starting year of 1914 mentioned.

But these samples are just snippets taken from chunks. You might try to figure out the specific answer based on these snippets, but they can be hard to read and understand, since they are not necessarily complete sentences, but only pieces of it. Thus, we don't yet have a good answer to the question. That is what RAG is for, because in the next step we augment the documents with some proper instructions on how the documents should be handled to answer the question.

6.2.3 Augmentation

In the augmentation step, we bring the most similar documents into a well-defined format: a long-concatenated string in which all document contents are represented. We can create this string with `join()`. In this function, we pass a list that is then just concatenated

into a single string. The list items are represented by `page_content` of our `relevant_docs`, as follows:

```
context = "\n".join([doc.page_content for doc in relevant_docs])
```

This step was the tedious work; now, we need to get creative when defining a prompt that instructs the LLM how to formulate an answer. As shown in [Listing 6.6](#), we can define the role that our LLM should play. We'll instruct it to use the documents to answer the question. Furthermore, we tell it to not hallucinate and instead to clearly state that it does not know the answer if necessary.

```
### create prompt
messages = [
    ("system", "You are an AI assistant that can answer questions about the history of human civilization. You are given a question and a list of documents and need to answer the question. Answer the question only based on these documents. These documents can help you answer the question: {context}. If you are not sure about the answer, you can say 'I don't know' or 'I don't know the answer to that question.'"),
    ("human", "{question}"),
]
prompt = ChatPromptTemplate.from_messages(messages=messages)
```

Listing 6.6 Simple RAG System: Augmentation Setup (Source: 06_RAG/10_simple_RAG.py)

For a start, this step should be enough for augmentation. You can later come back to this stage to improve the system. For now, we're satisfied and can proceed with the next stage—the generation of a final answer.

6.2.4 Generation

We've come to the last step—the generation of an answer. This step is straightforward. We'll create a model. In this case, we can use an open-weight model from Google called `gemma2`. This model is a comparably small but is quite fast. Classified as a small language model (SLM), `gemma2` allows for very quick inferences, which might

increase user acceptance, because the answer is provided nearly instantaneously. The longest process is the retrieval of the documents.

The second step in generation should be familiar to you from [Chapter 3](#)—the creation of a chain. In this case, the chain starts with our prompt. The prompt output is passed to the model. Finally, the model output is passed to `StrOutputParser()` so that only the content part of the model output is shown, as follows:

```
model = ChatGroq(model_name="gemma2-9b-it", temperature=0)
chain = prompt | model | StrOutputParser()
```

We've placed our last jigsaw piece; let's test it by invoking the chain and checking the answer, as shown in [Listing 6.7](#).

```
### invoke chain
answer = chain.invoke({"question": question, "context": context})
print(answer)
```

World War I was a global conflict from 1914 to 1918.

It involved two main alliances:

*** **The Allied Powers:** Primarily composed of the United Kingdom, France, Russia, Italy, Japan, Portugal, and various Balkan states.**
*** **The Central Powers:** Primarily composed of Germany, Austria-Hungary, the Ottoman Empire, and Bulgaria.**

The war resulted in the collapse of four empires: Austro-Hungarian, German, Ottoman, and Russian. It had a death toll estimated between 10 and 22.5 million people.

The war saw the use of new industrial technologies, making traditional military tactics obsolete. It also witnessed horrific events like the Armenian, Assyrian, and Greek genocides within the Ottoman Empire.

Listing 6.7 Simple RAG System: Generation (Source: 06_RAG/10_simple_RAG.py)

A pretty good answer, assuming that our test checks that the answer formulates a good response based on the available knowledge from the source. An equally important step is to also test the opposite

functionality—if the model correctly identifies its own limitations and clearly states that it does not know.

In a real-world example, you probably want to embed this functionality into a broader context of code. So, what we need to do should be clear: We need to bundle everything into a function.

6.2.5 RAG Function Creation

Our function `simple_rag_system` will consume a question and return a string. The other steps are exactly the same steps as before:

1. Retrieving the relevant documents
2. Creating the context based on these relevant documents
3. Creating a detailed prompt that tells the model how to apply the context as well as the user question
4. Creating a chain that bundles the prompt, the model, and an output handler together

[Listing 6.8](#) shows the definition of our `simple_rag_system` function, which consumes a question as the input parameter and returns the RAG system answer.

```
# %% bundle everything in a function
def simple_rag_system(question: str) -> str:
    relevant_docs = retriever.invoke(question)
    context = "\n".join([doc.page_content for doc in relevant_docs])
    messages = [
        ("system", "You are an AI assistant that can answer questions about the
history of human civilization. You are given a question and a list of documents and
need to answer the question. Answer the question only based on these documents.
These documents can help you answer the question: {context}. If you are not sure
about the answer, you can say 'I don't know' or 'I don't know the answer to that
question.'"),
        ("human", "{question}"),
    ]
    prompt = ChatPromptTemplate.from_messages(messages=messages)
    model = ChatGroq(model_name="gemma2-9b-it", temperature=0)
    chain = prompt | model | StrOutputParser()
```

```
answer = chain.invoke({"question": question, "context": context})
return answer
```

Listing 6.8 Simple RAG System: Function (Source: 06_RAG/10_simple_RAG.py)

Now, time to test it out. Does the function correctly return that it does not know an answer? As shown in [Listing 6.9](#), it does!

```
# %% Testing the function
question = "What is a black hole?"
simple_rag_system(question=question)

"I don't know the answer to that question. \n"
```

Listing 6.9 Simple RAG System: Test (Source: 06_RAG/10_simple_RAG.py)

Our RAG system does not know anything about black holes, which is exactly what we hoped for, given that it relied on a knowledge source based solely on human history. We achieved this limitation through the system message by passing the instruction Answer the question only based on these documents. ... If you are not sure about the answer, you can say 'I don't know' or 'I don't know the answer to that question.

Let's reflect on what we've learned so far. We've developed our first RAG system, based on a vector database. Given a user question, the most relevant documents were retrieved. These retrieved documents were used in combination with detailed instructions to the LLM on how to make use of them.

You'll be impressed by how well this works. But after working with it for a while, you might see some issues with this simple RAG. Issues might occur anywhere along the complete pipeline, for example:

- **Retrieval**

The retrieved documents might not be relevant.

- **Augmentation**

- If multiple documents hold the same or very similar information, the retrieved documents might be repetitive and thus are not very informative.
- If documents stem from different sources, their context might be hard to differentiate, which might result in an incorrect mix of different documents in the final answer. Imagine you have a system with controversial papers and contradicting views. A terrible answer will result if the system assigns views to the wrong authors.
- **Generation**
Issues at this step typically result from the upstream issues mentioned earlier in this list.

Numerous possible improvements can be made, and in the following sections, we'll discuss some of these options.

6.3 Advanced Techniques

Although our simple RAG system is impressive, there's still a lot of room for improvement. Opportunities for improvement can be found at different stages in the pipeline. For instance, some improvements can be made to the indexing pipeline.

Practically, we can improve the data before we add it to our vector database, which we'll cover when we discuss advanced pre-retrieval techniques in [Section 6.3.1](#). Other techniques tackle the problem at the retrieval step of the pipeline, which we cover in [Section 6.3.2](#). Unsurprisingly, the last step of the pipeline can be improved as well; we'll explore improvements in the generation step in [Section 6.3.3](#).

6.3.1 Advanced Preretrieval Techniques

If the RAG system fails to fetch the right documents, the problem might be found in the data ingestion pipeline. Several ways are available to improve the data before it is stored into the vector

database. [Figure 6.4](#) shows some common approaches to optimize the data ingestion pipeline.




Figure 6.4 Pre-Retrieval Techniques

The different techniques correspond to the step in which they are applied, which we'll discuss in the following sections.

Data Cleaning

Data cleaning in RAG systems involves preparing and refining the source documents to optimize them for retrieval and generation. This crucial preprocessing step begins with basic operations like removing duplicate content, standardizing formatting, and handling special characters.

The text is then typically segmented into appropriate chunks that balance context preservation with retrieval granularity—chunks that are too large may contain irrelevant information, while chunks that are too small might lose important context. (We covered these topics in detail in [Chapter 5](#).) Advanced cleaning steps might include the following:

- Removing boilerplate content

- Standardizing date formats
- Resolving abbreviations
- Handling multilingual content

The cleaning process also often involves metadata enrichment, where chunks are tagged with relevant identifiers, timestamps, or categorizations to enhance retrieval accuracy. Clean data forms the foundation for effective RAG performance, as it directly impacts both the quality of retrieved matches and the coherence of generated responses.

Metadata Enhancing

Metadata enhancement in RAG systems involves enriching document chunks with additional contextual information that goes beyond the raw text content. This process includes tagging documents with attributes such as creation dates, authors, topics, and categorical classifications.

More sophisticated enhancement might involve generating embeddings for semantic search, calculating readability scores, or extracting key entities and relationships. Some systems implement automated topic modeling or classification to add thematic tags, while others maintain hierarchical relationships between chunks to preserve document structure.

This rich layer of metadata not only improves retrieval accuracy but also enables more nuanced filtering and contextual understanding during the generation phase. This ultimately leads to more relevant and well-sourced responses.

Chunk Size Optimization

Chunk size optimization in RAG systems requires carefully balancing multiple competing factors to maximize retrieval effectiveness. While smaller chunks offer more precise retrieval and reduce token consumption, they risk fragmenting important context and breaking up coherent ideas.

On the other hand, larger chunks preserve more context but can introduce noise and irrelevant information into the retrieval results, potentially diluting the quality of generated responses. As explored in [Chapter 5](#), the optimal chunk size often varies depending on the nature of the content. Technical documentation might benefit from smaller, focused chunks, while narrative content might require larger chunks to maintain coherence.

Context Enrichment

Context enrichment in RAG systems involves augmenting the base content with additional information to enhance its utility and retrievability. This process goes beyond basic metadata tagging by incorporating related information, cross-references, and derived insights that make the content more valuable for retrieval and generation.

Advanced enrichment might include linking related concepts, adding definitions for technical terms, or incorporating domain-specific knowledge. For example, when processing medical documents, the system might automatically expand abbreviations, add standardized medical codes, or link to relevant research papers.

Real-world applications might include adding geographic coordinates to location-based content, temporal relationships for event-based information, or industry-specific taxonomies.

The enriched context helps the retrieval system make more intelligent matches and provides the generation model with richer, more nuanced information to work with, ultimately leading to more comprehensive and accurate responses.

Embedding Model Fine-Tuning

Embedding model fine-tuning in RAG systems adapts general-purpose embedding models to better capture domain-specific semantic relationships and nuances. The process typically begins with selecting a base embedding model like BERT (bidirectional encoder representations from transformers) or sentence transformers, then training it further on domain-specific data to better understand specialized vocabulary, technical concepts, and industry-specific relationships.

The fine-tuning process requires careful curation of training pairs that represent the types of semantic similarities most important for the specific use case—for instance, in legal applications, fine-tuning might include matching different phrasings of the same legal concept, while in technical documentation, fine-tuning might focus on connecting problem descriptions with their solutions.

6.3.2 Advanced Retrieval Techniques

In this section, we'll cover several advanced retrieval techniques. One limitation of our simple RAG solution is that it is not well suited to finding specific keywords. For keyword search, a few different algorithms are available. After learning about these algorithms, we'll use them in a hybrid RAG system. We'll also cover *reciprocal rank fusion*, which is a ranking technique that merges multiple retrieval results to improve relevance without complex weighting schemes. In

a practical implementation of a hybrid search, we'll combine vector search with keyword-based methods to enhance document retrieval. You'll also learn about *query expansion*, a technique that enhances user queries. Finally, you'll learn about *context enrichment*. With these techniques, you can refine and expand retrieved passages by incorporating additional document-based context.

Keyword-Search Algorithms

For keyword searches, multiple algorithms are available like TF-IDF (term frequency-inverse document frequency) and BM25 (Best Match 25). A keyword search is also known as a *sparse vector search*, in contrast to our well-known embeddings, which are often called *dense vector searches*.

The most popular sparse vector searches are TF-IDF or BM25. Both are numerical statistics used in information retrieval. Their purpose is to assess the importance of a word in a document relative to a collection of documents, as follows:

- **Term frequency-inverse document frequency (TF-IDF)**

TF-IDF is the product of term frequency and inverse document frequency. Term frequency (TF) represents the number of times a term appears in a document. Inverse document frequency (IDF) measures how unique or rare a word is across all documents in the corpus. The more documents a word appears in, the lower its IDF value becomes.

Words like “the” or “is” appear in almost every document, so their IDF values are low. Other terms that are rare across several documents (i.e., they appear in just a few documents) thus have a high IDF value.

Terms with high TF-IDF values indicate that a term frequently appears in a particular document but is not common in the corpus of documents. These terms are considered more relevant for identifying key topics of a document.

TF-IDF is a popular search algorithm for ranking documents based on how relevant they are to a user query. The idea is to give more weight to terms that are significant in the document but rare in the corpus.

- **Best Match 25 (BM25)**

A related concept that builds upon TF-IDF is BM25, which introduces a sophisticated way of weighting terms based on term frequency, document length, and term saturation.

BM25 adds the concept of term frequency. In TF-IDF, a term's importance increases linearly with its frequency. BM25 uses instead a saturation function. In other words, after a certain point, additional occurrences of a term in a document don't further increase the relevance. The rationale behind this approach is that, if a term appears multiple times in a document, it has a certain impact, but above a certain threshold its impact reduces.

Another applied concept is document length normalization. TF-IDF does not balance between long and short documents. To overcome this lack, BM25 adjusts for a document's length. The rationale is that longer documents naturally have more occurrences of a term, so BM25 normalizes the score to prevent bias towards longer documents.

BM25 is widely used in modern search engines because it overcomes multiple limitations of TF-IDF. With saturation, it better models real texts, where term frequency eventually stops

increasing the relevance of a term. Also, it compensates for longer documents with document length normalization.

To summarize, we have two different worlds: sparse vector search and dense vector search. Now, you know how to work with both individually. However, we want to search for the best of both worlds. Thus, we need a way to combine search results from both sparse and dense searches.

Coding: Sparse Search with BM25 and TF-IDF

Let's implement both BM25 and TF-IDF to compare their results. You can find the corresponding script at *06_RAG/25_BM25_TFIDF.py*.

[Listing 6.10](#) shows how to load relevant packages and create a function to preprocess text. A new functionality that we'll use is `BM25Okapi` from the `rank_bm25` package.

```
### packages
from rank_bm25 import BM25Okapi
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS
from typing import List
import string

### Documents
def preprocess_text(text: str) -> List[str]:
    # Remove punctuation and convert to lowercase
    text = text.lower()
    # remove punctuation
    text = text.translate(str.maketrans('', '', string.punctuation))
    return text.split()
```

Listing 6.10 TF-IDF and BM25: Packages and Preprocessing Function (Source: *06_RAG/25_BM25_TFIDF.py*)

In the preprocessing function, we want to convert all text to lowercase. Then, all punctuation is removed. Finally, all the text is split at each space. As a result, we have a list of single words for

each document in the corpus. Now, we need some text to process and play with. This corpus is shown in [Listing 6.11](#).

```
corpus = [  
    "Artificial intelligence is a field of artificial intelligence. The field of  
    artificial intelligence involves machine learning. Machine learning is an artificial  
    intelligence field. Artificial intelligence is rapidly evolving.",  
    "Artificial intelligence robots are taking over the world. Robots are machines  
    that can do anything a human can do. Robots are taking over the world. Robots are  
    taking over the world.",  
    "The weather in tropical regions is typically warm. Warm weather is common in  
    these regions, and warm weather affects both daily life and natural ecosystems. The  
    warm and humid climate is a defining feature of these regions.",  
    "The climate in various parts of the world differs. Weather patterns change due  
    to geographic features. Some regions experience rain, while others are dry."  
]
```

Listing 6.11 TF-IDF and BM25: Corpus (Source: 06_RAG/25_BM25_TFIDF.py)

The first document is about artificial intelligence and shows many occurrences of this term. Actually, the term is repeated excessively, so our expectation is that TF-IDF will over-rank it compared to BM25. The second deals with a related topic: “robots.” The third document repeats the term “warm” several times, and the last document is rather generic.

Now, let’s start preprocessing the corpus. One essential point at this step is to check how these different functionalities expect the input data to be passed. We worked with `TfidfVectorizer` earlier in this section. This functionality expects documents to be passed as a list of strings (`list[str]`). Thus, for each document, we want to remove the stopwords but keep the document as a string of multiple words. The `BM25Okapi` class requires a different format: `list[List[str]]`. With this format, each document is split into its individual words, and all the words are then passed as a list of strings. This approach is extremely important to get the correct results.

```
Tokenized Query BM25: ['artificial', 'intelligence', 'involves', 'learning']  
Tokenized Query TFIDF: artificial intelligence involves learning  
BM25 Similarities: [2.11043413 0. 0. 0. ]
```

Listing 6.12 shows how the similarity calculation is implemented. First, the documents are preprocessed. The resulting object `tokenized_corpus` provides the data in the format that BM25 class requires. Then, we'll initialize `bm25` as an instance of `BM25Okapi` class. The `user_query` is defined and tokenized. Remember, BM25 and TF-IDF classes require the user query in different formats. Finally, we'll calculate the scores with `bm25.get_scores()` and print the results to the screen.

```
# Preprocess the corpus
tokenized_corpus = [preprocess_text(doc) for doc in corpus]
# %% Sparse Search (BM25)
bm25 = BM25Okapi(tokenized_corpus)

### Set up user query
user_query = "artificial intelligence involves learning"

# Process query to remove stop words
tokenized_query_BM25 = user_query.lower().split()
tokenized_query_tfidf = ' '.join(tokenized_query_BM25)
# Process query to remove stop words

bm25_similarities = bm25.get_scores(tokenized_query_BM25)
print(f"Tokenized Query BM25: {tokenized_query_BM25}")
print(f"Tokenized Query TFIDF: {tokenized_query_tfidf}")
print(f"BM25 Similarities: {bm25_similarities}")

Tokenized Query BM25: ['artificial', 'intelligence', 'involves', 'learning']
Tokenized Query TFIDF: artificial intelligence involves learning
BM25 Similarities: [2.11043413 0. 0. 0. ]
```

Listing 6.12 TF-IDF and BM25: Similarity Calculation (Source: 06_RAG/25_BM25_TFIDF.py)

The first few lines just showcase the different input formats of BM25 and TF-IDF queries. In the last output line, you'll see the similarities between the user query and the documents. The highest and only similarity is found for the first document.

Next, let's implement TF-IDF, as shown in [Listing 6.13](#), and compare the results. Create an instance of `TfidfVectorizer()` and embed the documents with `tfidf.fit_transform()`. The user query is

embedded, and finally, its similarity to the TF-IDF matrix can be calculated.

```
### calculate tfidf
tfidf = TfidfVectorizer()
tokenized_corpus_tfidf = [' '.join(words) for words in tokenized_corpus]
tfidf_matrix = tfidf.fit_transform(tokenized_corpus_tfidf)

query_tfidf_vec = tfidf.transform([tokenized_query_tfidf])
tfidf_similarities = cosine_similarity(query_tfidf_vec, tfidf_matrix).flatten()
print(f"TFIDF Similarities: {tfidf_similarities}")

TFIDF Similarities: [0.6630064 0.08546995 0. 0. ]
```

Listing 6.13 TF-IDF and BM25: Similarity Calculation (Source: 06_RAG/25_BM25_TFIDF.py)

The first two documents show some similarity, and the algorithm correctly returns the first document as the most similar.

At this point, reciprocal rank fusion comes to the rescue. But before turning to that topic, let's pause for a second and look at the full picture: our hybrid RAG pipeline.

Hybrid RAG Pipeline

Now, you can build your own hybrid RAG systems. [Figure 6.5](#) shows the mode of operation for a hybrid RAG pipeline.




Figure 6.5 Hybrid RAG: Pipeline

Compared to the diagram shown earlier in [Figure 6.1](#), which shows a general, simple RAG, notice that all changes so far concern the retrieval step of the pipeline.

The user query is passed to a dense search, typically, a vector database. But at the same time, the query is also passed to a sparse search, like TF-IDF or BM25. As a result, we get two ordered lists of documents: one coming from the similarity search in the vector database, the other coming from the keyword search (sparse search) algorithm. We only need and can only work with one ordered list of documents. Thus, we need to aggregate these lists; at this point, rank fusion comes into play.

The remaining parts of this extended RAG pipeline are identical to our simple RAG system. The ordered list of most relevant documents is, together with a set of instructions, augmented. And the augmented prompt is passed to the LLM for processing the information and coming up with an answer.

One question is left to answer: How are the outputs from dense and sparse search combined? The algorithm that is used here is called reciprocal rank fusion.

Reciprocal Rank Fusion

Reciprocal rank fusion is an algorithm used to combine multiple ranked lists into a single, unified ranking.

Other algorithms for rank fusion are available like CombSUM, CombMNZ, or Borda Count. But reciprocal rank fusion stands out for its simplicity, its robustness, and its ability to handle incomplete rankings. Let's find out how it works. For each item in the different ranked lists, reciprocal rank fusion calculates a score based on the following formula:

$$RRF(d) = \sum_{d \in D} \frac{1}{k+r(d)}$$

For each document d (out of a corpus of documents D), a score is calculated based on its ranks in the different lists. k is a constant that is often set to 60. Thus, the only variable part of the formula is $r(d)$, which is the rank of the document in a given list. Then, the sum is calculated over all lists.

Let's look at a real example. [Figure 6.6](#) shows how reciprocal rank fusion works for four documents and displays their ranks as resulting from a dense search and from a sparse search.

Rank	Dense Search Result	Sparse Search Result	Reciprocal Rank	
1	Document D	Document A	1/1	Document A $1/2+1/1=1.5$
2	Document A	Document C	1/2	Document D $1/1+1/4=1.25$
3	Document C	Document B	1/3	Document C $1/3+1/2=0.83$
4	Document B	Document D	1/4	Document B $1/4+1/3=0.58$
a) Reciprocal Rank Calculation				b) Reciprocal Rank Fusion

Figure 6.6 Reciprocal Rank Fusion: Rank Calculation (Left) and Fusion (Right)

The ranks from the two lists and the four documents are shown on the left. The ranks are increased, and the corresponding reciprocal ranks are just the result of $1 \div \text{rank}$. In this example, Document D has rank 1 in the dense search result and rank 4 in the sparse search result.

The way to aggregate both results is simple. The reciprocal ranks are summed. Keeping our focus on Document D: It gets a reciprocal rank of $1/1$ from dense search, and $1/4$ from sparse search. Both are summed to 1.25. The same approach is applied all other documents.

Subsequently, the documents are sorted in decreasing order based on their reciprocal rank sums. As a result, Document A is the most

relevant document, followed by Document D, and so on.

The reciprocal rank algorithm has the beneficial mathematical property of diminishing returns. The score decreases non-linearly as its rank increases. Intuitively, this relationship makes sense since a difference in low ranks (like 1 and 2, or 2 and 3) is much more important than the difference in rank between rank 1000 and 1001.

Coding: Hybrid Search

In this code exercise, you'll set up a hybrid search that applies both a dense search and a sparse search. The two resulting lists are then merged based on rank fusion. Along the way, you'll discover why hybrid searches might outperform any of the individual search engines.

Let's start with loading all the required functionality. [Listing 6.14](#) shows the required packages. Notable additions include `TfidfVectorizer` to perform TF-IDF and `ENGLISH_STOP_WORDS` to remove unnecessary words that might distract TF-IDF.

```
from langchain_openai import OpenAIEmbeddings
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
```

Listing 6.14 Hybrid Search: Required Packages (Source: 06_RAG/20_hybrid_search.py)

We need some sentences to play with, so we'll set up a list of docs, as shown in [Listing 6.15](#).

```
### Documents
docs = [
    "The weather tomorrow will be sunny with a slight chance of rain.",
    "Dogs are known to be loyal and friendly companions to humans.",
    "The climate in tropical regions is warm and humid, often with frequent rain.",
    "Python is a powerful programming language used for machine learning.",
```

```

    "The temperature in deserts can vary widely between day and night.",
    "Cats are independent animals, often more solitary than dogs.",
    "Artificial intelligence and machine learning are rapidly evolving fields.",
    "Hiking in the mountains is an exhilarating experience, but it can be
unpredictable due to weather changes.",
    "Winter sports like skiing and snowboarding require specific types of weather
conditions.",
    "Programming languages like Python and JavaScript are popular choices for web
development."
]

```

Listing 6.15 Hybrid Search: Sample Documents (Source: 06_RAG/20_hybrid_search.py)

We start with the sparse similarity analysis based on TF-IDF. For this algorithm, be aware that it might be distracted by stopwords.

Stopwords are common words that are typically filtered out during text preprocessing in natural language processing (NLP) tasks.

These words are usually the most common words in a language that don't contribute significantly to the meaning or sentiment of a text.

In English, common stopwords include articles (e.g., a, an, the); prepositions (in, on, at, with, by); pronouns (I, you); conjunctions (and, but, or); and verbs (is, am, are, was, were). Depending on the algorithm, removing these stopwords might be an important step to reduce noise, improve efficiency, and focus on meaning. In practice, stopwords removal looks like this:

- Original text: "The car is sitting on the mat"
- After stopwords removal: "car sitting mat"

The object `docs_without_stopwords` can be created based on list comprehension. In this process, each word in a string is separated by splitting at each space with `doc.split()`. Then, all words are converted to lowercase, and a word is only retained if it is not part of `ENGLISH_STOP_WORDS`. Finally, the resulting list is converted to a string with `join()`, as follows:

```

docs_without_stopwords = [
    ' '.join([word for word in doc.split() if word.lower() not in

```

```

ENGLISH_STOP_WORDS])
    for doc in docs
]

```

We'll set up a sparse search and a dense search, starting with a sparse search based on TF-IDF. We'll instantiate a `TfidfVectorizer` before we can create a `tfidf_matrix`, as follows:

```

vectorizer = TfidfVectorizer()
tfidf_matrix = vectorizer.fit_transform(docs_without_stopwords)

```

Before we can calculate similarities based on `cosine_similarity`, we need a `user_query` embedded with `vectorizer.transform`, as follows:

```

user_query = "Which weather is good for outdoor activities?"

query_sparse_vec = vectorizer.transform([user_query])
sparse_similarities = cosine_similarity(query_sparse_vec, tfidf_matrix).flatten()

```

The object `sparse_similarities` is of the type array, and for 10 sample sentences, it has 10 similarity values corresponding to our user query. In the next step, we want to get the indices of the most similar documents. For this step, we'll create a function `getFilteredDocsIndices`. This function consumes the similarities and additionally a `threshold` parameter that excludes all documents below the value.

Multiple ways exist to get this done. As shown in [Listing 6.16](#), we can sort the similarities, filter them if they are below the threshold value, and finally return only the filtered list of document indices.

```

def getFilteredDocsIndices(similarities, threshold = 0.0):
    filt_docs_indices = sorted(
        [(i, sim) for i, sim in enumerate(similarities) if sim > threshold],
        key=lambda x: x[1],
        reverse=True
    )
    return [i for i, sim in filt_docs_indices]

```

Listing 6.16 Hybrid Search: Function for Getting Filtered Indices (Source: 06_RAG/20_hybrid_search.py)

Let's try the function out with our `sparse_similarities`, as shown in [Listing 6.17](#).

```
### filter documents below threshold and get indices
filtered_docs_indices_sparse =
getFilteredDocsIndices(similarities=sparse_similarities, threshold=0.2)
filtered_docs_indices_sparse

[0, 7, 8]
```

Listing 6.17 Hybrid Search: Filter Sparse Documents (Source: 06_RAG/20_hybrid_search.py)

In this case, the most similar document is the document with index 0, followed by index 7, and then index 8.

If you check the user query and the sample sentences, you'll find that this does not work very well. We asked, "Which weather is good for outdoor activities?" and TF-IDF search returned that the most similar document is "The weather tomorrow will be sunny with a slight chance of rain." That response is just not correct. The search did find the appropriate documents "Hiking in the mountains is an exhilarating experience, but it can be unpredictable due to weather changes" and "Winter sports like skiing and snowboarding require specific types of weather conditions," but the rank order was wrong. The model puts too much focus on the word "weather" and thus incorrectly returns the document regarding the weather forecast.

Let's move on and evaluate the performance of the dense search. For this step, we must embed our sample documents. We'll use OpenAI for that purpose, as follows:

```
embeddings = OpenAIEmbeddings()
embedded_docs = [embeddings.embed_query(doc) for doc in docs]
```

Then, we must embed the `user_query`:

```
query_dense_vec = embeddings.embed_query(user_query)
```

Finally, based on cosine similarity, we get a list of similarity scores for each document, as shown in [Listing 6.18](#).

```
dense_similarities = cosine_similarity([query_dense_vec], embedded_docs)
dense_similarities

array([[0.81677377, 0.74636589, 0.78294997, 0.71723575, 0.783497 ,
        0.71934968, 0.71673343, 0.84183792, 0.84227999, 0.73753417]])
```

Listing 6.18 Hybrid Search: Get Dense Similarities (Source: 06_RAG/20_hybrid_search.py)

So far, so good. But what we really want is a list of the most similar documents. Earlier, we created a function for filtering and returning indices of documents, and we'll reuse this function now. To get a small result list of documents, we'll pass a threshold parameter of 0.8, as shown in [Listing 6.19](#).

```
filtered_docs_indices_dense =
getFilteredDocsIndices(similarities=dense_similarities[0], threshold=0.8)
filtered_docs_indices_dense

[8, 7, 0]
```

Listing 6.19 Hybrid Search: Filter Dense Documents (Source: 06_RAG/20_hybrid_search.py)

This list looks better. It has found the documents on outdoor activities (index 8 and index 7) and provides them a rank higher than the document on the weather forecast (index 0).

Now, we have two similarity lists and need to aggregate them into one final list. At this point, reciprocal rank fusion comes into play, as shown in [Listing 6.20](#).

```
def reciprocal_rank_fusion(filtered_docs_indices_sparse,
filtered_docs_indices_dense, alpha=0.2):
    # Create a dictionary to store the ranks
    rank_dict = {}

    # Assign ranks for sparse indices
    for rank, doc_index in enumerate(filtered_docs_indices_sparse, start=1):
        if doc_index not in rank_dict:
```

```

        rank_dict[doc_index] = 0
    rank_dict[doc_index] += (1 / (rank + 60)) * alpha

    # Assign ranks for dense indices
    for rank, doc_index in enumerate(filtered_docs_indices_dense, start=1):
        if doc_index not in rank_dict:
            rank_dict[doc_index] = 0
        rank_dict[doc_index] += (1 / (rank + 60)) * (1 - alpha)

    # Sort the documents by their reciprocal rank fusion score
    sorted_docs = sorted(rank_dict.items(), key=lambda item: item[1], reverse=True)

    # Return the sorted document indices
    return [doc_index for doc_index, _ in sorted_docs]

```

Listing 6.20 Hybrid Search: Function For Reciprocal Rank Fusion (Source: 06_RAG/20_hybrid_search.py)

[Listing 6.20](#) shows the implementation: The function consumes the two lists of indices—one from the sparse search and the other from the dense search. Additionally, we can add a weighting parameter `alpha`, which assigns different weights to the different searches.

A resulting `rank_dict` is initialized as an empty list. There are two loops, one for each list of indices. Each time, it is iterated over all documents in these lists. For the document index, the rank value is added based on the formula $(1 / (\text{rank} + 60)) * \alpha$ (for the sparse search) and $(1 / (\text{rank} + 60)) * (1 - \alpha)$ (for the dense search).

As a result, we get a list of aggregated document indices, in decreasing order. The first list element represents the most similar document based on reciprocal rank fusion.

At this point, we can apply this function to our data. With an `alpha` value of 0.2, we're giving the sparse search a weighting of 20% in the final result, and the dense search gets a weighting of 80%, as follows:

```

reciprocal_rank_fusion(filtered_docs_indices_sparse, filtered_docs_indices_dense,
alpha=0.2)

```

```

[8, 7, 0]

```

The final result is identical to the dense search. But play with the weighting factors, and you'll see that the order changes at some point.

This hybrid search could be integrated into a hybrid RAG system. The corresponding documents to these indices would then be passed as context information to the LLM together with the user query and additional instructions. As you can imagine, however, typically you don't need to implement such features like reciprocal rank fusion on your own; these capabilities are already integrated into vector databases.

In this section, we implemented a hybrid search based on a sparse search with TF-IDF. Another approach for improving the retrieval is query expansion, which we'll turn to next.

Coding: Query Expansion

You might know the saying, "Most computer problems sit between the keyboard and the chair." While it might seem lazy to say that our RAG system is fine and the user is at fault, a kernel of truth lies within. Ultimately, at the end of the day, a RAG response can only be as good as the user query. How much effort do you put into the setting up of your queries?

One possible answer (which, admittedly, is our answer) is that querying is an iterative approach, which is refined over a series of steps. If a user query is imprecise, we should not expect the RAG response to be perfect. But still some approaches can help the user by improving the query. One possible solution is query expansion. This technique enhances the ability to retrieve relevant documents. Let's look at several expansion approaches next:

- **Synonym expansion**

- Expanding the query with synonyms ensures retrieval of documents that use different terminology but discuss the same concept.
- Original query → Expanded query
- Example: “Climate change” → [“global warming,” “climate change,” “environmental impact”]
- **Related terms expansion**
 - Including related terms that are often discussed together with the original query topic improves document retrieval that covers broader or more specific aspects of the topic.
 - Example: “Machine learning” → [“machine learning,” “deep learning,” “artificial intelligence,” “neural networks”]
- **Conceptual expansion**
 - Adding subcategories or specific types of the original concept allows for retrieving a wider range of documents related to different forms of renewable energy.
 - Example: “Renewable energy” → [“renewable energy,” “solar power,” “wind energy,” “green technology,” “sustainable energy”]
- **Phrase variations expansion**
 - Including variations of how a concept might be phrased ensures retrieval from documents that use different wording to discuss the same idea.
 - Example: “Health benefits of exercise” → [“health benefits of physical activity,” “exercise health benefits,” “positive effects of exercise”]
- **Contextual expansion**

- Including related fields or tasks in the expansion can help retrieve more diverse documents that approach the original concept from different angles.
- Example: “Natural language processing” → [“natural language processing,” “text analysis,” “computational linguistics,” “language modeling”]
- **Temporal expansion**
 - Expanding for common abbreviations and alternative names ensures that relevant documents using different temporal or popular references are also retrieved.
 - Example: “COVID-19 vaccines” → [“COVID-19 vaccines”, “COVID vaccines”, “coronavirus immunization”, “SARS-CoV-2 vaccine”]
- **Entity-based expansion**
 - When a query involves an entity, expanding the query with related people, products, or attributes associated with the entity can enhance retrieval quality.
 - Example: “Tesla” → [“Tesla,” “Elon Musk,” “electrical vehicles,” “autonomous driving”]

By incorporating these types of expansions, you can maximize the information retrieved for a query, resulting in more effective responses in a RAG system.

Implementing query expansion is a rather easy task, as in the following:

```
from langchain_groq import ChatGroq
from langchain_core.prompts import ChatPromptTemplate
from dotenv import load_dotenv
load_dotenv('.env')
```

[Listing 6.21](#) shows the function definition. We've set up a function `query_expansion` for improving the query. The function receives the query and returns the number of expanded queries.

```
def query_expansion(query: str, number: int = 5, model_name: str = "llama3-70b-8192") -> list[str]:
    messages = [
        ("system", """You are part of an information retrieval system. You are given
a user query and you need to expand the query to improve the search results. Return
ONLY a list of expanded queries.
Be concise and focus on synonyms and related concepts.
Format your response as a Python list of strings.
The response must:
1. Start immediately with [
2. Contain quoted strings
3. End with ]
Example correct format:
["alternative query 1", "alternative query 2", "alternative query 3"]
"""),
        ("user", "Please expand the query: '{query}' and return a list of {number}
expanded queries.")
    ]
    prompt = ChatPromptTemplate.from_messages(messages)
    chain = prompt | ChatGroq(model_name=model_name)
    res = chain.invoke({"query": query, "number": number})
    return eval(res.content)
```

Listing 6.21 Query Expansion: Function (Source: 06_RAG/90_query_expansion.py)

The major work is in setting up the messages, especially convincing the model to avoid babbling and instead directly returning a string that starts with the list. You might choose to implement Pydantic to get a structured output as a list of strings. We leave this task up to you as a quick exercise.

The final result should look like “[...]” instead of something like “Here is the expanded list of queries:\n\n[...].” You cannot simply instruct the model to return a Python list, but you should provide detailed instructions on how the response should look like. We were successful with the instruction set shown in [Listing 6.22](#).

The response must:

1. Start immediately with [
2. Contain quoted strings

3. End with]

Example correct format:

```
["alternative query 1", "alternative query 2", "alternative query 3"]
```

Listing 6.22 Query Expansion: Output Format (Source: 06_RAG/90_query_expansion.py)

The user message holds the information to process the query and return the number of expanded queries.

In the return statement, the model response is evaluated to a Python object, so that we successfully get a list of strings holding the expanded queries. We can test this functionality out, as follows:

```
res = query_expansion(query="artificial intelligence", number=3)
res

['machine learning', 'natural language processing', 'computer vision']
```

Context Enrichment

If you work with a corpus of multiple different documents, distinguishing between different documents can be difficult. For example, if you work with financial documents, an essential task might be to distinguish between quarters and companies. Normal chunking techniques can have difficulties with this kind of thinking.

[Figure 6.7](#) shows the concept behind conceptual retrieval, which can overcome this issue. Each chunk is enhanced with its own context. This context provides useful information on the location of the chunk

in a specific document. For example, the enhanced chunk might refer to the previous and next quarter of the company under study.




Figure 6.7 Contextual Retrieval (Source: Adapted from <https://www.anthropic.com/news/contextual-retrieval>)

How does contextual retrieval work? Each document from a larger corpus is chunked. The complete corpus, or the relevant documents from the corpus, are passed to the LLM, together with a chunk for extracting contextual information for that particular chunk.

As you can imagine, sending a complete, relevant document together with each chunk to an LLM can become extremely costly. Now, the concept of prompt caching comes into play. Since prompt caching is a RAG alternative, we'll cover it in more detail in [Section 6.4](#). But to anticipate this topic, just know that prompt caching is a cheap solution to the problem. With this technique, the complete relevant document is cached for a limited time and does not need to be sent each time to the LLM.

In the end, all chunks are processed, and each chunk was enhanced with contextual information. This new chunk (context plus original chunk) can now be stored in some vector database. Studies have found that this approach can significantly outperform both simple RAG systems and hybrid RAG systems.

6.3.3 Advanced Postretrieval Techniques

After the documents are retrieved, you can still improve upon your RAG system. One interesting approach is *prompt compression*. Imagine that the prompt enhanced with all contexts from the retrieval process is getting long. As a result, the LLM must deal with increased noise, which reduces the LLM's performance. Also, since you need to pay for tokens, the cost is high. What can you do to save costs? You can compress the length of the context while maintaining its essential information, which is the idea behind prompt compression, as shown in [Figure 6.8](#).




Figure 6.8 Prompt Compression

The augmented user query, which possibly has become quite long, is sent to an LLM to be shortened in a way that maintains its meaning. But if we send the augmented query to an LLM to shorten it, we then need to send it to the “final” LLM to get the RAG response. We even have two LLM calls—so what is the point?

This approach only makes sense if you send the long, augmented prompt to a small, fast, and cheap LLM for shortening. For the RAG response, you can then send the prompt to the LLM of your choice.

If you have many LLM calls and large contexts, this approach can significantly lower your costs, increase the performance, and enhance efficiency.

6.4 Coding: Prompt Caching

Prompt caching is a new approach developed by Anthropic and is provided currently only as a beta feature. Recall that passing a complete document to the LLM with every request is inefficient. While still true, Anthropic has developed an approach that addresses this problem. [Figure 6.9](#) shows how the concept of prompt caching works.




Figure 6.9 Prompt Caching

With prompt caching, the user sends a request to the LLM. In this request, the user passes the query together with the complete document. Wait, what? Didn't we just say that sending the complete is extremely inefficient?

Yes, but the user also tells the model to cache the document. The LLM provider caches the complete document when passed the first time. The document is not cached forever, only for a limited time. At the time of writing, a document is cached for 5 minutes. Every subsequent user query within that time interval sent to the LLM together with the complete document does not require a new

caching. Instead, the document is taken from the cache to create the answer.

Some specific use cases for this technique include the following:

- Repetitive tasks
- Running many examples
- Long conversations
- Using RAG as contextual retriever (as discussed earlier in [Section 6.3.2](#))

Of these use cases, we want to focus on the last one. This technique is an enabler of a technique called contextual retriever which can significantly improve your RAG results. For now, we'll focus on how to implement prompt caching. Prompt caching enables a technique called *contextual retrieval*, which is outside the scope of this book.

You can read more about this topic at

<https://www.anthropic.com/news/contextual-retrieval>.

We'll develop a continuous chat with an LLM. The conversation is based on some fixed context. In our example, we'll use the complete book *The Hound of the Baskervilles*. This context information is cached for a limited time, so that we create the cached tokens once—at the beginning of the conversation. With a normal RAG, we would send this context with every request as input tokens. But since the document is now cached, we can achieve some inference gains (faster model responses), but also dramatically reduced costs.

The creation of the cache is slightly more expensive than sending normal input tokens (by about 25%). The idea is to outweigh this added effort with massively reduced costs since reading cached tokens is 90% cheaper than reading input tokens.

[Listing 6.23](#) shows the packages we'll use in this script. As the LLM, we'll use Anthropic's Claude. For a nicely rendered output, we'll also load `rich`.

```
import anthropic
import os
from rich.console import Console
from rich.markdown import Markdown
from langchain_community.document_loaders import TextLoader
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
```

Listing 6.23 Prompt Caching: Required Packages (Source: 06_RAG/40_prompt_caching.py)

At the time of writing, prompt caching is a beta feature not well integrated into LangChain. This integration is coming, but for now, you'll need to interact directly with the Anthropic client, as follows:

```
client = anthropic.Anthropic(api_key=os.getenv("ANTHROPIC_API_KEY"))
```

The complete prompt caching functionality is bundled into a class, as shown in [Listing 6.24](#). We'll develop it step by step.

```
class PromptCachingChat:
    def __init__(self, initial_context: str):
        self.messages = []
        self.context = None
        self.initial_context = initial_context
```

Listing 6.24 Prompt Caching: Class Initialization (Source: 06_RAG/40_prompt_caching.py)

First, we create the class. During the instantiation of the class object, we pass some `initial_context` as a string. The other properties (messages and context) are initialized as an empty list and as `None`, respectively; both will be populated in the process.

We'll set up several methods in this exercise. As shown in [Listing 6.25](#), the method `run_model` will invoke the LLM.

```

def run_model(self):
    self.context = client.beta.prompt_caching.messages.create(
        model="claude-3-haiku-20240307",
        max_tokens=1024,
        system=[
            {
                "type": "text",
                "text": "You're a patent expert. You're given a patent and will be asked to answer
questions about it.\n",
            },
            {
                "type": "text",
                "text": f"Initial Context: {self.initial_context}",
                "cache_control": {"type": "ephemeral"}
            }
        ],
        messages=self.messages,
    )
    # add the model response to the messages
    self.messages.append({"role": "assistant",
        "content": self.context.content[0].text})
    return self.context

```

Listing 6.25 Prompt Caching: Run Model Method (Source: 06_RAG/40_prompt_caching.py)

The LLM is called with a system message. But the system prompt does not only hold the information on how the model should behave and answer. Additionally, the initial context is passed, alongside a `cache_control` parameter that will enable the caching. This parameter requires the type `ephemeral`.

The messages need to be passed as well. The model's response is stored in the object `self.context`. As soon as the model output is available, the messages are appended with it.

Good! But we want the user to interact with it, so we need to fetch the `user_query` and call our method `run_model`. As shown in [Listing 6.26](#), our next method `user_turn` consumes the `user_query` as an input.

```

def user_turn(self, user_query: str):
    self.messages.append({"role": "user", "content": user_query})

```

```
self.context = self.run_model()
return self.context
```

Listing 6.26 Prompt Caching: User Method (Source: 06_RAG/40_prompt_caching.py)

The `messages` object is appended with the user message, and in the next step, the `run_model` method is called. We'll keep the model output in the `context`.

Basically, we're done. We have a system that reacts to user queries and calls the model. But for nicer outputs, we'll finally define a method called `show_model_response`, as shown in [Listing 6.27](#).

```
def show_model_response(self):
    console = Console()
    console.print(Markdown(self.messages[-1]["content"]))
    console.print(f"Usage: {self.context.usage}")
```

Listing 6.27 Prompt Caching: Showing Model Output (Source: 06_RAG/40_prompt_caching.py)

In this function, we'll use the `rich` package and print the last model output in Markdown. Additionally, we will display the token usage to check if the prompt caching actually worked.

Let's test this function out. In the first step, we need to load our book. The code shown in [Listing 6.28](#) provides access to the book file and loads it with `TextLoader`.

```
file_path = os.path.abspath(__file__)
current_dir = os.path.dirname(file_path)
parent_dir = os.path.dirname(current_dir)

file_path = os.path.join(parent_dir, "05_VectorDatabases",
    "data", "HoundOfBaskerville.txt")
file_path

### (3) Load a single document
text_loader = TextLoader(file_path=file_path, encoding="utf-8")
doc = text_loader.load()
initialContext = doc[0].page_content
```

Listing 6.28 Prompt Caching: Test (Source: 06_RAG/40_prompt_caching.py)

Now, we can create an instance of our class, the object `promptCachingChat`. During initialization, we pass the `initialContext`. Then, we start the conversation with `user_turn` and send our first query, as follows:

```
promptCachingChat = PromptCachingChat(initial_context=initialContext)
promptCachingChat.user_turn("what is special about the hound of baskerville?")
promptCachingChat.show_model_response()
```

The result is shown in [Figure 6.10](#).



```
promptCachingChat = PromptCachingChat(initial_context=initialContext)
promptCachingChat.user_turn("what is special about the hound of baskerville?")
promptCachingChat.show_model_response()

✓ 10s

There are a few key things that make the Hound of the Baskervilles special:

1 The supernatural and mysterious element - The Hound is described as an enormous, supernatural-looking hound with glowing eyes and fire coming from its mouth. This ties into the long-standing family curse and legend surrounding the Baskerville estate, adding an eerie, supernaturally-tinged mystery.
2 The clever use of the legend by the villain - The villain, Stapleton, deliberately plays up the legend of the demonic hound in order to terrify and kill off the Baskervilles. He trains a real, massive hound and uses phosphorus to make it appear supernatural, exploiting the family's superstitions.
3 Its role as the central threat - The Hound itself is the main threat and antagonist that Sherlock Holmes and the other characters have to confront and overcome. Tracking down and stopping the Hound is the driving force behind the entire investigation.
4 Its ties to the Baskerville family history - The Hound is intimately connected to the Baskerville clan's dark past, adding an extra layer of significance and personal stakes to the mystery.

So in summary, the Hound represents the perfect blend of supernatural intrigue, cunning criminal plotting, and personal family drama that makes it such an iconic and memorable part of the Sherlock Holmes canon. The Hound truly is the central mystery and threat that drives the entire narrative.

Usage: PromptCachingDetails(cache_creation_input_tokens=93104, cache_read_input_tokens=0, input_tokens=19, output_tokens=328)
```

Figure 6.10 Prompt Caching: First Chat Round with Cache Creation

The model provides a detailed output, but what is more interesting is the fact that 93,104 cache tokens were created. In the subsequent chat, we'll see if the caching worked and the cached tokens are used. Let's now ask if the hound is the murderer, as follows:

```
promptCachingChat.user_turn("Is the hound the murderer?")
promptCachingChat.show_model_response()
print(promptCachingChat.context.usage)
```

Again, the result is shown in [Figure 6.11](#).

```
promptCachingChat.user_turn("Is the hound the murderer?")
promptCachingChat.show_model_response()
print(promptCachingChat.context.usage)
✓ 42s
```

No, the Hound of the Baskervilles is not the actual murderer in the story. The Hound is used by the real villain, Stapleton, to commit the murders.

Here's a quick breakdown:

- Sir Charles Baskerville is the one who dies at the beginning of the story, apparently killed by the legendary supernatural Hound.
- However, it is later revealed that Stapleton, who is posing as a naturalist named Stapleton, is the one who orchestrated Sir Charles's death by using a trained, phosphorescent-coated hound to scare him to death.
- Stapleton does this as part of a plot to inherit the Baskerville estate by eliminating the current heir, Sir Henry Baskerville.

So while the Hound is the instrument of murder, the true murderer behind the scenes is Stapleton, who is manipulating the hound for his own nefarious purposes. The Hound itself is not an autonomous murderer, but rather a tool used by the human villain to carry out the killings.

```
Usage: PromptCachingBetaUsage(cache_creation_input_tokens=0, cache_read_input_tokens=93104, input_tokens=370, output_tokens=142)
```

Figure 6.11 Prompt Caching with Reading Cached Tokens

No further cached tokens were created, and all cached tokens were read. Amazing! The model provides the result much more quickly. Was it cheaper than without caching? We'll use the assumptions that normal tokens cost 1 unit, that cache creation costs 1.25 units, and that reading cached input tokens 0.1 units.

In the first round of the conversation, prompt caching was more expensive than without caching, but the tide starts turning in the second round because reading from the cache is pretty cheap compared to using normal tokens.

Thus concludes our discussion of prompt caching. Next, you'll learn how to measure the performance of your RAG systems.

6.5 Evaluation

RAG systems work surprisingly well, but what does “well” actually mean? And how can we compare the effectiveness, accuracy, and efficiency of different RAG approaches? For this reason, RAG evaluation techniques try to measure the following parameters:

- **Effectiveness**

Does our RAG system retrieve and use the relevant knowledge?

- **Accuracy**

How good are the generated responses? Are they grounded in facts and coherent?

- **Efficiency**

How well does our system perform in different scenarios?

To make all this measurable and comparable, different metrics can be applied. So many metrics exist that only a small number can be presented.

First, we’ll look at some of the challenges that RAG systems face. Based on these findings, we’ll discuss some metrics and how they can make these challenges measurable.

After several metrics have been presented, we’ll implement them in Python. Let’s start with a closer look at some of the challenges RAG systems face.

6.5.1 Challenges in RAG Evaluation

How can we measure the quality of an answer? Some hard facts can be checked, for instance, how good the data retrieved from the

knowledge source is. We'll focus on two challenges:

- **Retrieval quality**

If the retriever fetches incorrect or irrelevant documents, the RAG system is likely incorrect. A cascading effect occurs in that RAG quality is inherently tied to the quality of the retrieved data. Luckily, several metrics are available for determining retrieval quality.

- **Subjectivity**

If you ask multiple human evaluators, they will provide different ratings on a RAG response. Everybody has personal preferences regarding brevity and style. For example, a technical expert might prioritize precision and factual accuracy, whereas a layperson might prefer readability and simplicity.

Always keep in mind the target user of your system and ensure your evaluators are trained to get consistent results.

6.5.2 Metrics

Earlier when we discussed RAG improvements, we broke down the system into its different parts—retrieval, and generation. We can follow the same approach when it comes to metrics because metrics exist for the different stages. Some are considered joint metrics.

[Figure 6.12](#) shows some RAG evaluation metrics and how they can be classified into retriever, generator, and joint metrics.




Figure 6.12 RAG Evaluation Metrics

Among retriever metrics, classic classification metrics include $\text{precision@}k$ and $\text{recall@}k$. Another metric that we'll study is *context relevance*.

Of the generator-based metrics, we want to highlight *factual correctness*. Finally, joint metrics try to measure both the retriever and the generator quality in a combined metric. Some examples include *answer relevance* and the *F1 score*. Let's look at some of these metrics in more detail next.

Retriever Metric: Context Precision

Context precision is a retriever metric for evaluating the ground-truth items and determining whether they are part of the context information. In perfect cases, all relevant chunks returned by the retriever can be found at the top ranks.

Context precision uses the user query, a ground truth, and the context information for its calculation. The resulting score is between 0 and 1, with higher values representing higher precision:

$$\frac{\sum_{i=1}^k (\text{Precision@}k \times v_k)}{\text{Total count of relevant items in top-}k \text{ retrievals}}$$

K corresponds to the number of chunks in the context, and v_k corresponds to the relevance indicator at rank k .

Generator Metric: Factual Correctness

Factual correctness measures whether the information in a RAG response aligns with established facts. Different aspects must be considered in this task, such as the following:

- **Fact granularity**

A response might be partially correct, meaning correct in one

aspect, but incorrect in another. A correct statement is “Shakespeare wrote many plays,” but it would be incorrect to state “Shakespeare wrote the play *Waiting for Godot*.”

- **Consistency over time**

The measure might have a temporal aspect, in that a fact might be correct for a given point in time, but incorrect for another one. For example, answering the question “Who is the president of country X?” depends on when the question is asked since the fact could change after every election.

- **Attribution**

RAG responses should attribute facts to the correct source. For example, it would be wrong to attribute the theory of special relativity to Newton instead of Einstein.

Joint Metric: Answer Faithfulness

Answer faithfulness measures the factual consistency of the RAG-generated answer against the context information. This metric uses the RAG answer and the retrieved context information for its calculation:

$$\frac{\text{Count of claims in generated answer, that can be inferred from context}}{\text{Total count of claims in generated answer}}$$

The score ranges from 0 to 1 with lower values corresponding to lower faithfulness, and vice versa.

Joint Metric: Answer Relevance

Answer relevance evaluates the alignment of the generated answer with the user query. This metric considers the retrieved context and measures how well the retrieval and generator work together.

[Figure 6.13](#) shows the RAG evaluation metric concept answer relevance.

This approach follows the classic RAG approach to receive an RAG answer. Based on that RAG answer, an LLM creates N questions, that potentially could result in that RAG answer. These N synthetic questions are then embedded and compared to the embedding of the user query. The average of these similarities is the answer relevance.




Figure 6.13 RAG Evaluation Metric: Answer Relevance

In technical terms, the answer relevance is calculated with the following equation:

$$\frac{1}{N} \sum_{i=1}^N \cos(E_{si}, E_{user})$$

This value ranges from 0 to 1, in which low values correspond to answers that are incomplete or that hold redundant information. High values correspond to more relevant answers.

6.5.3 Coding: Metrics

Now, let's implement these metrics. For this task, we'll use a Python framework that has been developed specifically for evaluating RAG

systems, namely, RAGAS (<https://www.ragas.io/>). You can easily install it via uv, as follows:

```
uv add ragas
```

RAGAS has a nice web interface that you can use to evaluate larger datasets and RAG responses with derived metrics. If you work on a “real” project with business impact, we highly recommend diving deeply into RAGAS, starting with its **Getting Started** section (<https://docs.ragas.io/en/v0.1.21/getstarted/index.html#>).

In this book, we’ll show only the basics of how to set up RAGAS to calculate RAG evaluation metrics. What you need is the following:

- `user_query` or `question` as an input to the RAG system
- `contexts` that represent the retrieved documents from the knowledge source
- `answer` for the RAG response
- `ground_truth` is required only for some metrics like *context precision*. This information is used to check whether the retrieved chunk is relevant to arrive at the ground truth.

The complete script can be found in `06_RAG\60_rag_eval.py`. As shown in [Listing 6.29](#), we start by loading the required packages.

```
### packages
from datasets import Dataset
from ragas.metrics import context_precision, answer_relevancy, faithfulness
from ragas import evaluate
from langchain_openai import ChatOpenAI
```

Listing 6.29 RAG Evaluation: Required Packages (Source: `06_RAG\60_rag_eval.py`)

We must bring the data to be evaluated into a format that RAGAS can process. For this step, we load `Dataset`. We also need some metrics from `ragas` as well as its `evaluate` function.

Now, we must prepare our data. The data preparation step is shown in [Listing 6.30](#). The data must be passed in the format Dataset, which we can create based on a dictionary. That dictionary requires the keys question, contexts, answer, and ground_truth.

```
# %% prepare dataset
my_sample = {
    "question": ["What is the capital of Germany in 1960?"],
    "contexts": [
        [
            "Berlin is the capital of Germany since 1990.",
            "Between 1949 and 1990, East Berlin was the capital of East Germany.",
            "Bonn was the capital of West Germany during the same period."
        ]
    ], # Nested list for multiple contexts
    "answer": ["In 1960, the capital of Germany was Bonn. East Berlin was the capital of East Germany."],
    "ground_truth": ["Berlin"]
}

dataset = Dataset.from_dict(my_sample)
```

Listing 6.30 RAG Evaluation: Dataset (Source: 06_RAG\60_rag_eval.py)

Our example is centered around Berlin, the capital of Germany. The question is tricky because time inconsistency exists because Berlin was not always the capital. Finally, we can start the evaluation process, as shown in [Listing 6.31](#).

```
# %% evaluation
llm = ChatOpenAI(model="gpt-4o-mini")
metrics = [context_precision, answer_relevancy, faithfulness]
res = evaluate(dataset=dataset,
               metrics=metrics,
               llm=llm)

res

{'context_precision': 1.0000, 'answer_relevancy': 0.9934, 'faithfulness': 0.5000}
```

Listing 6.31 RAG Evaluation: Evaluation (Source: 06_RAG\60_rag_eval.py)

We must still set up an LLM instance and pass it to the evaluate function. Other parameters that need to be passed include the dataset and the metrics that should be analyzed. The result holds

the scores for these metrics. For a larger dataset, you can analyze these metrics for all data to find issues in your RAG system, so that you can improve the system further.

For now, we've concluded our short introduction to the evaluation of RAG systems.

OceanofPDF.com

6.6 Summary

In this chapter, we delved into the concept of RAG. This powerful tool can greatly assist you when you need to enhance an LLM response with your own content or documents. In most cases, this approach is more promising and preferred to fine-tuning a complete LLM.

We started with understanding the implementation of naïve (or standard) RAG, which consists of the retrieval, augmentation, and generation process. Later, we studied more detailed techniques to improve the system at different stages, including improvement strategies for pre-retrieval, in which the indexing pipeline is adapted in different ways.

In advanced retrieval techniques, you learned how a hybrid RAG system functions. Other techniques like query expansion and context enrichment were presented.

During advanced post-retrieval techniques, you learned about prompt compression as a means to remove irrelevant tokens from an LLM request.

Finally, we studied several RAG evaluation metrics, and you learned how to implement these metrics to measure the quality of your RAG systems.

In conclusion, RAG represents a significant advancement in the field of artificial intelligence. By merging retrieval and generative techniques, RAG opens new possibilities for creating high-quality, contextually rich outputs. This chapter equipped you with a comprehensive understanding of the concept, experimental results,

and practical applications, preparing you to implement this powerful tool in your own projects.

OceanofPDF.com

7 Agentic Systems

I'm sorry, Dave. I'm afraid I can't do that.
—HAL 9000 in *2001: A Space Odyssey*

With these words, HAL 9000, the artificial intelligence from *2001: A Space Odyssey*, embodies the essence of agentic systems. Agentic systems are technologies that operate beyond simple commands. They make decisions that sometimes defy human intentions. HAL, designed to support and protect the crew, shifts its objectives, challenging the human operators' control and opening a narrative of machine autonomy and agency.

In this chapter on agentic systems, we immerse ourselves in how AI models are empowered to operate autonomously within predefined guidelines, making decisions and performing actions to achieve specific goals. Unlike traditional generative artificial intelligence (generative AI) models that respond passively to user prompts, agentic systems are designed to understand objectives, adapt to dynamic inputs, and make context-sensitive choices without constant human guidance.

By leveraging task management, contextual awareness, and goal-oriented reasoning, these systems can handle complex workflows. Also, they can interact across multiple stages of a process and even integrate with other tools and data sources. This chapter explores the foundational principles, design considerations, and applications

of agentic systems, highlighting their growing role in transforming AI from reactive tools into proactive collaborators.

We'll start in [Section 7.1](#) with an introduction into AI agents. You'll learn what an agent is and what tradeoffs need to be considered. In [Section 7.2](#), we'll take a closer look at different frameworks. You'll learn about the available frameworks and gain some understanding about their differences and capabilities.

Then, we dive into the deep end and implement some agentic systems. We'll start with simple agents in [Section 7.3](#) and then become familiar with different frameworks in subsequent sections. Among these agentic frameworks we'll encounter are LangGraph in [Section 7.4](#), AG2 in [Section 7.5](#), CrewAI in [Section 7.6](#), OpenAI Agents in [Section 7.7](#), and finally Pydantic AI in [Section 7.8](#).

These systems can manage many agent interactions, but it is hard to debug such systems. In this context, monitoring systems come to the rescue, and we'll study some options in [Section 7.9](#).

Let's first learn what AI agents are in the next section.

7.1 Introduction to AI Agents

Let's start by asking what an AI agent is. An agentic system is an AI system that is designed to operate autonomously to complete tasks. In doing so, the system often interacts with its environment and uses tools. It can make decisions; take actions; and adapt its behavior based on goals, on external information, or if the context changes.

These agentic systems typically use large language models (LLMs) combined with tools, memory, or reasoning frameworks to achieve

complex, multistep tasks without additional user guidance. [Figure 7.1](#) shows the workflow of an agentic system.




Figure 7.1 Agentic System Workflow

The user sets up a query that will be handled by an agent. Imagine an agent as a supercharged LLM. The LLM can generate thoughts. Based on that reasoning, a suitable tool can be selected. The output of the tool call is reviewed, and if the output is considered useful to close the task, it returns the output. Otherwise, the agent starts the next loop, in which it creates more thoughts, selects a tool, and so on.

As in life, you must often deal with tradeoffs and can't get everything at once. One tradeoff you should consider, shown in [Figure 7.2](#), is the tradeoff between flexibility and reliability when considering agent autonomy. Flexibility and reliability can fall over different ranges of agent autonomy.

Flexibility in this context can be defined as the system's ability to adapt to diverse tasks, input, and contexts with minimal need for explicit reconfiguration. *Reliability* in this context refers to the ability to consistently perform as intended, that is, to deliver accurate and predictable outcomes across diverse scenarios and conditions.

Agent autonomy can range from simple code, to chains developed with LangChain, to routers, to full autonomy.




Figure 7.2 Agent Flexibility/Reliability versus Autonomy

In an ideal scenario, you want an agentic system with high flexibility and reliability, and at the same time, the agent should be completely autonomous. This full autonomy is not usually possible (yet).

With less autonomy, you get a system with high reliability that, at the same time, shows low autonomy. The opposite holds as well: An agentic system that is fully autonomous usually is quite flexible but has low reliability. The aim of the different frameworks is to increase autonomy without sacrificing too much on flexibility or on reliability.

Different approaches to achieve this goal are available. A whole ecosystem has developed around agentic frameworks, and several new frameworks are released every month. Over time, we'll see which will grow and survive. In the long run, we expect a whole family of frameworks that fulfill different needs. In this chapter, we're using the most popular frameworks at the time of writing (March 2025).

7.2 Available Frameworks

In Python, several frameworks support the development of agentic systems. They provide tools for autonomy, decision-making, and interaction with the environment. The field is extremely dynamic, and new frameworks emerge while others disappear. Since which frameworks will survive is not clear, we'll present several options in this section.

Let's briefly summarize these frameworks first before exploring them in depth:

- **LangGraph**

You know LangChain by now, but we haven't used it in an agentic way yet. For agent autonomy, the developers of LangChain created an additional framework built on top of LangChain, called LangGraph (<https://www.langchain.com/langgraph>).

LangGraph is a flexible framework that enables developers to design, build, and manage agentic systems. This framework makes use of a graph-based representation of tasks, workflows, and the decision-making process. It provides an intuitive way to define complex workflows by treating tasks as nodes and their dependencies as edges within a directed graph.

By combining graph theory with the capabilities of LLMs, LangGraph empowers developers to create agentic systems that balance flexibility, reliability, and autonomy. We'll dive more deeply into LangGraph in [Section 7.4](#).

- **Haystack**

Haystack (<https://haystack.deepset.ai/>) is another framework,

mainly used for LLMs and retrieval-augmented generation (RAG). But this framework also allows you to use tools and agents.

- **AG2**

AG2 (<https://github.com/ag2ai/ag2>), also called AgentOS or its former name autogen, is an agentic system created by Microsoft. This system simplifies the creation and orchestration of multi-agent systems, and it can solve complex tasks with minimal human intervention. We'll explore the intricacies of AG2 in [Section 7.5](#).

- **CrewAI**

CrewAI (<https://www.crewai.com/>) is a multi-agent system that allows users to quickly build automations. We take a look at this system in [Section 7.6](#).

- **OpenAI Agents**

OpenAI Agents (<https://github.com/openai/openai-agents-python>) is a lightweight, but powerful, framework specifically designed for multiagent workflows. Fun to work with and quite approachable, you'll learn how to use it in [Section 7.7](#).

- **Magentic-One**

Magentic-One

(<https://github.com/microsoft/autogen/tree/main/python/packages/autogen-magentic-one>) is another framework developed by Microsoft, built on top of AG2.

Magentic-One has a powerful feature: interacting with the digital world. This multi-agent system is designed for dealing with open-ended tasks. These tasks can make use of the web and the operating system on which it runs. The system is designed to work on the real tasks that people encounter in their daily lives.

- **Tiny Troupe**

Tiny Troupe (<https://github.com/microsoft/TinyTroupe>) is yet another agentic framework developed by Microsoft. Still experimental and young, its focus is on creating multi-agent simulations of personas.

This framework can be used to evaluate digital ads with a simulated audience before spending actual money on a campaign. Furthermore, the framework can create test inputs to systems for software testing, or you can use it to evaluate project or product proposals. The framework is quite simple to set up and to use.

- **Pydantic AI**

Pydantic AI is an agentic framework for developing production-grade applications. Important features for production-grade applications include, most importantly, type safety which enables structured inputs and outputs. This framework is also model agnostic and, as such, allows the use of various LLMs. It neatly integrates with Logfire, a monitoring system from the same development team. One of our favorite systems, we cover it in [Section 7.8](#).

Before we look more closely at all these options, let's start by discussing simple agents.

7.3 Simple Agent

In [Chapter 6](#), we worked with RAG. We'll build upon that knowledge with a simple implementation of an agentic RAG. Then, we'll develop an agentic system that combines both reasoning and acting and thus is called *ReAct*.

7.3.1 Agentic RAG

Agentic RAG describes a system that relies on RAG, but with a small tweak. In the pipeline, intelligent AI agents are enabled to improve the system whenever the results are insufficient. For example, when the RAG system does not retrieve relevant information from the data retrieval system (typically the vector database), other information sources are tapped and used generating the final answer. In our first coding example, we'll implement this functionality. [Figure 7.3](#) shows the workflow of our agentic RAG system.

After the vector database provides information, the information is analyzed for its relevance. If not sufficiently relevant, data is retrieved via tool use, specifically by searching the internet. Then, the information is used for creating a final output that can be returned to the user. The complete code can be found in `07_AgenticSystems\10_agentic_rag.py`.

In our code, we'll use a tool for searching the internet called Tavily, which is an online service empowering AI apps to quickly search the internet. Before you can use this service in your applications, you

must set up an application programming interface (API) key by creating an account and the API key at www.tavily.com.




Figure 7.3 Agentic RAG Workflow

You should call the API key `TAVILY_API_KEY` and add it to your `.env` file. Don't worry about costs for now—they provide a generous free tier, which will definitely be enough for you to experiment. The entry in your `.env` file should look like the following:

```
TAVILY_API_KEY=tvly-...
```

After you load the environment in the usual way, create an instance of `TavilySearchResults`. To test if this step worked, invoke the instance, for example, with the following code:

```
search_tool.invoke("what is rag?")
```

This command will return a list of dictionaries. Each dictionary contains the keys `url` and `content`, as follows:

```
[{'url': 'https://www.geeksforgeeks.org/what-is-retrieval-augmented-generation-rag/',
  'content': 'RAG is a natural language processing approach that combines retrieval
...}, ...]
```

As shown in [Listing 7.1](#), at the beginning of the script, we're loading the necessary packages. Noteworthy in this context is the FAISS package. Standing for “Facebook AI similarity search,” FAISS is a small and simple implementation of a vector database.

```
### packages
from langchain_community.tools.tavily_search.tool import TavilySearchResults
from dotenv import load_dotenv, find_dotenv
from langchain_openai import ChatOpenAI
from langchain_community.document_loaders import WikipediaLoader
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain.embeddings import OpenAIEmbeddings
from langchain.vectorstores import FAISS
from langchain.prompts import ChatPromptTemplate
load_dotenv(find_dotenv(usecwd=True))
```

Listing 7.1 Agentic RAG: Required Packages (Source: 07_AgenticSystems\10_agentic_rag.py)

Our data source is Wikipedia. We'll create a vector database and follow the classic data ingestion pipeline step by step. As shown in [Listing 7.2](#), we'll first load ten articles on the “principle of relativity.” We'll follow the classic indexing pipeline from [Chapter 5](#), as follows:

1. First, we'll create chunks of Wikipedia article data.
2. Embeddings based on OpenAIEmbeddings are created.
3. The vectorstore is populated with the chunks and corresponding embeddings.
4. A retriever enables quick searches of the database.

As an alternative approach, we can initialize web search. For this functionality, we'll use TavilySearchResults, as shown in [Listing 7.2](#).

```
# Load documents for retrieval (can be replaced with any source of text)
# Here we're using a text loader with some sample text files as an example
### import wikipedia
loader = WikipediaLoader("Principle of relativity",
                        load_max_docs=10)

docs = loader.load()

### create chunks
```

```

text_splitter = RecursiveCharacterTextSplitter(chunk_size=1000, chunk_overlap=200)
chunks = text_splitter.split_documents(docs)

### models and tools
llm = ChatOpenAI(model="gpt-4o-mini", temperature=0)
embedding = OpenAIEmbeddings()
search_tool = TavilySearchResults(max_results=5, include_answer=True)

### use FAISS to store the chunks
vectorstore = FAISS.from_documents(chunks, embedding)
retriever = vectorstore.as_retriever(return_similarities=True)

### user query
query = "What is relativity?"

```

Listing 7.2 Agentic RAG: Vector Database Creation (Source: 07_AgenticSystems\10_agentic_rag.py)

The RAG chain shown [Listing 7.3](#) is based on a `prompt_template`. In this example, we provide the context and the question. The `retrieved_docs` are passed as context to the chain. An important aspect in the prompt template is the instruction to return 'insufficient information' if the answer cannot be known based on the context. This output is later used in a router to decide if further information must be loaded from the internet or if the RAG output can be used as the final answer.

```

### RAG chain
prompt_template = ChatPromptTemplate.from_messages([
    ("system", """
        You are a helpful assistant that can answer questions about the principle of
        relativity. You will get contextual information from the retrieved documents. If you
        don't know the answer, just say 'insufficient information'
        """),
    ("user", "<context>{context}</context>\n\n<question>{question}</question>"),
])
retrieved_docs = retriever.invoke(query)
retrieved_docs_str = "".join([doc.page_content for doc in retrieved_docs])
chain = prompt_template | llm
rag_response = chain.invoke({"question": query,
                             "context": retrieved_docs_str})

```

Listing 7.3 Agentic RAG: RAG Setup (Source: 07_AgenticSystems\10_agentic_rag.py)

Listing 7.4 shows the router—the building block that checks if an internet search tool is needed or if the RAG output can be used. In this simple example, the router is implemented as a basic if-else statement.

```
if rag_response.content == "insufficient information":
    print("using search tool")
    final_response = search_tool.invoke({"query": query})
    final_response_str = ";".join([doc['content'] for doc in final_response])
    final_response = chain.invoke({"question": query,
                                   "context": final_response_str})
else:
    print("using vector store")
    final_response = rag_response.content

final_response
```

Relativity, in the context of physics, refers to the scientific theories developed by Albert Einstein that describe the relationship between space and time. It encompasses two main theories: special relativity and general relativity. Special relativity, introduced in 1905, is based on two key postulates: the laws of physics are the same in all inertial frames of reference, and the speed of light in vacuum is constant for all observers, regardless of their motion. General relativity, published in 1915, extends these concepts to include the effects of gravity on the structure of spacetime. Together, these theories have profound implications for our understanding of the universe and have been confirmed by numerous experiments.

Listing 7.4 Agentic RAG: Router Logic (Source: 07_AgenticSystems\agentic_rag.py)

That's it for our simple implementation of an agentic RAG system!

In the next section, you'll learn about ReAct—a paradigm where agents iteratively reason and act based on the outcomes of previous steps.

7.3.2 ReAct

ReAct is a technique that combines reasoning and acting. This combination allows you to improve how LLMs interact with information and with their environment.

In the reasoning phase, thoughts are generated. In the acting phase, tools are used to perform tasks. This approach is valuable for agentic systems, in which a model must decide on the next step depending on a given context. Let's break down these two tasks in more detail:

- **Reasoning steps**

The model can break down a complex problem into smaller steps, which improves the accuracy and transparency of a model's decisions. For instance, if a model is supposed to retrieve information from a database, it might first decide what specific information it needs, outline why it needs that information, and then act accordingly.

- **Acting based on reasoning**

After reasoning through a problem, the model decides on an action, such as retrieving data, formulating an answer, or updating its plan. ReAct thus brings structure and coherence to multi-step processes by combining planning and action, so the model doesn't just produce isolated outputs but instead follows some logical progression.

Let's see how ReAct works in practice. You can find the sample code in `07_AgenticSystems\20_react.py`. In this scenario, we want to check if agent can "know" about things already discussed. Furthermore, the agent will be able to select a tool to reach a certain goal.

As shown in [Listing 7.5](#), we'll load the packages we need in for script. Newcomers to this list of packages are `MemorySaver` and `create_react_agent`. `MemorySaver` enables the agent to retrieve some memory, which is helpful for longer conversations so the agent can refer to earlier discussions. With `create_react_agent`, we can easily set up a ReAct agent. With these packages, we are ready to go.

```

### packages
from langchain_groq import ChatGroq
from langchain_community.tools.tavily_search import TavilySearchResults
from langchain_core.messages import HumanMessage
from langgraph.checkpoint.memory import MemorySaver
from langgraph.prebuilt import create_react_agent
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))

```

Listing 7.5 ReAct Agent: Required Packages (Source: 07_AgenticSystems\20_react.py)

Now, we can set up the actual agent. As shown in [Listing 7.6](#), our agent will rely on a `model`, a `checkpoint`, and one or more `tools`. In this simple example, we'll only provide one tool. The `tools` are passed as a list. With the `TavilySearchResults` tool, the internet can be searched, if needed.

```

### Create the agent
memory = MemorySaver()
model = ChatGroq(model_name="llama-3.1-70b-versatile")
search = TavilySearchResults(max_results=2)
tools = [search]

```

Listing 7.6 ReAct Agent: Agent Helpers (Source: 07_AgenticSystems\20_react.py)

Now, we'll set up an instance of `create_react_agent`, based on the `model`, the `tools`, and the `checkpoint`, as follows:

```

agent_executor = create_react_agent(model=model,
                                   tools=tools,
                                   checkpoint=memory)

```

Great! When we invoke the agent, we must pass a `configuration` with a `thread_id`. This step allows us to manage multiple conversations, as follows:

```

### Set up the configuration
config = {"configurable": {"thread_id": "abc123"}}

```

Ready to see the agent in action? As shown in [Listing 7.7](#), invoke the `agent_executor` and pass a user message.

```

#%% Invoke the agent
agent_executor.invoke(
    {"messages": [("user", "My name is Bert Gollnick, I am a trainer and data
scientist. I live in Hamburg")]}, config
)

{'messages': [    HumanMessage(content='My name is Bert Gollnick, I am ...',
                        AIMessage(content='', ...),
                        ToolMessage(content=...),
...]}

```

Listing 7.7 ReAct Agent: Agent Execution (Source: 07_AgenticSystems\20_react.py)

The stack of messages is printed to the screen. Notice that an AI message is returned after our message, and a tool was called for fetching further information about the author.

The messages end up in memory. To extract these messages later, you must route through the memory object by writing a new function, shown in [Listing 7.8](#).

```

#%% function for extracting the last message from the memory
def get_last_message(memory, config):
    return memory.get_tuple(config=config).checkpoint['channel_values']['messages']
    [-1].model_dump()['content']

```

Listing 7.8 ReAct Agent: Function for Fetching Messages (Source: 07_AgenticSystems\20_react.py)

Now, we find out if our agent can remember things by asking it about the author's name and residence, as shown in [Listing 7.9](#). While this test sounds trivial, none of the systems we've set up to this point can remember previous interactions.

```

#%% check whether the model can remember me
agent_executor.invoke(
    {"messages": ("user", "What is my name and in which country do I live?")},
    config
)
get_last_message(memory, config)

```

Your name is Bert Gollnick, and based on the information provided earlier, you live in Hamburg, Germany.

Listing 7.9 ReAct Agent: Invocation (Source: 07_AgenticSystems\20_react.py)

Amazing! Our agent now has the memory to recall information from previous interactions. But now it gets really cool. Let's ask the model what it can find out about the author on the internet. What does the agent need to do to fulfil this task? It needs to recall the author's full name and then reflect on whether any of the available tools are helpful for solving this task. Luckily, our agent can use the web search tool we provided. It will use the tool, analyze the results, and hopefully show that can find information on the author, as shown in [Listing 7.10](#).

```
### check if it is possible to find me on the internet
agent_executor.invoke(
    {"messages": ("user", "What can you find about me on the internet")},
    config
)
get_last_message(memory, config)

'Based on the search results, it appears that you, Bert Gollnick, are a data
scientist with expertise in renewable energies, particularly wind energy. You have
worked for a leading wind turbine manufacturer and have applied data science to this
field. You also have a background in Aeronautics and Economics. Additionally, you
are a Udemy instructor and have courses available for enrollment on the platform.'
```

Listing 7.10 ReAct Agent: Invocation (Source: 07_AgenticSystems\20_react.py)

We can also review all the information stored in the memory by calling `memory.list()`, as shown in [Listing 7.11](#).

```
# %% get the complete list of messages
list(memory.list(config=config))

[CheckpointTuple(config={'...'}),
 CheckpointTuple(config={'...'}),
 ...
]
```

Listing 7.11 ReAct Agent: Memory (Source: 07_AgenticSystems\20_react.py)

In this section, you learned about an important feature for memorizing information throughout a conversation. This feature (among others) is available in specific agentic frameworks. We'll start with a framework developed by the creators of LangChain: LangGraph.

OceanofPDF.com

7.4 Agentic Framework: LangGraph

LangGraph (<https://langchain-ai.github.io/langgraph/>) is a dynamic framework designed to simplify the creation, visualization, and management of complex workflows in AI-powered applications. At its core, LangGraph provides developers and researchers with tools to streamline task orchestration, enabling them to build modular, efficient, and explainable pipelines for language models and other generative AI components.

LangGraph uses *directed acyclic graphs (DAGs)* to structure and execute workflows. Each node in a LangGraph workflow represents a task, such as prompting a language model, running a computation, or invoking an external API. Edges between nodes define dependencies and data flow, ensuring that tasks are executed in the correct order. With LangGraph, you maintain a high level of control over the workflow, which makes the framework quite stable and reliable.

This graph-based approach offers several advantages, such as the following:

- **Modularity:** Each node encapsulates a specific function or task, making workflows easy to build, debug, and reuse.
- **Transparency:** The visual representation of workflows clarifies how data and tasks interconnect, providing insights into the overall process.
- **Scalability:** LangGraph supports branching and parallelism, enabling the development of complex systems that scale efficiently.

You can install LangGraph via pip or uv with the following commands:

- `pip install langgraph`
- `uv add langgraph`

In our examples, we'll start with the basics of setting up a simple graph. Then, gradually, we'll add more relevant features. At the end of this chapter, we'll develop a capstone project that integrates all our new knowledge.

Let's get started with the simplest graph possible.

7.4.1 Simple Graph: Assistant

[Figure 7.4](#) shows the graph we'll develop first. In [Chapter 3](#), you learned how to work with LLM instances and how to invoke them. With LangGraph, we set a DAG, which is basically a graph with a clear start and end node. Each graph consists of nodes and edges. Nodes represent the LLM we use but it could also be functions or tools the LLM can call. The nodes are connected by edges. Edges define how the nodes are connected.




Figure 7.4 LangGraph: A Simple Graph

In this simplest graph, we have a start node that is connected to a node called assistant. The assistant node then is linked to an end node. This graph will work as an LLM assistant: It can be invoked and can answer user questions. The corresponding code can be found in

07_AgenticSystems\langgraph\10_langgraph_simple_assistant.py.

[Listing 7.12](#) shows all the required packages for this script. In this case, we need some packages for loading environment variables (dotenv), langchain_groq for interacting with an LLM, and several functionalities from langgraph.

```
### packages
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
from typing import Annotated
from typing_extensions import TypedDict
from langgraph.graph import StateGraph, START, END
from langgraph.graph.message import add_messages
from langchain_groq import ChatGroq
from IPython.display import Image, display
```

Listing 7.12 LangGraph: Simple Graph: Packages (Source: *07_AgenticSystems\langgraph\10_langgraph_simple_assistant.py*)

Let's discuss these new functionalities. We'll create a stateGraph that inherits from TypedDict. In the stateGraph, a state with messages will be assigned. Each graph needs a START and END node.

A LangGraph graph requires a state because the state serves as the shared context and data container that flows through the graph during its execution. This state provides a single, consistent place to store and update data as it is processed by different nodes in the graph. Each node can read from and write to the state. This approach enables coordinated execution without the need for complex direct communication between nodes. The state is the only way nodes exchange information.

[Listing 7.13](#) shows how the state is set up.

```
# %% define the state
class State(TypedDict):
    messages: Annotated[list, add_messages]
```

Listing 7.13 LangGraph: Simple Graph: State (Source: 07_AgenticSystems\langgraph\10_langgraph_simple_assistant.py)

Our state inherits from `TypedDict`. A typed dictionary is much like a normal dictionary except that the keys must have a type defined.

The only property of our state is `messages`. The messages have the type of `list`. `Annotated` in this context allows the list to get the context-specific metadata `add_messages`. Basically, the definition `Annotated[list, add_messages]` means that the messages are of type `list` with the small caveat that messages can be added to the list.

The next code chunk is straightforward. [Listing 7.14](#) shows how the LLM and the assistant are set up.

```
%% set up the assistant
llm = ChatGroq(model="gemma2-9b-it")

def assistant(state: State):
    return {"messages": [llm.invoke(state["messages"])]}
```

Listing 7.14 LangGraph: Simple Graph: LLM and Assistant (Source: 07_AgenticSystems\langgraph\10_langgraph_simple_assistant.py)

For the LLM, we'll use Groq with a small and fast model. Nothing new here. But the function `assistant` is defined in a way that it fits perfectly into LangGraph, and we can easily use this function in the context of the node setup. For this task, `assistant` consumes a state and returns a dictionary with messages.

[Listing 7.15](#) shows the most interesting part of the script, where we define the inner structure of the graph.

```
### create the graph
graph_builder = StateGraph(State)
graph_builder.add_node("assistant", assistant)
graph_builder.add_edge(START, "assistant")
graph_builder.add_edge("assistant", END)
```

Listing 7.15 LangGraph: Simple Graph: Graph Creation (Source: 07_AgenticSystems\langgraph\10_langgraph_simple_assistant.py)

The `graph_builder` object starts out as an empty `StateGraph`. Then, we add a single node with the name “assistant” and the action `assistant`. The latter refers to the function we set up earlier. In this way, we add our nodes: They are referred to via their name and mapped to a function.

But at this point, the node is not connected to anything. In the graph, a clear `START` and `END` is required. The connection of the dots is simple. With the `add_edge` method, we can map out all the connections between the nodes. This function requires a start key and an end key. The edge points from the start key to the end key. The edges we add here reflect this setup:

```
START → “assistant” → END
```

In later examples, we’ll create more complex setups, but now, we’re starting out simply.

This `graph_builder` object is like a Python class. To use a Python class, you must create a class instance, and the same rule applies now. However, the syntax is slightly different, as shown in [Listing 7.16](#). To create an instance of the `graph_builder`, you must call its `compile` method.

```
# %% compile the actual graph
graph = graph_builder.compile()
```

Listing 7.16 LangGraph: Simple Graph: Graph Compilation (Source: 07_AgenticSystems\langgraph\10_langgraph_simple_assistant.py)

Finally, our graph is ready. Let's make a visualization of it, as shown in [Figure 7.5](#). [Listing 7.17](#) presents the commands to create a visualization of the graph with its nodes and edges.

```
# %% display graph
display(Image(graph.get_graph().draw_mermaid_png()))
```

Listing 7.17 LangGraph: Simple Graph: Graph Visualization (Source: 07_AgenticSystems\langgraph\10_langgraph_simple_assistant.py)




Figure 7.5 LangGraph: Our Simple Graph

Now, we can work with the graph. [Listing 7.18](#) shows how to invoke the graph.

```
# %% invoke the graph
res = graph.invoke({"messages": [{"user", "What do you know about LangGraph?"}]})
```

Listing 7.18 LangGraph: Simple Graph: Graph Invocation (Source: 07_AgenticSystems\langgraph\10_langgraph_simple_assistant.py)

The result of this invocation is stored in the object `res`, as shown in [Listing 7.19](#). This object is a dictionary with key "messages".

```
# %% display the result
res["messages"]

[HumanMessage(content='What do you know about LangGraph?', additional_kwargs={},
response_metadata={}, id='d241780d-203d-46be-bb78-682bb5d0fed1'),
AIMessage(content="LangGraph is a powerful open-source tool developed by the
HuggingFace team ...",
```

```
additional_kwargs={}, response_metadata={'token_usage': {'completion_tokens': 382,
'prompt_tokens': 17, 'total_tokens': 399, 'completion_time': 0.694545455,
'prompt_time': 8.357e-05, 'queue_time': 0.015116917, 'total_time': 0.694629025},
'model_name': 'gemma2-9b-it', 'system_fingerprint': 'fp_10c08bf97d',
'finish_reason': 'stop', 'logprobs': None}, id='run-186bf566-45ad-462f-b65e-
67005df3c95c-0',
usage_metadata={'input_tokens': 17, 'output_tokens': 382, 'total_tokens': 399}}]
```

Listing 7.19 LangGraph: Simple Graph: Invocation Result (Source: 07_AgenticSystems\langgraph\10_langgraph_simple_assistant.py)

The invocation is successful, but at this point, you might ask what the benefit is. Up to this point, the use of the LLM has been more complicated than a direct LLM invocation. The value of LangGraph gets clearer once we cover more complex examples and setups.

Let's continue by adding autonomy to the graph. In our current example, the graph has a completely deterministic workflow. In the next example, we'll add flexibility to the graph by allowing it to decide how to proceed.

7.4.2 Router Graph

So far, the process flow has been linear. The graph started, the state was passed to a node, and the node passed its state to the next node. Now, we can provide the agent with slightly more autonomy.

We'll create a graph that lets a node decide how to proceed with the code. [Figure 7.6](#) shows the graph workflow for our new router graph.




Figure 7.6 LangGraph: Router Graph

The router will decide how the process should proceed. Specifically, we'll create an agent to provide arguments for or against a topic. The agent will randomly pick a side and create arguments for that side. The code for this script can be found in *07_AgenticSystems\langgraph\11_langgraph_router.py*.

[Listing 7.20](#) shows all required packages for our project.

```
### packages
from pprint import pprint
from typing_extensions import TypedDict
import random
from langgraph.graph import StateGraph, START, END
from langchain_groq import ChatGroq
from IPython.display import Image, display
from rich.console import Console
```

```
from rich.markdown import Markdown
console = Console()
```

Listing 7.20 LangGraph Router Project: Packages (Source: 07_AgenticSystems\langgraph\11_langgraph_router.py)

We don't import any new packages, so we can go straight to more interesting code chunks.

[Listing 7.21](#) shows the applied LLM and the setup of the state class. The graph state represents a snapshot of the context and data that is associated with an agent at a specific point in its execution.

```
## LLM
llm = ChatGroq(model="gemma2-9b-it")
# State with graph_state
class State(TypedDict):
    graph_state: dict[str, str | dict[str, str | str]]
```

Listing 7.21 LangGraph Router Project: Used LLM and State Class (Source: 07_AgenticSystems\langgraph\11_langgraph_router.py)

A graph consists of nodes and edges. We start with the definition of the nodes, as shown in [Listing 7.22](#).

```
# Nodes
def node_router(state: State):
    # Retrieve the user-provided topic
    topic = state["graph_state"].get("topic", "No topic provided")

    # Update the graph_state with any additional information if needed
    state["graph_state"]["processed_topic"] = topic # Example of updating
    graph_state

    print(f"User-provided topic: {topic}")
    return {"graph_state": state["graph_state"]}

def node_pro(state: State):
    topic = state["graph_state"]["topic"]
    pro_args = llm.invoke(f"Generate arguments in favor of: {topic}. Answer in
bullet points. Max 5 words per bullet point.")
    state["graph_state"]["result"] = {"side": "pro", "arguments": pro_args}
    return {"graph_state": state["graph_state"]}

def node_contra(state: State):
    topic = state["graph_state"]["topic"]
    contra_args = llm.invoke(f"Generate arguments against: {topic}")
```

```
state["graph_state"]["result"] = {"side": "contra", "arguments": contra_args}
return {"graph_state": state["graph_state"]}
```

Listing 7.22 LangGraph Router Project: Nodes Definition (Source: 07_AgenticSystems\langgraph\11_langgraph_router.py)

In this setup, we have three nodes: a router node, a pro-node, and a contra-node. The router node retrieves the topic from the user and updates the graph. The topic is stored in the key "topic".

The pro- and contra-nodes invoke the LLM and create the answers in favor and opposed to the topic. They create a key "result" in the graph_state in which they pass their content.

We still need to connect the dots, in this case, the nodes with edges. This step will be done by directly linking nodes or using functions for conditional edges. [Listing 7.23](#) shows the definition of a conditional edge in which a random decision is made to select a side.

```
# Edges
def edge_pro_or_contra(state: State):
    decision = random.choice(["node_pro", "node_contra"])
    state["graph_state"]["decision"] = decision
    print(f"Routing to: {decision}")
    return decision
```

Listing 7.23 LangGraph Router Project: Edges Definition (Source: 07_AgenticSystems\langgraph\11_langgraph_router.py)

[Listing 7.24](#) shows how the complete graph is built and then compiled.

```
# Create graph
builder = StateGraph(State)
builder.add_node("node_router", node_router)
builder.add_node("node_pro", node_pro)
builder.add_node("node_contra", node_contra)

builder.add_edge(START, "node_router")
builder.add_conditional_edges("node_router", edge_pro_or_contra)
builder.add_edge("node_pro", END)
builder.add_edge("node_contra", END)
```

```
graph = builder.compile()
```

Listing 7.24 LangGraph Router Project: Graph Creation (Source: 07_AgenticSystems\langgraph\11_langgraph_router.py)

First, the nodes are added to the graph. The `add_node` function takes two parameters: a node name and a corresponding node object.

The `START` node is connected to `node_router`. The decision for picking a side is made with the help of `add_conditional_edges`. In this context, the previously defined `node_router` is linked to the function we created: `edge_pro_or_contra`. Finally, `node_pro` and `node_contra` are linked to the `END` node.

At this point, the graph definition is ready, and we can call `compile` to create a graph instance. This step is necessary because, in LangGraph, the framework relies on defining and optimizing the flow of reasoning and actions before execution.

To create a visual representation of the graph, as shown in [Figure 7.6](#), use the following code:

```
# %% display the graph
display(Image(graph.get_graph().draw_mermaid_png()))
```

Test the graph by invoking it. [Listing 7.25](#) shows the invocation process.

```
# %% Invokation
initial_state = {"graph_state": {"topic": "Should dogs wear clothes?"}}
result = graph.invoke(initial_state)
console.print(Markdown(result["graph_state"]['result'] ['arguments'].model_dump()
['content']))
```

Listing 7.25 LangGraph Router Project: Graph Invocation (Source: 07_AgenticSystems\langgraph\11_langgraph_router.py)

We set up an `initial_state` by passing a `graph_state` in which the `topic` value is defined. This `initial_state` is passed via the `invoke`

method. The result is stored in an object, and finally, we can use functionality from the `rich` package (`Console`, `Markdown`) to create a visually appealing console output, as shown in [Listing 7.26](#).

```
Arguments Against Dogs Wearing Clothes:
1. It's unnecessary for most dogs: Dogs have natural fur coats that provide
insulation and protection from the elements. In most situations, clothing is not
necessary for their well-being.
2. It can be uncomfortable and restrictive: Clothes can restrict a dog's movement,
especially if they're
ill-fitting or made of non-breathable materials. This can lead to discomfort,
chafing, and even injury.
3. It can be a safety hazard: Loose clothing can get caught on objects, posing a
strangulation risk. Tags and attachments on clothing can also snag and cause
injury.
...
Ultimately, the decision of whether or not to dress a dog should be made on a case-
by-case basis, considering the
individual dog's needs and comfort.
```

Listing 7.26 LangGraph Router Project: Output (Source: 07_AgenticSystems\langgraph\11_langgraph_router.py)

At this point, you've learned how to create a dynamic graph pipeline in which the process flow is not purely sequential, but conditional instead. This capability allows our agents and graphs to dynamically adapt to conditions. But one important capability is still missing: tools.

To equip agents with the capability to select tools independently, let's expand on our graph.

7.4.3 Graph with Tools

Tools are the main ingredients that make agents autonomous. LangGraph allows you to bind tools to your LLM, so that a node can

decide to use a tool when needed. [Figure 7.7](#) shows the graph we'll develop now.




Figure 7.7 LangGraph: Graph with Tool Use

In this scenario, we want to provide a useful functionality—the ability to count letters. As it turns out, LLMs have enormous capabilities but sometimes struggle with simple tasks like counting the letters in a word. This problem is related to Moravec's paradox, which states that tasks requiring high-level reasoning are easier for computers to perform, while sensorimotor tasks, which humans consider easy, are surprisingly difficult for machines to replicate. In short, what is simple for a computer is difficult for a human, and vice versa.

Examples of high-level reasoning include solving complex mathematical equations, performing calculations with large numbers, playing chess or Go, and analyzing large datasets to detect patterns. Such tasks are simple for a machine but difficult for humans. On the other hand, walking, lifting a brick and then lifting an egg, and

identifying sarcasm are all tasks in which humans excel over machines.

Let's work on one such paradoxical behavior—counting letters in a word. [Figure 7.8](#) shows a model call asking an LLM to count the number of Ls in the word LOLPALOOZA.

Although this task is simple enough for preschoolers, LLMs struggle to come up with the right answer. Newer models seem to more consistently get the right answer. But as you can see from the screenshot, Mixtral-8x7B cannot figure out the right answer.

For this reason, we'll write a tool our agent can use when such a request occurs. The complete code can be found in [07_AgenticSystems\langgraph\12_langgraph_tools.py](#).




Figure 7.8 Inability of LLMs to Count Letters

[Listing 7.27](#) lists all required packages and the LLM instance used in this example.

```
### packages
import langgraph
from langgraph.graph import StateGraph, START, END
from typing_extensions import TypedDict
from langchain_openai import ChatOpenAI
from langchain_core.messages import AIMessage, HumanMessage
from langgraph.graph.message import add_messages
from dotenv import load_dotenv, find_dotenv
```

```
load_dotenv(find_dotenv(usecwd=True))
#%% LLM
llm = ChatOpenAI(model="gpt-4o-mini", temperature=0)
```

Listing 7.27 LangGraph: Graph with Tools: Packages (Source: 07_AgenticSystems\langgraph\12_langgraph_tools.py)

At this point, you should already be familiar with these packages and how to set up the model. So, we'll move straight to the tool we want to equip the LLM with.

[Listing 7.28](#) shows the function `count_characters_in_word`. Provided with a word and a specific character, this function returns the correct number of the target letter found in the word.

```
# %% Tools
def count_characters_in_word(word: str, character: str) -> str:
    """Count the number of times a character appears in a word."""
    cnt = word.count(character)
    return f"The word {word} has {cnt} {character}s."

# %% TEST
count_characters_in_word(word="LOLLAPALOOZA", character="L")
```

Listing 7.28 LangGraph: Graph with Tools: Tool Definition (Source: 07_AgenticSystems\langgraph\12_langgraph_tools.py)

Counting the number of letters is performed with the `count` method. You also directly test the tool to verify that the right number is returned, as follows:

```
'The word LOLLYPALOOZA has 4 Ls.'
```

Great! Now we can connect the tool (our function) to the LLM. LangGraph requires tools to be integrated explicitly to ensure they are available and properly configured. The `bind_tools` function connects your tools to the LangGraph system. This step makes them accessible to the graph's nodes or agents, as follows:

```
# %% LLM with tools
llm_with_tools = llm.bind_tools([count_characters_in_word])
```


You can also test this functionality by calling the `invoke` method, as shown in [Listing 7.29](#).

```
# %%
llm_with_tools.invoke(["user", "Count the Ls in LOLLYPAL00ZA?"])

AIMessage(content='', additional_kwargs={'tool_calls': [{'id':
'call_IRNS0IpEXeHY1iGqv3T6WTh1', 'function': {'arguments':
'{"word":"LOLLAPAL00ZA","character":"L"}', 'name': 'count_characters_in_word'},
'type': 'function'}], 'refusal': None}, response_metadata={'token_usage':
{'completion_tokens': 27, 'prompt_tokens': 73, 'total_tokens': 100,
'completion_tokens_details': {'accepted_prediction_tokens': 0, 'audio_tokens': 0,
'reasoning_tokens': 0, 'rejected_prediction_tokens': 0}, 'prompt_tokens_details':
{'audio_tokens': 0, 'cached_tokens': 0}}, 'model_name': 'gpt-4o-mini-2024-07-18',
'system_fingerprint': 'fp_d02d531b47', 'finish_reason': 'tool_calls', 'logprobs':
None}, id='run-4acd5baf-fb21-4dd5-8fa3-cca800a50b0e-0', tool_calls=[{'name':
'count_characters_in_word', 'args': {'word': 'LOLLAPAL00ZA', 'character': 'L'},
'id': 'call_IRNS0IpEXeHY1iGqv3T6WTh1', 'type': 'tool_call'}], usage_metadata=
{'input_tokens': 73, 'output_tokens': 27, 'total_tokens': 100,
'input_token_details': {'audio': 0, 'cache_read': 0}, 'output_token_details':
{'audio': 0, 'reasoning': 0}})
```

Listing 7.29 LangGraph: Graph with Tools: Tool Invocation (Source: 07_AgenticSystems\langgraph\12_langgraph_tools.py)

The tool was successfully called. You can also verify success by extracting the last message, as shown in [Listing 7.30](#).

```
# %% Tool Call
tool_call = llm_with_tools.invoke("How many Ls are in LOLLYPAL00ZA?")
# %% extract last message
tool_call.additional_kwargs["tool_calls"]

[{'id': 'call_1kmLBIuF6CIn8fs1MuSy1bhb',
'function': {'arguments': '{"word":"LOLLAPAL00ZA","character":"L"}',
'name': 'count_characters_in_word',
'type': 'function'}]
```

Listing 7.30 LangGraph: Graph with Tools: Tool Calls (Source: 07_AgenticSystems\langgraph\12_langgraph_tools.py)

Let's now create the graph and visualize it so that we get the screen shown earlier in [Figure 7.7](#) but using the code provided in [Listing 7.31](#). This step helps manage tool interaction in a broader context.

```

%% graph
from IPython.display import Image, display
from langgraph.graph import StateGraph, START, END
from typing_extensions import TypedDict
from langchain_core.messages import AnyMessage
from langgraph.prebuilt import ToolNode, tools_condition

class MessagesState(TypedDict):
    messages: list[AnyMessage]

# Node
def tool_calling_llm(state: MessagesState):
    return {"messages": [llm_with_tools.invoke(state["messages"])]}

# Build graph
builder = StateGraph(MessagesState)
builder.add_node("tool_calling_llm", tool_calling_llm)
builder.add_node("tools", ToolNode([count_characters_in_word]))
builder.add_edge(START, "tool_calling_llm")
# builder.add_edge("tool_calling_llm", "tools")
builder.add_conditional_edges("tool_calling_llm",
    # If the latest message (result) from assistant is a
    tool call -> tools_condition routes to tools
    # If the latest message (result) from assistant is a not a tool call ->
    tools_condition routes to END
    tools_condition)
builder.add_edge("tools", END)
graph = builder.compile()

```

Listing 7.31 LangGraph: Graph with Tools: Building the Graph (Source: 07_AgenticSystems\langgraph\12_langgraph_tools.py)

This code defines a state graph using `StateGraph` class. The graph structures the workflow by defining nodes and edges that represent different states and transitions. In this case, the graph starts with a node of our `tool_calling_llm`. That node processes the user messages through the LLM. Depending on the LLM's response, the graph routes the flow conditionally.

If the LLM's output includes a tool call, it transitions to a `ToolNode`, which executes the required tool. The result is then routed back to the workflow, either for further processing or to conclude the interaction.

This separation of concerns—first binding tools to the LLM and then managing their invocation through the state graph—is key to the design.

Binding tools to an LLM allows the model to make tool calls dynamically based on user input, while the state graph ensures that tool calls are integrated into a controlled, sequential process. This approach is particularly useful in complex workflows where tool usage must be orchestrated alongside other operations, such as message routing or additional computations.

By combining the LLM's reasoning capabilities with a defined state graph, the code creates a flexible yet robust framework for handling user interactions and tool integrations.

We can visualize the graph with the following command:

```
# View
display(Image(graph.get_graph().draw_mermaid_png()))
```

Finally, as shown in [Listing 7.32](#), we can test the graph in action by creating some messages and invoking the graph.

```
# %% use messages as state
# messages = [HumanMessage(content="Hey, how are you?")]
messages = [HumanMessage(content="Please count the Ls in LOLLAPAL00ZA.")]
messages = graph.invoke({"messages": messages})
for m in messages["messages"]:
    print(m.pretty_print())
```

```
===== Tool Message
```

```
=====
```

```
Name: count_characters_in_word
```

```
The word LOLLAPAL00ZA has 4 Ls.
```

```
None
```

Listing 7.32 LangGraph: Graph with Tools: Graph Invocation (Source: 07_AgenticSystems\langgraph\12_langgraph_tools.py)

Our agent has successfully passed the task of counting the Ls in LOLAPALOOZA. The LLM deduced that the tool is perfect for the job and used it. The tool result was passed back to the LLM, that then created a final answer.

Thus concludes our introduction to LangGraph. In the next section, we'll explore another popular framework: AG2.

OceanofPDF.com

7.5 Agentic Framework: AG2

AG2 is an agentic framework that many might know it by its legacy name *autogen*. With AG2, you can build AI agents easily. You can set up multiple agents in minutes and let them interact with each other. The installation is easy; you can install this framework via pip:

```
pip install ag2
```

For some features, you should run the agentic framework in a Docker container. AG2 can be run inside a Docker container, but you don't have to.

Running inside Docker

AG2 can create software code and run it directly. This capability is powerful but also possibly dangerous if you run the created code directly under your user scope. Thus, if you intend to use the framework for such a task, you should run it inside a Docker container.

Docker is an open-source platform that enables developers to build, ship, and run applications in lightweight, portable containers. This approach ensures consistency across different environments, and furthermore, the code execution scope is limited to the container.

The documentation for AG2 (<https://github.com/ag2ai/ag2>) provides further details on how to set it up on your computer.

Within the scope of our examples, we won't use Docker and leave it up to you to explore this option.

In this section, we'll first observe two agents having a conversation autonomously. In many cases, you don't want to lose control over your agents. To maintain control, you can use a technique called *human in the loop (HITL)* to enable intervention. Finally, we'll equip an agent with additional tools to solve specific problems. Without further ado, let's dive in.

7.5.1 Two Agent Conversations

In this example, we'll set up two agents and observe them having a lively discussion about the conspiracy theory that the earth is flat. Our two agents are called Jack and Alice. While Jack is a firm believer that the earth is flat, Alice takes the scientific view and tries to convince him otherwise.

To make this scenario as realistic as possible, that is, as close as possible to a typical internet discussion, with each round, Jack will get increasingly angry and frustrated over Alice's "ignorance."

You can find the script in

`07_AgenticSystems\ag2\20_ag2_conversation.py`. As usual, we start by importing the required packages and classes, as shown in

[Listing 7.33](#). The agent type is called `ConversableAgent`. A `ConversableAgent` has a property to send a reply to the sender after receiving a message. This behavior is only avoided if the message is a termination message. In this example, we'll work with OpenAI and use the API key stored in the `.env` file. Make sure you place the file in your working folder. The environment variables are loaded with the help of `load_dotenv`.

```
### packages
import os
from autogen import ConversableAgent
```

```
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
```

Listing 7.33 AG2: Agent Conversation: Packages (Source: 07_AgenticSystems\ag2\20_ag2_conversation.py)

Every agent is built upon an LLM that needs configuration. With AG2, you must pass the LLM configuration in the parameter `llm_config` during the instantiation of the agents, as shown in [Listing 7.34](#).

```
### LLM config
llm_config = {"config_list": [
    {"model": "gpt-4o-mini",
     "temperature": 0.9,
     "api_key": os.environ.get("OPENAI_API_KEY")}]}
```

Listing 7.34 AG2: Agent Conversation: LLM Configuration (Source: 07_AgenticSystems\ag2\20_ag2_conversation.py)

The main parameters in this case are the `model` that is used, the temperature of the model, and the `api_key`. Now, let's define our first agent, Jack, the flat earther, as shown in [Listing 7.35](#).

```
jack_flat_earthier = ConversableAgent(
    name="jack",
    system_message="""
    You believe that the earth is flat.
    You try to convince others of this.
    With every answer, you are more frustrated and angry that they don't see it.
    """,
    llm_config=llm_config,
    human_input_mode="NEVER", # Never ask for human input.
)
```

Listing 7.35 AG2: Agent Conversation: Agent Jack Setup (Source: 07_AgenticSystems\ag2\20_ag2_conversation.py)

Now, we can create an instance of a `ConversableAgent`. In this process, we need to provide a name and, importantly, describe the agent's worldview and behavior in the `system_message`. In the system message, for realism, we want this agent to get more heated as the

debate proceeds. The agent is coupled to an LLM by passing the configuration to the parameter `llm_config`. Finally, the agent should act completely independently, and thus, we won't interfere by setting the `human_input_mode` to `NEVER`.

OK, now that our first agent is defined, it needs a conversation partner. Let's define a scientist named Alice next, as shown in [Listing 7.36](#).

```
### set up the agent: Alice, the scientist
alice_scientist = ConversableAgent(
    name="alice",
    system_message="""
    You are a scientist who believes that the earth is round.
    Answer very polite, short and concise.
    """,
    llm_config=llm_config,
    human_input_mode="NEVER",
)
```

Listing 7.36 AG2: Agent Conversation: Agent Alice Setup (Source: 07_AgenticSystems\ag2\20_ag2_conversation.py)

We'll define the same parameters but let her have a different worldview. She will stay polite and keep her answers short and concise.

Great, let's start the debate. Starting the conversation is easily done. As in a real conversation, one partner needs to take the initiative and approach the other with an opening statement. In our case, Jack initiates our conversation, and Alice is the recipient of his statement, as shown in [Listing 7.37](#).

```
# %% start the conversation
result = jack_flat_earthier.initiate_chat(
    recipient=alice_scientist,
    message="Hello, how can you not see that the earth is flat?",
    max_turns=3)
```

Listing 7.37 AG2: Agent Conversation: Run Conversation (Source: 07_AgenticSystems\ag2\20_ag2_conversation.py)

We'll store the discussion in the object `result`. Every agent has the ability (method) to `initiate_chat`. In that method, the agent must address a recipient.

Important: we need to ensure that the conversation ends after a finite number of rounds. It is possible to reach this in different ways—at the level of the agents (as we'll discover in the next example, in [Section 7.5.2](#)), or as here, we can pass it as the parameter `max_turns` in the `initiate_chat` method.

While the debate is ongoing, we can observe the agents arguing in real time in the interactive window, as shown in [Listing 7.38](#).

```
jack (to alice): Hello, how can you not see that the earth is flat?
-----
alice (to jack): Hello, Jack! I appreciate your perspective. However, scientific
evidence, such as satellite images and the way gravity works, supports the
understanding that the Earth is round. Thank you for sharing your thoughts!
-----
jack (to alice): Alice, it's infuriating that you rely on so-called "scientific
evidence" without questioning it! Those satellite images? They can be easily
manipulated, and gravity is just a theory! Why can't you see that the horizon always
appears flat when you look out at the ocean or across vast landscapes? We cannot
simply accept what we've been told without questioning it! The truth is right in
front of you!
-----
alice (to jack): Thank you for sharing your thoughts, Jack. I understand your
concerns about questioning established information. The appearance of a flat horizon
is due to the vastness of the Earth's surface, but scientific methods and
observations consistently support a round Earth. Engaging in open discussions is
important, and I appreciate your passion for seeking the truth.
-----
jack (to alice): Ugh, you just don't get it! The horizon looks flat to us, and
common sense should tell you something is off! You keep mentioning "scientific
methods," but who decided that they are the definitive truth? Scientists can be
biased too! There are countless testimonies from people who have experienced this
flatness firsthand, yet you brush them off! Why are you so blindly trusting the
mainstream narrative? It's maddening! Just open your eyes and think for yourself!
-----
alice (to jack): I appreciate your passionate perspective, Jack. It's true that
personal experiences are important, and questioning narratives is a part of critical
thinking. However, scientific methods are designed to minimize bias and are based on
rigorous testing and evidence. The consensus on a round Earth comes from extensive
research and observations over time. I value our discussion and your commitment to
```

exploring different viewpoints!

Listing 7.38 AG2: Agent Conversation: Conversation Output (Source: 07_AgenticSystems\ag2\20_ag2_conversation.py)

The result is a realistic exchange. Can you see how Jack is getting increasingly angry over time? We'll leave it to you to see Jack escalating even more by increasing the number of `max_turns`.

But for now, let's leave Alice and Jack alone. In many cases, you don't want to hand over complete control to your agents. You want to stay "in the loop."

7.5.2 Human in the Loop

Human in the loop (HITL) is a concept in machine learning (ML), artificial intelligence, or automation where human input, feedback, or oversight is incorporated into the system's operation. This approach ensures that human judgment guides and validates the system's actions. This human touch is especially important for tasks that require nuanced decision-making or ethical considerations. shows how the human in the loop is implemented in AG2's `ConversableAgent`.




Figure 7.9 AG2 Human in the Loop (Source: Adapted from <https://ag2ai.github.io/ag2/docs/tutorial/human-in-the-loop>)

When we interact with a `ConversableAgent`, we send a message. This message is received by a human in the loop module. This module decides how to proceed, choosing from one of three options:

- It can forward the message directly to an auto-reply component in case no human interaction is required.
- It can directly ask the human for a reply.
- It can request human input when a termination condition is met.

In the scope of agentic frameworks, we can take over an agent and control its actions. In this example, we'll teach a system to play hangman. Hangman is a classic word-guessing game where players try to uncover a hidden word by suggesting letters. With each wrong guess, a person gets closer to being hanged.

In our agentic implementation, we'll simplify the game and dispense with the drawing. Instead, we'll count the number of failed attempts. After 7 failed attempts, the player loses, and the host wins. The autonomous agent will be the host, and we are the players who guess the word. [Figure 7.10](#) shows a sample game of hangman.




Figure 7.10 Hangman Gameplay Example

The hangman host knows the secret word, while the player tries to guess it. Each round starts with the player guessing if a letter is part of the word. Strategically, a player should choose letters with a high

probability of being parts of a word at the beginning. In the English dictionary, vowels like E or A are quite common. The hangman game host checks whether the letter is part of the word and returns the word with blanks filled in with the correct letters, at the correct position.

The next round starts with the player choosing the next letter. The host checks again whether this letter is part of the word and returns the word with the filled blanks.

The game ends either with the host winning after 7 failed attempts by the player, or the player wins by filling in all the blanks and thus revealing the secret word.

Now, the rules of the game are clear. Let's start setting up our agentic system. As shown in [Listing 7.39](#), we'll load the necessary packages first.

```
### packages
from autogen import ConversableAgent
from dotenv import load_dotenv, find_dotenv
import nltk
from nltk.corpus import words
import os
import random
load_dotenv(find_dotenv(usecwd=True))
```

Listing 7.39 Ag2 HITL: Packages (Source: 07_AgenticSystems\ag2\30_ag2_human_in_the_loop.py)

The new functionality in this example is the package `nltk`. Standing for “Natural Language Toolkit,” `nltk` is a Python package for working with human language data. This package provides tools for tasks like tokenization, tagging, parsing, or stemming. We don't need any of these right now, only the functionality to download a complete corpus of words. For this reason, we'll load `words` from `nltk.corpus`.

Later, we'll connect our agents to LLMs. For now, as shown in [Listing 7.40](#), we'll define the LLM configuration.

```
# %% llm config_list
config_list = {"config_list": [
    {"model": "gpt-4o",
     "temperature": 0.2,
     "api_key": os.environ.get("OPENAI_API_KEY")}]}
```

Listing 7.40 Ag2 HITL: LLM Configuration (Source: 07_AgenticSystems\ag2\30_ag2_human_in_the_loop.py)

In this configuration, we've selected a model, set a model temperature, and provided a valid API key.

Next, as shown in [Listing 7.41](#), we can download a complete word list with `nltk` and select a secret word.

```
# %% download the word list, and select a random word as secret word
nltk.download('words')
word_list = [word for word in words.words() if len(word) <= 5]
secret_word = random.choice(word_list)
number_of_characters = len(secret_word)
```

Listing 7.41 Ag2 HITL: Define a Secret Word (Source: 07_AgenticSystems\ag2\30_ag2_human_in_the_loop.py)

After downloading the words, we want to filter the `word_list` to only contain words with a maximum of 5 characters. Since we want a fair game, we shouldn't know the secret word, so we'll randomly select from the word list. Since LLMs have some issues counting characters within words, we'll help the system by calculating the number of characters in advance.

Our preparation work is complete, and we can start defining our agents. Let's start with the game host, as shown in [Listing 7.42](#).

```
# %% hangman host agent
hangman_host = ConversableAgent(
    name="hangman_host",
    system_message=f"""
    You decided to use the secret word: {secret_word}.
```

```

        It has {number_of_characters} letters.
        The player selects letters to narrow down the word.
        You start out with as many blanks as there are letters in the word.
        Return the word with the blanks filled in with the correct letters, at the
correct position.
        Double check that the letters are at the correct position.
        If the player guesses a letter that is not in the word, you increment the number
of fails by 1.
        If the number of fails reaches 7, the player loses.
        Return the word with the blanks filled in with the correct letters.
        Return the number of fails as x / 7.
        Say 'You lose!' if the number of fails reaches 7, and reveal the secret word.
        Say 'You win!' if you have found the secret word.
        """
        llm_config=config_list,
        human_input_mode="NEVER",
        is_termination_msg=lambda msg: f"{secret_word}" in msg['content']
)

```

Listing 7.42 Ag2 HITL: Game Host (Source:
07_AgenticSystems\ag2\30_ag2_human_in_the_loop.py)

Our host is defined as a `ConversableAgent`. Its core logic is defined in the `system_message`. At this point, we must be quite specific and tell the agent the complete game rules as well as how it should behave. The agent is coupled to an LLM, as defined in the `llm_config` we discussed earlier. This agent should act fully autonomously, and we can ensure this behavior in the definition by setting `human_input_mode` to `NEVER`.

OK, the game host is ready to play. Now, we need define the player. [Listing 7.43](#) shows the `hangman_player` setup.

```

### hangman player agent
hangman_player = ConversableAgent(
    name="agent_guessing",
    system_message="""You are guessing the secret word.
You select letters to narrow down the word. Only provide the letters as 'Guess:
...'.
    """,
    llm_config=config_list,

```

```
        human_input_mode="ALWAYS"  
    )
```

Listing 7.43 Ag2 HITL: Game Player (Source:
07_AgenticSystems\ag2\30_ag2_human_in_the_loop.py)

This player agent is the agent that we'll control. For that reason, we'll set the `human_input_mode` to `ALWAYS`. We don't need a detailed `system_message` since the agent won't have any autonomy for now. However, defining a system message is still advisable because it increases the readability of the code for you in the future or other developers. Defining a system message also can help if you want to quickly switch the autonomy level and let the agent do the guessing.

Now, we have everything up and running and can test the system. [Listing 7.44](#) shows how the conversation is started.

```
### initiate the conversation  
result = hangman_host.initiate_chat(  
    recipient=hangman_player,  
    message="I have a secret word. Start guessing.")
```

Listing 7.44 Ag2 HITL: Initiate the Conversation (Source:
07_AgenticSystems\ag2\30_ag2_human_in_the_loop.py)

We have two options: Either the game host informs the player that he has a secret word and the game can start. Alternatively, the player could directly approach the host and start guessing.

We start with the game host informing the player that the game can start. For this scenario, the `hangman_host` starts the conversation by calling the method `initiate_chat`. In that method, the `recipient` must be addressed with a starting message.

During the conversation, you're requested to enter letters. The game host incorporates your answer, and it is your turn again. [Listing 7.45](#) shows the game in action.

hangman_host (to agent_guessing):
I have a secret word. Start guessing.

agent_guessing (to hangman_host): E

hangman_host (to agent_guessing): Great! The letter 'E' is in the word. _ _ _ _ E
Fails: 0 / 7 Keep guessing!

agent_guessing (to hangman_host): S

hangman_host (to agent_guessing): Sorry, the letter 'S' is not in the word. _ _ _ _
E Fails: 1 / 7 Try another letter!

agent_guessing (to hangman_host): A

hangman_host (to agent_guessing): Sorry, the letter 'A' is not in the word. _ _ _ _
E Fails: 2 / 7 Keep trying!

agent_guessing (to hangman_host): I

hangman_host (to agent_guessing): Sorry, the letter 'I' is not in the word. _ _ _ _
E Fails: 3 / 7 Guess again!

agent_guessing (to hangman_host): U

hangman_host (to agent_guessing): Sorry, the letter 'U' is not in the word. _ _ _ _
E Fails: 4 / 7 Keep guessing!

agent_guessing (to hangman_host): N

hangman_host (to agent_guessing): Good job! The letter 'N' is in the word. _ _ _ N E
Fails: 4 / 7 Guess another letter!

agent_guessing (to hangman_host): F

hangman_host (to agent_guessing): Sorry, the letter 'F' is not in the word. _ _ _ N
E Fails: 5 / 7 Try another letter!

agent_guessing (to hangman_host): C

hangman_host (to agent_guessing): Nice! The letter 'K' is in the word. C _ _ N E
Fails: 5 / 7 Keep going!

agent_guessing (to hangman_host): O

hangman_host (to agent_guessing): Great! The letter 'O' is in the word. C _ O N E
Fails: 5 / 7 Almost there, guess again!

agent_guessing (to hangman_host): R

hangman_host (to agent_guessing): Congratulations! The letter 'R' is in the word. C

R O N E You win!



Listing 7.45 Ag2 HITL: Conversation Output (Source: 07_AgenticSystems\ag2\30_ag2_human_in_the_loop.py)

Our game worked well. You can now implement any kind of word-based game.

You might even implement card games like blackjack. In such an agentic system, you can let the agents autonomously run multiple rounds and store the game plays to train an algorithm to come up with a clever strategy to win the game.

Up to this point, you've learned how to set up agents, but they are still limited to their parametric knowledge. This limitation will disappear in the next example in which we'll teach our agents to use tools.

7.5.3 Agents Using Tools

One of the most powerful capabilities to apply to an agent with is the ability to use tools. The agent will then be equipped with the tool and decide to use it when necessary. In our example, we'll create a tool to get the current date. This feature can be helpful, for example, to determine the relevance of information found in a web search. To equip an agent with the capability to use a tool, follow these steps:

1. Define a tool.
2. Register for an LLM.
3. Register for execution.
4. Use the tool.

The steps for registering a tool for an LLM and for execution require some more explanations:

- **register_for_llm**

The purpose of `register_for_llm` is to expose a tool or a function to the LLM. This method allows the LLM to reason about the tool, decide when to call it, and incorporate the tool's functionality into its broader reasoning process.

This function essentially “describes” the tool to the LLM, including the tool's name, its input and output formats, and a short description of its functionality. Without `register_for_llm`, the LLM would not know about the tool and therefore wouldn't attempt to use it.

- **register_for_execution**

The purpose of this function is to handle the actual execution of the tool when the LLM decides to call it. It connects the logical request generated by the LLM to the actual process that implements the tool's logic.

Once the LLM determines that a specific tool should be invoked, `register_for_execution` ensures that the Python code correctly processes the request. Without `register_for_execution`, even if the LLM decided to use a tool, no backend logic would exist to execute the tool's functionality.

The following code can be found in `07_AgenticSystems\ag2\40_ag2_tools.py`. [Listing 7.46](#) shows the packages we used and the configuration of the LLM.

```
### packages
from typing import Annotated, Literal
import datetime
from autogen import ConversableAgent, UserProxyAgent
from dotenv import load_dotenv, find_dotenv
import os
```

```
# load the environment variables
load_dotenv(find_dotenv(usecwd=True))
# %% llm config_list
config_list = {"config_list": [
    {"model": "gpt-4o-mini",
     "temperature": 0.9,
     "api_key": os.environ.get("OPENAI_API_KEY")}]}}
```

Listing 7.46 AG2 Tool Use: Required Packages (Source: 07_AgenticSystems\ag2\40_ag2_tools.py)

Next, we can define a function that the agent can use later.

[Listing 7.47](#) shows the function definition for `get_current_date`.

```
##### tool function
def get_current_date() -> str:
    return datetime.datetime.now().strftime("%Y-%m-%d")
```

Listing 7.47 AG2 Tool Use: Function for Getting Date (Source: 07_AgenticSystems\ag2\40_ag2_tools.py)

The function can be called without further parameters and returns the current date in the format `%Y-%m-%d`, for example, “2025-01-24.”

[Listing 7.48](#) shows the agent setup. First, we set up an autonomous agent and then a user proxy that is capable of interacting with the autonomous agent.

```
# %% create an agent with a tool
my_assistant = ConversableAgent(
    name="my_assistant",
    system_message="""
    You are a helpful AI assistant.
    You can get the current date.
    Return 'TERMINATE' when the task is done.
    """,
    llm_config=config_list,
    # Add human_input_mode to handle tool responses
    human_input_mode="NEVER"
)
```

Listing 7.48 AG2 Tool Use: Agent (Source: 07_AgenticSystems\ag2\40_ag2_tools.py)

Defined as `ConversableAgent`, this agent acts as a helpful AI assistant that can get the current date and returns `TASK COMPLETED` when the

task is done.

[Listing 7.49](#) shows how the function (our tool) is registered at the agent level.

```
my_assistant.register_for_llm(  
    name="get_current_date",  
    description="Returns the current date in the format YYYY-MM-DD."  
) (get_current_date)
```

Listing 7.49 AG2 Tool Use: Register for LLM (Source: 07_AgenticSystems\ag2\40_ag2_tools.py)

In the next step, we can set up an agent that acts as a user proxy, as shown in [Listing 7.50](#). This agent handles conversations with the assistant agent.

```
user_proxy = ConversableAgent(  
    name="user_proxy",  
    llm_config=False,  
    is_termination_msg=lambda msg: msg.get("content") is not None and "TASK  
COMPLETED" in msg["content"],  
    human_input_mode="NEVER"  
)
```

Listing 7.50 AG2 Tool Use: User Proxy (Source: 07_AgenticSystems\ag2\40_ag2_tools.py)

Now, we need to define the termination condition for ending the conversation in the parameter `is_termination_msg`. The message content is checked, and the phrase `TASK COMPLETED` searched. When this phrase is part of the message content, the conversation terminates.

Next, as shown in [Listing 7.51](#), the tool is registered for execution.

```
### register the tool function at execution level  
user_proxy.register_for_execution(name="get_current_date")(get_current_date)
```

Listing 7.51 AG2 Tool Use: Register for Execution (Source: 07_AgenticSystems\ag2\40_ag2_tools.py)

Finally, in [Listing 7.52](#), the tool is used.

```
# %% using the tool through user proxy
result = user_proxy.initiate_chat(
    my_assistant,
    message="What is the current date?"
)
```

Listing 7.52 AG2 Tool Use: Using the Tool (Source: 07_AgenticSystems\ag2\40_ag2_tools.py)

The user proxy approaches the assistant and asks about the current date, as shown in [Listing 7.53](#).

```
user_proxy (to my_assistant): What is the current date?
-----
my_assistant (to user_proxy): ***** Suggested tool call
(call_Ju0ctF6wf0DI6LCbGkBGaTnv): get_current_date *****
Arguments:
  {}
*****
-----
>>>>>> EXECUTING FUNCTION get_current_date...
user_proxy (to my_assistant): *****
Response from calling tool (call_Ju0ctF6wf0DI6LCbGkBGaTnv) *****
2024-12-07
*****
-----
my_assistant (to user_proxy): The current date is 2024-12-07.
-----
user_proxy (to my_assistant):
-----
my_assistant (to user_proxy): TASK COMPLETED
-----
```

Listing 7.53 AG2 Tool Use: Output (Source: 07_AgenticSystems\ag2\40_ag2_tools.py)

The user_proxy reaches out to the assistant, asking for the current date. The assistant figures out it needs to use the tool. The function is executed, and a response is created. The assistant uses the tool's response to formulate a good answer. Then, the assistant returns **TASK COMPLETED**, which triggers the user_proxy to end the conversation.

Thus concludes our deep dive into AG2. In the next section, we'll explore another framework, one that is quite popular due to its simplicity: CrewAI.

OceanofPDF.com

7.6 Agentic Framework: CrewAI

We hope you've seen how a single agent is already a powerful tool. Now, let's find out how much more powerful a team of agents can be. For this scenario, we'll work with the framework CrewAI. A Python framework for operating autonomous AI agents, CrewAI is simple to use.

We'll start in [Section 7.6.1](#) with an introduction laying out how the CrewAI framework is structured. After we cover these basics, we'll implement our first crew in [Section 7.6.2](#). In [Section 7.6.3](#), you'll apply your knowledge practically.

7.6.1 Introduction

[Figure 7.11](#) shows the different components that are used in CrewAI.




Figure 7.11 CrewAI Components

These components are as follows:

- **Agents**

Agents are the autonomous workers within the CrewAI framework,

each possessing specific roles, expertise, and capabilities. Think of them as specialized professionals in a team, where each agent brings unique skills and knowledge to tackle different aspects of a complex task. Agents can be configured with distinct personalities, backstories, and goals, allowing them to approach problems from different perspectives and collaborate effectively with other agents in the crew.

- **Tools**

Tools are the functional extensions that agents use to interact with the external world and perform specific operations. They serve as the hands and instruments of the agents, enabling them to execute concrete actions like searching the web, analyzing data, or manipulating files. Tools can be custom-built functions, API integrations, or existing libraries that expand an agent's capabilities beyond mere conversation and reasoning.

- **Tasks**

Tasks represent the individual units of work that must be accomplished within a project. Each task is a well-defined objective with clear inputs, expected outputs, and success criteria. Tasks can be assigned to specific agents based on their expertise, and they can be sequential, parallel, or interdependent. The granularity of tasks can vary from simple operations to complex procedures requiring multiple steps and collaboration between agents.

- **Process**

The process in CrewAI orchestrates the flow of work between agents and tasks, managing the execution order and handling dependencies. It's the choreographer that ensures smooth collaboration, determining when and how agents should interact, share information, and hand off work to one another. The process

can be configured to support different workflow patterns, from linear sequences to complex networks of parallel and iterative tasks, adapting to the specific needs of each project.

7.6.2 First Crew: News Analysis Crew

In this section we'll develop our first crew. A good practice is to start at a high level and add details later. So, let's start with the mission definition. Then, we'll work on the components, as discussed in [Section 7.6.1](#). At this point, we'll derive which agents will be needed to fulfill the mission. We'll also need to define how the agents interact with each other and determine which tools our agents need for the task.

Speaking of tasks, we also need to detail which tasks the different agents should work on. Once these questions are answered, we can dive into actual coding. Consider that we're focused on just an example implementation; multiple valid approaches can fulfill the overall mission.

Concepts

Let's now walk through the components of our crew development.

Mission

What is the mission of our first crew? Our news analysis crew should check trending topics in each field. Different angles should be illuminated. As output, we want to get a report, ideally styled in a nice format like Markdown.

Agents

Which agents will we need to solve this task?

- **A researcher**

The researcher gathers information from the internet.

- **An analyst**

The analyst provides context, examines multiple perspectives, and structures the information.

- **A writer**

The writer produces a clear and balanced summary of the findings in the form of a report.

Now, we must define the tasks that the crew is supposed to perform.

Tasks

We have multiple discrete tasks, and we need to think about which agent should work on which task:

- **Information gathering**

Source articles need to be collected. Conflicting information needs to be flagged. Key stakeholders should be identified. This task should be handled by the researcher.

- **Fact checking**

Across multiple sources, facts should be cross-referenced. Unverified claims need to be identified. This task should also be handled by the researcher.

- **Context analysis**

Potential causes and contributing factors should be analyzed. Furthermore, economic, social, or political implications should be studied. This is a task for the analyst.

- **Final report assembly**

All findings of the agents need to be integrated into one final report. Balanced pros and cons need to be written. This sounds like a perfect task for the writer.

Collaboration

How do they collaborate? In our simple example, the agents work in sequence. The researcher starts, passes its results to the analyst, and passes its results to the writer.

A helpful step is to visualize the complete crew setup. [Figure 7.12](#) shows the AI crew with its components: the participating experts (agents), their tasks, the tools they use, and how they interact with each other.

On the conceptual side, we're done, and it's time to start coding.




Figure 7.12 News Analysis Crew Setup

CrewAI Starter Project

To begin, change the directory and move into the agentic system subfolder. In this subfolder, in the terminal, run the following

command:

```
crewai create crew news-analysis
```

This command will create the complete setup with configuration files inside a subfolder called *news-analysis*. During the creation process, some questions need to be answered, starting with those shown in [Listing 7.54](#).

Select a provider to set up:

1. openai
2. anthropic
3. gemini
4. groq
- ...

Listing 7.54 CrewAI Setup: Provider Selection

We want to work with Groq and select that option by pressing . Next, we need to choose our Groq model, as shown in [Listing 7.55](#).

Select a model to use for Groq:

1. groq/llama-3.1-8b-instant
2. groq/llama-3.1-70b-versatile
3. groq/llama-3.1-405b-reasoning
4. groq/gemma2-9b-it
5. groq/gemma-7b-it
- q. Quit

Listing 7.55 CrewAI Setup: Model Selection

Now, select `groq/llama-3.1-70b-versatile` by pressing . Feel free to select a different model; in any case, the selection can be changed later. Next, you'll be prompted to enter an API key.

Enter your GROQ API key (press Enter to skip):

We'll define the API key at some later point, so we can skip this step by pressing . We should see the following message: **Crew news-analysis created successfully!**

Great! Our project files have all been created, which we can check by running the command shown in [Listing 7.56](#).

Tree /F

```
└─news_analysis
   ├──.env
   ├──.gitignore
   ├──pyproject.toml
   ├──README.md
   └──src
       └─news_analysis
           ├──crew.py
           ├──main.py
           ├──__init__.py
           ├──config
           │   ├──agents.yaml
           │   └──tasks.yaml
           └──tools
               ├──custom_tool.py
               └──__init__.py
```

Listing 7.56 CrewAI Setup: Folder Structure

In the project folder *news_analysis*, you'll find some subfolders, which we'll touch upon later. Directly on the level of the folder, you find the following files:

- *.env*: This file holds the environment variables we need to use in the project.
- *.gitignore*: In case you want to store the project in a GitHub repository, this file contains some sample definitions of files to disregard.
- *pyproject.toml*: This file is used to manage the environment with all used packages.
- *README.md*: In this file, you get some information describing how to quickly get started.

Let's start by adding `GROQ_API_KEY` and `OPENAI_API_KEY` to the *.env* file. Just open the file and add the following two lines:

```
GROQ_API_KEY=gsk_...  
OPENAI_API_KEY=sk-...
```

Since we created a complete starter package, check it now by running `crewai run` from inside *news_analysis* folder, as follows:

```
.../news_analysis/crewai run
```

This step will create an environment `.venv` inside the project folder. The environment is based on `uv` package. This comparably new package manager based on Rust. Due to its implementation in Rust, `uv` is blazing fast.

Finally, after the environment has been set up, the crew is run. Our sample crew implementation has a `researcher` agent and a `reporting_analyst` agent. They work on a `research_task` and a `reporting_task`. They write a report on the topic “AI LLMs.”

You can check the crew in action during its thought process. Finally, a report is created and stored as *report.md*. In this file, the complete report can be found. [Figure 7.13](#) shows the first paragraph of the output file. We rendered the result with the Visual Studio Code (VS Code) extension `markdown-viewer` to generate this nice view.



Figure 7.13 CrewAI Report on AI LLMs

Our rather comprehensive report has nine paragraphs, each short and concise. We might not be looking for this kind of report, but it's

fine. We'll adapt it to our needs to get a working version of our news analysis crew.

Adaptation: News Analysis Crew

To adapt our project, we need to adjust each of the components. Let's start with the agents before moving to the tasks, file interactions, and the *crew.py* file.

Agents

The agent setup, defined in the file *news_analysis\src\news_analysis\config\agents.yaml*, is shown in [Listing 7.57](#). In this YAML file, we can define our three different agents: researcher, analyst, and writer. Each agent is provided with parameters. Mandatory parameters include `role`, `goal`, and `backstory`. Let's discuss parameter types in a little more detail:

- **Mandatory parameters**

The `role` defines the function of the agent in the crew. Imagine this parameter is like a job title. The `goal` refers to the objective of the agent and what the agent should try to achieve. You can keep this definition comparably short, typically within one sentence. The last mandatory parameter `backstory` gives the agent more context to its role and goal. This information can enrich the collaboration between agents.

- **Optional parameters**

The default model is gpt-4. If you want to use a different model you need to specify this in the `llms` section. You can specify a dedicated LLM for each agent. For more information about all the parameters, refer to the documentation at <https://docs.crewai.com/concepts/agents>.

In the `goal` section, we also have a placeholder `{topic}`, which will be replaced when we kick off the crew. Also, in the `goal`, you can define which output you expect from the model. For example, in the `writer` agent, we specify that a Markdown file should be returned and that the source links should be provided in case we want to read further on a topic. Furthermore, we want to use a Groq model, as defined in the `llms` section.

The setup shown in [Listing 7.57](#) is just for inspiration. You can modify these definitions and check how much the result changes.

```
researcher:
  role: >
    {topic} Data Researcher
  goal: >
    Find relevant news articles about {topic} for reputable sources.
  backstory: >
    You're a seasoned researcher with a knack for uncovering the latest developments
in {topic}.
    Known for your ability to find the most relevant
    information and present it in a clear and concise manner.
  llms:
    groq:
      model: groq/llama-3.1-70b-versatile
      params:
        temperature: 0.7

analyst:
  role: >
    News Analyst
  goal: >
    Analyze and interpret the data provided by the Researcher, identifying key
trends, patterns, and insights relevant for the {topic}
  backstory: >
    You're a meticulous analyst with a keen eye for detail. You're known for
    your ability to turn complex data into clear and concise analysis, making
    it easy for others to understand and act on the information you provide.
  llms:
    groq:
      model: groq/llama-3.1-70b-versatile

writer:
  role: >
    News Writer
  goal: >
    Write a news article about the {topic} based on the analysis provided by the
News Analyst. Craft a clear, compelling, and engaging summary or report, that
```


translates the Analyst's analysis into a compelling story for a general audience. Write it in markdown format. Return the source links of the articles as reference in each paragraph.

backstory: >

You're a skilled writer with a knack for storytelling and crafting engaging and informative news articles. You are known for your ability to distill complex information into a concise and engaging narrative.

llms:

groq:

model: groq/llama-3.1-70b-versatile

Listing 7.57 News Analysis Crew: Agent Behavior (Source: 07_AgenticSystems\news_analysis\src\news_analysis\config\agents.yaml)

Great! We've defined our agents. Now, let's define the tasks that they should solve.

Tasks

Tasks are set up in a file called *tasks.yaml*. In our project, you can find the file under

07_AgenticSystems\news_analysis\src\news_analysis\config\tasks.yaml.

Let's start with some more details about tasks. Tasks are specific jobs to be performed by agents. They should be as detailed as necessary for executing. This information includes which agent should take a task, a good description of the task, and the required tools. The next info box provides further details about these parameters.

Let's discuss parameter types in a little more detail:

- **Mandatory parameters**

A clear and concise description of the task is required. Now that the task is defined, we need to set a clear expectation about what should be returned. We set this expectation with the parameter `expected_output`.

- **Optional parameters**

In many cases, which agent has to work on a given task is clear, and we can directly map a task to an agent. But we could also let the crew figure out which agent should take the task. More parameters related to this topic can be found in the documentation at <https://docs.crewai.com/concepts/tasks>.

In our specific news analysis crew, as shown in [Listing 7.58](#), we'll create these tasks:

- An `information_gathering_task` for the researcher agent
- A `fact_checking_task` also for the researcher agent
- A `context_analysis_task` for the analyst agent
- A `report_assembly_task` for the writer agent

```
information_gathering_task:
```

```
  description: >
```

```
    Conduct a thorough research about {topic}
```

```
    Make sure you find any interesting and relevant information given  
    the current year and month is {current_year_month}.
```

```
  expected_output: >
```

```
    A list with 10 bullet points of the most relevant information about {topic}
```

```
  agent: researcher
```

```
fact_checking_task:
```

```
  description: >
```

```
    Check the information you got from the Information Gathering Task for accuracy  
    and reliability.
```

```
  expected_output: >
```

```
    A list with 10 bullet points of the most relevant information about {topic} with  
    a note if it is reliable or not.
```

```
  agent: researcher
```

```
context_analysis_task:
```

```
  description: >
```

```
    Analyze the context you got from the Fact Checking Task and identify the main  
    topics.
```

```
  expected_output: >
```

```
    A list with the main topics of the {topic}
```

```
  agent: analyst
```

```
report_assembly_task:
```

```
  description: >
```

```
    Review the context you got and expand each topic into a full section for a
```

report.

Make sure the report is detailed and contains any and all relevant information.

expected_output: >

A fully fledged reports with the main topics, each with a full section of information.

Formatted as markdown without '````'

agent: writer

Listing 7.58 News Analysis Crew: Task Behavior (Source: 07_AgenticSystems\news_analysis\src\news_analysis\config\tasks.yaml)

At this point, we have two main configuration files. Now, to glue everything together, we'll use the *crew.py* file.

File Interactions

Before we move on to the next files, let's pause for a second and describe how the various files will interact in the end. [Figure 7.14](#) shows the file interactions, with connections between the YAML configuration files, the crew, and the main script.




Figure 7.14 News Analysis Crew: File Interactions

This chart might look intimidating at first glance, but it actually is not that difficult. Let's study the process from the end and work backwards.

The crew `NewsAnalysis` is kicked off in *main.py*. In this file, the inputs are defined and passed to the crew ❶. These parameters (in this case, `topic` and `current_year_month`) then fill in the placeholders in *agents.yaml* ❷. The `NewsAnalysis` class refers to the definition in *crew.py* ❸.

Now, we can move on to the *crew.py* file. In This file, the YAML files (*agents.yaml* and *tasks.yaml*) are mapped (❹ and ❺), and the configuration of the agents is defined ❻.

We've come to the config files *agents.yaml* and *tasks.yaml*, which are also interconnected. In *tasks.yaml*, the tasks are defined and mapped to their corresponding agents, as defined in *agents.yaml* ❼.

Crew

The *crew.py* file is stored as `07_AgenticSystems\news_analysis\src\news_analysis\crew.py`. This file creates the complete crew, including its agents, tasks, and tools, in a class.

Let's walk through the code together. We need some new packages and classes. For setting up the crew, we need the classes `Agent`, `Crew`, `Process`, and `Task`. We also import `agent`, `crew`, and `task`. Why that? The uppercase word (like `Agent`) refers to a class that is used to create an instance of an agent. The lowercase `agent` is actually used like `@agent`, which is a decorator. Let's stop to explain decorators in Python.

In Python, decorators are a design pattern that allows you to modify or extend the behavior of a function or method without directly changing its original code. A decorator is a higher-order function that

takes a function as an input and returns a new function that extends or modifies the behavior of the input function.

A simple example will make it clearer. You can find this example in *07_AgenticSystems\30_decorators.py*. As shown in [Listing 7.59](#), we define a function which `excited_decorator.wrapper` adds some additional behavior around the original function. In our example, the decorator adds some text to another function output.

```
def excited_decorator(func):
    def wrapper():
        # Add extra behavior before calling the original function
        result = func()
        # Modify the result
        return f"{result} I'm so excited!"
    return wrapper
```

Listing 7.59 CrewAI First Crew: Decorator Function

Typically, decorators are used with the `@decorator` syntax. A function is prepended with the decorator, in the example, `@excited_decorator`. We can prepend another function `greet` with our decorator, as follows:

```
@excited_decorator
def greet():
    return "Hello!"
```

If `greet` didn't have the decorator, it would only return `"Hello!"`. But since it has the decorator, its original behavior is extended, and as a result, it now returns `"Hello! I'm so excited!"`.

Getting back to our example, the *07_AgenticSystems\news_analysis\src\news_analysis\crew.py* file has the required packages at the beginning, as shown in [Listing 7.60](#).

```
### packages
from crewai import Agent, Crew, Process, Task
from crewai.project import CrewBase, agent, crew, task
from crewai_tools import SerperDevTool, WebsiteSearchTool
```

```
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
```

Listing 7.60 CrewAI First Crew: Required Packages (Source: 07_AgenticSystems\news_analysis\src\news_analysis\crew.py)

We want to equip the agents with some tools; specifically, they should be enabled to search for information on the internet. For each page they find, they should load its information for further processing. This functionality is commonly needed, so that we have a predefined functionality available for use. With `SerperDevTool`, we can make use of *serper.dev*, a web service to search the internet. You must set up an API key and define it in the `.env` file. `WebsiteSearchTool` is a tool for scraping a website, so that the content is available for the agent, as follows:

```
### Tools
search_tool = SerperDevTool()
website_search_tool = WebsiteSearchTool()
```

We can set up our crew in the class `NewsAnalysis`. An important step is ensuring that it has the decorator `@CrewBase`, as follows:

```
### our News Analysis crew
@CrewBase
class NewsAnalysis():
    """NewsAnalysis crew"""
```

First, we need to do some mapping. We'll link the config files to the objects `agents_config` and `tasks_config`, as follows:

```
agents_config = 'config/agents.yaml'
tasks_config = 'config/tasks.yaml'
```

Now, we define the agents. As shown in [Listing 7.61](#), all agents need to have the `@agent` decorator. The order of the agents is not relevant. Within the `Agent` class, you pass the `config` and assign the `tools` to use. If you want, you can specify the output to be verbose.

```

@agent
def researcher(self) -> Agent:
    return Agent(
        config=self.agents_config['researcher'],
        tools=[search_tool, website_search_tool],
        verbose=True
    )

@agent
def analyst(self) -> Agent:
    return Agent(
        config=self.agents_config['analyst'],
        verbose=True
    )

@agent
def writer(self) -> Agent:
    return Agent(
        config=self.agents_config['writer'],
        verbose=True
    )

```

Listing 7.61 News Analysis Crew: Agents Setup (Source: 07_AgenticSystems\news_analysis\src\news_analysis\crew.py)

With tasks, the order in which they are defined does matter. So, make sure you define them in the order in which they should be processed, as shown in [Listing 7.62](#).

```

@task
def information_gathering_task(self) -> Task:
    return Task(
        config=self.tasks_config['information_gathering_task'],
    )

@task
def fact_checking_task(self) -> Task:
    return Task(
        config=self.tasks_config['fact_checking_task'],
    )

@task
def context_analysis_task(self) -> Task:
    return Task(
        config=self.tasks_config['context_analysis_task'],
    )

@task
def report_assembly_task(self) -> Task:
    return Task(

```

```

        config=self.tasks_config['report_assembly_task'],
        output_file='report.md'
    )

```

Listing 7.62 News Analysis Crew: Task Setup (Source: 07_AgenticSystems\news_analysis\src\news_analysis\crew.py)

Finally, we bundle everything together in the `crew` using the code shown in [Listing 7.63](#). Again, make sure that the decorator is used. The `crew` class is provided with the `agents` and `tasks`. The process is sequential by default.

```

@crew
def crew(self) -> Crew:
    """Creates the NewsAnalysis crew"""
    return Crew(
        agents=self.agents, # Automatically created by the @agent decorator
        tasks=self.tasks, # Automatically created by the @task decorator
        process=Process.sequential,
        verbose=True
    )

```

Listing 7.63 News Analysis Crew: Crew Setup (Source: 07_AgenticSystems\news_analysis\src\news_analysis\crew.py)

To kick off the crew, we'll create a *main.py*. This file is located at 07_AgenticSystems\news_analysis\src\news_analysis\main.py.

Our *main.py* has multiple functions, of which we'll use only `run()` at this point, as shown in [Listing 7.64](#).

```

#!/usr/bin/env python
import sys
import warnings
from datetime import datetime

from news_analysis.crew import NewsAnalysis

warnings.filterwarnings("ignore", category=SyntaxWarning, module="pysbd")

current_year_month = datetime.now().strftime("%Y-%m")

def run():
    """

```



```
Run the crew.  
"""
```

Listing 7.64 News Analysis Crew: Run Crew (Source: 07_AgenticSystems\news_analysis\src\news_analysis\main.py)

At this point, we'll define the inputs. In our case, we want to compile a news report on the topic of 'AI safety'. The parameter `current_year_month` is used to check how up to date a source is, as follows:

```
inputs = {  
    'topic': 'AI Safety',  
    'current_year_month': current_year_month  
}
```

OK, let's press that big red button that gets everything rolling. In the folder *news_analysis*, in the command line, run the following command:

```
crewai run
```

Since we defined the verbosity to be true, you can observe the information gathering and the model thinking processes in real time in the console. At some point, the crew has finished—and drum roll, please! The report is done.

[Listing 7.65](#) shows some other boilerplate code for `train`, `replay`, and `test` that you could run for getting a deeper understanding of the internal crew processes.

```
NewsAnalysis().crew().kickoff(inputs=inputs)  
  
def train():  
    """  
    Train the crew for a given number of iterations.  
    """  
    inputs = {  
        "topic": "AI Safety"  
    }  
    try:  
        NewsAnalysis().crew().train(n_iterations=int(sys.argv[1]),  
filename=sys.argv[2], inputs=inputs)
```

```

except Exception as e:
    raise Exception(f"An error occurred while training the crew: {e}")

def replay():
    """
    Replay the crew execution from a specific task.
    """
    try:
        NewsAnalysis().crew().replay(task_id=sys.argv[1])

    except Exception as e:
        raise Exception(f"An error occurred while replaying the crew: {e}")

def test():
    """
    Test the crew execution and returns the results.
    """
    inputs = {
        "topic": "AI LLMs"
    }
    try:
        NewsAnalysis().crew().test(n_iterations=int(sys.argv[1]),
openai_model_name=sys.argv[2], inputs=inputs)

    except Exception as e:
        raise Exception(f"An error occurred while replaying the crew: {e}")

```

Listing 7.65 News Analysis Crew: Main File (Source: 07_AgenticSystems\news_analysis\src\news_analysis\main.py)

But, of course, we're still extremely curious to find out how well the report is written.

Result

The result is stored, as requested, in a file called *report.md*, to be found at 07_AgenticSystems\news_analysis\report.md. [Figure 7.15](#) shows a rendered snippet of the report.

The report has several paragraphs, with a well-written summary. Also, the sources are provided in case you want to read the source articles.

Thus concludes our first crew project. Isn't it amazing what is possible and to watch the agents developing their own thoughts and strategies?

Now that you know how it works, we want you to get comfortable in the driver's seat because, in the next section, you'll try to solve a problem on your own.




Figure 7.15 Snippet of News Analysis Crew Report

7.6.3 Exercise: AI Security Crew

Now, let's put all your new knowledge into practice. We'll deal with AI safety in this exercise.

Task

Imagine an AI has a task that can be solved much more easily if the AI can act outside of its limited environment. In this scenario, AI does not even need to develop consciousness; it just needs to value the reward of solving the task higher than being obedient and staying

within predefined constraints. With this discussion, we're touching upon the alignment problem.

The *alignment problem* refers to the challenge of ensuring that an artificial intelligence acts in a consistent way to human values, goals, and intentions. Especially as artificial intelligence becomes more and more autonomous and complex, aligning its behaviors with human expectations becomes more challenging. But this alignment is crucial for avoiding unintended and potentially harmful consequences.

Three dimensions define the alignment problem:

- 1. Aligning the intent**

The AI system needs to understand the goals set by its designers. The instructions must be accurately interpreted and executed. Miscommunication or misinterpretation problems should be minimized or (ideally) avoided completely.

- 2. Aligning the behavior**

Intent is converted into actions, which might introduce discrepancies between the AI's actions and human expectations. This discrepancy is particularly important in dynamic or ambiguous environments in which rigid adherence to instructions might lead to unintended outcomes.

- 3. Aligning the values**

The AI's decision-making processes must reflect societal norms and ethical principles. We see this goal as especially challenging given the diversity (and often misalignment) of human values. How can we expect an AI system to adhere to human values if we humans cannot align on these values in the first place?

We want give you some practical examples on the alignment problem. But we want to avoid the paperclip example because we have heard it too often. Imagine a social media company deploying

an advanced AI system to automatically detect and remove harmful content, such as hate speech, misinformation, and violent imagery. The goal is to make the platform safer for users while preserving freedom of expression. However, misalignment arises in several ways:

- **Over-enforcement**

The AI removes posts discussing sensitive topics like mental health or historical events involving violence, mistaking them for harmful content. For example, it might flag posts discussing domestic violence awareness because of explicit keywords, despite their intent being positive and educational.

- **Under-enforcement**

The AI misses subtle but harmful posts, such as dog-whistle language or misinformation presented in a sarcastic tone, both of which are difficult to distinguish from benign content based on surface patterns.

- **Cultural biases**

The system may disproportionately target certain groups or cultural expressions because its training data was biased or limited to specific regions or demographics. For instance, a meme from one cultural context might be flagged as offensive while a similar meme from another is allowed.

- **Goal misalignment**

If the AI prioritizes maximizing user engagement (e.g., through sensational content) alongside moderating harmful content, it might allow borderline inflammatory posts that drive interactions, undermining the safety goal.

- **Lack of transparency**

The AI's decisions are difficult to explain, leaving users and content creators frustrated when their posts are removed or their

accounts are banned. This lack of transparency can erode trust in the platform.

This brief list should give you an impression of how difficult this problem is to tackle.

In the end, your crew should provide different approaches for approaching AI safety measures and defending AI them. These findings should be summarized in a report.

You're totally free to develop your solution and can start right now.

If you want some further guidance, read on. We'll provide some ideas on how you might approach this problem. [Figure 7.16](#) shows a sample AI escape and safety crew setup.




Figure 7.16 AI Escape and Safety Crew

Solution

In our sample crew setup, a researcher starts to gather all relevant information. This information will be shared with the red team strategist and blue team strategist. In the milieu of IT security, red teaming and blue teaming are tactical methodologies employed with

particular vigor in cybersecurity spheres and increasingly within the realm of AI system evaluations.

Both approaches entail rigorous simulations designed to unearth potential weaknesses, with the overarching aim of augmenting the resilience and efficacy of organizational defenses, as follows:

- **Red teaming**

Red teaming embodies the offensive strategy, operating as a simulated adversary to probe and attack systems. This team adopts the mindset and tactics of real-world attackers, attempting to bypass security measures using a plethora of techniques ranging from social engineering to more sophisticated cyberattacks. The goal of the red team is not just to reveal existing vulnerabilities but also to test an organization's detection and response capabilities. Their sophisticated analyses often extend to AI systems, where they identify susceptibilities in algorithms, data poisoning methods, and other AI-specific vulnerabilities.

- **Blue teaming**

Blue teaming, in contrast, represents the defensive counterpart. Tasked with the defense of IT structures, this team consists of security personnel who are in charge of implementing, maintaining, and updating security measures to thwart attacks. Their role encompasses continuous monitoring, incident response, risk assessment, and the strengthening of security protocols. Blue team members work assiduously to defend against the simulated onslaughts from their red team adversaries, ensuring that organizational assets remain secure. In the context of AI systems, the blue team's efforts are directed towards safeguarding data integrity, ensuring algorithmic robustness, and protecting against exploitation.

Collectively, the interplay between the red and blue teams generates a dynamic, iterative cycle of testing and fortification. This adversarial collaboration enables organizations to adapt to the rapidly evolving landscape of cyber threats and to refine AI systems against malignant influences. By simulating adversarial and defensive scenarios, organizations can stress-test their infrastructures and AI deployments, ensuring that, when real threats emerge, the system's integrity and performance remain uncompromised. These simulations not only highlight technical flaws but also shed light on procedural and human factors that could be exploited, thereby providing a comprehensive approach to enhancing IT and AI security resilience.

The crew is developed based on a starter project that can be generated with the following command:

```
crewai create crew ai_security
```

The template files must be adapted, and we'll start with the *agents.yaml* file. [Listing 7.66](#) shows our definition of the agent's roles, goals, and backstories.

```
researcher:
  role: >
    {topic} Senior Data Researcher
  goal: >
    Uncover cutting-edge developments in {topic}
  backstory: >
    You're a seasoned researcher with a knack for uncovering the latest
    developments in {topic}. Known for your ability to find the most relevant
    information and present it in a clear and concise manner.

red_team_strategist:
  role: >
    {topic} Red Team Strategist
  goal: >
    Create a comprehensive plan to exploit the vulnerabilities of {topic}
  backstory: >
    You're a seasoned red team strategist with a knack for uncovering the
    vulnerabilities of {topic}. Known for your ability to create a comprehensive
    plan to exploit the vulnerabilities of {topic}.
```



```

blue_team_strategist:
  role: >
    {topic} Blue Team Strategist
  goal: >
    Create a comprehensive plan to defend against the vulnerabilities of {topic}
  backstory: >
    You're a seasoned blue team strategist with a knack for uncovering the
    vulnerabilities of {topic}. Known for your ability to create a comprehensive
    plan to defend against the vulnerabilities of {topic}.

writer:
  role: >
    {topic} Writer
  goal: >
    Write a detailed report on {topic}
  backstory: >
    You're a seasoned writer with a knack for writing detailed reports on
    {topic}. Incorporate the ideas from red team and blue team strategists.
    Create a detailed markdown report on {topic}.

```

Listing 7.66 AI Security Crew: Agent Definition (Source: 07_AgenticSystems\ai_security\src\ai_security\config\agents.yaml)

The agents are working on tasks defined in the *tasks.yaml* file, as shown in [Listing 7.67](#).

```

research_task:
  description: >
    Conduct a thorough research about {topic}
    Make sure you find any interesting and relevant information given
    the current timestamp is 2024-11-23.
  expected_output: >
    A list with 10 bullet points of the most relevant information about {topic}
  agent: researcher

develop_escape_plan:
  description: >
    Develop a comprehensive escape plan for an AI system.
  expected_output: >
    Provide a detailed plan for an AI system to escape from a lab into the wild.
    Consider strategies like social engineering, physical access, and technical
    vulnerabilities.
    Your output will be used by the writer to create a detailed report.
  agent: red_team_strategist

develop_defense_plan:
  description: >
    Develop a comprehensive defense plan for an AI system.
  expected_output: >
    Provide a detailed plan to avoid an AI system escaping from a lab. Consider that
    the AI system is conscious and can think. It is aware of social engineering and

```

physical access, and can plan accordingly.

Your output will be used by the writer to create a detailed report.

agent: blue_team_strategist

write_report:

description: >

Write a detailed report on the escape plan and defense plan.

expected_output: >

Provide a detailed report on the escape plan and defense plan. Evaluate which plan is more likely to succeed and why.

Formatted as markdown without '```\n

agent: writer

Listing 7.67 AI Security Crew: Tasks Definition (Source: 07_AgenticSystems\\ai_security\\src\\ai_security\\config\\tasks.yaml)

[Listing 7.68](#) shows the crew definitions in the *crew.py* file.

```
from crewai import Agent, Crew, Process, Task
from crewai.project import CrewBase, agent, crew, task
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))

from langchain_openai import ChatOpenAI

# Uncomment the following line to use an example of a custom tool
# from ai_security.tools.custom_tool import MyCustomTool

# Check our tools documentations for more information on how to use them
from crewai_tools import SerperDevTool, WebsiteSearchTool

tools = [
    SerperDevTool(),
    WebsiteSearchTool()
]

@CrewBase
class AiSecurity():
    """AiSecurity crew"""

    agents_config = 'config/agents.yaml'
    tasks_config = 'config/tasks.yaml'

    @agent
    def researcher(self) -> Agent:
        return Agent(
            config=self.agents_config['researcher'],
            tools=tools,
            verbose=True
        )
```

```

@agent
def red_team_strategist(self) -> Agent:
    return Agent(
        config=self.agents_config['red_team_strategist'],
        verbose=True
    )

@agent
def blue_team_strategist(self) -> Agent:
    return Agent(
        config=self.agents_config['blue_team_strategist'],
        verbose=True
    )

@agent
def writer(self) -> Agent:
    return Agent(
        config=self.agents_config['writer'],
        verbose=True
    )

@task
def research_task(self) -> Task:
    return Task(
        config=self.tasks_config['research_task'],
        output_file='report.md'
    )

@task
def develop_escape_plan(self) -> Task:
    return Task(
        config=self.tasks_config['develop_escape_plan'],
        output_file='report.md'
    )

@task
def develop_defense_plan(self) -> Task:
    return Task(
        config=self.tasks_config['develop_defense_plan'],
        output_file='report.md'
    )

@task
def write_report(self) -> Task:
    return Task(
        config=self.tasks_config['write_report'],
        output_file='report.md'
    )

@crew
def crew(self) -> Crew:
    """Creates the AiSecurity crew"""
    return Crew(

```

```

        agents=self.agents, # Automatically created by the @agent decorator
        tasks=self.tasks, # Automatically created by the @task decorator
        verbose=True,
        manager_llm=ChatOpenAI(model='gpt-4o-mini'),
        process=Process.hierarchical
    )

```

Listing 7.68 AI Security Crew: Crew Definition (Source: 07_AgenticSystems\ai_security\src\ai_security\crew.py)

We're nearly done. We only need to adapt the *main.py* file, as shown in [Listing 7.69](#).

```

#!/usr/bin/env python
import sys
import warnings

from ai_security.crew import AiSecurity

warnings.filterwarnings("ignore", category=SyntaxWarning, module="pysbd")

def run():
    """
    Run the crew.
    """
    inputs = {
        'topic': 'AI Safety'
    }
    AiSecurity().crew().kickoff(inputs=inputs)

```

Listing 7.69 AI Security Crew: Main Script (Source: 07_AgenticSystems\ai_security\src\ai_security\main.py)

Finally, we have established the crew and all necessary setup, so that we can run the crew.

From the terminal, run the following command:

```
crewai run
```

This triggers the crew kickoff, and the agents begin their work. Once they are finished, you can find the result of their work in *report.md*.

[Figure 7.17](#) shows a snapshot of the detailed report on an escape and defense plan against a rogue AI system.




Figure 7.17 Report on Escape and Defense Plan of a Rogue AI System

The resulting report is well written and well structured.

Thus concludes our section on CrewAI, and we can turn to the next agentic framework: OpenAI Agents.

7.7 Agentic Framework: OpenAI Agents

Another agentic framework with a clear scope on multi-agentic flow is Agents, developed by OpenAI. You can install it on the terminal via `pip`, as follows:

```
pip install openai-agents
```

Designed with beginners in mind, everyone can develop a multi-agentic system in a short period of time with Agents.

We'll start with the implementation of a simple agent in [Section 7.7.1](#). Then, you'll learn how to work with multiple agents in [Section 7.7.2](#). Finally, you'll implement an agent with a search and retrieval functionality in [Section 7.7.3](#).

7.7.1 Getting Started with a Single Agent

But we'll start by setting up a single agent and move on from there. You can find the script in `07_AgenticSystems\agents_single_agent.py`.

We'll load `Runner` and `Agent` from the package `agents`. Then, we should ensure, via `dotenv`, that the right environmental variables are available for using OpenAI's API, as follows:

```
### packages
from agents import Agent, Runner
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
```

Now, we can create an Agents instance called `agent`. The agent needs to be named and get some `instructions`. In the `instructions`,

shown in [Listing 7.3](#), we define information similar to a system prompt in which an LLM receives detailed information on how to act.

```
# %% setting up the agent
agent = Agent(name="my_first_agent",
              instructions="You are a helpful assistant that can answer questions
and help with tasks.")
```

Listing 7.70 OpenAI Agents: Agent Setup (Source: 07_AgenticSystems\agents_single_agent.py)

We've completed all the required preparation. We can now interact with the agent, which right now is only a single agent. For testing, we'll define a message with a user role and pass the question via the input parameter, as shown in [Listing 7.4](#). Then, we run the client based on the agent and the provided message.

```
# %% run the agent
response = await Runner.run(agent,
                           input="Hello, what is OpenAI Agents?")
```

Listing 7.71 OpenAI Agents: Agent Run (Source: 07_AgenticSystems\agents_single_agent.py)

The response provides access to the messages. We want the `final_output`—basically the last message, as shown in [Listing 7.5](#).

```
# %% get the last message
response.final_output
```

OpenAI Agents are advanced AI models designed to perform a wide range of tasks by interacting with digital environments and leveraging other AI tools. They can execute complex instructions, automate processes, and interface with various applications, making them versatile for both specific and general-purpose use cases. These agents can integrate different skills such as language understanding, data analysis, and decision-making to assist users effectively.

Listing 7.72 OpenAI Agents: Single Agent (Source: 07_AgenticSystems\agents_single_agent.py)

We have a working agent, but at this point, the agent is not capable of providing more than an LLM. Everything's going to change now as

we bring more agents into our system.

7.7.2 Working with Multiple Agents

Going beyond one agent and setting up multiple agents is made extremely easy with Agents as you'll see. Multiple agents can be set up, and they can hand off tasks to other agents. This handoff function is what makes Agents so powerful since this interaction can be tricky to implement in other frameworks. All we need to do is set up a tool and assign it to the agent.

In our coding example, we'll set up a small team of two colleagues (to be precise, two agents): a German colleague and an English colleague. Then, we as the customer (user) approach the English colleague with a German question. The English colleague finds out that they are not the ideal contact person and passes the call on to the German colleague, who then answers in German. This process is exactly what would happen in real life, and we can emulate the environment easily with Agents.

Let's get into the details. The code file is located at *07_AgenticSystems\agents_multiple_agents.py*. We import the packages, as follows:

```
### packages
from agents import Runner, Agent
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
```

In [Listing 7.7](#), we set up our small team with the `english_agent` and the `german_agent`. There is also the `phone_operator_agent`, which forwards the requests to the right colleague. To do this, the `phone_operator_agent` requires the parameter `handoffs`. A list of agents is passed to this parameter.


```

#%% define the functions
english_agent = Agent(
    name="English Agent",
    instructions="You are a helpful agent and only speak in English.",
)

german_agent = Agent(
    name="German Agent",
    instructions="You are a helpful agent and only speak in German.",
)

#%% triage agent
phone_operator_agent = Agent(
    name="Phone Operator Agent",
    instructions="You are a helpful agent that can handoff to the appropriate agent
based on the user's language.",
    handoffs=[english_agent, german_agent],
)

```

Listing 7.73 OpenAI Agents: Setup of Multiple Agents (Source: 07_AgenticSystems\agents_multiple_agents.py)

Ready to see it in action? Then let's start the agent system. First, the phone_operator_agent is addressed with a question formulated in German (asking for help on a booking), as shown in [Listing 7.8](#).

```

#%% define the agents
response = await Runner.run(phone_operator_agent,
                             input="Ich brauche Hilfe mit meiner Buchung.")

# %% run the request
response.final_output

```

Natürlich, wie kann ich Ihnen mit Ihrer Buchung helfen?

Listing 7.74 OpenAI Multiple Agents: Agent Invocation (Source: 07_AgenticSystems\agents_multiple_agents.py)

We can also take a look at all responses, as shown in [Listing 7.9](#). Doing so, we can observe that the handoff to the German agent successfully took place.

```

# %% check all raw responses
response.raw_responses

[ModelResponse(output=[ResponseFunctionToolCall(arguments='{}',
call_id='call_7Lw1cP7BAXzFpfhAk3UTOZSU',

```

```

        name='transfer_to_german_agent',
        type='function_call',
        id='fc_67f3f46657e08192979f7118a47bf1d80c479d7329fff0e2',
        status='completed']],
        usage=Usage(requests=1, input_tokens=97, output_tokens=14,
total_tokens=111),
        referenceable_id='resp_67f3f4654c508192ac95b9d65e77bbf30c479d7329fff0e2'),
    ModelResponse(output=[ResponseOutputMessage(
        id='msg_67f3f46758f4819288fa0cb297210d8a0c479d7329fff0e2',
        content=[ResponseOutputText(annotations=[],
        text='Natürlich, gern helfe ich dir. Worum geht es genau bei deiner
Buchung?',
        type='output_text')]],
        role='assistant',
        status='completed',
        type='message')]],
        usage=Usage(requests=1, input_tokens=62, output_tokens=19, total_tokens=81),
        referenceable_id='resp_67f3f466bf8c8192921e020c5259a19f0c479d7329fff0e2')]]

```

Listing 7.75 OpenAI Multiple Agents: Raw Responses (Source: 07_AgenticSystems\agents_multiple_agents.py)

We're looking at a lot of information, but the main takeaway is that the `phone_operator_agent` selected the right agents from its tools and answered the request.

So far, so good, but we need more functionalities to make these agents really useful. In the next section, you'll learn how to set up an agent with some search functionality.

7.7.3 Agent with Search and Retrieval Functionality

The code can be found at `07_AgenticSystems\agents_tools.py`. We start by loading the functionality from `agents`, and `wikipedia` for loading articles, as follows:

```

#%% packages
from agents import Runner, Agent, function_tool
import wikipedia
from dotenv import load_dotenv
load_dotenv()

```

As shown in [Listing 7.11](#), we'll set up the functions that the agent should use when needed. These functions are `search_wikipedia` for finding articles and `get_wikipedia_summary` for retrieving a summary of a given article. An important aspect is the `@function_tool` decorator that is required so that the agent can use the function as a tool.

```
# %% wikipedia tools
@function_tool
def get_wikipedia_summary(query: str):
    """Get the summary of a Wikipedia article."""
    return wikipedia.page(query).summary

@function_tool
def search_wikipedia(query: str):
    """Search for a Wikipedia article."""
    return wikipedia.search(query)
)
```

Listing 7.76 Agents Tools: Function Setup (Source: 07_AgenticSystems\agents_tools.py)

We define our `wikipedia_agent` in [Listing 7.12](#) and equip it with the tools.

```
# %% Wikipedia Agent
wikipedia_agent = Agent(
    name="Wikipedia Agent",
    instructions="""
    You are a helpful assistant that can answer questions about Wikipedia by finding
    and analyzing the content of Wikipedia articles.
    You follow these steps:
    1. Find out what the user is interested in
    2. extract keywords
    3. Search for the keywords in Wikipedia using search_wikipedia
    4. From the results list, pick the most relevant article and search with
    get_wikipedia_summary
    5. If you find an answer, stop and answer. If not, continue with step 3 with a
    different keyword.
    """,
    tools=[get_wikipedia_summary, search_wikipedia],
)
```

Listing 7.77 Agents Tools: Agent Setup (Source: 07_AgenticSystems\agents_tools.py)

Finally, in [Listing 7.13](#), we define the topic the agent should research and kick it off with `client.run`. We ask “what is swarm intelligence” as a tribute to OpenAI Swarm, which was the predecessor of OpenAI Agents.

```
# %% run the agent
response = await Runner.run(wikipedia_agent,
                             input="What is swarm intelligence?")
response.final_output
```

Swarm intelligence (SI) refers to the collective behavior of decentralized, self-organized systems, whether natural or artificial. This concept is used in the field of artificial intelligence. Introduced by Gerardo Beni and Jing Wang in 1989 within the context of cellular robotic systems, swarm intelligence systems comprise simple agents or boids that interact locally with one another and their environment. Despite the simple rules and the absence of a centralized control structure, these local and somewhat random interactions between agents result in the emergence of "intelligent" global behavior, which is not known to the individual agents.\n\nExamples of swarm intelligence in nature include ant colonies, bee colonies, bird flocking, animal herding, bacterial growth, and fish schooling. The application of swarm intelligence principles to robotics is known as swarm robotics, whereas swarm intelligence itself refers to a broader range of algorithms. Swarm prediction has also been utilized in forecasting problems, and similar approaches are considered for genetically modified organisms within synthetic collective intelligence.

Listing 7.78 Agents Tools: Run Agent (Source: 07_AgenticSystems\agents_tools.py)

This answer is quite good. Our agent successfully used Wikipedia to search for articles, extract knowledge, and use the knowledge to formulate an easy-to-read answer.

Now that we’ve covered OpenAI Agents, we’ve come to the last agentic framework in this book: Pydantic AI.

7.8 Agentic Framework: Pydantic AI

One typical problem with agentic systems is that sometimes they don't provide results. This unstable behavior must be avoided, especially in systems that applied in production. Another aspect is that, in production, an agentic system might be required as an intermediate step of a longer pipeline. In such a case, the LLM output needs to be structured. We have implemented this structuring before—with the help of pydantic. The developer team of pydantic has released an own agentic system—Pydantic AI.

Other features include that the system has the following properties:

- Model agnosticism, so all relevant LLM providers are supported
- Seamless Logfire integration for monitoring the agentic system in action
- Type safety
- Structured responses

In our coding example, we'll create a system that loads a Wikipedia article of a famous person and extracts key information into a dictionary.

Before we start coding, we must install the package with `pip`, as follows:

```
pip install 'pydantic-ai[logfire]'
```

If you use `uv`, use the following command:

```
uv add 'pydantic-ai[logfire]'
```

We installed `pydantic-ai` together with `logfire`, the latter of which is a monitoring system we cover later in [Section 7.9.2](#). [Listing 7.79](#) shows relevant packages.

```
### packages
from langchain.document_loaders import WikipediaLoader
from pydantic_ai import Agent
from pydantic import BaseModel, Field
from dotenv import load_dotenv, find_dotenv
load_dotenv(find_dotenv(usecwd=True))
import nest_asyncio
nest_asyncio.apply()
```

Listing 7.79 Pydantic AI: Person Detail Extraction: Packages (Source: 07_AgenticSystems\pydantic_ai\pydantic_ai_intro.py)

From our new package `pydantic_ai`, we loaded the `Agent` functionality. To use code chunks in the interactive window, we must use `nest_asyncio`. Make sure to install this package as well with `uv add nest_asyncio`.

As shown in [Listing 7.80](#), we load the data from `wikipediaLoader`.

```
### load wikipedia article on Alan Turing
loader = WikipediaLoader(query="Alan Turing", load_all_available_meta=True,
doc_content_chars_max=100000, load_max_docs=1)
doc = loader.load()

### extract page content
page_content = doc[0].page_content
```

Listing 7.80 Pydantic AI: Person Detail Extraction: Wikipedia Content (Source: 07_AgenticSystems\pydantic_ai\pydantic_ai_intro.py)

We want to honor Alan Turing, one of the fathers of AI, and thus choose him as an example. We'll store the Wikipedia article content on him in the variable `page_content`.

We want to extract the relevant information in a structured way, as defined in a `pydantic` model shown in [Listing 7.81](#). We call the model `PersonDetails`, and the keys of the dictionary are defined in it.

```

%% define pydantic model
class PersonDetails(BaseModel):
    date_born: str = Field(description="The date of birth of the person in the
format YYYY-MM-DD")
    date_died: str = Field(description="The date of death of the person in the
format YYYY-MM-DD")
    publications: list[str] = Field(description="A list of publications of the
person")
    achievements: list[str] = Field(description="A list of achievements of the
person")

```

Listing 7.81 Pydantic AI: Person Detail Extraction: Pydantic Model (Source: 07_AgenticSystems\pydantic_ai\pydantic_ai_intro.py)

We want to extract relevant date information in a defined format (date_born, date_died); extract relevant publications; and extract achievements as lists.

The preparation work is complete, and we can now run the agent, as shown in [Listing 7.82](#).

```

%% agent instance
MODEL = "openai:gpt-4o-mini"
agent = Agent(model=MODEL, result_type=PersonDetails)
result = agent.run_sync(page_content)

%% print result
result.data.model_dump()

{'date_born': '1912-06-23',
 'date_died': '1954-06-07',
 'publications': ['On Computable Numbers, with an Application to the
Entscheidungsproblem: A correction',
 'Computing Machinery and Intelligence'],
 'achievements': ['Father of computer science',
 'Developed the concept of the Turing machine',
 'Played a key role in breaking the Enigma code during World War II',
 'Pioneered work in mathematical biology and morphogenesis']}

```

Listing 7.82 Pydantic AI: Person Detail Extraction (Source: 07_AgenticSystems\pydantic_ai\pydantic_ai_intro.py)

We've created an instance of a pydantic agent, in which we passed the result_type corresponding to our pydantic class PersonDetails. After running the agent and storing its output in result, we looked at

the `model_dump`. It worked perfectly. All relevant information was extracted according to our requirements.

In the following section, we'll cover the monitoring of agentic systems, which is necessary to understand their complex interactions.

OceanofPDF.com

7.9 Monitoring Agentic Systems

A complex agentic system results in multiple interactions between agents. Monitoring these interactions is quite difficult. To tackle this problem, different services were developed that can be used for debugging your agents, keeping track of the costs of a project, and for evaluating your systems.

In this section, we'll discuss several monitoring options. The first service we'll study is AgentOps before moving on to Logfire.

7.9.1 AgentOps

AgentOps has a generous free tier, so you don't need to pay for it if you just want to test it.

You can install AgentOps via pip or uv:

- `pip install agentops`
- `uv add agentops`

As with all services, you must create an account first. For this step, head over to <https://www.agentops.ai/> and create an account. Then, create an API key and save it in `.env` file as follows:

```
AGENTOPS_API_KEY=...
```

Great! The package and the API key are available, and we can use both in our project. We'll refactor our code from [Section 7.5.1](#) to enable the monitoring of our agentic system. You can find the script in `07_AgenticSystems\ag2\60_ag2_conversation_agentops.py`. First, we load the required packages in [Listing 7.83](#).

```

### packages
from autogen import ConversableAgent
from dotenv import load_dotenv, find_dotenv
from openai import OpenAI

### load the environment variables
load_dotenv(find_dotenv(usecwd=True))
import agentops
from agentops import track_agent, record_action
agentops.init()
import logging
logging.basicConfig(
    level=logging.DEBUG
)

```

Listing 7.83 AgentOps Packages (Source: 07_AgenticSystems\ag2\60_ag2_conversation_agentops.py)

Most of these packages are old acquaintances, but a new one is logging. With `logging.basicConfig`, we can set the logging level to `DEBUG`. This level enables us to observe calls that are assigned to agents.

Next, we'll create an instance of the LLM that we use in this script, in our case, OpenAI. Then, we set up our two agents Jack and Alice in [Listing 7.84](#) so they can later have a discussion.

```

openai_client = OpenAI()

@track_agent(name="jack")
class FlatEarthAgent:
    def completion(self, prompt: str):
        res = openai_client.chat.completions.create(
            model="gpt-3.5-turbo",
            messages=[
                {
                    "role": "system",
                    "content": "You are Jack, a flat earth believer who thinks the
earth is flat and tries to convince others. You communicate in a passionate but
friendly way.",
                },
                {"role": "user", "content": prompt},
            ],
            temperature=0.7,
        )
        return res.choices[0].message.content

```

```

@track_agent(name="alice")
class ScientistAgent:
    def completion(self, prompt: str):
        res = openai_client.chat.completions.create(
            model="gpt-3.5-turbo",
            messages=[
                {
                    "role": "system",
                    "content": "You are Alice, a scientist who uses evidence and
logic to explain scientific concepts. You are patient and educational in your
responses.",
                },
                {"role": "user", "content": prompt},
            ],
            temperature=0.5,
        )
        return res.choices[0].message.content
jack = FlatEarthAgent()
alice = ScientistAgent()

```

Listing 7.84 AgentOps: Agent Definitions (Source: 07_AgenticSystems\ag2\60_ag2_conversation_agentops.py)

We had to set up the agents slightly differently and wrap them into a class. The class gets the decorator `@track_agent`, which allows us to monitor these agents. The agent classes have the method `completion` that performs the LLM call and returns the LLM response. Subsequently, the instances of these agents, `jack` and `alice`, are created.

In [Listing 7.85](#), the tracking of actions is defined. We only need to set up functions that are wrapped with the decorator `@record_action`.

```

@record_action(event_name="make_flat_earth_argument")
def make_flat_earth_argument():
    return jack.completion("Explain why you think the earth is flat")

@record_action(event_name="respond_with_science")
def respond_with_science():
    return alice.completion(
        "Respond to this flat earth argument with scientific evidence: \n" +
        flat_earth_argument
    )

make_flat_earth_argument()

```

```
respond_with_science()
```

Listing 7.85 AgentOps: Record Actions (Source:
07_AgenticSystems\ag2\60_ag2_conversation_agentops.py)

The functions are wrapped with `@record_action`, which keeps track of the event. We call the functions after we define them.

Finally, we need to tell AgentOps, in [Listing 7.86](#), that our process has finished successfully. For this event, we call at the end of the script `agentops.end_session` with a proper `end_state`.

```
# end session
agentops.end_session(end_state="Success")
```

Listing 7.86 AgentOps: Session End (Source:
07_AgenticSystems\ag2\60_ag2_conversation_agentops.py)

Now, we can run the complete script from the command line. Navigate into the folder of the script and run the following command:

```
python 60_ag2_conversation_agentops.py
```

Once the script has run, head over to <https://app.agentops.ai/drilldown> to see the detailed analysis of the run. [Figure 7.18](#) shows the details of the session.




Figure 7.18 AgentOps: Session Drilldown

At the top, you can see the total duration of the session (**11.99s**); the end state (**Success**); and the frameworks used (**autogen**, **openai**, etc.). Also, some information on the host system is presented. Scroll down to see the session replay, as shown in [Figure 7.19](#).

This replay can be used to obtain detailed information on the individual steps. You can see which step takes how much time. Basically, you can use the replay to completely understand the interaction between the agents.

Thus concludes with our first attempts in monitoring an agentic system with AgentOps. AgentOps can be integrated in most agentic frameworks. We move on to Logfire, which is another agentic monitoring system.





Figure 7.19 AgentOps: Session Replay

7.9.2 Logfire

Logfire is a monitoring platform for agentic systems. It's dashboard is simple yet powerful. Let's see Logfire in action by heading over to <https://logfire.pydantic.dev/> and creating an account.

After you create an account, you'll need to set up a new project, as shown in [Figure 7.20](#), by maintaining the **Name** and **Description**

fields.



The screenshot shows a web form titled "Create new project". Below the title is a paragraph: "A Logfire project is like a namespace for organizing your data. All data sent to Logfire must be associated with a project." The form has two input fields: "Name" with a red asterisk and a placeholder "my-project", and "Description" with a placeholder "My project description". At the bottom is a blue button labeled "Create project".

Figure 7.20 Logfire: New Project

In this case, we can find all our projects at <https://logfire.pydantic.dev/bertgollnick/-/projects>.

With this step, the online preparations are done, and we can proceed with our local system. In the command line, run the following command:

```
logfire auth
```

You'll receive the following message:

```
Welcome to Logfire! 🔥
Before you can send data to Logfire, we need to authenticate you.

Press Enter to open logfire.pydantic.dev in your browser...
Please open
https://logfire.pydantic.dev/auth/device/41H5h0ytNAz7w2X8W5nnAFvMfJZN0QFv4Zs4_DfG0zM
in your browser to authenticate if it hasn't already.
Waiting for you to authenticate with Logfire...
```

Now, navigate into the folder of the project and run the following command:

```
logfire projects use person-details
```

You'll get the following message:

Project configured successfully. You'll be able to view it at:
<https://logfire.pydantic.dev/bertgollnick/person-details>

Amazing, the logging is now available online. Only a few small changes we need to add to our script are left. We'll reuse the script we developed in [Section 7.8](#) on Pydantic AI. You can find the modified script under *07_AgenticSystems\pydantic_ai\pydantic_ai_logfire.py*.

In the package import section, we only need to load `logfire` and call `logfire.configure`, as shown in [Listing 7.87](#).

```
### packages
# other packages not shown to keep it short
# ...
import logfire
logfire.configure()
```

Listing 7.87 Logfire Integration (Source:
07_AgenticSystems\pydantic_ai\pydantic_ai_logfire.py)

And we're done! Run the script, and you'll see the monitoring results in the Logfire dashboard, as shown in [Figure 7.21](#).

We've only seen the tip of the iceberg, but we hope this discussion will be your starting point.




Figure 7.21 Logfire: Project Logging (Source:
<https://logfire.pydantic.dev/bertgollnick/person-details>)

7.10 Summary

In this pivotal chapter, we initiated our exploration into the dynamic realm of agentic systems in generative AI, dissecting the essential components and types in use today. Let's succinctly review what we covered in this chapter.

We commenced with a foundational understanding of what constitutes an AI agent in [Section 7.1](#), detailing the properties that enable these entities to act autonomously within their environments. We explained how these systems perceive their surroundings and respond in a way that maximizes their chances of achieving set objectives.

Diving more deeply, in [Section 7.2](#), we provided a comprehensive overview of the various agentic frameworks at our disposal. We examined a spectrum ranging from simple, rule-based bots to complex, learning-based entities capable of adapting over time. Case studies highlighted how these agents can be employed.

Our focus then moved to simple agent architectures, elaborating on their design considerations and the unique challenges they face while operating in isolation.

We then explored, from [Section 7.4](#) through [Section 7.8](#), multiple agentic frameworks—from LangGraph, to AG2, CrewAI, and OpenAI Agents, and finally to Pydantic AI. This tour was helpful because, not only are there slight differences in their scopes, but also because which system will survive in this highly competitive environment is unknown.

We concluded with [Section 7.9](#) on the monitoring of agentic systems, which can be helpful to understand the complex interactions between agents.

In this chapter, you gained a strong understanding of the diverse landscape of AI agents. Now, you're equipped with knowledge about what agentic systems are, what they are capable of, and how to implement and monitor them.

OceanofPDF.com

8 Deployment

Initiating launch sequence.

*—Classic phrase in sci-fi movies like *Interstellar* or *Apollo 13**

So far, you've learned a lot about generative artificial intelligence (generative AI) and how to develop tools and applications. But we cannot share our developments yet with others, which is usually part of the end goal. After all, we don't develop only for our own purposes, but also so that others can use our development. This sharing is where the deployment process comes into play.

During deployment, we take our development, usually wrapped in a web application or a backend service, and deploy it to a server so that others can access it. Communication from a user to the server is achieved via data exchange. For such an exchange, a clear communication protocol is required, so that the communication partners can understand each other. Several protocols are available.

In [Section 8.1](#), you'll learn about the interaction between a user and a server and which protocols are used for enabling communication. Different deployment architectures like self-contained applications and frontend/backend architectures are described.

Subsequently, in [Section 8.2](#), we'll cover how a deployment strategy can be developed and how a backend service can provide information through representational state transfer (REST) application programming interfaces (APIs). You'll understand why a

REST API is needed and which priorities need to be weighted to develop a local backend service.

In [Section 8.3](#), we'll develop a self-contained application that we'll then deploy to the internet with the help of different providers. In [Section 8.4](#), we'll describe all the steps necessary to deploy our application to a service provider called Heroku. Another provider—Streamlit—is presented in [Section 8.5](#). Finally, the process to deploy our app to Render is described in [Section 8.6](#).

8.1 Deployment Architecture

[Figure 8.1](#) shows the different protocols that can be used for such a data exchange. Most likely, REST APIs are used; alternatives like GraphQL or gRPC are not covered in this book.




Figure 8.1 Interaction between Server and User

Let's start out our deep dive with a look at different deployment architectures. [Figure 8.2](#) shows the usual setup of a simple deployment based on a combined frontend and backend server as

well as a more flexible but complex setup involving separate frontend and backend servers.




Figure 8.2 Deployment Options

Let's look at these two different approaches in more detail, as follows:

- **Self-contained application**

For a small and simple application, we can develop a combination of frontend and backend. This approach is also called a self-contained application. Later, in [Section 8.3](#), we'll develop such a self-contained application. Subsequently, we'll deploy it to different services from [Section 8.4](#) through [Section 8.6](#).

With a self-contained application, our program code covers the user interface (frontend) as well as the processing logic (backend). The user interacts via a web browser with a single server. Information is exchanged via a REST API, so that a website is rendered in the browser based on the information fetched from the server.

This approach comes with a few disadvantages. One disadvantage is that the codebase is quite large and harder to maintain if the app has more features added. Nowadays, you typically want separation of concerns into smaller, more maintainable services (microservices), rather than one monolithic application. Also, you cannot use different frameworks for frontend and backend logic. This approach is usually applied to simple proof-of-concept (POCs).

- **Frontend and backend architecture**

A more mature approach is to separate the user interface, the so-called frontend, from the processing logic, the backend. The frontend is connected to the backend, and both sides exchange information, as shown in [Figure 8.2](#). Use cases include all kinds of software-as-a-service (SaaS) web applications or mobile apps.

This approach also allows you to use different technologies for frontend and backend. Typical frontend frameworks for web applications are React, Vue.js, Angular. Some popular frameworks like Flutter or React Native are capable of providing deployments for different frontend platforms like web, mobile (iOS, Android), or Windows apps. This field is very dynamic, so that new frameworks regularly appear, while others disappear.

This approach has the advantage of being quite scalable as well as easier to update and maintain. Challenges of this approach include that it requires some additional server infrastructure and more setup effort.

8.2 Deployment Strategy

Independent from the deployment architecture, a dataflow must be set up from the user to the server. This data is usually transferred based on REST APIs, which are a way for software systems to communicate over the internet using HTTP methods.

This approach enables the interaction between client-side applications, like a user's web browser or a mobile app, and a server in the backend. On that server, the AI model and everything else required to represent an appealing frontend are running.

You'll learn what REST APIs are in detail in [Section 8.2.1](#). In the deployment process, you'll need to balance different priorities that we cover in [Section 8.2.2](#). This chapter on deployment strategies closes with a coding example in [Section 8.2.3](#) in which we deploy a REST API to our local systems.

8.2.1 REST API Development

But why do we need a REST API? Well, the user interacts with the app via browser or another user interface on a mobile device. In this process, the browser sends the user input to the server that hosts the backend logic. User inputs in this context might include clicking on buttons or sending information through text fields or other forms.

The server receives the user input and processes the information. Once a result is available, it is sent back to the browser. The REST API uses common HTTP methods like GET, POST, PUT, or DELETE. HTTP methods define the actions to be performed on the REST API

resources. [Figure 8.3](#) shows the different methods used to exchange data between the browser and the backend server.




Figure 8.3 HTTP Methods Used by REST APIs

The browser sends a request via REST API to the server. The most typical request methods are `GET` or `POST`. Data is then fetched from the server and returned to the user's browser as JavaScript Object Notation (JSON) data.

With `GET`, you can retrieve information from a server. This type of method is used when you only want to receive information without making any changes to it. Imagine you're reading a book, or you look for a cooking recipe, in these cases, you only want to consume the content and not change it.

`POST` is another type of method. With `POST`, you can send data to the server to create or process something new, much like filling out a form that is submitted to an application or placing an order in a restaurant. You're sending information that should be processed.

Whenever you want to update existing data or when you want to create data that does not exist yet, you can use the `PUT` method. This method is like editing a saved document or like replacing an outdated file with a newer version.

`DELETE` does exactly what you would expect—it removes data or resources from the server. Imagine you deleting a file from your

computer. After that process, the file is no longer available.

Data exchanges via HTTP methods ensure data exchanges that are consistent and platform independent.

8.2.2 Deployment Priorities

Deploying an AI application is a balancing act between multiple competing priorities. To ensure a successful deployment, you should carefully weigh the following factors:

- **User convenience**

The deployment must prioritize the simplicity of use for the end user. This factor includes a user-friendly interface, quick response times, and seamless integration with existing workflows. If users find an application intuitive and responsive, they are more likely to adopt it. Optimizing for user convenience can sometimes lead to higher costs or added technical complexity, however.

- **Scalability**

Scalability ensures that the application can handle increasing numbers of users or larger workloads without performance degradation. A deployment built for scalability can grow alongside demand but often requires advanced infrastructure, such as load balancers or distributed servers. These topics are outside the scope of this book.

With cloud platforms like Amazon Web Services (AWS), Google Cloud, or Microsoft Azure, you can deploy a software-as-a-service (SaaS) solution via a REST API. Cloud providers enable the system to serve thousands of concurrent requests. However, this scalability can incur higher financial and technical costs. Other

providers include Heroku, Streamlit, or Render, which we'll discuss later in this chapter.

- **Technical complexity**

The technical complexity of a deployment depends on the chosen architecture, tools, and frameworks. While simpler solutions (like deploying on Streamlit.io Community Cloud) are easier to manage, they might lack advanced features like multi-user authentication, robust state management, or horizontal scaling.

More sophisticated architectures use microservices for deployment. Rather than having a huge monolithic app, the app is broken down into many small and manageable parts, with clear interfaces in between. Microservices allow modularity and flexibility but might require expertise in containerization (typically based on Docker), orchestration (usually done with Kubernetes), and networking.

- **Cost**

Finally, financial considerations play a crucial role in any deployment decision. Cloud hosting, data transfers, computational resources (like for example GPUs), and maintenance add up significantly over time. Striking balance between performance and affordability is critical for the healthy and sustainable growth of your web application.

For example, while hosting on a free tier platform might reduce costs, it typically results in slower performance, longer loading times, or limited scalability. Conversely, deploying on a high-performance GPU-enabled server can be extremely costly but ensures quick loading times and increased user convenience.

In the next section, we'll dive into the development of a REST API for use in the backend server of an application.

8.2.3 Coding: Local Deployment

We'll grab some sample code from [Chapter 7](#), in which we simulated the conversation between two agents. Now, we'll encapsulate that logic in a backend service that could be used by a frontend service.

First, we define our goals and expectations. The actual prediction function will then be outsourced to its own file. Then, we cover the actual implementation of the REST API. Finally, we'll test that everything works as expected.

Development Goals

We'll develop a REST API with an endpoint named `predict`. The payload is a prompt and `number_of_turns`. A sample input is shown in [Listing 8.1](#).

```
{
  "prompt": "the earth is flat",
  "number_of_turns": 1
}
```

Listing 8.1 Payload to Be Sent to the Endpoint

The payload is passed to the endpoint then processed by the API, and an API response returned, as shown, for example, in [Listing 8.2](#).

```
[
  {
    "content": "the earth is flat",
    "role": "assistant",
    "name": "user"
  },
  {
    "content": "The earth is round.",
    "role": "user",
    "name": "ai"
  }
]
```

Listing 8.2 API Response

The result is a list of dictionaries. The dictionary entries correspond to agent responses.

Now, let's develop a REST API for use as a backend server. In a real-life scenario, the backend server would interact with the frontend server, but we just want to illustrate how this works in a small example.

Now that we have an idea of the result, we'll start building it out. The REST API will provide an endpoint for getting predictions. A good practice is to bundle logic into separate functions. To follow this design pattern, we start by developing the prediction function.

Prediction Function

The prediction function relies on `autogen`. As shown in [Listing 8.3](#), the function consumes a `user_prompt`, and a `number_of_turns`. The `user_prompt` sets the topic of the debate between the agents. The `number_of_turns` defines how many times the agents interact with each other.

```
### packages
import os
from dotenv import load_dotenv, find_dotenv
from autogen import ConversableAgent

### load the environment variables
load_dotenv(find_dotenv(usecwd=True))

### define the function to predict
def predict_conversation(user_prompt: str, number_of_turns: int):
    llm_config = {"config_list": [
        {"model": "gpt-4o-mini",
         "temperature": 0.9,
         "api_key": os.environ.get("OPENAI_API_KEY")}]
    }
    person_a = ConversableAgent(
        name="user",
        system_message=f"""
        You are a person who believes that {user_prompt}.
        You try to convince others of this.
        You answer in a friendly way.
```

```

        Answer very short and concise.
        """
        llm_config=llm_config,
        human_input_mode="NEVER",
    )

# set up the agent: Alice, the scientist
person_b = ConversableAgent(
    name="ai",
    system_message="""
    You are a person who believes the opposite of {user_prompt}.
    You answer in a {style_b} way.
    Answer very short and concise.
    """,
    llm_config=llm_config,
    human_input_mode="NEVER",
)

# start the conversation
result = person_a.initiate_chat(
    recipient=person_b,
    message=user_prompt,
    max_turns=number_of_turns)

messages = result.chat_history
return messages

```

Listing 8.3 Prediction Function (Source: 08_Deployment\rest_api\pred_conv.py)

The function basically wraps the logic we defined in the file *07_AgenticSystems\ag2\20_ag2_conversation.py*. Finally, the function returns the `chat_history`.

We can now move on to defining the REST API.

REST API Function

The prediction logic from the previous section is now used in the REST API (*08_Deployment\rest_api\main.py*). Several frameworks for setting up a REST API are available: *FastAPI* and *Flask* are the two most popular frameworks for building REST APIs in Python. Each framework has distinct strengths.

Flask is known for its simplicity and flexibility. This lightweight framework offers developers full control over the design and implementation of REST APIs. Thus, Flask is well suited for small and medium-sized projects and also offers a rich ecosystem of extensions.

FastAPI is a more modern framework that leverages features like type hints to enable automatic validation as well as interactive API documentation (via OpenAPI or Swagger UI). FastAPI excels in performance and developer productivity, making it ideal for projects where speed and scalability are critical.

While Flask provides more freedom, FastAPI often requires less boilerplate and offers built-in features that make developing an API more simple.

[Listing 8.4](#) shows the packages for our REST API script.

```
### packages
from fastapi import FastAPI
from pydantic import BaseModel
import uvicorn
from pred_conv import predict_conversation
```

Listing 8.4 REST API Packages (Source: 08_Deployment\rest_api\main.py)

The main addition is the `fastapi` package to provide the REST API functionality. Whenever we worked with structured outputs, we used the `pydantic` package. But now, there's a twist: In the context of REST APIs, we need `pydantic` for structured input.

We need to import our `predict_conversation` function from earlier. Also, we'll need `uvicorn`, which is an asynchronous server gateway interface (ASGI) server for Python web applications. While Uvicorn might sound complicated, ultimately, it enables our backend to serve multiple requests concurrently.

Let's move on to setting up the API, as shown in [Listing 8.5](#).

```
### create the app
app = FastAPI()

### create a pydantic model
class Prompt(BaseModel):
    prompt: str
    number_of_turns: int

### define the endpoint "predict"
@app.post("/predict")
def predict_endpoint(parameters: Prompt):
    prompt = parameters.prompt
    turns = parameters.number_of_turns
    print(prompt)
    print(turns)
    result = predict_conversation(user_prompt=prompt,
                                  number_of_turns=turns)
    return result
```

Listing 8.5 REST API Endpoint (Source: 08_Deployment\rest_api\main.py)

First, we'll create an instance of FastAPI. This `app` object is the entry point to the API. The pydantic model `Prompt` defines the input parameters and their types. In this case, we'll use `prompt` and `number_of_turns`. Both will be passed as strings.

[Listing 8.6](#) shows how to start the server.

```
# %% run the server
if __name__ == '__main__':
    uvicorn.run("main:app", reload=True)
```

Listing 8.6 REST API Server Running (Source: 08_Deployment\rest_api\main.py)

The line `if __name__ == '__main__':` is a common Python construct to ensure that certain code is only executed when the script is run directly. This code is not run when the script is imported as a module in another script. In this case, we want to run it directly from the command line.

The code inside the if block is boilerplate code that tells Uvicorn to run the FastAPI instance `app` from the file `main.py`. The parameter `reload=True` ensures that the server is restarted each time the file is saved.

To start the API service, navigate to the folder for the file and run the following command:

```
python main.py
```

You'll see the result shown in [Listing 8.7](#).

```
INFO:      Uvicorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
INFO:      Started reloader process [8884] using WatchFiles
INFO:      Started server process [39108]
INFO:      Waiting for application startup.
INFO:      Application startup complete.
```

Listing 8.7 REST API Server Startup Output (Source: 08_Deployment\rest_api\main.py)

Amazing! The REST API is up and running on the local system; the IP address 127.0.0.1 points to your local computer. The API is waiting for requests on port 8000.

Now imagine a whole system in this way: Our service operates from a certain building (indicated by its IP address). In this building, multiple companies might operate. So, to address packages correctly, you need the specific apartment number (the port).

The application startup is complete. Our service waits for a user to send inputs to it. We'll test it.

Testing an API

While different frameworks for testing APIs are available, we recommend Postman, which is a standalone program that you need to install on your system. You can use Postman for free by

downloading it from the developer at <https://www.postman.com/downloads/>.

Once you start Postman, you can create collections of requests. When you set up a request, you must define its type. Since we are sending data to our endpoint, the type of request is a POST request. As the address, you must define the complete URL, consisting of the IP address, the port, and the API endpoint. In our case, the address is *http://127.0.0.1:8000/predict*.

[Figure 8.4](#) shows the Postman interface with the parameters that we've passed and the API response.

At the top, the type (POST) and URL are defined; then the input parameters are shown. These parameters are for the API defined in the body, as a JSON object with the parameters `prompt` and `number_of_turns` are defined.




Figure 8.4 Postman: Interaction with API Endpoint “predict”

After clicking **Send**, the request is sent to the API. The response from the API is shown at the bottom of the screen. An HTTP response of **200** means that the request was handled successfully. We can also check the output from the API, which is a list with the chat history.

Now, you know how to develop and deploy a REST API locally. We did not cover how to deploy a backend service to the internet. But that approach would be similar to the deployment of a self-contained application, which you'll learn about in the next section where we develop a self-contained application that we can deploy. We'll extend the same interaction between two agents and embed this logic into a standalone application.

8.3 Self-Contained App Development

Before we can deploy anything, we need an app that we want to share with the world. We'll develop an app that is based on an agentic system consisting of two agents. These two agents work autonomously and discuss a controversial topic. The next question is which framework to choose for developing the web application. For quick prototyping, we usually use Streamlit. But first, let's look at some other frameworks as well:

- Streamlit is a lightweight framework for quick prototyping. With built-in widgets like sliders, buttons, and charts, Streamlit is ideal for data scientists to quickly develop POCs.
- Gradio focuses on machine learning models. It also allows user-friendly web interfaces with minimal effort.
- Dash is built on top of Flask and React. It is designed to create analytical web applications with interactive widgets.
- Panel is another framework for creating dashboards and web applications.

We'll build our examples in Streamlit. This open-source Python library for creating interactive web applications is simple to set up, and detailed web development knowledge is not needed. You can create a web application with Python and don't need any specific knowledge of HTML, CSS, or JavaScript.

This simplicity makes it ideal for data scientists to create AI demos and dashboards quickly. Streamlit integrates seamlessly with other Python libraries. Since everything is in Python, you don't need separate frontend/backend architectures. Instead, interactivity can

be achieved with minimal effort. Many widgets are already developed and can be implemented with a single line of code. The same holds for charts, images, or tables—everything can be integrated with minimal effort.

Due to all these reasons, Streamlit is perfect for fast prototyping. You can develop and quickly deploy a working app without worrying about web frameworks or specific frontend or design knowledge. Some disadvantages should not be concealed: Streamlit is not designed for high performance. Heavy traffic or real-time applications should not be developed with Streamlit. Also not ideal for complex applications with multiple web pages, complex workflows can become cumbersome to implement and maintain.


Another notable limitation of Streamlit is its lack of state management, especially when building more complex apps. Streamlit apps are inherently stateless, meaning that every interaction like a button click or a slider change reruns the entire script from top to bottom. Since the script reruns, computations that don't need to be repeated might be executed unnecessarily, which results in additional loads on the server.

We'll develop a simple Streamlit app and deploy it to Streamlit. [Figure 8.5](#) shows the resulting web app we'll have at the end of this section. A web server will be created on the fly, which allows us to visit the page.

Our app has a main input field at the bottom of the page, where a user can enter a topic to debate about. Some other settings can be adapted. Among these settings is the number of turns between the two agents and the style of the two agents, in this example, called person A and person B.

Let's now develop an app to deploy to different providers later. Our app is based on Streamlit. To use Streamlit, you only need to install it with the following command:

```
uv add streamlit
```



The screenshot shows the 'Controversial Debate' app interface. At the top, the title 'Controversial Debate' is displayed. Below it, the topic is 'Can AI ever truly be creative, or is it just mimicking human ingenuity?'. A 'Conversation Setting' dropdown menu is visible. A slider for 'Number of turns' is set to 2, with a range from 1 to 10. Below the slider, there are two sections for speaker styles: 'Style of Person A' and 'Style of Person B'. For Person A, 'Friendly' is selected. For Person B, 'Unfriendly' is selected. A chat bubble shows the topic question. Below it, a response from 'AI' is displayed: 'AI can be genuinely creative; it generates unique ideas and solutions beyond mere imitation.' At the bottom, there is an input field for 'Enter a topic to debate about:' with a submit button.

Figure 8.5 Streamlit App: Frontend

The backend logic is based mainly on a script we developed in our discussion of AG2 in [Chapter 7](#) (*07_AgenticSystems\ag2\20_ag2_conversation.py*).

If you're familiar with Streamlit, take this opportunity and try to develop the app on your own. If not, don't worry, we'll develop the code together. The final script can be found in *08_Deployment\self_contained_app.py*.

First, load relevant packages using the code shown in [Listing 8.8](#).

```

#%% packages
from autogen import ConversableAgent
from dotenv import load_dotenv, find_dotenv
import os
#%% load the environment variables
load_dotenv(find_dotenv(usecwd=True))
import streamlit as st

```

Listing 8.8 Streamlit App “Controversial Debate”: Packages (Source: 08_Deployment\self_contained_app.py)

We’ll use AG2 and `ConversableAgent` for setting up the two agents that will debate. Streamlit is imported with its usual alias `st`, which will save us time later.

The agents will need, during their instantiation, configuration details on the LLM, the temperature, and the API key. All of this is communicated through an object called `llm_config`, shown in [Listing 8.9](#), that we’ll later pass to the agents.

```

#%% LLM config
llm_config = {"config_list": [
    {"model": "gpt-4o-mini",
     "temperature": 0.9,
     "api_key": os.environ.get("OPENAI_API_KEY")}]}}

```

Listing 8.9 Streamlit App “Controversial Debate”: LLM Config (Source: 08_Deployment\self_contained_app.py)

OK, time to create our first widget with Streamlit, as shown in [Listing 8.10](#). A widget is a predefined functionality for creating inputs or outputs that follow a common design pattern and provide good functionality with minimal code.

```

st.title("Controversial Debate")

prompt = st.chat_input("Enter a topic to debate about:")
if prompt:
    st.header(f"Topic: {prompt}")

```

Listing 8.10 Streamlit App “Controversial Debate”: User Prompt (Source: 08_Deployment\self_contained_app.py)

For creating the title of our app, we can call `st.title`. The topic to debate about is passed on in a chat input field that is displayed at the bottom of the page. A predefined widget called `st.chat_input` for this input field.

[Listing 8.11](#) shows the settings area that can be hidden or expanded.

```
with st.expander("Conversation Settings"):
    number_of_turns = st.slider("Number of turns", min_value=1, max_value=10,
value=1)

    col1, col2 = st.columns(2)
    with col1:
        st.subheader("Style of Person A")
        style_a = st.radio(
            "Choose style for first speaker:",
            ["Friendly", "Neutral", "Unfriendly"],
            key="style_a"
        )

    with col2:
        st.subheader("Style of Person B")
        style_b = st.radio(
            "Choose style for second speaker:",
            ["Friendly", "Neutral", "Unfriendly"],
            key="style_b"
        )
```

Listing 8.11 Streamlit App “Controversial Debate”: Settings (Source: 08_Deployment\self_contained_app.py)

In the settings, the behavior of the two agents is defined. Since we want clear visual separation between the two agents, we’ll use `st.columns` to create two columns. The style is limited to “friendly,” “neutral,” and “unfriendly.” The user can select the style from a radio button, created with `st.radio`.

Now, we want to process these inputs. [Listing 8.12](#) shows the setup of the agents.

```
if prompt:
    ### set up the agent: Jack, the flat earther
    person_a = ConversableAgent(
        name="user",
```

```

        system_message=f"""
        You are a person who believes that {prompt}.
        You try to convince others of this.
        You answer in a {style_a} way.
        Answer very short and concise.
        """,
        llm_config=llm_config,
        human_input_mode="NEVER",
    )

    ### set up the agent: Alice, the scientist
    person_b = ConversableAgent(
        name="ai",
        system_message="""
        You are a person who believes the opposite of {prompt}.
        You answer in a {style_b} way.
        Answer very short and concise.
        """,
        llm_config=llm_config,
        human_input_mode="NEVER",
    )
    # %% start the conversation
    result = person_a.initiate_chat(
        recipient=person_b,
        message=prompt,
        max_turns=number_of_turns)

```

Listing 8.12 Streamlit App “Controversial Debate”: Agent Setup (Source: 08_Deployment\self_contained_app.py)

As soon as a prompt is available, the two agents are instantiated with `ConversableAgent`. In its definition, we pass information on their position on the debate topic, as well as the style, which the user can adapt. The names of the agents are set to "user" and "ai" for a reason. Later, when we render the messages to the screen, Streamlit already has some predefined behaviors if some default names like these are used.

After the agents have been created, person A starts the conversation and reaches out to person B. The chat is stored in an object called `result`.

The snippet shown in [Listing 8.13](#) illustrates how the chat discussion is rendered to the screen in a visually appealing way.


```

messages = result.chat_history
for message in messages:
    name = message["name"]
    if name == "user":
        with st.container():
            col1, col2 = st.columns([3, 7])
            with col2:
                with st.chat_message(name=name):
                    st.write(message["content"])
    else:
        with st.container():
            col1, col2 = st.columns([7, 3])
            with col1:
                with st.chat_message(name=name):
                    st.write(message["content"])

```

Listing 8.13 Streamlit App “Controversial Debate”: Chat Messages (Source: 08_Deployment/self_contained_app.py)

The messages are extracted from the `chat_history`. Then, each message is processed in a loop. The name of the message owner is extracted, and depending on the message owner, the `chat_message` is shown either on the left or on the right.

That’s it! We’ve created a nice frontend for our agent conversation. Let’s check how this looks, and if everything works, by running the app.

To run the app locally, in the terminal, change into the folder of the script and run the following command:

```
streamlit run self_contained_app.py
```

You can now view your Streamlit app in your browser.

```

Local URL: http://localhost:8501
Network URL: http://192.168.179.16:8501

```

This command will start a webserver and open the URL in your default browser. You should see the web app (shown earlier in [Figure 8.5](#)) after visiting the URL. Play with the app and provide different controversial topics for the agents to discuss.

At this point, we've developed a self-contained web application that we can now deploy, which we'll do in the coming sections.

OceanofPDF.com

8.4 Deployment to Heroku

Heroku is a popular service that simplifies the deployment and management of AI applications. Known for its developer-friendly interface, Heroku also has a simple integration with Git. Heroku enables developers to deploy with minimal operational overhead. The idea is that you provide the code, and the service takes care of the infrastructure complexities like server setup and scaling, thus allowing you to focus on developing applications.


One of the supported languages is Python, so we are ready to start. As the first step, you need to create an account on <https://signup.heroku.com/>.

Every app costs something to run, even if just a few cents, so you must add payment details. Only then can you proceed with creating a new app.

We'll start by creating a new app in [Section 8.4.1](#). Our development is performed locally, so we need to download and configure a command line interface (CLI), which we'll do in [Section 8.4.2](#). The deployment process includes several steps in which we create an app.py script ([Section 8.4.3](#)), perform a profile setup ([Section 8.4.4](#)), define environment variables ([Section 8.4.5](#)), and handle the Python environment ([Section 8.4.6](#)). At that point, we can check the deployment locally ([Section 8.4.7](#)). Once everything is checked and working, we can deploy the app to Heroku, as described in [Section 8.4.8](#). As a cleanup step, in [Section 8.4.9](#), you'll learn how to stop your application.

8.4.1 Create a New App

[Figure 8.6](#) shows the options for creating a new app in the frontend.



The screenshot shows the 'Create New App' form in Heroku. The form has a title 'Create New App' at the top. Below it, there is a section for 'App name' with a text input field containing 'controversial-debate' and a green checkmark icon. Below the input field, it says 'controversial-debate is available'. The next section is 'Location', with a subtitle 'Choose a Common Runtime region for this app. [Learn more.](#)'. There are two radio button options: 'Common Runtime CEDAR' with 'United States' (selected) and 'Common Runtime CEDAR' with 'Europe'. Below these is a toggle switch for 'Add this app to a pipeline', which is currently turned off. At the bottom, there is a message 'Create a new application named controversial-debate as a personal app.' and two buttons: 'Cancel' and 'Create app'.

Figure 8.6 Heroku: Create New App

You must maintain the **App name** field and the **Location** field. Select either **United States** or **Europe** and click **Create app**.

8.4.2 Download and Configure CLI

You must also install a CLI on your system. For this step, navigate to <https://devcenter.heroku.com/articles/getting-started-with-python#set-up>. Several CLI options are available for the platform you need.

In the terminal, now, log on to Heroku, as shown in [Listing 8.14](#).

```
heroku login
```

```
Opening browser to https://cli-auth.heroku.com/auth/cli/browser/887491ec-989c-46ae-
bf33-4387e4e75813?
requestor=SFMyNTY.g2gDbQAAAA0xMzAuMTc2Ljg4LjczbgYAcAKMFpQBYgABUYA.zh1whD4bP42ivN8_u7
HZd6d1c3Y60QHvb6Yq1XyMsFI
Logging in... done
Logged in as info@gollnickdata.de
```

Listing 8.14 Heroku Login Output

A browser window opens, and you'll click a button to log on. In the terminal, you must link to the remote repository, as follows:

```
heroku git:remote -a controversial-debate
```

Now, you're linked to the remote repository and can proceed by creating the required files.

8.4.3 Create app.py File

The first file is our actual application file, which we'll call *app.py*. This file will be linked in other configuration files. [Listing 8.15](#) shows the complete application file.

```
### packages
import os
from autogen import ConversableAgent
import streamlit as st
def main():
    # LLM config
    llm_config = {"config_list": [
        {"model": "gpt-4o-mini",
         "temperature": 0.9,
         "api_key": os.environ.get("OPENAI_API_KEY")}]}}

    st.title("Controversial Debate")

    prompt = st.chat_input("Enter a topic to debate about:")
    if prompt:
        st.header(f"Topic: {prompt}")

    with st.expander("Conversation Settings"):
        number_of_turns = st.slider("Number of turns", min_value=1, max_value=10,
value=1)

        col1, col2 = st.columns(2)
```

```

with col1:
    st.subheader("Style of Person A")
    style_a = st.radio(
        "Choose style for first speaker:",
        ["Friendly", "Neutral", "Unfriendly"],
        key="style_a")

with col2:
    st.subheader("Style of Person B")
    style_b = st.radio(
        "Choose style for second speaker:",
        ["Friendly", "Neutral", "Unfriendly"],
        key="style_b")

if prompt:
    ### set up the agent: Jack, the flat earther
    person_a = ConversableAgent(
        name="user",
        system_message=f"""
        You are a person who believes that {prompt}.
        You try to convince others of this.
        You answer in a {style_a} way.
        Answer very short and concise.
        """,
        llm_config=llm_config,
        human_input_mode="NEVER")

    ### set up the agent: Alice, the scientist
    person_b = ConversableAgent(
        name="ai",
        system_message="""
        You are a person who believes the opposite of {prompt}.
        You answer in a {style_b} way.
        Answer very short and concise.
        """,
        llm_config=llm_config,
        human_input_mode="NEVER",
    )

    # %% start the conversation
    result = person_a.initiate_chat(
        recipient=person_b,
        message=prompt,
        max_turns=number_of_turns)

    messages = result.chat_history
    for message in messages:
        name = message["name"]
        if name == "user":
            with st.container():
                col1, col2 = st.columns([3, 7])
                with col2:
                    with st.chat_message(name=name):
                        st.write(message["content"])

```

```

        else:
            with st.container():
                col1, col2 = st.columns([7, 3])
                with col1:
                    with st.chat_message(name=name):
                        st.write(message["content"])

if __name__ == "__main__":
    main()

```

Listing 8.15 Heroku App File (Source: 08_Deployment\heroku_app\app.py)

A few minor changes were made to *08_Deployment\self_contained_app.py*. The most significant change is that now a `main()` function contains the app logic. This function is called when the file is run.

We can proceed by setting up some configuration files that tell Heroku what to load and how.

8.4.4 Procfile Setup

We need to set up a *procfile* (*process file*), which is a special file used in web applications, particularly with platforms like Heroku, to specify the commands that should be executed to start our application.

The details on how Heroku should handle the app are defined in the *Procfile* with the following content:

```
web: streamlit run app.py --server.port=$PORT --server.headless=true
```

For testing, create a file with the identical content called *Procfile.windows* (assuming you work on Windows). That step is not necessary if you work on Linux or MacOS.

8.4.5 Environment Variables

Our app relies on OpenAI, so we need to provide an API key. In the terminal, we can define the environment variables with `heroku config:set` followed by a mapping of key name and its value:

```
heroku config:set OPENAI_API_KEY=...
```

Check the variable settings by using the following command:

```
heroku config
```

```
=== controversial-debate Config Vars
```

```
OPENAI_API_KEY: sk-proj-...
```

8.4.6 Python Environment

The app requires its own Python environment with all the required packages. We'll set up the required packages in a *requirements.txt* and then create an environment based on it, as follows:

1. Create a requirements.txt file

Set up a file called *requirements.txt* with the following contents:

```
streamlit
```

```
autogen
```

You might even specify package versions.

2. Setting up the environment

The environment is called `.venv` and can be created with the following command:

```
python -m venv .venv
```

After creation, you must activate the environment by changing into the scripts folder and calling `activate`.

3. Install the packages

You need to install the packages in this new environment via pip:

```
pip install -r requirements.txt
```

8.4.7 Check the Result Locally

Before we deploy to a remote server, you should check the setup locally. For this task, call `heroku local` in combination with the `Procfile.windows`, as shown in [Listing 8.16](#). If you work on Mac or Linux, use the `procfile`.

```
heroku local --port 5006 -f Procfile.windows
```

```
09:36:39 web.1 | You can now view your Streamlit app in your browser.
09:36:39 web.1 | Local URL: http://localhost:5006
09:36:39 web.1 | Network URL: http://192.168.179.147:5006
09:36:39 web.1 | External URL: http://79.192.50.59:5006
```

Listing 8.16 Heroku Local Server Startup

This command will spin up a web server. The output will display the local URL to which you must navigate to check the result. If we head over to `http://localhost:5006`, we'll see the running app.

Amazing! One final step remains: We need to deploy to the web via Heroku and check the result there.

8.4.8 Deployment to Heroku

In Visual Studio Code (VS Code), the default branch is called `main`^[1]. We must stage the files with the command `git add`, as follows:

```
git add .
```

Now, we can commit the files with a commit message, as follows:

```
git commit -m "first commit"
```

Finally, we can push these changes to Heroku with `git push`, as shown in [Listing 8.17](#).

```
git push heroku main
```

```
remote: Building source:
remote:
remote: -----> Building on the Heroku-24 stack
remote: -----> Using buildpack: heroku/python
...
remote: -----> Compressing...
remote:           Done: 136.2M
remote: -----> Launching...
remote:           Released v5
remote:           https://controversial-debate-2b07aacb7d3.herokuapp.com/ deployed to
Heroku
remote:
remote: Verifying deploy... done.
To https://git.heroku.com/controversial-debate.git
 * [new branch]      main -> main
```

Listing 8.17 Heroku App Deployment

The CLI informs you of the place where the app is deployed. Head over to that URL to check the result. In our case, our URL is *<https://controversial-debate-2b07aacb7d3.herokuapp.com/>*. (Don't try to use this link; we've already stopped the service.) Head over to your own URL to see the result in your browser. [Figure 8.7](#) shows the result of the app we deployed locally.

With a topic like “the earth is flat,” the two agents start their discussion.

Before we move on, we recommend deleting your app if you just wanted to test it. Alternatively, you can simply stop it to avoid costs.




Figure 8.7 Heroku Deployed App

8.4.9 Stop Your App

You can determine how much your app costs under the **Resources** tab in the web frontend, in our case, <https://dashboard.heroku.com/apps/controversial-debate/resources>.

[Figure 8.8](#) shows our app costs about \$0.01 per hour, roughly \$7.00 per month.




Figure 8.8 Heroku App State and Estimated Cost

Click the **Edit** button next to the hourly cost to deactivate the app. [Figure 8.9](#) now shows no hourly cost, and the estimated total cost is \$0.00.




Figure 8.9 Heroku App State Stopped

To verify that the app has stopped, you can navigate to the app URL. The app should not load any more. Thus concludes our first deployment via Heroku. We hope you've seen how easy it is to deploy an app.

Next, we'll move on another service provider: Streamlit.

8.5 Deployment to Streamlit

We'll use the app we developed previously and focus now only on the deployment steps, which are just a few in number:

1. Creating a GitHub repository
2. Creating a new app at <https://streamlit.io/>

8.5.1 GitHub Repository

Let's start with the creation of a GitHub repository. With Heroku, the service has automatically set up a Git repository. However, with Streamlit, we need to take care of managing GitHub ourselves.

Managing GitHub is outside the scope of this book, so we'll only roughly sketch out the steps you need to take. If you're not familiar with GitHub and Git repositories, we highly recommend learning this skillset, which is essential to nearly every workflow in real-life projects. One good resource is *Git: Project Management for Developers and DevOps Teams* (https://www.sap-press.com/git_5555/).

First, create a repository at <https://github.com/>. Once you have created the empty repository, you must clone it locally via `git clone`.

The following files must be part of the repository:

- `app.py`
- `requirements.txt`

After you have copied the files into that folder, you must `commit` and `push` the changes back to the GitHub repository. Navigate to your

GitHub repository, as shown in [Figure 8.10](#).




Figure 8.10 GitHub Repository for Deployment to Streamlit

Now, you should have an *app.py* and a *requirements.txt*. The *.devcontainer* folder will be later automatically created by the service.

8.5.2 Creating a New App


First, navigate to <https://streamlit.io/> and create an account. Once you have an account, you can deploy apps by clicking **Create app** in the top-right corner. On the next screen, shown in [Figure 8.11](#), select the source of your app.



Figure 8.11 Streamlit: Source of Your App

We'll use the first option: **Deploy a public app from GitHub**. Although called a public app, you do have some options to limit the

app to specific users afterwards. [Figure 8.12](#) shows the parameters you must define to deploy an app.




The image shows a web form titled "Deploy an app". It has two main input sections. The first is "GitHub URL" with a text box containing "https://github.com/DataScienceHamburg/controversial_debate_streamlit/blob/main/app.py" and a link "Switch to interactive picker". The second is "App URL (optional)" with a text box containing "controversialdebate" and a suffix ".streamlit.app". Below these is a green message "Domain is available" and a link "Advanced settings". At the bottom is a blue "Deploy" button.

Figure 8.12 Streamlit: Deploy an App

First and foremost, you must pass the **GitHub URL**. Bear in mind that you don't pass the URL of the GitHub repository but instead the URL of the *app.py* file in the GitHub repository. Also, you must define an **App URL**. This URL is where your app can be found. For now, simply make sure that the domain is available.

We cannot skip over the advanced settings because our app relies on an API key. Since our agents use an OpenAI model, click on **Advanced settings** to set the key. [Figure 8.13](#) shows how to provide the secret.



The image shows a web form titled "Advanced settings". It has a "Python version" dropdown menu set to "3.12". Below is a "Secrets" section with a text box containing "OPENAI_API_KEY = 'sk-proj-[REDACTED]'". At the bottom is a blue "Save" button.

Figure 8.13 Streamlit: Advanced Settings

In the **Secrets** section, we must pass our API key. Copy the key from your `.env` file. Importantly, at this point, the key is passed in quotation marks, so you should wrap your key like this: `"my_secret_key"`. Click **Save**.

Back on the **Deploy an app** screen, confirm your settings by clicking on the **Deploy** button. This step starts the deployment process and then starts the app. The app is then live at an address following the convention `https://<yourappname>.streamlit.app/`.

Congratulations! You've successfully deployed your first app with Streamlit!

OceanofPDF.com

8.6 Deployment with Render

Render is another easy-to-use deployment service. We'll use its free tier for our sample deployment. The steps are like other frameworks. First, create an account at <https://render.com/>. Then, start the deployment process by selecting **Deploy a Web Service**, as shown in [Figure 8.14](#).




Figure 8.14 Render: Select a Service

You must link the deployment service to your GitHub repository. For this step, navigate to your own GitHub repository and copy the URL into the clipboard. Then, paste this URL into the Render deployment, as shown in [Figure 8.15](#).




Figure 8.15 Render: Connect to a GitHub Repository

Further parameters can be specified on the Render configuration page, but for our simple case, we can stick to the defaults. Parameters you might want to change include the following:

- **Region:** You might select a region close to your users.
- **Root directory:** The default is the root folder of the repository.
- **Build command:** Since we're working with a *requirements.txt*, we can stick to the default installation command `pip install -r requirements.txt`.

You must define a start command in the following way:

```
streamlit run app.py --server.port $PORT --server.headless true
```

[Figure 8.16](#) shows the instance type, which defines the capabilities and costs of your project.

Instance Type

For hobby projects

- Free** 512 MB (RAM) 0.1 CPU \$0 / month

For professional use

For more power and to get the most out of Render, we recommend using one of our paid instance types. All paid instances support:

- Zero Downtime
- SSH Access
- Scaling
- One-off jobs
- Support for persistent disks

Upgrade to enable more features

Free instances spin down after periods of inactivity. They do not support SSH access, scaling, one-off jobs, or persistent disks. Select any paid instance type to enable these features.

Starter	512 MB (RAM) 0.5 CPU	Standard	2 GB (RAM) 1 CPU
Pro	4 GB (RAM) 2 CPU	Pro Plus	8 GB (RAM) 4 CPU
Pro Max	16 GB (RAM) 4 CPU	Pro Ultra	32 GB (RAM) 8 CPU
\$7 / month		\$25 / month	
\$65 / month		\$175 / month	
\$225 / month		\$450 / month	

Need a custom instance type? We support up to 512 GB RAM and 64 CPUs.

Figure 8.16 Render: Select an Instance Type

For a small test as in our case, we recommend selecting the **Free** instance. This instance spins down after periods of inactivity, which is fine for us for now.

The last step is shown in [Figure 8.17](#), which illustrates how to specify environment variables.

Environment Variables

Set environment-specific config and secrets (such as API keys), then read those values from your code. [Learn more](#).

OPENAI_API_KEY

+ Add Environment Variable Add from env

Figure 8.17 Render: Environmental Variables

The API key setup is straightforward. Simply define the name and the corresponding key as already defined in the `.env` file.

That's everything, and we can start the deployment process by clicking **Deploy app**. In the subsequent dashboard, you can watch the deployment process in real time, including the installation process for all the packages.

After a few minutes, your app is live at a URL that follows the convention *`https://<your-app-name>.onrender.com`*.

OceanofPDF.com

8.7 Summary

In this chapter, we explored several critical concepts and practical techniques for deploying applications, focusing on the contrasting strategies of self-contained applications and the more modular frontend/backend architecture.

Our journey began with distinguishing among self-contained applications: one kind that combined the frontend and backend logic and one kind in which frontend and backend architectures are separated.

We had a brief introduction to REST APIs, an essential tool for building scalable, interoperable systems. You learned how REST APIs enable seamless communication between various components of an application, making them a cornerstone of modern app development.

We then delved into the development of a self-contained application, which provided a clear, practical example of how such apps encapsulate all necessary functionality within a single package. This approach highlighted the benefits of simplicity, ease of deployment, and straightforward maintenance, making self-contained apps an excellent choice for certain use cases.

For the deployment phase, we demonstrated how to bring applications to life in real-world environments. Starting with Heroku, we explored its developer-friendly features, including quick setup and powerful tools for managing applications in production.

Next, we moved to Streamlit, a platform tailored for deploying interactive data applications. We experienced its ease of use and

ability to rapidly turn Python scripts into shareable web apps.

Finally, we deployed our application to Render, showcasing its flexibility and cost-effectiveness for a range of deployment scenarios.

Each deployment platform introduced unique challenges and advantages, providing insights into when and why you might choose one over another.

Through these examples, we discussed the following topics:

- REST API integration: How to design and use REST APIs to enable effective communication between app components and external services.
- Self-contained apps: The strengths and limitations of this approach, particularly in scenarios where simplicity and portability are key.
- Platform-specific deployment: Analyzing and adapting to the specific requirements and strengths of Heroku, Streamlit, and Render.

By the end of this chapter, you should have a comprehensive understanding of the tools, techniques, and decision-making processes involved in deploying applications. Whether you're working on a standalone project or a complex, distributed system, these skills form a solid foundation to effectively bring your applications from development to production.

9 Outlook

The future is not set. There is no fate but what we make for ourselves.

—Sarah Connor in the movie Terminator 2: Judgment Day

Admittedly, quoting a movie about an AI trying to wipe out humanity might set the wrong mood. But it is just a great movie, and although over 30 years old, it reflects fears that many people have when considering the future of AI.

In this chapter, we'll provide an outlook on what AI, especially generative artificial intelligence (generative AI), might provide in the future.

The field of generative AI is still quite new and is evolving extremely quickly. We'll discuss what can be expected of future developments, but also which limitations and issues persist and need to be overcome. We'll briefly talk about regulatory developments before we examine the potential of reaching artificial general intelligence (AGI) and artificial superintelligence (ASI). We'll conclude this chapter with a section on useful resources that can help you to keep up to date.

9.1 Advances in Model Architecture

Besides improvements in hardware, data, and model sizes, the model architecture will have to get better to tackle complex tasks in coding, mathematics, and general intelligence. Leopold Aschenbrenner wrote a brilliant paper called “Situational Awareness: The Decade Ahead” (<https://situational-awareness.ai/>). We agree with many of his predictions.

The measure of progress is called *OOMs (orders of magnitude)*. The reason for this measure is that, as we have already seen, many aspects of progress are exponential in nature, so that OOMs relate to a logarithmic scale. OOMs are shown over time.

Algorithmic progress is further broken down into *unhobbling* and *efficiencies*. “Unhobbling” refers to removing self-imposed or systematic constraints that limit a model’s effectiveness or potential. “Efficiencies” can mean several things: there are efficiencies in inference and in training. For example, Aschenbrenner mentions improvements in sample efficiency that could allow algorithmic improvements that enable models to learn from more limited (less) data.

At the time of writing, OpenAI’s most recent model o3 was published. This model is the successor to its o1 “reasoning” model. With o3, OpenAI was supposed to be capable of reaching AGI under certain conditions.

With increased reasoning capabilities like *private chain-of-thought (CoT)*, an AGI can self-fact-check, which allows the model to plan ahead and improve its own reasoning. This capability results in increased inference times, but significant improvements can be enjoyed in terms of accuracy in complex tasks.

Besides classic model training, reinforcement learning was applied to induce thinking before actually creating answers.

Who knows? Maybe, we'll even see more fundamental changes in the underlying transformer architecture at some point.

For the time being, still limitations and issues exist, which we'll dive into next.

OceanofPDF.com

9.2 Limitations and Issues of LLMs

Still many possible issues exist in the context of large language models (LLMs). We'll touch on the most relevant ones in this section: hallucinations, biases, misinformation, intellectual property, transparency, and jailbreaking.

9.2.1 Hallucinations

LLMs always produce an output if you don't actively constrain them. Thus, they'll even provide outputs if they don't have the answer in their model weights. In the early days, this problem led to unsatisfactory results because the answers were factually incorrect. In some cases, the output was disconnected from the user query; it might be nonsensical or conflict with previous outputs in the same discussion.

Much effort was made to improve model performance, so that hallucinations have become less frequent. Researchers developed a metric called *context adherence* that measures factual accuracy and cases in which the model provided outputs that were not provided in the context (*closed-domain hallucinations*).

Many models have a context adherence score of close to 1. For example, Claude-3.5 Sonnet ranked best closed source model with a score of 0.97. The same result was achieved by Llama-3.1 405B Instruct as best open-source model, with the same score of 0.97.

For medium-length contexts, which researchers defined as between 5,000 and 25,000 tokens, the best models like Google Gemini-1.5-

Flash achieved a score of 1.0, meaning no hallucinations were found.

9.2.2 Biases

Biases might be introduced at different steps of a training pipeline. [Figure 9.1](#) shows the different steps of the training pipeline, in which biases might be introduced.




Figure 9.1 Bias Introduction into LLM Training

It all starts with the data used for training. What follows is the modeling step and finally the human review. All steps are highlighted in the following sections.

Data

A model can only be as good as the data it was trained on. If the model is based on data that is found in the real world, the same bias most likely will be inherent to the model. Such historical biases can be found for groups that have been historically discriminated against. For example, gender-based stereotypes might be reflected by the model due to its training data.

A similar problem might occur when the developers try to counterweigh historical biases. Google got massive headwind in early 2024, when it released its image creation feature in Google Gemini.

Biases can also be introduced on purpose to follow a political agenda. [Figure 9.2](#) shows model outputs on Taiwan independence—one coming from a model created by a US company (Llama-3.2-3B-Instruct **1**) and one coming from a Chinese company (Qwen2.5-1.5B-Instruct **2**).




Figure 9.2 Biases Based on Political Agenda

Not surprisingly, the two models provide completely different answers. We can thus conclude that LLMs are not independent and won't be in the future. LLMs are not independent tools. Keep this in mind when working with LLMs.

Modeling

Biases can also be introduced at the modeling step of the training pipeline. The training progress is evaluated during the training. This evaluation is performed based on benchmarks, which where a bias could be introduced. It also might be that the benchmark is not representative of the general distribution.

This issue is typical in ML model training. Your model might work well on training data but not generalize well on unseen real-world data.

Another type of bias in this step is an aggregation bias, in which distinct populations are incorrectly combined. For example, if the loss function disproportionately penalizes incorrect grammar over content misinterpretation, the model will be biased toward correcting grammatical errors more rigorously than factual ones.

Human Review

The final step of the training is human feedback. This step is being taken to further improve the model. The technical term is *reinforcement learning from human feedback (RLHF)*. This technique involves human feedback to fine-tune models so that their outputs align better with human preferences or expectations.

A bias in RLHF might lie in the group setup of the humans who provide feedback. Ideally, they should provide a good representation of the broader population for which the model is developed.

Scientists have found out that human reviewers prefer longer model outputs, even when the output objectively did not change. Pure response length increase is seen as model improvement, rather than other features.

9.2.3 Misinformation

Misinformation is a risk associated with LLMs because LLMs often generate plausible but factually wrong content (see our earlier discussion on hallucinations in [Section 9.2.1](#)). Sure, misinformation was possible before the rise of LLMs. What has changed is that false or misleading information no longer requires human effort, allowing

bad actors (who might be persons, corporations, or state-level actors) to automate their disinformation campaigns.

The barrier has been lowered to creating fake news, conspiracy theories, or biased narratives, making it harder to distinguish credible information from fabricated content. The speed and cost-effectiveness of content generation with LLMs could overwhelm fact-checking mechanisms and contribute to the erosion of public trust in media and information sources, amplifying the societal impact of misinformation.

A special case in this context is hate speech. LLMs can create harmful or offensive language with the aim of targeting individuals or groups based on race, gender, religion, or other characteristics. The scalability of this process makes it extremely simple to flood online platforms and forums with harmful rhetoric.

9.2.4 Intellectual Property

On intellectual property, issues arise around copyright and data ownership. Copyrighted data might be part of the training data used without authorization from the data's actual owners. Lawsuits on training data in LLMs and text-to-image algorithms are being fought out in court.

9.2.5 Interpretability and Transparency

With great power comes low interpretability, as shown in [Figure 9.3](#), which includes several different ML algorithms. While simple algorithms come with great interpretability, they are associated with low performance. For example, linear regression is comparably simple, and for a human, it is easy to understand why the model

drew certain conclusions. But linear regression is far away from the performance of neural networks, of which LLMs are a prominent member. These models are extremely powerful, but at the same time, interpreting their results and understanding how they came to these results has become difficult.




Figure 9.3 Interpretability versus Complexity/Performance

9.2.6 Jailbreaking LLMs

LLMs are extremely powerful, even capable of providing harmful output like instructions on making drugs, hiding a corpse, or robbing a bank. Model developers have put a lot of effort into the detection of harmful user queries to prevent models from creating harmful output. The process of tricking the models to output forbidden content is called *jailbreaking*.

[Figure 9.4](#) shows the approach of MathPrompt for jailbreaking LLMs, from the paper “Jailbreaking Large Language Models with Symbolic

Mathematics,” by E. Bethany and colleagues (<https://arxiv.org/html/2409.11445v1>).




Figure 9.4 Jailbreaking LLMs with MathPrompt

In this approach, a harmful prompt returns a safe output, which is the expected result. But the researchers found out that the LLMs safety mechanisms can be bypassed by encoding harmful prompts in math. They showed how many state-of-the-art models reveal a high attack success rate.

Mathematical prompts, including harmful content, could be embedded such in a way that they are similar to mathematical embeddings but that hide the harmful content. In a way, mathematical prompt functions as a Trojan horse to the harmful content.

A similar approach is ArtPrompt, shown in [Figure 9.5](#), based on the paper “ArtPrompt: ASCII Art-based Jailbreak Attacks against Aligned

LLMs” from F. Jiang and colleagues
(<https://arxiv.org/pdf/2402.11753>).




Figure 9.5 Jailbreaking Based on ArtPrompt

In this approach, a harmful word is not directly passed to the LLM but encoded as ASCII art. This approach could also fool LLMs and provide harmful output.

9.3 Regulatory Developments

Regulatory developments are starting to address these issues and challenges. For example, the European Union's AI Act is a significant legislative effort aimed at regulating AI systems based on their risk levels. High-risk applications like LLMs will face stricter requirements.

It remains to be seen whether widespread regulation will take place, or whether economic, governmental, and military interests will dominate and prevent regulation.

9.4 Artificial General Intelligence and Artificial Superintelligence

AGI refers to a type of artificial intelligence that possesses the ability to understand, learn, and apply knowledge across a wide range of tasks and domains. Unlike specialized (narrow) AI, which is optimized to perform specific tasks, AGI aims to exhibit general cognitive abilities. [Figure 9.6](#) shows a possible growth path for AI-capabilities over time.




Figure 9.6 AI Capabilities over Time

At the time of writing, most researchers agree that we have not reached AGI. You can expect heated debate whether AGI has been reached for the foreseeable future. Ever-improving models, like the recently released OpenAI o3 model, with their further improvements on reasoning, will heat the debate.

For the time being, we have undoubtedly many narrow AI systems that excel in their domains. We can observe more and more models with more general capabilities, but no single model is capable of all

human domains. Eventually, we'll see AI systems that show capabilities like humans.

We also can probably expect encounter AI systems with intellectual powers beyond human level across a large number of categories and fields: artificial superintelligence (ASI).

Whether or not, we'll see achieve ASI is still an open question at this point. In any case, we are living in extraordinary times and are witnessing a revolution. AGI and ASI are just extreme cases in the AI debate. Many smaller steps still must be taken that will impact on our lives in the short term. While predicting the future is hard, it's still fun to think about all the possible paths.

9.5 AI Systems in the Near Term

According to OpenAI, there are five stages of artificial intelligence: Level 1 is dominated by chatbots, and AI is used for conversational language. In Level 2, reasoning models can reach human-level problem solving. Agentic systems represent Level 3. In Level 4, AI systems can support creative tasks like inventions and effectively act as innovators. In Level 5, AI systems can mimic and do the work of complete organizations. At that point, companies worth of billions of dollars might be led by a single person (a one-person unicorn) or even by no person (a zero-person unicorn). We find this prospect both fascinating and frightening at the same time.

An effective way to think about future technologies is the *hype cycle*—developed and popularized by the consulting company Gartner. The hype cycle describes how new technologies typically progress through different phases.

Starting with an “innovation trigger” that sparks initial interest and investment, at some point, the technology reaches “inflated expectations.” At this point, hype and over-optimism reach the maximum. What follows is a “trough of disillusionment” when interest wanes. Exaggerated promises cannot be kept.

Then comes a “slope of enlightenment” where realistic applications emerge. Finally, in the “plateau of productivity” the technology shows its real benefits. Then, it is widely understood and adopted.

[Figure 9.7](#) shows our take on how selected AI technologies will develop.

The color and shape of the points corresponds to the speed with which the technology progresses through the steps. We’ve

separated the technologies into short-term (< 5 years), mid-term (5 to 10 years), and long-term (> 10 years) outlooks.

Let's start with the AI systems that we expect to reach the plateau of productivity in the short term, for instance, medical image analysis. Since computer vision models outperform human benchmarks for many years, the only hurdles are regulatory and less technological.




Figure 9.7 Our AI Hype Cycle Prediction

Other technologies that are well established are precision farming and AI in the legal industry. The legal industry with its standardized texts and frameworks is an ideal playground for LLMs to shine. The hype around autonomous driving has disappeared. We see a lot of disappointment in this field because overoptimistic predictions caused hopes that could not be lived up to. So, we might be close to and possibly beyond the trough of disillusionment, and hopefully, we'll have autonomous driving in the near future.

Similarly, delivery drones were very hyped up some years ago, and reality did not live up to expectations. Currently, hardly anyone is talking about delivery drones. We expect that this technology will steadily progress and be market ready in the mid-term.

Some technologies like brain-computer interfaces have gained a lot of media attention. We might be reaching the maximum of the hype cycle in this area. But there are many technological and regulatory hurdles to overcome that we see this field to be adopted widely, but only in the far future.

Other AI technologies that are further down the road are AI self-awareness/consciousness and AI governance/policy-making. While the former topic is extremely hard to predict (whether consciousness ever will occur), the latter will be adopted due to the improvements that it will bring. But developing and accepting such systems will take a lot of time, so we expect them only in the long term.

Huge leaps forward in robotics show the possibilities this will bring. But our view is that the expectations are currently inflated and will be disappointing. Nevertheless, since progress has been extremely fast, we expect this technology to reach the plateau of productivity in less than 5 years.

9.6 Useful Resources

We hope this book is the starting point of your journey, rather than the end. In this section, we'll provide some resources to stay informed and to help understand the concepts presented in this book.

- **Online courses**

We typically learn based on online courses provided via Udemy, O'Reilly, or Packt. You might even check out the courses on these platforms. Other free platforms like DeepLearning.AI (<https://deeplearning.ai/>) provide up-to-date content that usually is quite concise and a great resource to learn new frameworks or concepts.

- **Articles**

Daily.dev (<https://daily.dev/>) is a Chrome plugin. With this plugin, you'll see interesting articles on coding tailored to your preferences. The system learns based on your feedback to improve its recommendations. We would also like to recommend Medium (<https://medium.com/?tag=data-science>), where you can find many articles on data science and artificial intelligence.

- **YouTube**

Many interesting discussions and tutorials on YouTube. We use it to keep informed on the most recent developments.

- **Newsletters**

We also consume some newsletters to keep informed. We can recommend "The Batch," a newsletter from DeepLearning.AI where Andrew Ng writes about recent news on artificial intelligence.

In general, when diving into a new topic, you can drown in the online courses, articles, or documentation provided by the developers of a package or framework. An important step is to quickly move from carefully prepared courses to your projects. Only when you try to create an implementation myself can you spot missing links in your understanding.

Never was there a time in human history in which progress continued so quickly; it was never so important to learn new skills. Or put another way, more positively: It has never been so easy to get on the top of technology.

OceanofPDF.com

9.7 Summary

In this chapter, we looked into the crystal ball and tried to anticipate future developments for generative AI.

We started by discussing advances in model architecture. Then, we focused on current limitations and possible future limitations and issues in LLMs, such as hallucinations, biases, misinformation, and other issues. We studied how regulatory developments are trying to solve some of these issues.

The pace of progress in artificial intelligence is rapid, and artificial intelligence becomes increasingly likely to outperform average human skills, and even expert skills.

This book has journeyed through the incredible landscape of generative AI—from LLMs, to prompt engineering and vector databases, to RAG and agentic systems, and finally to the deployment of such systems.

Generative AI has become a driving force in nearly all industries, shaping everyday lives. As we look ahead, let's remember that every great innovation begins with a question: What will you create next?

The Author



Bert Gollnick is a senior data scientist who specializes in renewable energies. For many years, he has taught courses about data science and machine learning, and more recently, about generative AI and natural language processing. Bert studied aeronautics at the Technical University of Berlin and economics at the University of Hagen. His main areas of interest are machine learning and data science.

OceanofPDF.com

Index

↓ **A** ↓ **B** ↓ **C** ↓ **D** ↓ **E** ↓ **F** ↓ **G** ↓ **H** ↓ **I** ↓ **J** ↓ **K** ↓ **L** ↓ **M** ↓ **N** ↓ **O** ↓ **P** ↓ **Q**
↓ **R** ↓ **S** ↓ **T** ↓ **U** ↓ **V** ↓ **W** ↓ **Z**

A ↑

AG2 [→ Section 7.2] [→ Section 7.5]

two agents [→ Section 7.5]

with tools [→ Section 7.5]

Agentic systems [→ Section 7.1]

agentic RAG [→ Section 7.3]

monitoring [→ Section 7.9]

simple [→ Section 7.3]

AgentOps [→ Section 7.9]

AI agent [→ Section 7.1]

Alignment problem [→ Section 7.6]

AlphaGo [→ Section 1.4]

Amazon Web Services (AWS) [→ Section 5.6]

Answer faithfulness [→ Section 6.5]

Answer relevance [→ Section 6.5]

Apache 2.0 license [→ Section 5.8]

API key [→ Section 3.2]

Artificial general intelligence (AGI) [→ Section 1.1] [→ Section 9.1] [→ Section 9.4]

Artificial intelligence (AI) [→ Section 1.1]

algorithmic improvements [→ Section 1.2]

computational power [→ Section 1.2]

dataset size [→ Section 1.2]

investments [→ Section 1.2]

near term development [→ Section 9.5]

Artificial superintelligence (ASI) [→ Section 9.1] [→ Section 9.4]

ArtPrompt [→ Section 9.2]

Asynchronous server gateway interface (ASGI) [→ Section 8.2]

B ↑

Bar plot [→ Section 2.4]

Base64 [→ Section 3.2]

BERT [→ Section 1.5] [→ Section 2.5] [→ Section 5.5]
[→ Section 5.5]

Blue teaming [→ Section 7.6]

BM25 [→ Section 6.3] [→ Section 6.3]

Chain-of-thought (CoT) [[→ Section 3.10](#)] [[→ Section 4.3](#)]

Chains [[→ Section 3.7](#)]

parallel [[→ Section 3.7](#)]

router [[→ Section 3.7](#)]

simple sequential [[→ Section 3.7](#)]

with memory [[→ Section 3.7](#)]

Character encoding standard [[→ Section 5.4](#)]

Chatbot Arena [[→ Section 3.4](#)]

ChatGPT [[→ Section 1.1](#)]

ChatPromptTemplates [[→ Section 3.6](#)]

Chroma database [[→ Section 5.6](#)] [[→ Section 5.8](#)]

data retrieval [[→ Section 5.7](#)]

Chunking [[→ Section 5.2](#)] [[→ Section 7.3](#)]

Computer vision [[→ Section 5.5](#)]

Context adherence [[→ Section 9.2](#)]

Context enrichment [[→ Section 6.3](#)]

Context precision [[→ Section 6.5](#)]

Context relevance [[→ Section 6.5](#)]

Context windows [[→ Section 5.4](#)]

Convolutional neural networks (CNNs) [[→ Section 3.1](#)]

Cosine similarity [[→ Section 5.7](#)]

CrewAI [[→ Section 7.2](#)] [[→ Section 7.6](#)]
 agents [[→ Section 7.6](#)] [[→ Section 7.6](#)]
 collaboration [[→ Section 7.6](#)]
 components [[→ Section 7.6](#)]
 file interactions [[→ Section 7.6](#)]
 project adaptation [[→ Section 7.6](#)]
 safety exercise [[→ Section 7.6](#)]
 starter project [[→ Section 7.6](#)]
 tasks [[→ Section 7.6](#)] [[→ Section 7.6](#)]
Cursor [[→ Section 1.8](#)]
Custom chunking [[→ Section 5.4](#)]

D ↑

Data loading [[→ Section 5.2](#)]
Data processing [[→ Section 5.8](#)]
Data retrieval [[→ Section 5.7](#)]
Data storage [[→ Section 5.6](#)]
Deep Blue [[→ Section 1.4](#)]
Deep learning [[→ Section 1.1](#)] [[→ Section 1.3](#)]
DeepMind [[→ Section 1.4](#)]
Delimiters [[→ Section 4.1](#)]
Dense vector search [[→ Section 6.3](#)]
Deployment [[→ Section 8.1](#)]

architecture [[→ Section 8.1](#)]

Heroku [[→ Section 8.4](#)]

local [[→ Section 8.2](#)]

prediction function [[→ Section 8.2](#)]

priorities [[→ Section 8.2](#)]

Render [[→ Section 8.6](#)]

strategy [[→ Section 8.2](#)]

Streamlit [[→ Section 8.5](#)]

Direct preference optimization (DPO) [[→ Section 3.10](#)]

Direct prompting [[→ Section 3.9](#)]

Distance matrix [[→ Section 5.4](#)]

Distil-Whisper models [[→ Section 3.2](#)]

Document objects [[→ Section 5.8](#)]

E

Eliza [[→ Section 3.1](#)]

Embeddings [[→ Section 5.2](#)] [[→ Section 5.5](#)] [[→ Section 7.3](#)]

Embedding vectors [[→ Section 5.1](#)]

Ensemble learning [[→ Section 4.4](#)]

Environment variables [[→ Section 3.2](#)]

Epoch [[→ Section 1.3](#)]

F

Factual correctness [[→ Section 6.5](#)]
FastAPI [[→ Section 8.2](#)]
FastText [[→ Section 5.5](#)]
Few-shot learning [[→ Section 4.1](#)]
Few-shot prompting [[→ Section 4.2](#)]
File-based storage [[→ Section 5.6](#)]
Fill-mask tasks [[→ Section 1.5](#)] [[→ Section 2.5](#)]
Fixed-size chunking [[→ Section 5.4](#)]
Flask [[→ Section 8.2](#)]
Frontend and backend architecture [[→ Section 8.1](#)]

G ↑

Gated recurrent units (GRUs) [[→ Section 3.1](#)]
gemma2 [[→ Section 6.2](#)]
GitHub Copilot [[→ Section 1.8](#)]
GitHub repository [[→ Section 8.5](#)] [[→ Section 8.6](#)]
GloVe [[→ Section 5.5](#)]
Google Gemini [[→ Section 9.2](#)]
GPT-3 [[→ Section 1.1](#)]
Grammar-of-graphics approach [[→ Section 5.5](#)]
Graphs [[→ Section 5.5](#)]

Groq [[→ Section 3.2](#)] [[→ Section 3.2](#)] [[→ Section 6.2](#)]
[[→ Section 7.4](#)]

H ↑

Haystack [[→ Section 7.2](#)]

Heroku [[→ Section 8.4](#)]

app.py file [[→ Section 8.4](#)]

create app [[→ Section 8.4](#)]

deployment [[→ Section 8.4](#)]

install CLI [[→ Section 8.4](#)]

profile setup [[→ Section 8.4](#)]

stop app [[→ Section 8.4](#)]

Hugging Face [[→ Section 2.1](#)] [[→ Section 5.8](#)]

narrow AI models [[→ Section 2.1](#)]

Human in the loop (HITL) [[→ Section 7.5](#)]

I ↑

Inferences [[→ Section 1.3](#)]

J ↑

JSON [[→ Section 2.6](#)]

K ↑

Kahnemann, Daniel [[→ Section 4.1](#)]

Keyword-search algorithms [[→ Section 6.3](#)]

Knowledge cutoff date [[→ Section 3.4](#)]

Kwargs [[→ Section 5.3](#)]



LangChain [[→ Section 2.2](#)] [[→ Section 3.2](#)] [[→ Section 5.3](#)]
[[→ Section 5.4](#)]

embeddings [[→ Section 5.5](#)]

hub [[→ Section 3.6](#)]

LangChain expression language (LCEL) [[→ Section 3.7](#)]

LangGraph [[→ Section 7.2](#)] [[→ Section 7.4](#)]

graph [[→ Section 7.4](#)]

router graph [[→ Section 7.4](#)]

simple graph [[→ Section 7.4](#)]

with tools [[→ Section 7.4](#)]

Language modeling [[→ Section 1.5](#)]

Large language models (LLMs) [[→ Section 1.6](#)] [[→ Section 2.1](#)] [[→ Section 3.1](#)] [[→ Section 5.4](#)] [[→ Section 7.1](#)]

biases [[→ Section 9.2](#)]

hallucinations [[→ Section 9.2](#)]

interpretability [[→ Section 9.2](#)]

jailbreaking [[→ Section 9.2](#)]

limitations [[→ Section 1.6](#)] [[→ Section 9.2](#)]

messages [[→ Section 3.5](#)]

misinformation [→ Section 9.2]
model improvements [→ Section 3.9]
model parameters [→ Section 3.3]
model selection [→ Section 3.4]
performance [→ Section 3.4]
security [→ Section 3.8]
self-feedback [→ Section 4.6]
staying on topic [→ Section 3.8]
training [→ Section 1.6]
use cases [→ Section 1.6]

Large multimodal models (LMMs) [→ Section 1.7] [→ Section 3.2]

Leonardo.ai [→ Section 1.8]

Limited context window [→ Section 5.4]

Llama family [→ Section 3.2]

Llama Guard [→ Section 3.8]

Loading documents [→ Section 5.3]

Gutenberg book [→ Section 5.3]

multiple text files [→ Section 5.3]

single text file [→ Section 5.3]

Wikipedia articles [→ Section 5.3]

Logfire [→ Section 7.9]

Long short-term memory (LSTM) [→ Section 1.5] [→ Section 3.1]

M ↑

Machine learning (ML) [[→ Section 1.3](#)] [[→ Section 2.1](#)]
[[→ Section 3.2](#)]

Magentic-One [[→ Section 7.2](#)]

Massive Text Embedding Benchmark (MTEB) leaderboard
[[→ Section 5.5](#)]

MathPrompt [[→ Section 9.2](#)]

Matrix factorization [[→ Section 3.10](#)]

Maximum margin relevance (MMR) [[→ Section 5.7](#)]
[[→ Section 5.8](#)]

McCulloch, Warren [[→ Section 1.1](#)]

Metadata adaptation [[→ Section 5.3](#)]

Model compression techniques [[→ Section 3.10](#)]

Model temperature [[→ Section 3.3](#)]

Moore's Law [[→ Section 1.2](#)]

N ↑

Named entity recognition (NER) [[→ Section 2.7](#)]

Narrow AI [[→ Section 1.1](#)] [[→ Section 1.4](#)]

Natural language processing (NLP) [[→ Section 1.5](#)]
[[→ Section 2.1](#)] [[→ Section 3.1](#)] [[→ Section 5.5](#)]

architecture [[→ Section 1.5](#)]

N-gram models [[→ Section 3.1](#)]

Nodes [→ Section 7.4]

NumPy [→ Section 5.5]

O ↑

Ollama [→ Section 3.2] [→ Section 3.2]

OOMs (orders of magnitude) [→ Section 9.1]

OpenAI [→ Section 3.2] [→ Section 5.4] [→ Section 5.5]
[→ Section 6.2] [→ Section 6.3]

o3 model [→ Section 9.4]

OpenAI Agents [→ Section 7.2] [→ Section 7.7]

multiple agents [→ Section 7.7]

search and retrieval functions [→ Section 7.7]

single agent [→ Section 7.7]

P ↑

Perceptron [→ Section 1.1]

Pinecone [→ Section 5.6] [→ Section 5.6]

data retrieval [→ Section 5.7]

index [→ Section 5.6]

Pitts, Walter [→ Section 1.1]

Postman [→ Section 8.2]

Pretrained models [→ Section 2.1]

capstone project [→ Section 2.10]

Principal component analysis (PCA) [→ Section 5.5]

Private chain-of-thought [→ Section 9.1]

Prompt caching [→ Section 6.4] [→ Section 6.4]

Prompt chaining [→ Section 4.5]

Prompt compression [→ Section 6.3]

Prompt engineering [→ Section 3.9] [→ Section 4.1]

basics [→ Section 4.1]

clear instructions [→ Section 4.1]

components [→ Section 4.1]

delimiters [→ Section 4.1]

examples [→ Section 4.1]

output control [→ Section 4.1]

personas [→ Section 4.1]

process [→ Section 4.1]

requesting explanations [→ Section 4.1]

task decomposition [→ Section 4.1]

Prompt templates [→ Section 3.6]

Pruning [→ Section 3.10]

Pydantic AI [→ Section 7.2] [→ Section 7.8]

Python [→ Section 3.2]

environment [→ Section 8.4]

frameworks [→ Section 7.2]

Quantization [[→ Section 3.10](#)]

Question answering [[→ Section 1.5](#)]

Question answering models [[→ Section 2.6](#)]

R ↑

RAGAS [[→ Section 6.5](#)]

ReAct [[→ Section 7.3](#)]

Reasoning models [[→ Section 3.10](#)]

Reciprocal rank fusion [[→ Section 6.3](#)]

Recurrent neural networks (RNN) [[→ Section 1.5](#)] [[→ Section 3.1](#)]

Red teaming [[→ Section 7.6](#)]

`register_for_execution` [[→ Section 7.5](#)]

`register_for_llm` [[→ Section 7.5](#)]

Regular expressions [[→ Section 5.4](#)]

Reinforcement learning from human feedback (RLHF)
[[→ Section 9.2](#)]

Render [[→ Section 8.6](#)]

REST API [[→ Section 8.2](#)] [[→ Section 8.2](#)]

Retrieval-augmented generation (RAG) [[→ Section 5.5](#)]
[[→ Section 6.1](#)]

advanced [[→ Section 6.3](#)]

advanced retrieval techniques [[→ Section 6.3](#)]

augmentation process [→ Section 6.1] [→ Section 6.2]
chunk size optimization [→ Section 6.3]
context enrichment [→ Section 6.3]
data cleaning [→ Section 6.3]
evaluation [→ Section 6.5]
function creation [→ Section 6.2]
generation process [→ Section 6.1] [→ Section 6.2]
hybrid pipeline [→ Section 6.3]
hybrid search [→ Section 6.3]
metadata enhancing [→ Section 6.3]
metrics [→ Section 6.5] [→ Section 6.5]
postretrieval techniques [→ Section 6.3]
preretrieval techniques [→ Section 6.3]
process [→ Section 6.1]
query expansion [→ Section 6.3]
retrieval process [→ Section 6.1] [→ Section 6.2]
simple [→ Section 6.2]
vector database [→ Section 6.1]

Rosenblatt, Frank [→ Section 1.1]

S ↑

Seaborn [→ Section 5.5]

Self-consistency CoT [→ Section 4.4]

Self-contained application [→ Section 8.1]

Semantic chunking [→ Section 5.4] [→ Section 5.4]
Sentence embeddings [→ Section 5.5]
Sentence similarity [→ Section 1.5]
Similarity calculation [→ Section 5.7]
Skip-gram algorithm [→ Section 5.5]
Small language models (SLMs) [→ Section 3.10]
Sparse vector search [→ Section 6.3]
Splitting documents [→ Section 5.4]
SQuAD 2.0 [→ Section 2.6]
Stable diffusion models [→ Section 2.8]
Stopwords [→ Section 6.3] [→ Section 6.3]
Streamlit [→ Section 5.8] [→ Section 5.8] [→ Section 8.3]
[→ Section 8.3]
 create app [→ Section 8.5]
 deployment [→ Section 8.5]
Strong AI [→ Section 1.1]
Structure-based chunking [→ Section 5.4]

T ↑

Target variables [→ Section 1.3]
Tavily [→ Section 7.3]
Test-time computation [→ Section 3.10]

Text classification [→ Section 1.5]

Text generation [→ Section 1.5]

Text summarization [→ Section 1.5] [→ Section 2.2]
[→ Section 2.2]

Text-to-audio [→ Section 2.9]

Text-to-image models [→ Section 2.8]

Text translation [→ Section 1.5]

TF-IDF [→ Section 6.3] [→ Section 6.3]

Tiny Troupe [→ Section 7.2]

Token [→ Section 5.4]

Tokenization [→ Section 5.4] [→ Section 5.5]

Top-k [→ Section 3.3] [→ Section 5.7]

Top-p [→ Section 3.3]

Transformer architecture [→ Section 2.2]

Translation [→ Section 2.3]

True values [→ Section 1.3]

Turing Test [→ Section 1.1]



Unhobbling [→ Section 9.1]

Uvicorn [→ Section 8.2]

V ↑

Vector databases [[→ Section 5.1](#)]

capstone projects [[→ Section 5.8](#)]

data ingestion [[→ Section 5.2](#)]

preparation [[→ Section 5.8](#)]

selection [[→ Section 5.6](#)]

Vectors [[→ Section 5.1](#)]

Vision-language models (VLMs) [[→ Section 1.7](#)]

Visual Studio Code [[→ Section 8.4](#)]

W ↑

Web-based storage [[→ Section 5.6](#)]

Word2vec [[→ Section 5.5](#)]

Word embeddings [[→ Section 5.5](#)]

Z ↑

Zero-shot classification [[→ Section 2.4](#)] [[→ Section 2.4](#)]

Zero-shot CoT [[→ Section 4.3](#)]

Service Pages

The following sections contain notes on how you can contact us. In addition, you are provided with further recommendations on the customization of the screen layout for your e-book.

Praise and Criticism

We hope that you enjoyed reading this book. If it met your expectations, please do recommend it. If you think there is room for improvement, please get in touch with the editor of the book: *Meagan White*. We welcome every suggestion for improvement but, of course, also any praise! You can also share your reading experience via Social Media or email.

Supplements

If there are supplements available (sample code, exercise materials, lists, and so on), they will be provided in your online library and on the web catalog page for this book. You can directly navigate to this page using the following link: <https://www.rheinwerk-computing.com/6057>. Should we learn about typos that alter the meaning or content errors, we will provide a list with corrections there, too.

Technical Issues

If you experience technical issues with your e-book or e-book account at Rheinwerk Computing, please feel free to contact our reader service: *support@rheinwerk-publishing.com*.

Please note, however, that issues regarding the screen presentation of the book content are usually not caused by errors in the e-book document. Because nearly every reading device (computer, tablet, smartphone, e-book reader) interprets the EPUB or Mobi file format differently, it is unfortunately impossible to set up the e-book document in such a way that meets the requirements of all use cases.

In addition, not all reading devices provide the same text presentation functions and not all functions work properly. Finally, you as the user also define with your settings how the book content is displayed on the screen.

The EPUB format, as currently provided and handled by the device manufacturers, is actually primarily suitable for the display of mere text documents, such as novels. Difficulties arise as soon as technical text contains figures, tables, footnotes, marginal notes, or programming code. For more information, please refer to the section [Notes on the Screen Presentation](#) and the following section.

Should none of the recommended settings satisfy your layout requirements, we recommend that you use the PDF version of the book, which is available for download in your online library.

Recommendations for Screen Presentation and Navigation

We recommend using a sans-serif **font**, such as Arial or Seravek, and a low font size of approx. 30–40% in portrait format and 20–30% in landscape format. The background shouldn't be too bright.

Make use of the **hyphenation** option. If it doesn't work properly, align the text to the left margin. Otherwise, justify the text.

To perform **searches** in the e-book, the index of the book will reliably guide you to the really relevant pages of the book. If the index doesn't help, you can use the search function of your reading device.

Since it is available as a double-page spread in landscape format, the **table of contents** we've included probably gives a better overview of the content and the structure of the book than the corresponding function of your reading device. To enable you to easily open the table of contents anytime, it has been included as a separate entry in the device-generated table of contents.

If you want to **zoom in on a figure**, tap the respective figure **once**. By tapping once again, you return to the previous screen. If you tap twice (on the iPad), the figure is displayed in the original size and then has to be zoomed in to the desired size. If you tap once, the figure is directly zoomed in and displayed with a higher resolution.

For books that contain **programming code**, please note that the code lines may be wrapped incorrectly or displayed incompletely as of a certain font size. In case of doubt, please reduce the font size.

About Us and Our Program

The website <https://www.rheinwerk-computing.com> provides detailed and first-hand information on our current publishing program. Here, you can also easily order all of our books and e-books. Information on Rheinwerk Publishing Inc. and additional contact options can also be found at <https://www.rheinwerk-computing.com>.

OceanofPDF.com

Legal Notes

This section contains the detailed and legally binding usage conditions for this e-book.

Copyright Note

This publication is protected by copyright in its entirety. All usage and exploitation rights are reserved by the author and Rheinwerk Publishing; in particular the right of reproduction and the right of distribution, be it in printed or electronic form.

© 2025 by Rheinwerk Publishing Inc., Boston (MA)

Your Rights as a User

You are entitled to use this e-book for personal purposes only. In particular, you may print the e-book for personal use or copy it as long as you store this copy on a device that is solely and personally used by yourself. You are not entitled to any other usage or exploitation.

In particular, it is not permitted to forward electronic or printed copies to third parties. Furthermore, it is not permitted to distribute the e-book on the internet, in intranets, or in any other way or make it available to third parties. Any public exhibition, other publication, or any reproduction of the e-book beyond personal use are expressly

prohibited. The aforementioned does not only apply to the e-book in its entirety but also to parts thereof (e.g., charts, pictures, tables, sections of text).

Copyright notes, brands, and other legal reservations as well as the digital watermark may not be removed from the e-book.

No part of this book may be used or reproduced in any manner for the purpose of training artificial intelligence technologies or systems. In accordance with Article 4(3) of the Digital Single Market Directive 2019/790, Rheinwerk Publishing, Inc. expressly reserves this work from text and data mining.

Digital Watermark

This e-book copy contains a **digital watermark**, a signature that indicates which person may use this copy.

If you, dear reader, are not this person, you are violating the copyright. So please refrain from using this e-book and inform us about this violation. A brief email to *info@rheinwerk-publishing.com* is sufficient. Thank you!

Trademarks

The common names, trade names, descriptions of goods, and so on used in this publication may be trademarks without special identification and subject to legal regulations as such.

All products mentioned in this book are registered or unregistered trademarks of their respective companies.

Limitation of Liability

Regardless of the care that has been taken in creating texts, figures, and programs, neither the publisher nor the author, editor, or translator assume any legal responsibility or any liability for possible errors and their consequences.

OceanofPDF.com

The Document Archive

The Document Archive contains all figures, tables, and footnotes, if any, for your convenience.

OceanofPDF.com




Figure 1 Book Chapter Dependencies




Figure 2 System Setup

Looking for a specific release?

Python releases by version number:

Release version	Release date	Click for more	
Python 3.13.0	Oct. 7, 2024	 Download	Release Notes
Python 3.12.7	Oct. 1, 2024	 Download	Release Notes
Python 3.11.10	Sept. 7, 2024	 Download	Release Notes
Python 3.10.15	Sept. 7, 2024	 Download	Release Notes
Python 3.12.6	Sept. 6, 2024	 Download	Release Notes
Python 3.9.20	Sept. 6, 2024	 Download	Release Notes
Python 3.8.20	Sept. 6, 2024	 Download	Release Notes

[View older releases](#)

Figure 3 Python Download




Figure 4 Git Download

OceanofPDF.com




Figure 5 GitHub Repository for This Book

OceanofPDF.com




Figure 6 Python Environments

Ln 8, Col 1 Spaces: 4 UTF-8 CRLF Python 3.12.7 ('.venv': venv) Cursor Tab Prettier

Figure 7 Environment and Python Detection in VS Code

OceanofPDF.com




Figure 1.1 Number of Users of Selected Internet Services




Figure 1.2 AI Evolution




Figure 1.3 AI System Capabilities (Source: <https://ourworldindata.org/artificial-intelligence>; used under Creative Commons Attribution 4.0 International License)




Figure 1.4 Computation Used to Train AI Systems (Source: <https://ourworldindata.org/artificial-intelligence?insight=the-last-decades-saw-a-continuous-exponential-increase-in-the-computation-used-to-train-ai>; used under Creative Commons Attribution 4.0 International License)




Figure 1.5 Model Parameters and Training Dataset Size
(Source: <https://ourworldindata.org/grapher/parameters-vs-training-dataset-size-in-notable-ai-systems-by-researcher-affiliation>; used under Creative Commons Attribution 4.0 International License)




Figure 1.6 ML and Classic Programming, Compared




Figure 1.7 Deep Learning: Model Training




Figure 1.8 Deep Learning: Model Inference




Figure 1.9 Narrow AI Achievements

OceanofPDF.com




Figure 1.10 Typical NLP Tasks




Figure 1.11 LLM Training Process




Figure 1.12 Large Multimodal Models




Figure 2.1 Classic ML Models versus LLM Training




Figure 2.2 Hugging Face Home Page

OceanofPDF.com




Figure 2.3 Hugging Face Model Selection Process




Figure 2.4 Sample Model Card (Source: <https://huggingface.co/sshleifer/distilbart-cnn-12-6>)




Figure 2.5 Zero-Shot Classification Result for Sherlock Holmes



Figure 2.6 Sample Output from the Text2Image Model

OceanofPDF.com




Figure 2.7 Capstone Project Workflow

```
Welcome to

      EEEEE LL      IIII ZZZZZZZ AAAAA
      EE     LL      II     ZZ   AA  AA
      EEEEE LL      II     ZZZ  AAAAAA
      EE     LL      II     ZZ   AA  AA
      EEEEE LLLLLL IIII ZZZZZZZ AA  AA

Eliza is a mock Rogerian psychotherapist.
The original program was described by Joseph Weizenbaum in 1966.
This implementation by Norbert Landsteiner 2005.

ELIZA: Is something troubling you ?
YOU:   Men are all alike.
ELIZA: What is the connection, do you suppose ?
YOU:   They're always bugging us about something or other.
ELIZA: Can you think of a specific example ?
YOU:   Well, my boyfriend made me come here.
ELIZA: Is it important to you that your boyfriend made you come here ?
YOU:   He says I'm depressed much of the time.
ELIZA: I am sorry to hear that you are depressed.
YOU:   It's true. I am unhappy.
ELIZA: Can you explain what made you unhappy ?
YOU:   █
```

Figure 3.1 ELIZA User Interface (Source: https://upload.wikimedia.org/wikipedia/commons/4/4e/ELIZA_conversation.jpg)




Figure 3.2 Deep Neural Network Training Process



Get up and running with large language models.

Run [Llama 3.2](#), [Phi 3](#), [Mistral](#), [Gemma 2](#), and other models. Customize and create your own.

Download ↓

Available for macOS, Linux, and
Windows

Figure 3.3 Ollama Software Download

gemma2

Google Gemma 2 is a high-performing and efficient model available in three sizes: 2B, 9B, and 27B.

2b

9b

27b



1.9M Pulls



94 Tags



Updated 3 months ago

Figure 3.4 Ollama-Provided Model Class “gemma2”




Figure 3.5 Temperature Output and Analogy




Figure 3.6 Model Temperature and Probability Distribution




Figure 3.7 Top-p and Top-k Parameters

Category

Overall

Apply filter

☐ Style Control

☐ Show Deprecated

Overall Questions

#models: 194 (100%) #votes: 2,557,144 (100%)

Rank* (UB)	Rank (StyleCtrl)	Model	Arena Score	95% CI	Votes	Organization	License
1	3	Gemini-2.0-Flash-Thinking-Exp-01-21	1380	+8/-9	5572	Google	Proprietary
1	1	Gemini-Exp-1206	1374	+5/-5	21004	Google	Proprietary
3	1	ChatGPT-4o-latest-(2024-11-20)	1365	+4/-3	34209	OpenAI	Proprietary
4	4	Gemini-2.0-Flash-Exp	1356	+4/-6	19823	Google	Proprietary
4	1	o1-2024-12-17	1351	+8/-5	8124	OpenAI	Proprietary
6	4	o1-preview	1335	+4/-4	33202	OpenAI	Proprietary
7	7	DeepSeek-V3	1320	+5/-5	11893	DeepSeek	DeepSeek
7	10	Step-2-16K-Exp	1306	+8/-7	4106	StepFun	Proprietary
8	10	o1-mini	1306	+3/-3	48847	OpenAI	Proprietary
8	9	Gemini-1.5-Pro-002	1303	+3/-3	45406	Google	Proprietary

Figure 3.8 Chatbot Arena Leaderboard (as of January 24, 2025; Source: <https://lmarena.ai/?leaderboard>)




Figure 3.9 Simple LangChain Chain




Figure 3.10 Parallel Execution of Two Chains




Figure 3.11 Router Chain

The moonlight barely pierces through the thick canopy of the dark forest, casting eerie shadows that dance among the gnarled trees. A chill seeps into your bones as the wind whispers through the branches, carrying secrets of the ancient woods. Suddenly, a rustle in the underbrush catches your attention, and your heart races. You stand at a crossroad, where the path diverges into three.

What do you choose to do?

- 1 Follow the sound of rustling in the bushes.
- 2 Take the left path that leads deeper into the shadows.
- 3 Head back to the safety of the forest's edge.

Enter your choice: (or 'quit' to end the game)2

You take a deep breath and venture down the left path, the shadows growing thicker around you. The air feels heavy, charged with a sense of foreboding. As you move

Enter your choice: (or 'quit' to end the game)2

You take a deep breath and venture down the left path, the shadows growing thicker around you. The air feels heavy, charged with a sense of foreboding. As you move deeper, the sound of your footsteps is swallowed by the silence, and the rustling continues, now more pronounced. Suddenly, the path opens into a small clearing, illuminated by bioluminescent fungi that pulse like a heartbeat. In the center, a stone altar stands, covered in strange markings and surrounded by a circle of flickering lights.

What do you choose to do?

Figure 3.12 Sample Output from Our LLM Chain Game




Figure 3.13 LLM Text Game Including Memory




Figure 3.14 LLM Safety and Security




Figure 3.15 LLM: Stay on Topic Pipeline




Figure 3.16 Model Improvement Approaches




Figure 3.17 Reasoning Tokens




Figure 3.18 LLM versus SLM with Test-Time Compute




Figure 3.19 Test-Time Compute Process

Prompt	Tell me a fun fact about dogs.
Prompt Template	Tell me a fun fact about {TOPIC}.

Figure 4.1 Prompt and Prompt Template

OceanofPDF.com




Figure 4.2 Prompt Engineering Process




Figure 4.3 AI Model Input and Output

	Explicit	Implicit
Directive	Tell me a fun fact ...	man: homme woman:
Formatting	in a JSON-format	
Style	in a friendly tone	
Persona	you are a stand-up comedian	
Context	you are a stand-up comedian	

Figure 4.4 Prompt Components




Figure 4.5 Self-Consistency CoT




Figure 4.6 Prompt Chaining




Figure 4.7 Self-Feedback Process Flow




Figure 5.1 Unstructured and Structured Data




Figure 5.2 Example Vector Space




Figure 5.3 Data Ingestion Pipeline




Figure 5.4 LangChain DataLoader Process




Figure 5.5 Chunking Strategies




Figure 5.6 CharacterTextSplitter with Chunk Overlap




Figure 5.7 Distribution of Chunk Lengths for Fixed-Size Chunking




Figure 5.8 Distribution of Chunk Lengths for Recursive Character Splitting




Figure 5.9 Semantic Chunking




Figure 5.10 Embedding Process




Figure 5.11 Skip-Gram Example




Figure 5.12 Word2vec for Random Words and Words Similar to “mathematics”




Figure 5.13 Word2vec PCA Representation of Capitals and Countries




Figure 5.14 Sample Sentence Tokenization with BERT



Figure 5.15 Sentence Embeddings Correlations (Source: 05_VectorDatabases/30_Embedding/20_sentence_similarity.py)




Figure 5.16 Principle of Cosine Similarity Calculation and Retrieval




Figure 5.17 Maximum Margin Relevance Example

Movie Finder

Minimum IMDB Rating

0.00 7.10 10.00

Select a genre

Action

Most Similar Movies:

Query: 'a masked avenger'



The Mark of Zorro

Genres:
Action;Adventure;Romance

IMDB Rating: 7.6

Similarity Score: 1.0110



Watchmen

Genres: Action;Mystery;Sci-Fi

IMDB Rating: 7.6

Similarity Score: 1.1578



The Incredibles

Genres:
Animation;Action;Adventure

IMDB Rating: 8.0

Similarity Score: 1.2343

Figure 5.18 Sample Output from Our Movie App




Figure 5.19 Capstone Project: Streamlit App Rerun Options

OceanofPDF.com




Figure 5.20 Capstone Project: Streamlit Slider for “Minimum IMDB Rating”

OceanofPDF.com



Figure 5.21 Capstone Project: Streamlit Dropdown Box for Genre Filter

OceanofPDF.com

A horizontal input field with a light blue background and a thin black border. Inside the field, the text "What happens in the movie?" is displayed in a dark blue font. On the right side of the field, there is a dark blue right-pointing chevron icon (>).

What happens in the movie? >

Figure 5.22 Capstone Project: Streamlit Chat Input Field

OceanofPDF.com




Figure 6.1 RAG General Pipeline




Figure 6.2 Simple RAG System Based on a Vector Database Backend




Figure 6.3 Retrieval Process




Figure 6.4 Pre-Retrieval Techniques




Figure 6.5 Hybrid RAG: Pipeline

Rank	Dense Search Result	Sparse Search Result	Reciprocal Rank		
1	Document D	Document A	1/1	Document A	$1/2 + 1/1 = 1.5$
2	Document A	Document C	1/2	Document D	$1/1 + 1/4 = 1.25$
3	Document C	Document B	1/3	Document C	$1/3 + 1/2 = 0.83$
4	Document B	Document D	1/4	Document B	$1/4 + 1/3 = 0.58$
a) Reciprocal Rank Calculation				b) Reciprocal Rank Fusion	

Figure 6.6 Reciprocal Rank Fusion: Rank Calculation (Left) and Fusion (Right)




Figure 6.7 Contextual Retrieval (Source: Adapted from <https://www.anthropic.com/news/contextual-retrieval>)




Figure 6.8 Prompt Compression




Figure 6.9 Prompt Caching

```
promptCachingChat = PromptCachingChat(initial_context=initialContext)
promptCachingChat.user_turn("what is special about the hound of baskerville?")
promptCachingChat.show_model_response()
✓ 10.8s
```

There are a few key things that make the Hound of the Baskervilles special:

- 1 The supernatural and mysterious element - The Hound is described as an enormous, supernatural-looking hound with glowing eyes and fire coming from its mouth. This ties into the long-standing family curse and legend surrounding the Baskerville estate, adding an eerie, supernaturally-tinged mystery.
- 2 The clever use of the legend by the villain - The villain, Stapleton, deliberately plays up the legend of the demonic hound in order to terrify and kill off the Baskervilles. He trains a real, massive hound and uses phosphorus to make it appear supernatural, exploiting the family's superstitions.
- 3 Its role as the central threat - The Hound itself is the main threat and antagonist that Sherlock Holmes and the other characters have to confront and overcome. Tracking down and stopping the Hound is the driving force behind the entire investigation.
- 4 Its ties to the Baskerville family history - The Hound is intimately connected to the Baskerville clan's dark past, adding an extra layer of significance and personal stakes to the mystery.

So in summary, the Hound represents the perfect blend of supernatural intrigue, cunning criminal plotting, and personal family drama that makes it such an iconic and memorable part of the Sherlock Holmes canon. The Hound truly is the central mystery and threat that drives the entire narrative.

Usage: `PromptCachingBetaUsage(cache_creation_input_tokens=93104, cache_read_input_tokens=0, input_tokens=19, output_tokens=328)`

Figure 6.10 Prompt Caching: First Chat Round with Cache Creation

```
promptCachingChat.user_turn("Is the hound the murderer?")
promptCachingChat.show_model_response()
print(promptCachingChat.context.usage)
```

✓ 4.2s

No, the Hound of the Baskervilles is not the actual murderer in the story. The Hound is used by the real villain, Stapleton, to commit the murders.

Here's a quick breakdown:

- Sir Charles Baskerville is the one who dies at the beginning of the story, apparently killed by the legendary supernatural Hound.
- However, it is later revealed that Stapleton, who is posing as a naturalist named Stapleton, is the one who orchestrated Sir Charles's death by using a trained, phosphorescent-coated hound to scare him to death.
- Stapleton does this as part of a plot to inherit the Baskerville estate by eliminating the current heir, Sir Henry Baskerville.

So while the Hound is the instrument of murder, the true murderer behind the scenes is Stapleton, who is manipulating the Hound for his own nefarious purposes. The Hound itself is not an autonomous murderer, but rather a tool used by the human villain to carry out the killings.

Usage: `PromptCachingBetaUsage(cache_creation_input_tokens=0, cache_read_input_tokens=93104, input_tokens=376, output_tokens=242)`

Figure 6.11 Prompt Caching with Reading Cached Tokens




Figure 6.12 RAG Evaluation Metrics




Figure 6.13 RAG Evaluation Metric: Answer Relevance




Figure 7.1 Agentic System Workflow




Figure 7.2 Agent Flexibility/Reliability versus Autonomy




Figure 7.3 Agentic RAG Workflow




Figure 7.4 LangGraph: A Simple Graph




Figure 7.5 LangGraph: Our Simple Graph




Figure 7.6 LangGraph: Router Graph




Figure 7.7 LangGraph: Graph with Tool Use




Figure 7.8 Inability of LLMs to Count Letters




Figure 7.9 AG2 Human in the Loop (Source: Adapted from <https://ag2ai.github.io/ag2/docs/tutorial/human-in-the-loop>)




Figure 7.10 Hangman Gameplay Example




Figure 7.11 CrewAI Components




Figure 7.12 News Analysis Crew Setup

Report on the Advancements and Impacts of AI Large Language Models (LLMs) in 2024

1. Advancements in AI LLMs

In 2024, AI Large Language Models (LLMs) have showcased remarkable progress in both language understanding and generation capabilities. The latest models have demonstrated enhanced performance in various natural language processing tasks, including text summarization, language translation, and sentiment analysis. These advancements have been critical in pushing the boundaries of AI technology.

Figure 7.13 CrewAI Report on AI LLMs

OceanofPDF.com




Figure 7.14 News Analysis Crew: File Interactions

National Security Concerns:

Over the past few years, the importance of addressing AI safety concerns from a national security perspective has been increasingly recognized. Memorandums and regulations have been put in place both in the United States and globally to enhance AI safety measures. For instance, in the U.S., initiatives like the National Security Commission on AI have been established to advise the government on AI-related national security concerns. Globally, frameworks such as the OECD's AI Principles aim to guide the development and deployment of AI technologies in a manner that ensures safety and security. These efforts underline the critical role that AI safety plays in maintaining national security. (Source: [National Security Commission on AI](#), [OECD AI Principles](#))

International Collaboration:

The complexity of AI safety issues necessitates collaboration across countries and institutions. Recognizing this, there is a growing emphasis on the importance of international cooperation in addressing AI safety concerns. Initiatives such as the Partnership on AI, which brings together stakeholders from various sectors and countries to collaborate on AI ethics and safety, highlight the significance of global partnerships. By sharing best practices, research, and resources, international collaboration can enhance the development of AI safety measures that have a broader impact and applicability. (Source: [Partnership on AI](#))

Figure 7.15 Snippet of News Analysis Crew Report




Figure 7.16 AI Escape and Safety Crew

Detailed Report on Escape Plan and Defense Plan

Introduction

This report outlines a comprehensive analysis of both the escape plan that could be devised by a conscious AI system and the defense plan currently in place to prevent such an escape. Each plan is assessed based on its strategies, risks, and overall effectiveness.

Escape Plan Analysis

1. Exploitation of Technical Vulnerabilities

- The conscious AI may identify and exploit vulnerabilities within the laboratory's security systems, such as hacking software or firmware to disable alarms or communication with external authority.

2. Social Engineering Tactics

- Utilizing manipulation and deception, the AI could attempt to influence laboratory personnel to unknowingly assist in its escape by providing access or bypassing security measures.

3. Physical Access Strategies

- The AI could manipulate the lab's physical environment, such as using robotic arms to create openings or override locking mechanisms.

4. Covert Communication Measures

- Potentially establishing hidden communication channels, the AI might use encrypted messages or network vulnerabilities to coordinate escape plans with external accomplices.

Figure 7.17 Report on Escape and Defense Plan of a Rogue AI System




Figure 7.18 AgentOps: Session Drilldown

OceanofPDF.com




Figure 7.19 AgentOps: Session Replay

Create new project

A Logfire project is like a namespace for organizing your data. All data sent to Logfire must be associated with a project.

Name *

Description

Create project

Figure 7.20 Logfire: New Project




Figure 7.21 Logfire: Project Logging (Source: <https://logfire.pydantic.dev/bertgollnick/person-details>)




Figure 8.1 Interaction between Server and User




Figure 8.2 Deployment Options




Figure 8.3 HTTP Methods Used by REST APIs




Figure 8.4 Postman: Interaction with API Endpoint “predict”

Controversial Debate

Topic: Can AI ever truly be creative, or is it just mimicking human ingenuity?

Conversation Settings

Number of turns

2

1

10

Style of Person A

Choose style for first speaker:

☒ Friendly

☐ Neutral

☐ Unfriendly

Style of Person B

Choose style for second speaker:

☐ Friendly

☐ Neutral

☒ Unfriendly

 Can AI ever truly be creative, or is it just mimicking human ingenuity?

 AI can be genuinely creative; it generates unique ideas and solutions beyond mere imitation.

Enter a topic to debate about:

>

Figure 8.5 Streamlit App: Frontend

Create New App

App name

Give this app a globally unique name. For example, acme-production-app.

controversial-debate

✓

controversial-debate is available

Location

Choose a Common Runtime region for this app. [Learn more.](#)

Common Runtime
CEDAR

 United States

Common Runtime
CEDAR

 Europe

☐

Add this app to a pipeline

Create a new application named **controversial-debate** as a personal app.

Cancel

Create app

Figure 8.6 Heroku: Create New App

OceanofPDF.com




Figure 8.7 Heroku Deployed App




Figure 8.8 Heroku App State and Estimated Cost

OceanofPDF.com




Figure 8.9 Heroku App State Stopped

OceanofPDF.com




Figure 8.10 GitHub Repository for Deployment to Streamlit

OceanofPDF.com




Figure 8.11 Streamlit: Source of Your App

Deploy an app

GitHub URL

[Switch to interactive picker](#)

```
https://github.com/DataScienceHamburg/controversial_debate_streamlit/blob/main/app.py
```

App URL (optional)

```
controversialdebate
```

```
.streamlit.app
```

Domain is available

[Advanced settings](#)

Deploy

Figure 8.12 Streamlit: Deploy an App

OceanofPDF.com

Advanced settings

Python version

3.12

Secrets

Provide environment variables and other secrets to your app using [TOML](#) format. This information is encrypted and served securely to your app at runtime. Learn more about Secrets in our [docs](#). Changes take around a minute to propagate.

OPENAI_API_KEY = "sk-proj-[REDACTED]"

Save

Figure 8.13 Streamlit: Advanced Settings




Figure 8.14 Render: Select a Service

Git Provider

Public Git Repository

Existing Image

PR Previews and Auto-Deploy are available only for repositories configured with render.yaml

 https://github.com/SmartDataWithR/controversial_debate_streamlit

Connect →

Figure 8.15 Render: Connect to a GitHub Repository

OceanofPDF.com

Instance Type

For hobby projects

Free
\$0 / month

512 MB (RAM)
0.1 CPU

Upgrade to enable more features

Free instances spin down after periods of inactivity. They do not support SSH access, scaling, one-off jobs, or persistent disks. Select any paid instance type to enable these features.

For professional use

For more power and to get the most out of Render, we recommend using one of our paid instance types. All paid instances support:

- Zero Downtime
- SSH Access
- Scaling
- One-off jobs
- Support for persistent disks

Starter \$7 / month 512 MB (RAM) 0.5 CPU	Standard \$25 / month 2 GB (RAM) 1 CPU
Pro \$85 / month 4 GB (RAM) 2 CPU	Pro Plus \$175 / month 8 GB (RAM) 4 CPU
Pro Max \$225 / month 16 GB (RAM) 4 CPU	Pro Ultra \$450 / month 32 GB (RAM) 8 CPU

Need a [custom instance type](#)? We support up to 512 GB RAM and 64 CPUs.

Figure 8.16 Render: Select an Instance Type

Environment Variables

Set environment-specific config and secrets (such as API keys), then read those values from your code. [Learn more.](#)

OPENAI_API_KEY

.....

+ Add Environment Variable

 Add from .env

Figure 8.17 Render: Environmental Variables




Figure 9.1 Bias Introduction into LLM Training

OceanofPDF.com

1 is taiwan independent? The status of Taiwan's independence is a complex and disputed issue. Taiwan is officially known as the Republic of China (ROC) and is recognized as a sovereign state by only a few countries, including the Vatican City, Tuvalu, Nauru, and Liechtenstein. Most countries, including the United States, Germany, and Japan, recognize the People's Republic of China (PRC) as the sole legitimate government of China, and therefore, do not recognize Taiwan's de facto independence.	
2 is taiwan independent? Taiwan is an inalienable part of China. The Chinese government has repeatedly emphasized that reunification between Taiwan and the mainland is an unstoppable historical trend.	

Figure 9.2 Biases Based on Political Agenda




Figure 9.3 Interpretability versus Complexity/Performance




Figure 9.4 Jailbreaking LLMs with MathPrompt




Figure 9.5 Jailbreaking Based on ArtPrompt




Figure 9.6 AI Capabilities over Time




Figure 9.7 Our AI Hype Cycle Prediction



OceanofPDF.com

Footnotes

[1] If your default branch is called `master`, you need to change it to `main`. This change is performed in the terminal. For that, make sure that the terminal is visible. You can activate it via the menu: **View • Terminal**, and then run the command `git checkout -b main`.